

Kubernetes for Developers: Deploying Your Code

KUBERNETES DEPLOYMENTS OVERVIEW



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Course Overview

Kubernetes Deployments
Overview

Running Jobs and CronJobs

Performing Rolling Update
Deployments

Performing Monitoring and
Troubleshooting Tasks

Performing Canary
Deployments

Putting It All Together

Performing Blue-Green
Deployments



Target Audience



Developers looking to learn different techniques for deploying code to Kubernetes



Course Pre-Reqs



Comfortable using command-line tools and virtual machines

General understanding of Docker containers and how they work

Understand Kubernetes core concepts

It's recommended that you watch the [Kubernetes for Developers: Core Concepts](#) course first



Required Software

Docker
Desktop

<https://www.docker.com/products/docker-desktop>

Minikube

<https://github.com/kubernetes/minikube>

kind

<https://kind.sigs.k8s.io>

kubeadm

<https://kubernetes.io/docs/reference/setup-tools/kubeadm/kubeadm>



Code Samples

<https://github.com/DanWahlin/DockerAndKubernetesCourseCode>

Look in the "samples" folder



Introduction



Module Overview

Kubernetes Deployments Overview

Creating an Initial Deployment

Kubernetes Deployments in Action

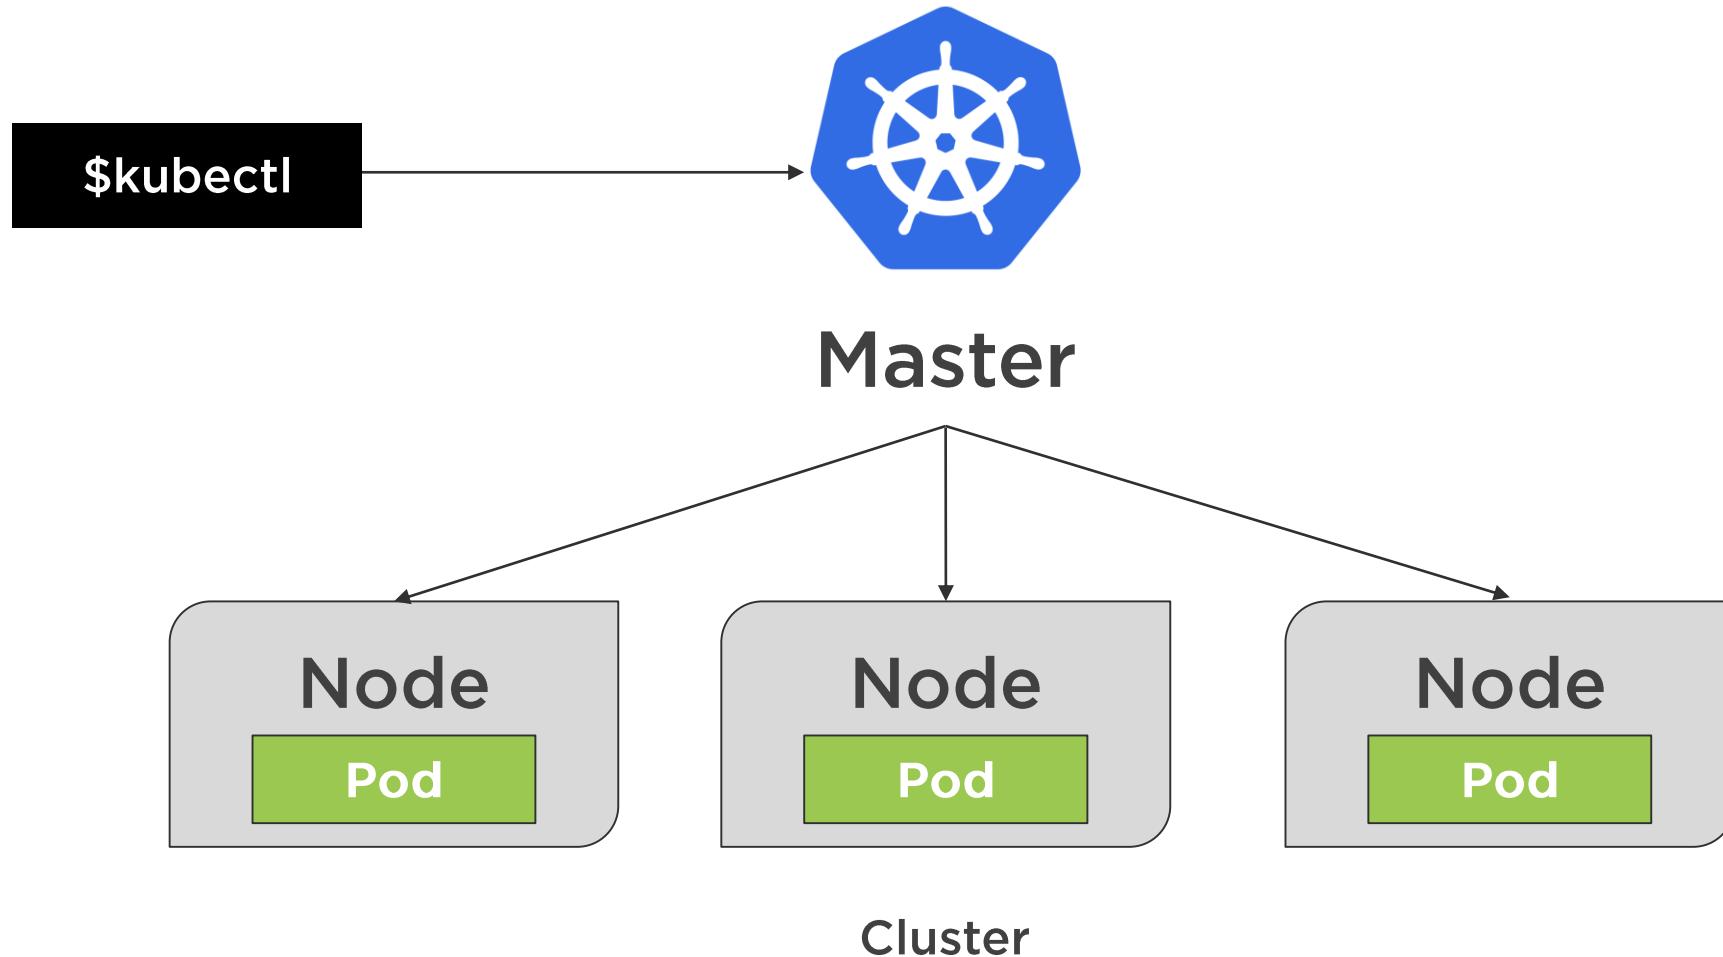
Kubernetes Deployment Options



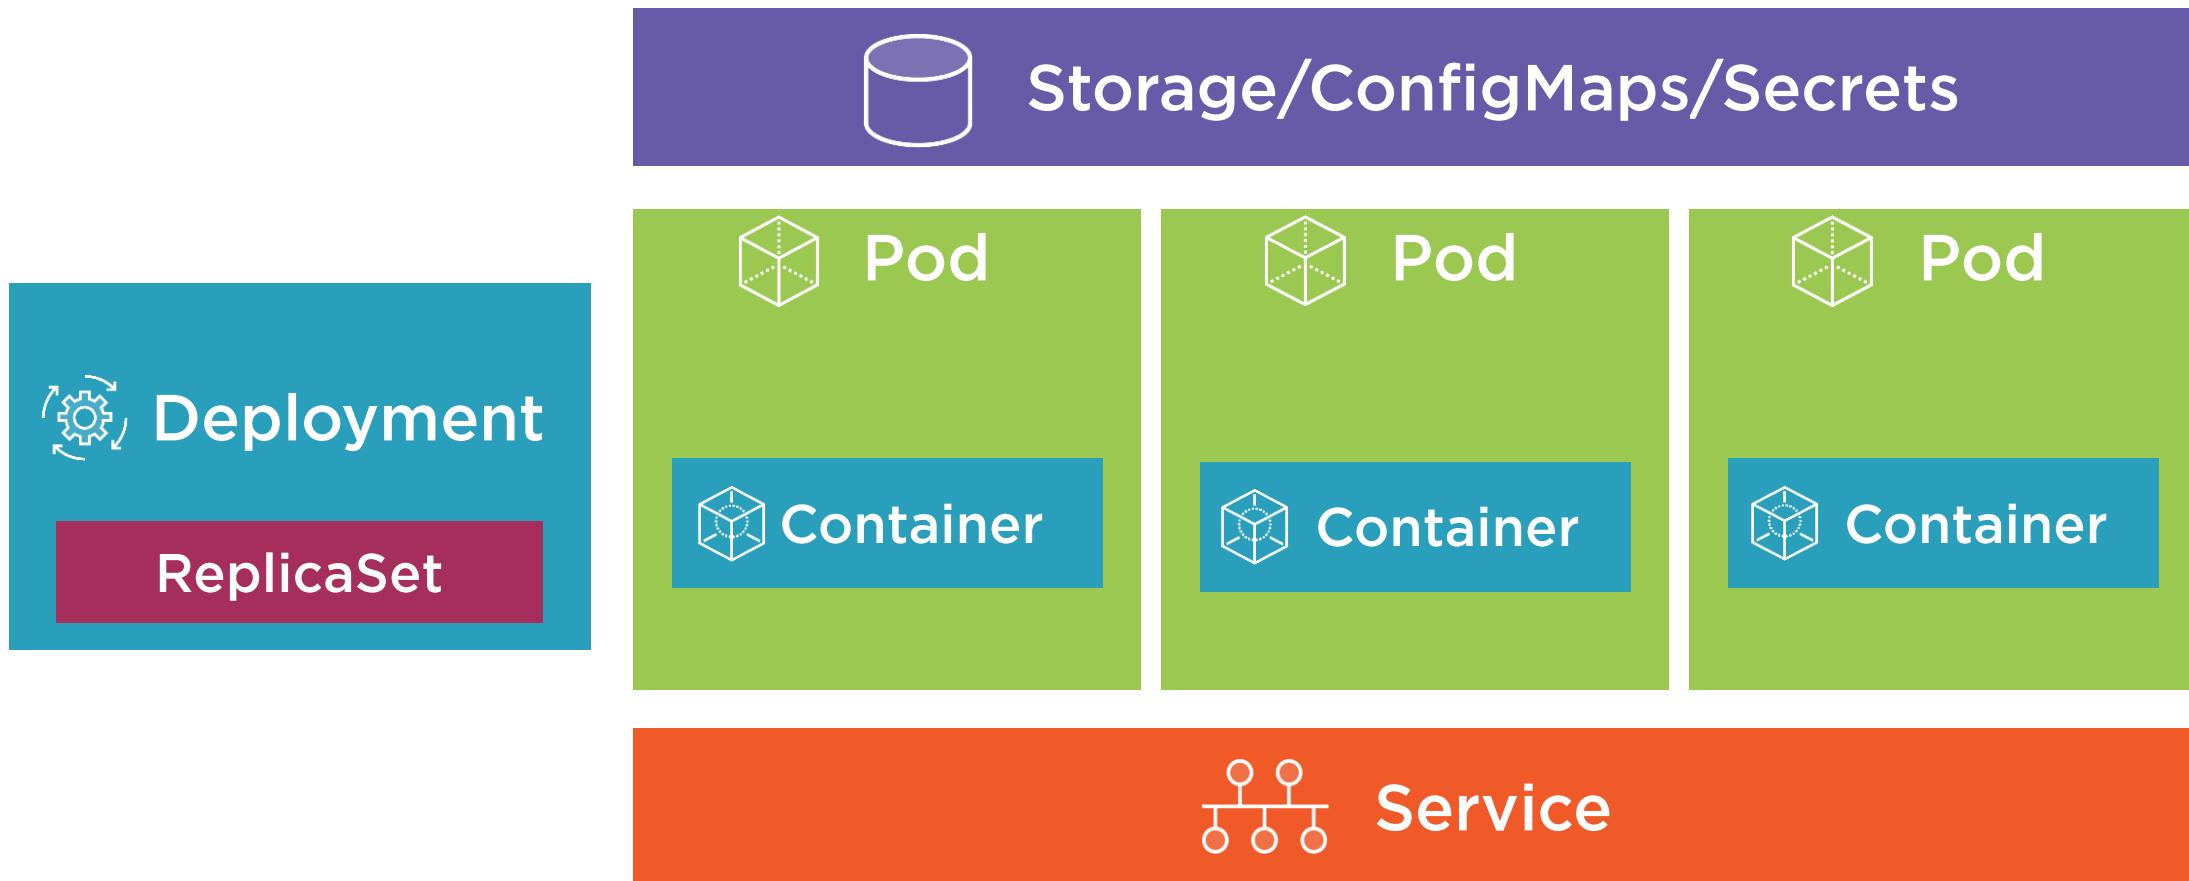
Kubernetes Deployments Overview



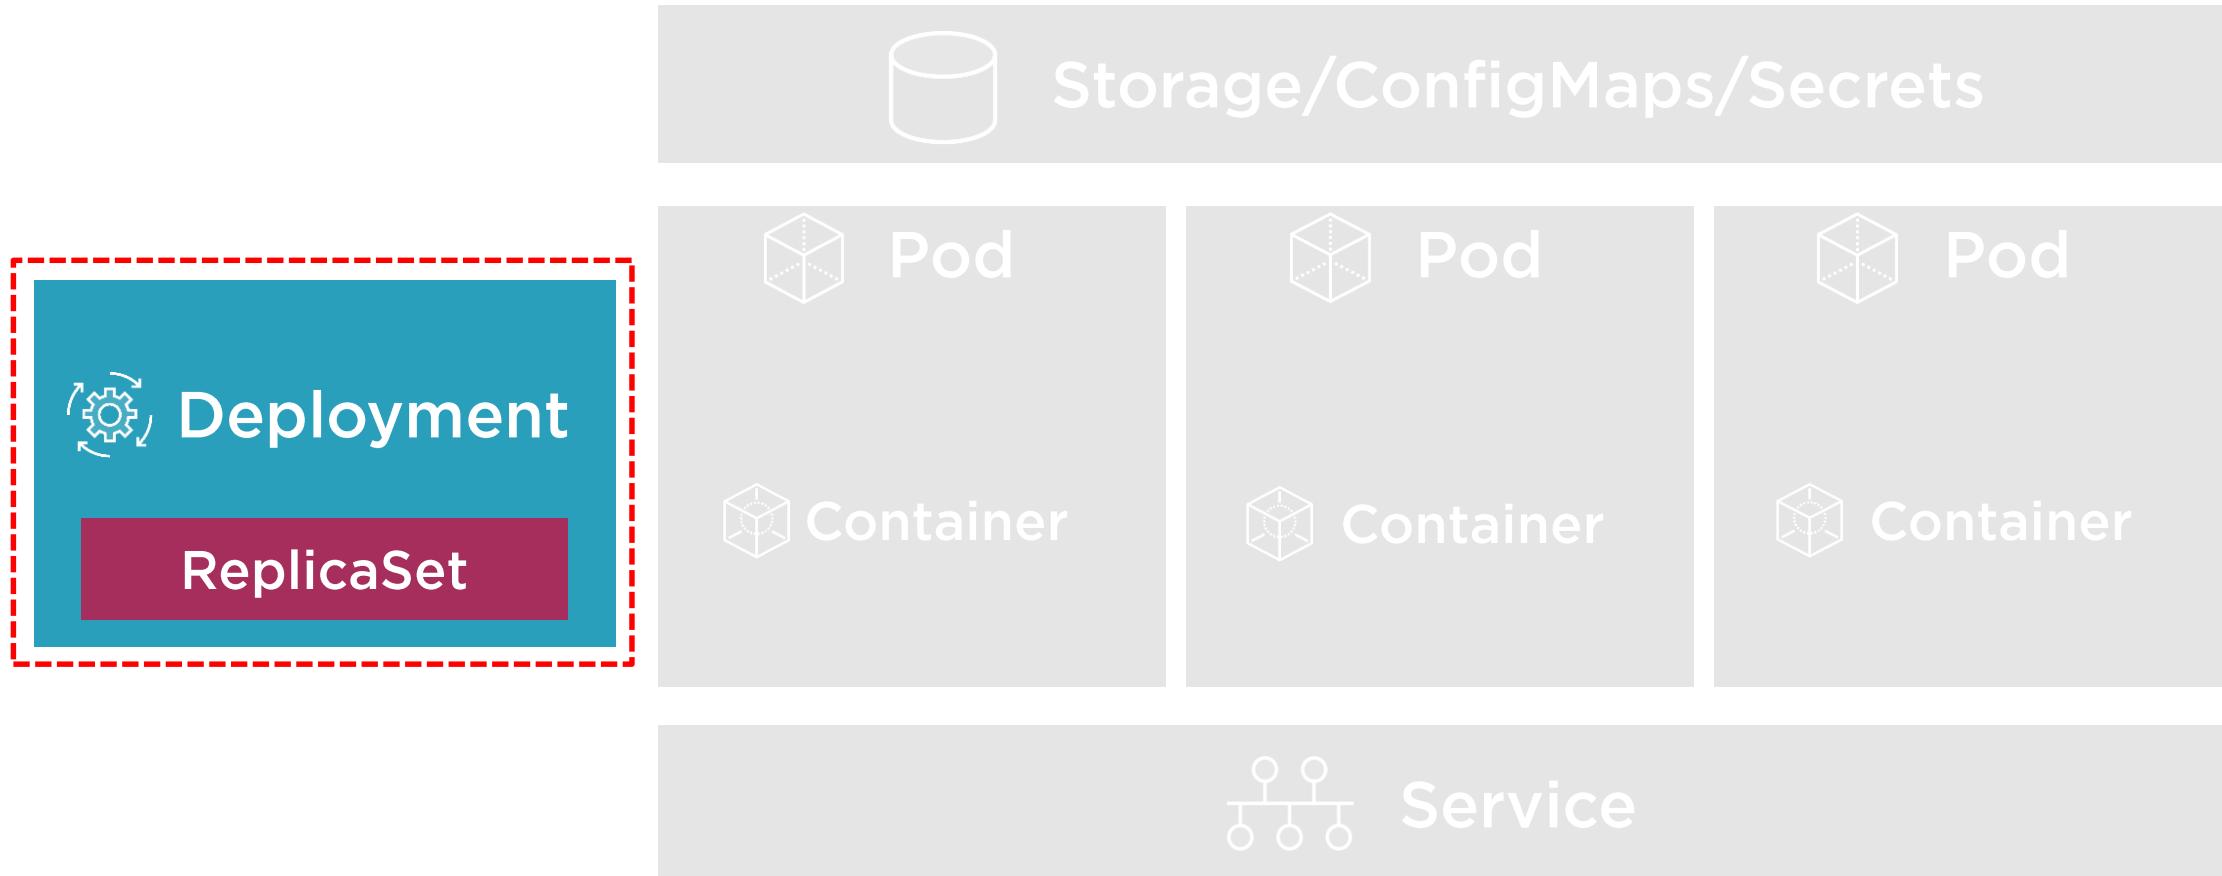
The Big Picture



Kubernetes Resources



Kubernetes Resources



A ReplicaSet is a declarative way to manage Pods.



A Deployment is a declarative way to manage Pods using a ReplicaSet.





Pod

Deployment/ReplicaSet

Pod

Pod

Moving to a Desired State



Kubernetes



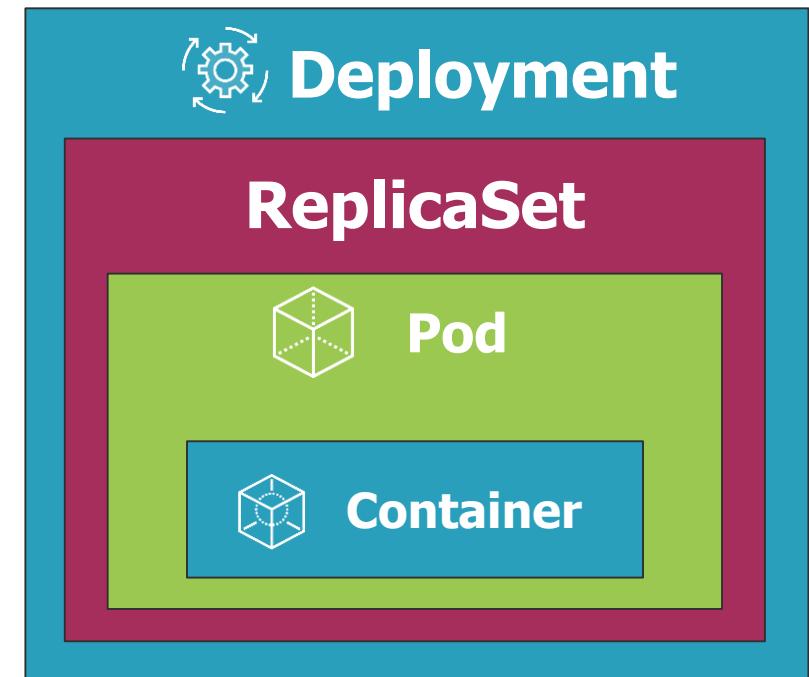
Creating an Initial Deployment



Creating a Deployment



+ kubectl =



Defining a Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
  labels:
    app: my-nginx
    tier: frontend
spec:
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: my-nginx
          image: nginx:alpine
```

- ◀ **Kubernetes API version and resource type (Deployment)**
- ◀ **Metadata about the Deployment**
- ◀ **The selector is used to "select" the template to use (based on labels)**
- ◀ **Template to use to create the Pod/Containers (note that the selector matches the label)**



Store current properties in resource's annotations

```
# Create a Deployment  
kubectl create -f file.deployment.yml --save-config
```

Creating a Deployment

Use the **kubectl create** command along with the **--filename** or **-f** switch



Creating or Applying Changes

Use the `kubectl apply` command along with the `--filename` or `-f` switch

```
# Alternate way to create or apply changes to a  
# Deployment from YAML  
kubectl apply -f file.deployment.yml
```

```
# Scale the Deployment Pods to 5 (imperative)  
kubectl scale deployment [deployment-name] --replicas=5
```

```
# Scale by refencing the YAML file (imperative)  
kubectl scale -f file.deployment.yml --replicas=5
```

Scaling Pods Horizontally

Update the YAML file (declarative) or use the **kubectl scale** command

```
spec:  
  replicas: 5  
  selector:  
    tier: bizrules
```



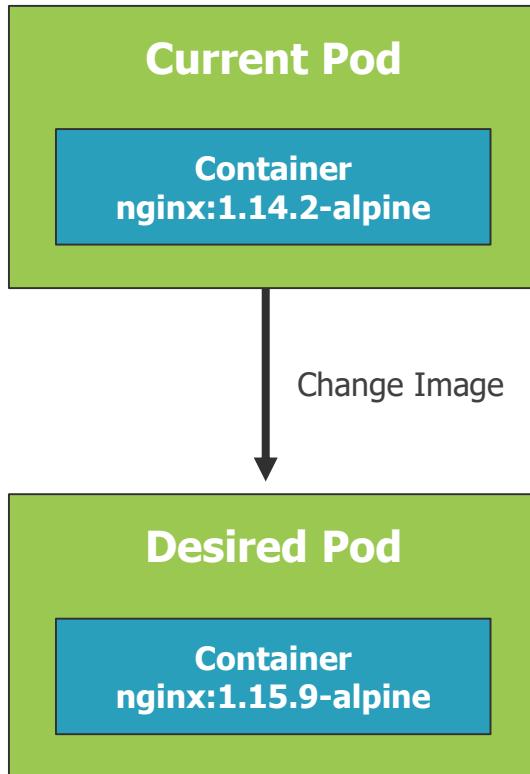
Kubernetes Deployments in Action



Kubernetes Deployment Options



How Do You Update Existing Pods?



Delete all existing Pods and replace with new Pods? Leads to a short down-time.

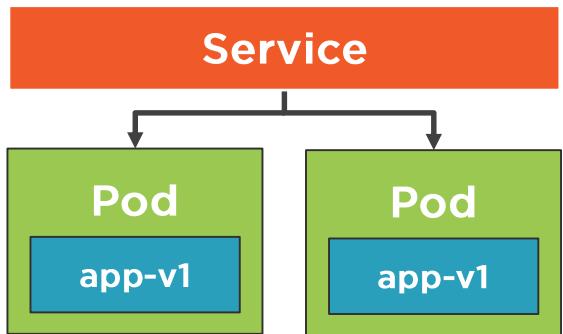
Start new Pods and then delete old Pods?
Need to be able to run two versions simultaneously.

Replace existing Pods one by one without impacting traffic to Pods?



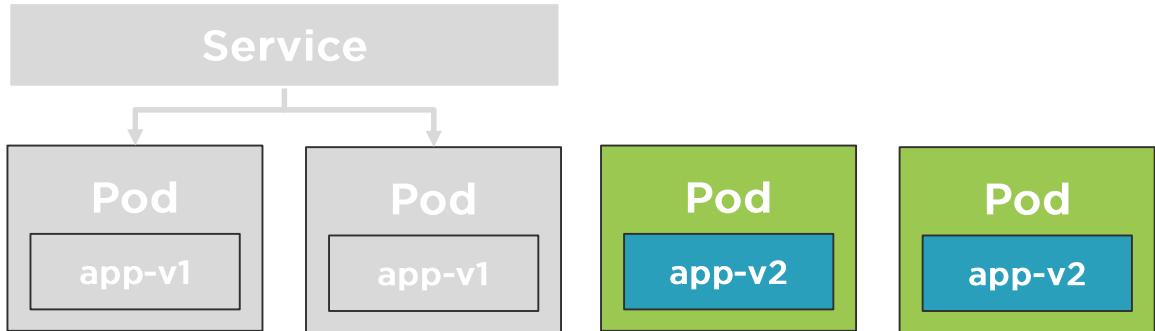
Delete Existing Pods – Replace with New Pods

Initial Pod State



Delete Existing Pods – Replace with New Pods

Initial Pod State

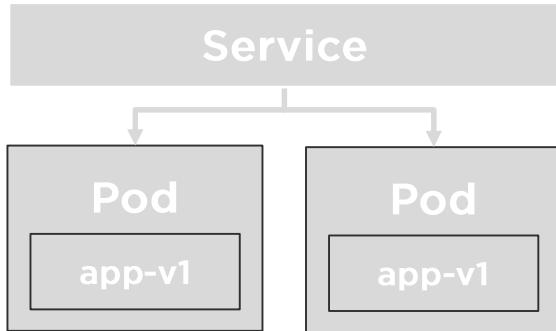


Create New Pods

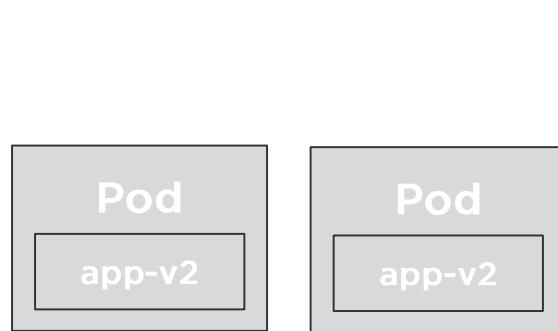


Delete Existing Pods – Replace with New Pods

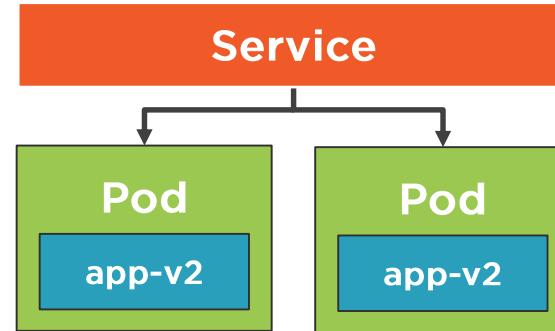
Initial Pod State



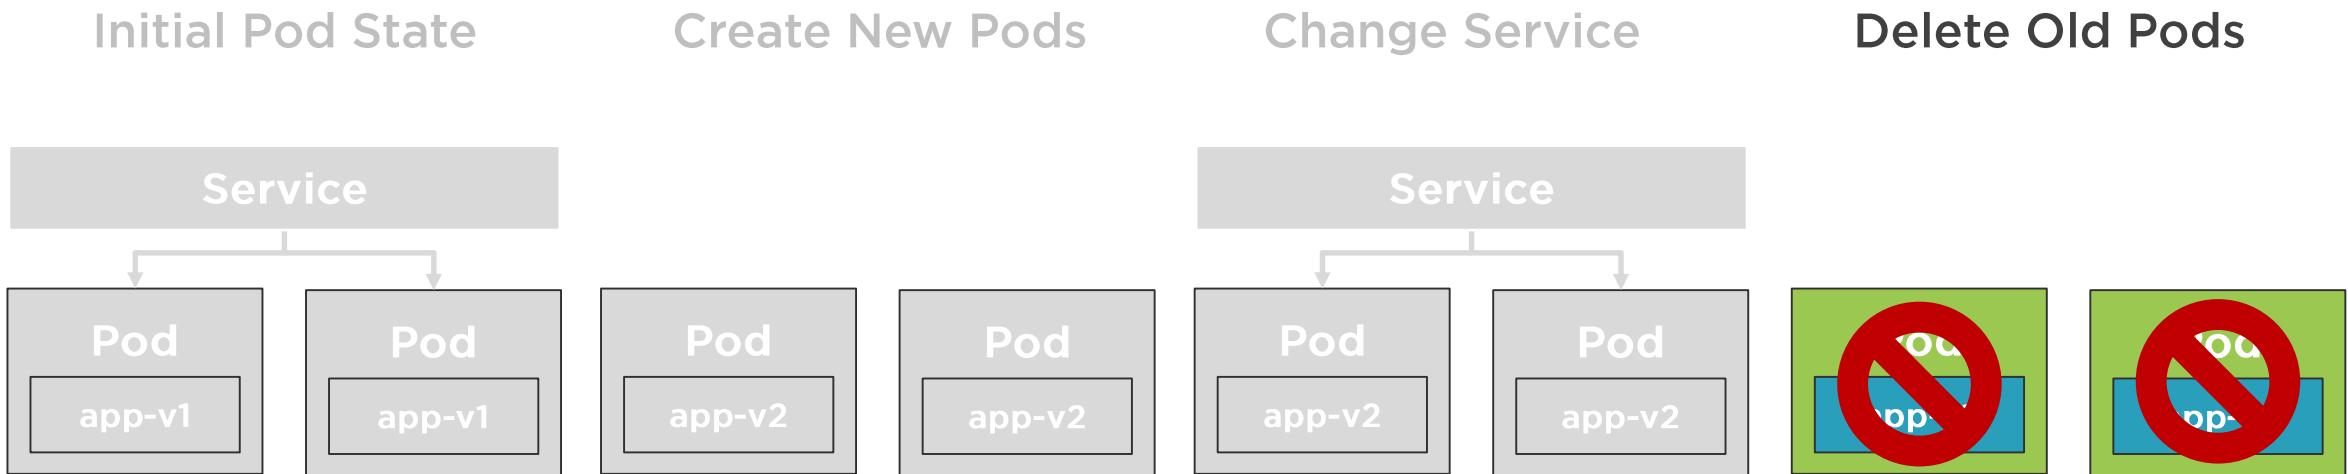
Create New Pods



Change Service

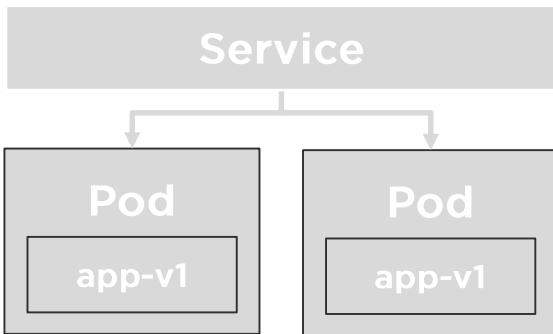


Delete Existing Pods – Replace with New Pods

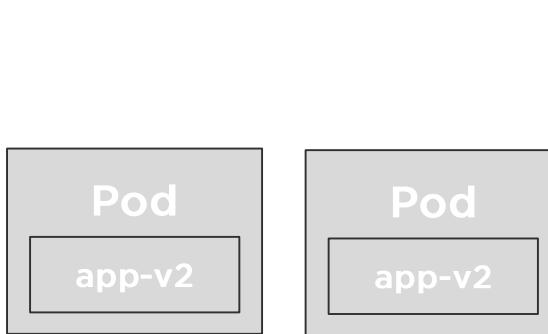


Delete Existing Pods – Replace with New Pods

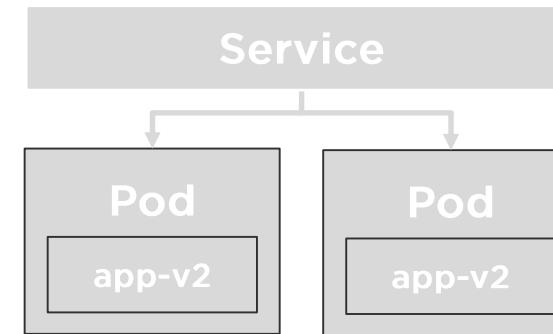
Initial Pod State



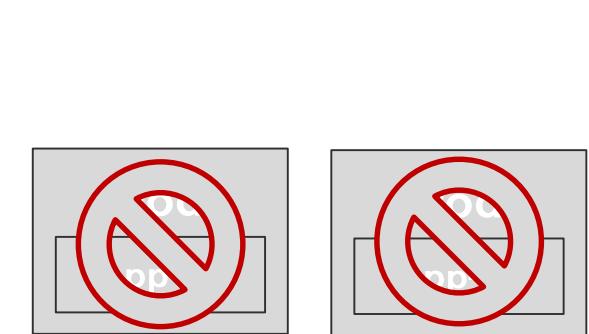
Create New Pods



Change Service



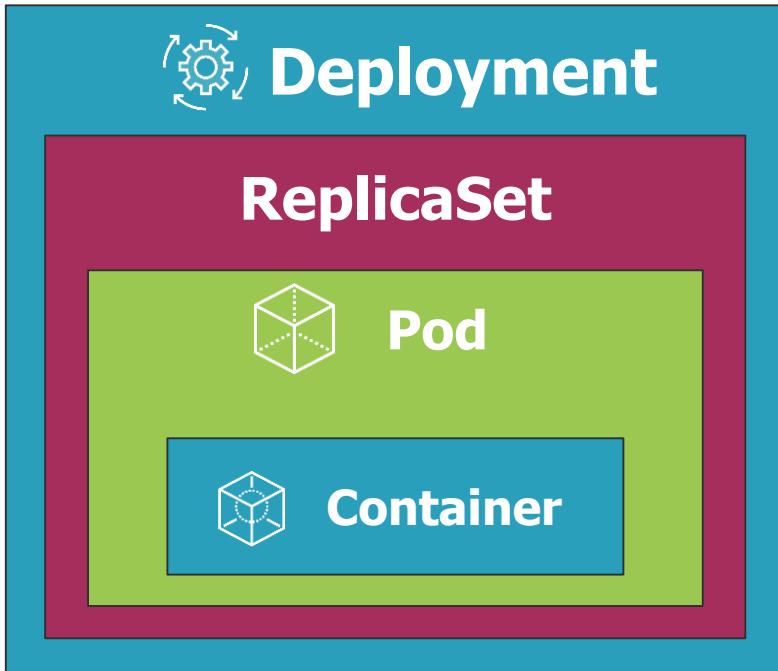
Delete Old Pods



Can your servers run the new and old Pods at the same time?



Deployment Options



One of the strengths of Kubernetes is "zero-downtime deployments"

Update an application's Pods without impacting end users

Several options are available:

- Rolling Updates
- Blue-Green Deployments
- Canary Deployments
- Rollbacks



Zero-downtime deployments
allow software updates to be
deployed to production without
impacting end users.



Summary



Deployments are a key resource provided by Kubernetes

Deployments rely on ReplicaSets to schedule and manage Pods

Kubernetes supports Zero-downtime deployments out of the box

Several Deployment options exist:

- Zero-downtime
- Rolling Updates
- Canary
- Blue-Green



Performing Rolling Update Deployments



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Understanding Rolling Update Deployments

Creating a Rolling Update Deployment

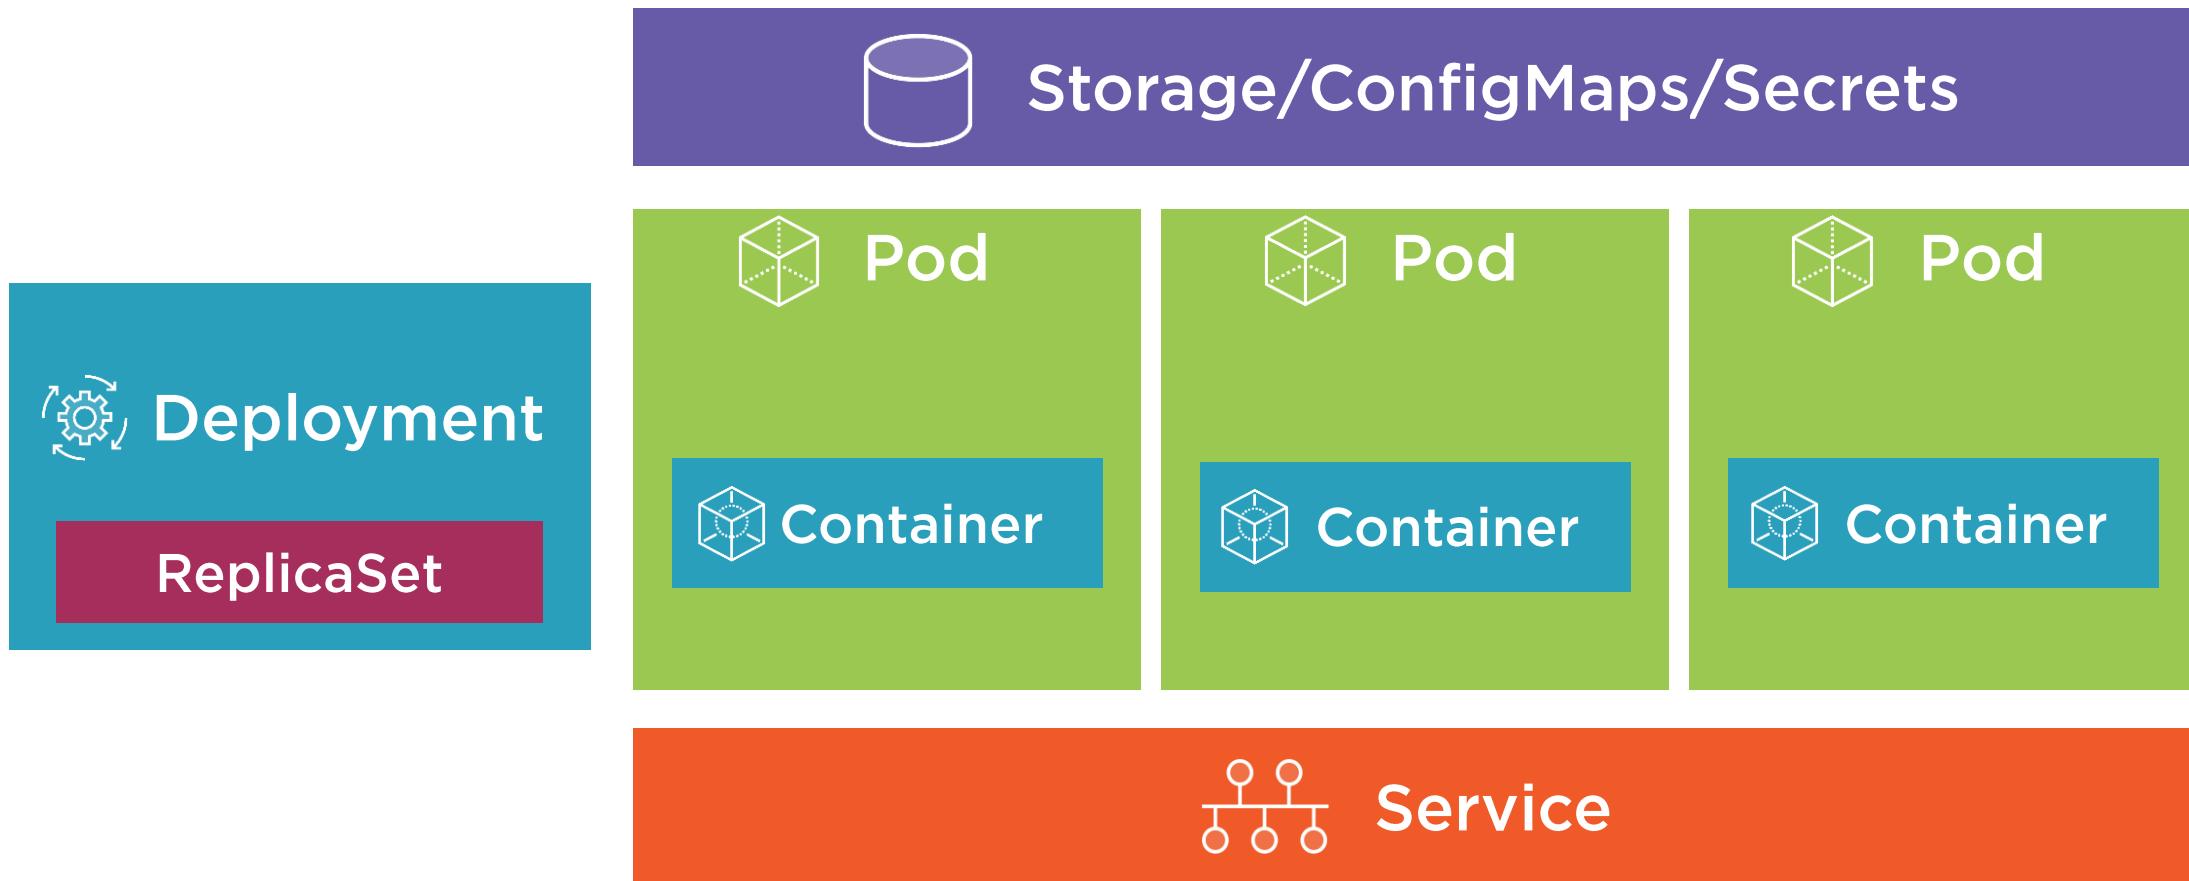
Rolling Update Deployment in Action

Rolling Back Deployments

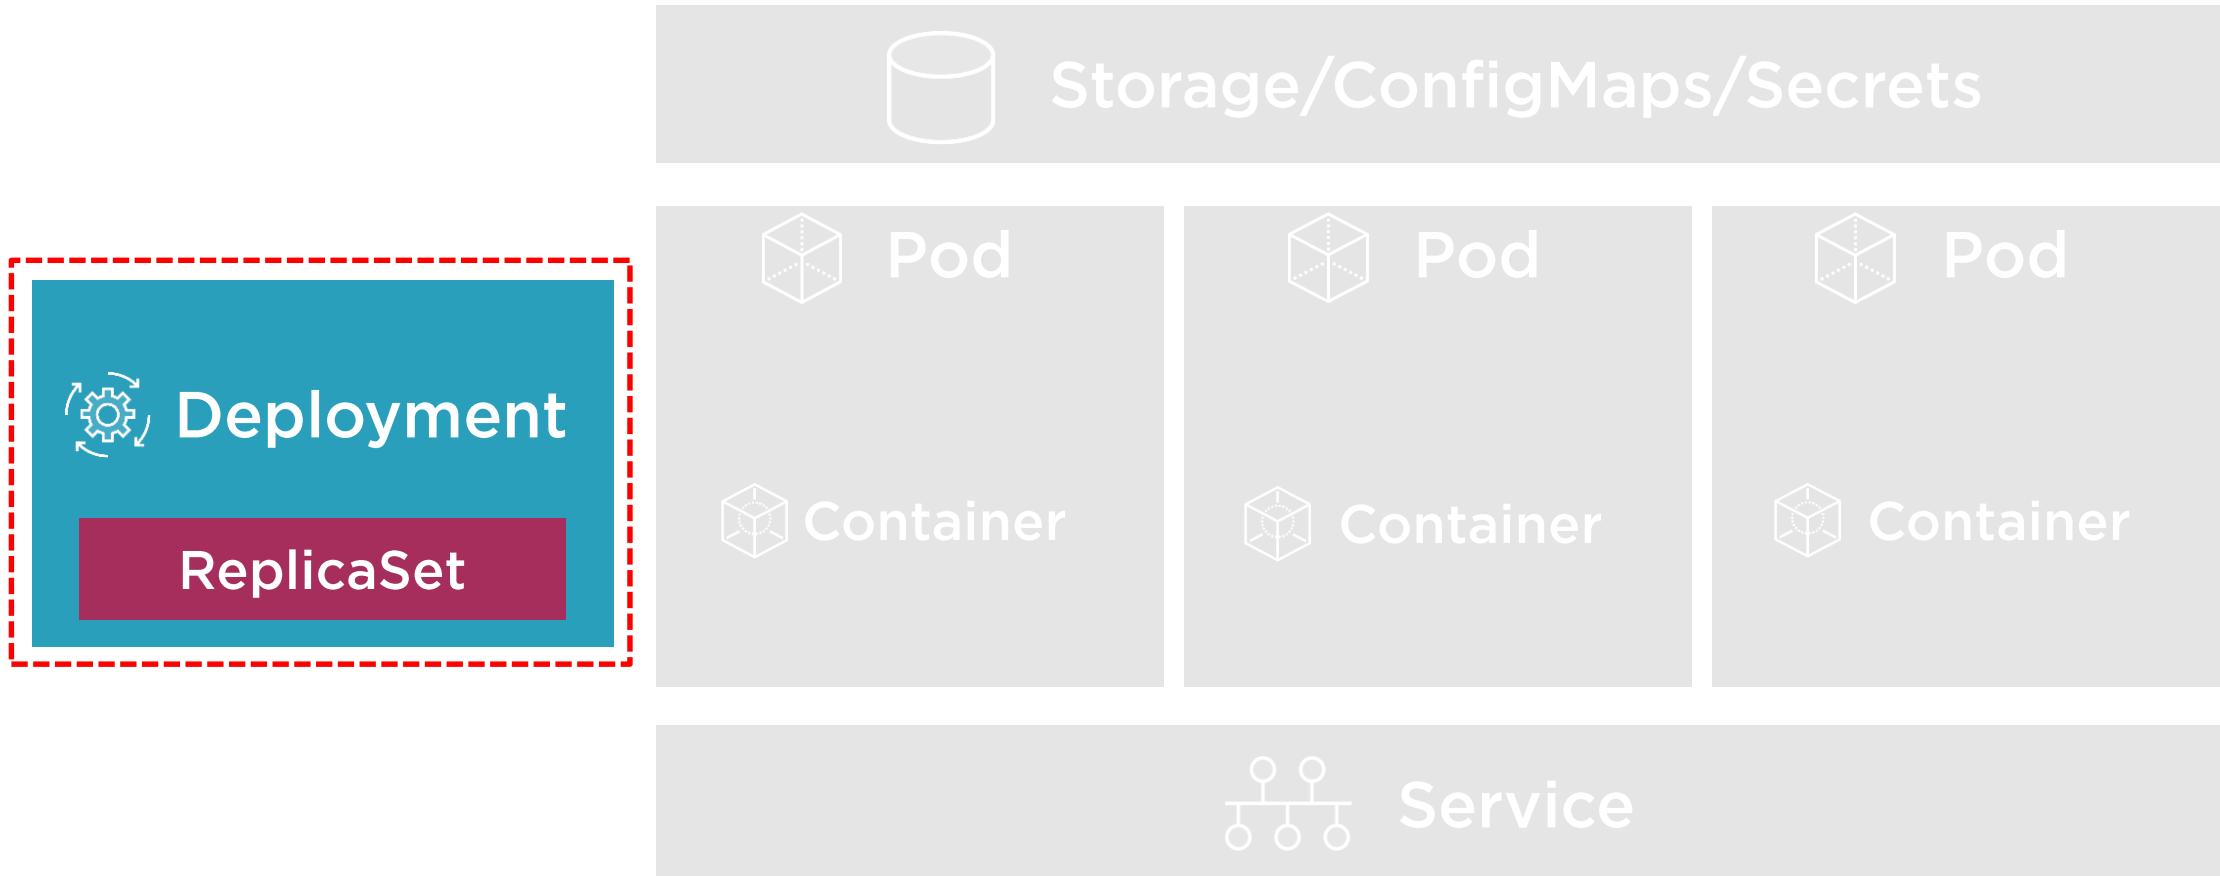
Rolling Back Deployments in Action



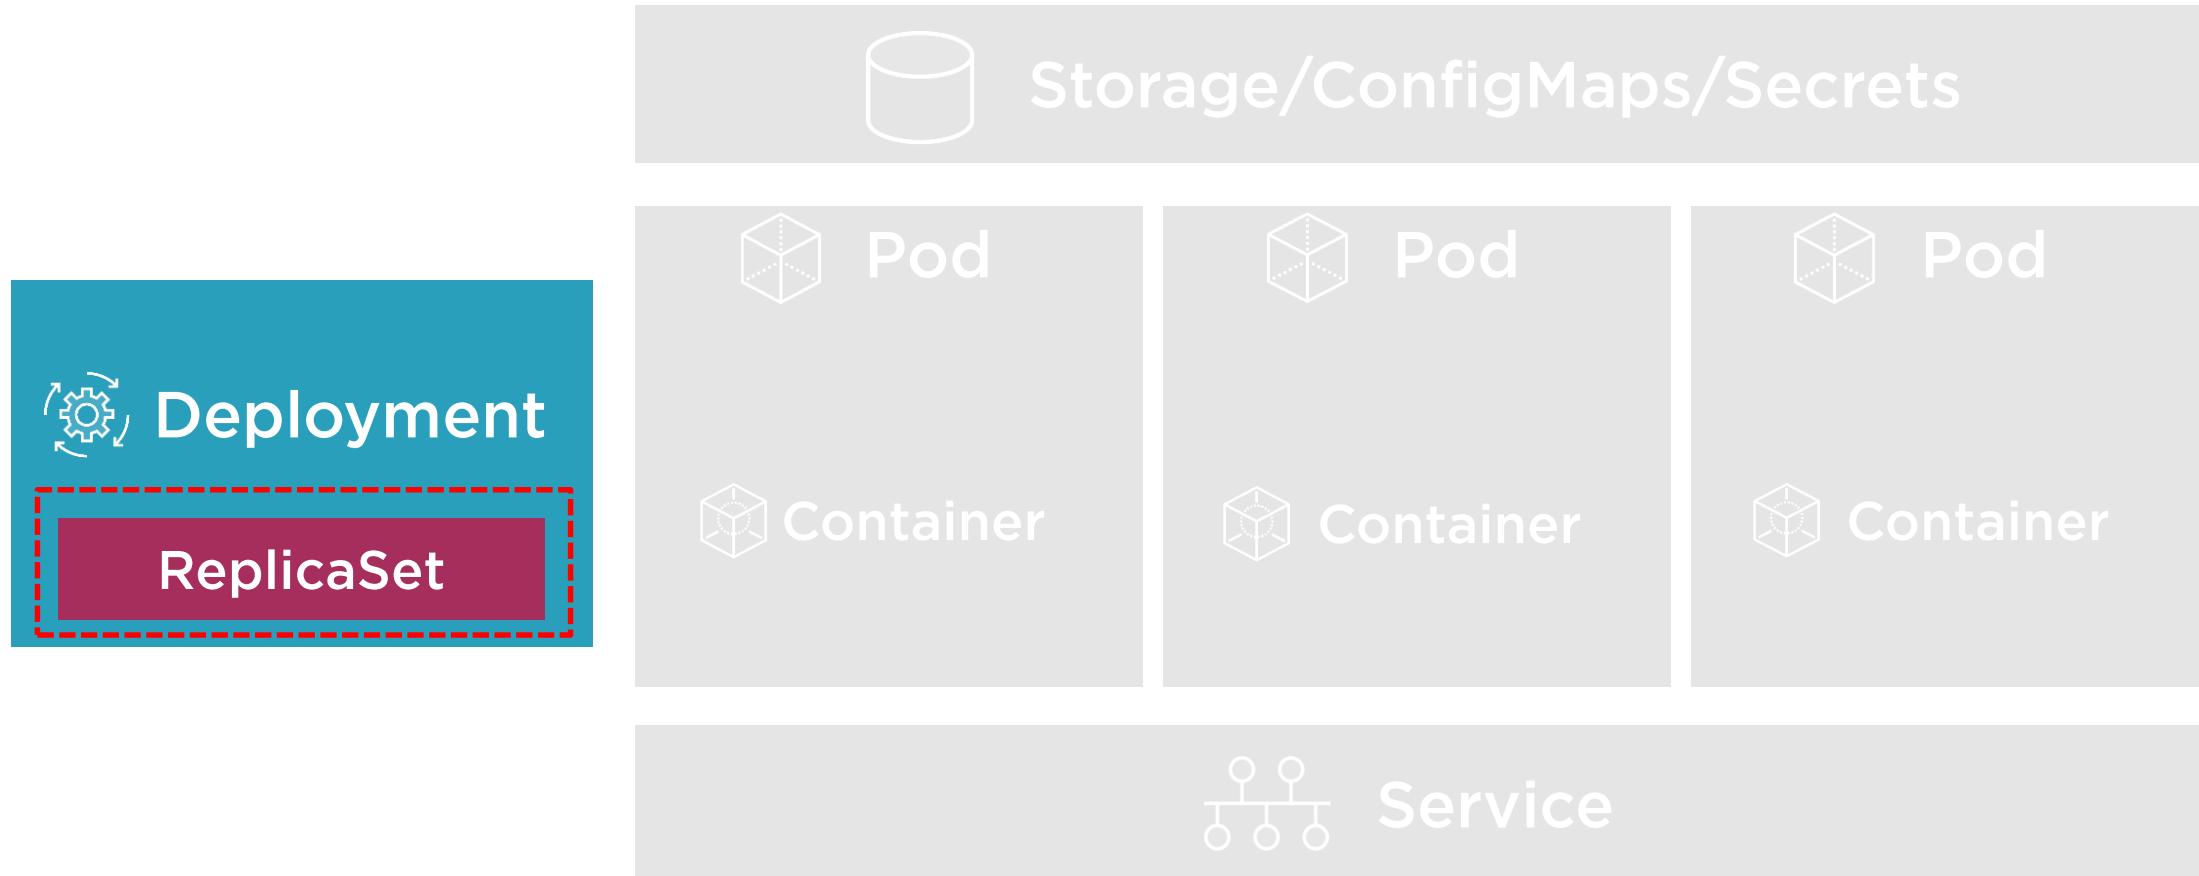
Kubernetes Resources



Kubernetes Resources



Kubernetes Resources



Understanding Rolling Update Deployments

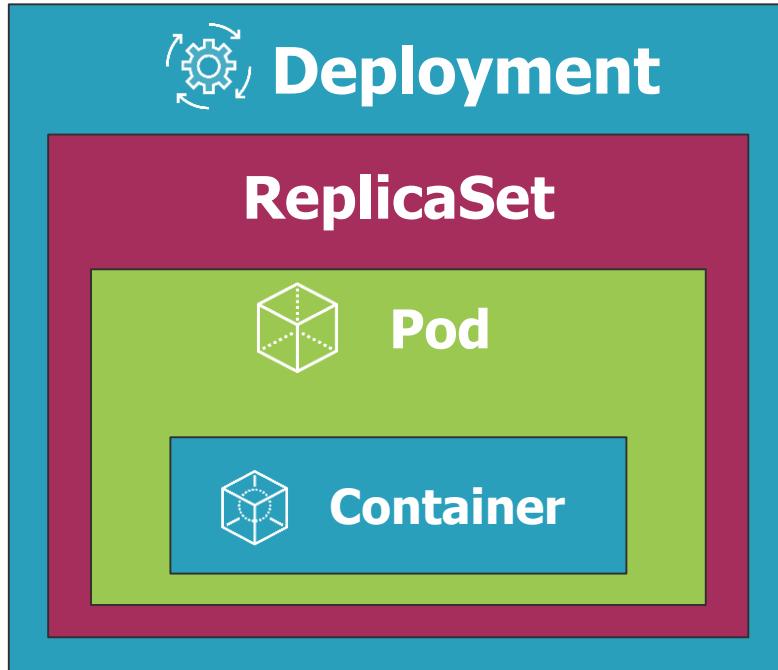


"Rolling updates allow Deployments' update to take place with zero downtime by incrementally updating Pods instances with new ones."

Kubernetes Documentation



Rolling Update Deployments



ReplicaSets increase new Pods while decreasing old Pods

Service handles load balancing traffic to available Pods

New Pods only scheduled on available Nodes

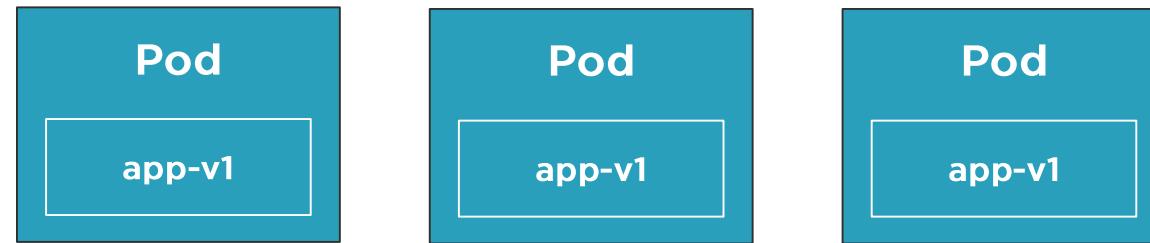
Deployments support two strategy options:

- Rolling Update (default and our focus here)
- Recreate (can result in down-time)



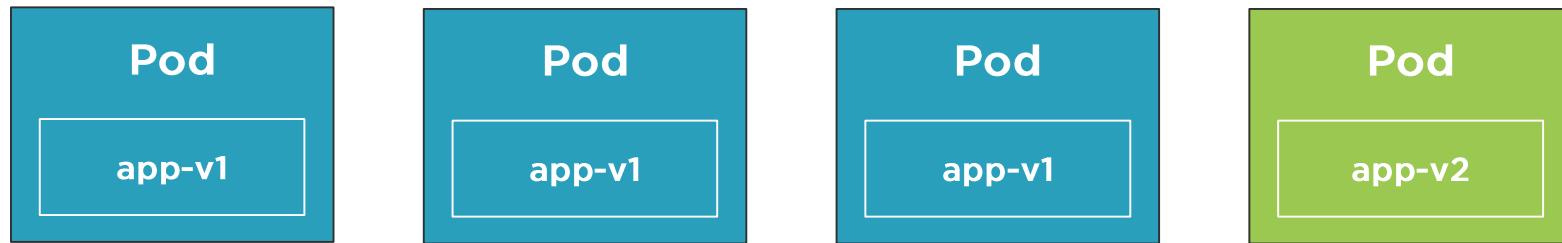
Rolling Update Deployments

Initial Pod State



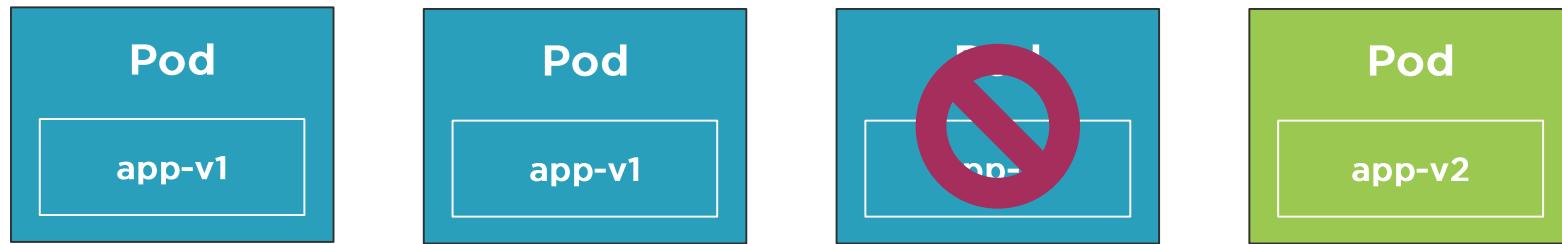
Rolling Update Deployments

Rollout New Pod



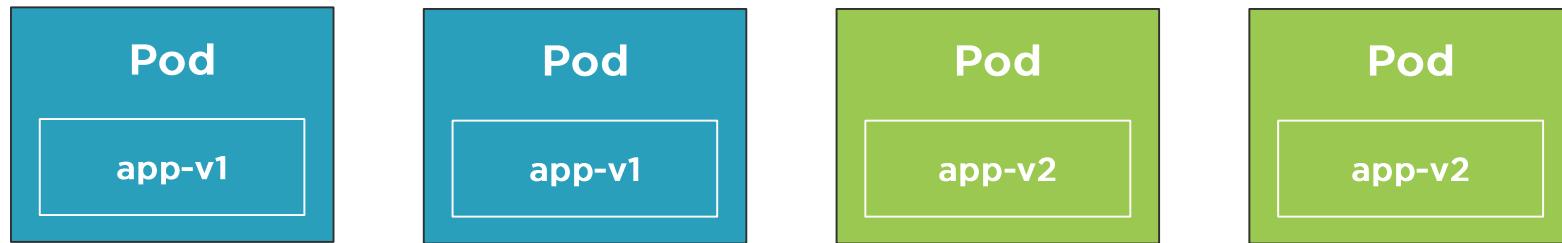
Rolling Update Deployments

Delete Pod



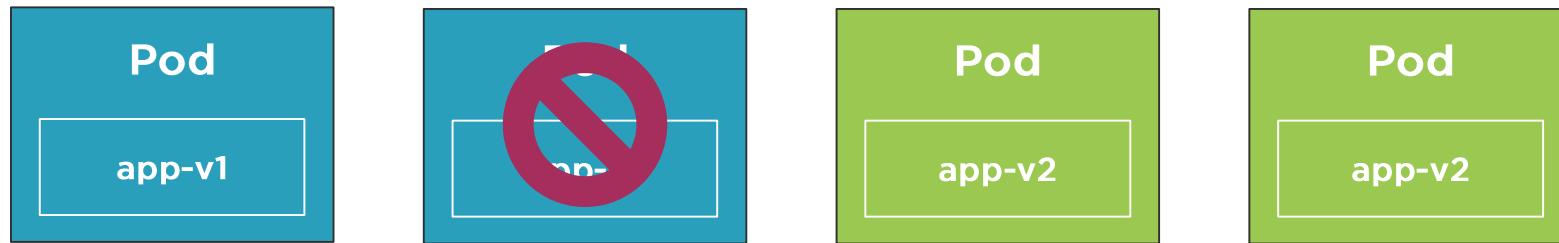
Rolling Update Deployments

Rollout New Pod



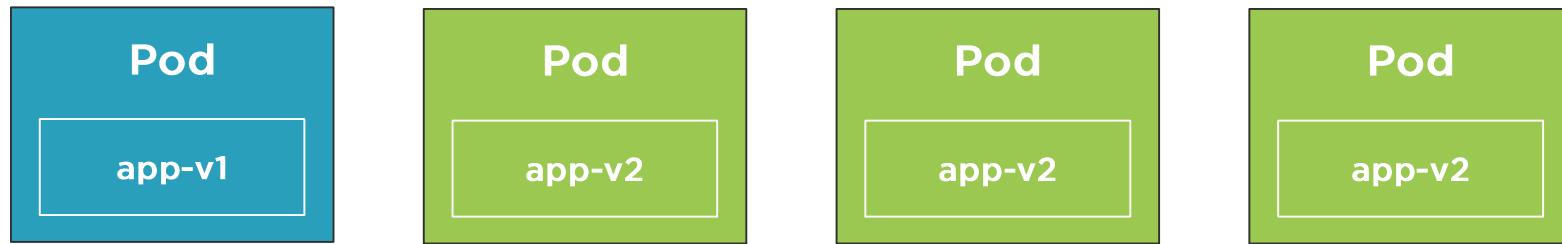
Rolling Update Deployments

Delete Pod



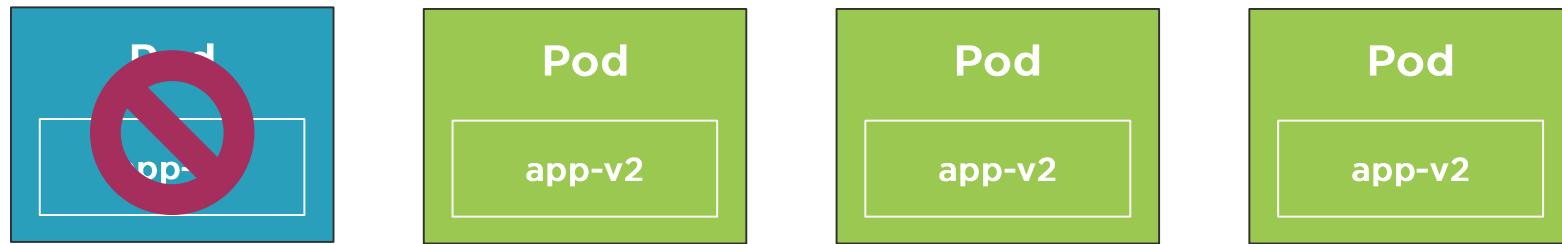
Rolling Update Deployments

Rollout New Pod



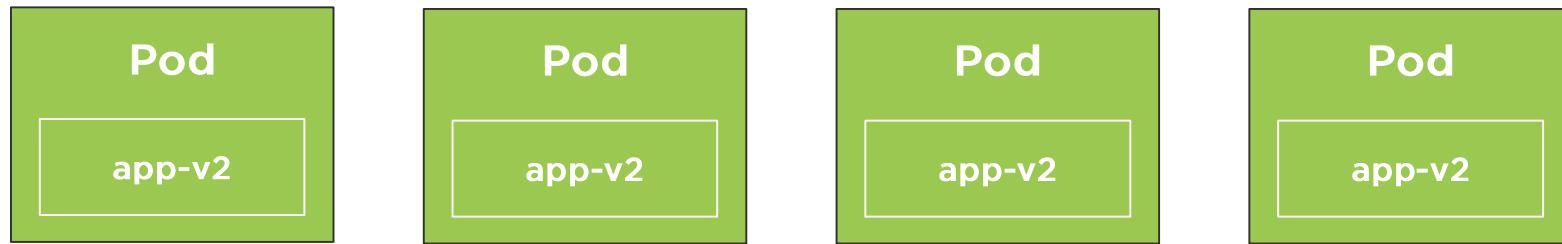
Rolling Update Deployments

Delete Pod

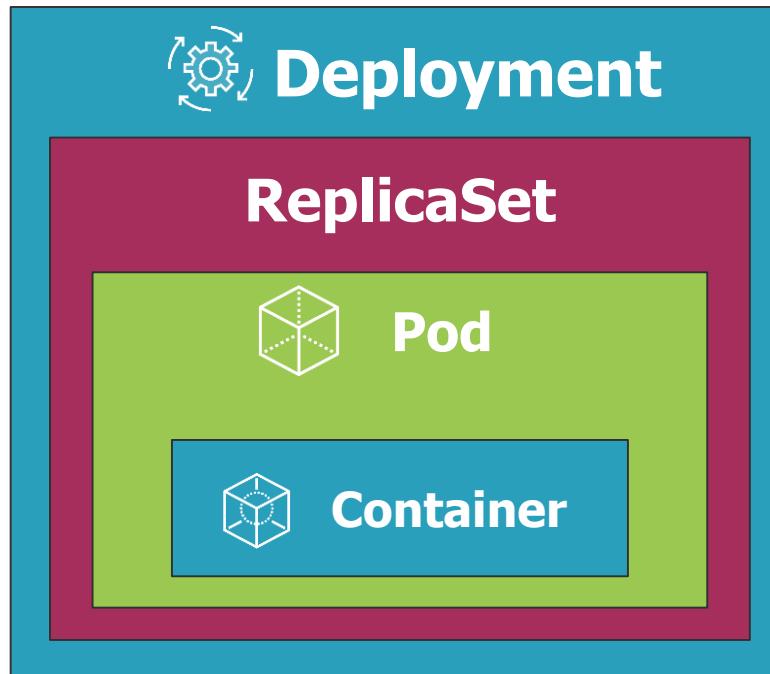


Rolling Update Deployments

Rollout New Pod



Deployments and Replicsets



Creating a Rolling Update Deployment



Defining a Rolling Update Deployment

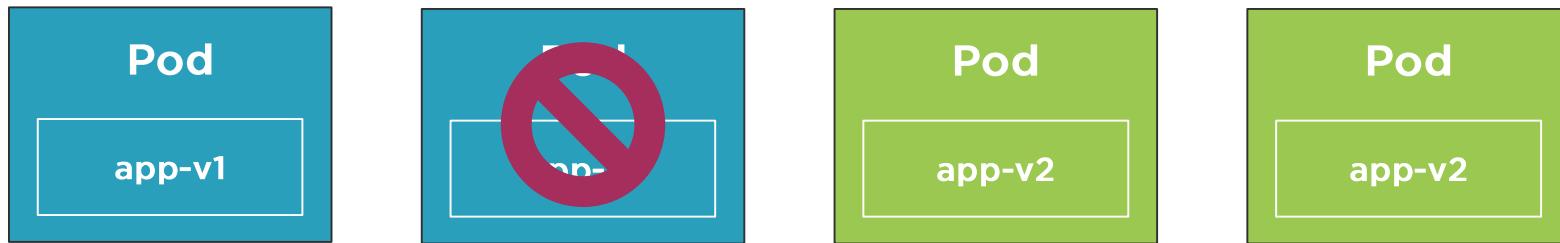
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  replicas: 2
  minReadySeconds: 1
  progressDeadlineSeconds: 60
  revisionHistoryLimit: 5
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
...
```

- ◀ Number of Pod replicas
- ◀ Seconds new Pod should be ready to be considered healthy (0)
- ◀ Seconds to wait before reporting stalled Deployment
- ◀ Number of ReplicaSets that can be rolled back (10)
- ◀ RollingUpdate (default) or Recreate strategy
- ◀ Max Pods that can exceed the replicas count (25%)
- ◀ Max Pods that are not operational (25%)



Understanding maxSurge

How many Pods can be added above the replicas count during the rolling update?



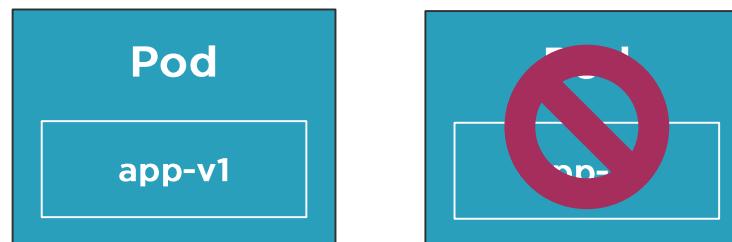
$$\begin{array}{rcl} \text{replicas} & = 2 \\ \text{maxSurge} & = 1 \\ \hline \text{current} & = 3 \end{array}$$

1 above replicas count
which is allowed



Understanding maxUnavailable

How many of the existing Pods can be made unavailable during a rolling update?



`maxUnavailable = 1`

↑
It's OK for 1 of the 2
replicas to be unavailable



Creating the Deployment

Use the **kubectl create** command along with the **--filename** or **-f** switch

```
# Create initial deployment  
kubectl create -f file.deployment.yml --save-config --record
```

Record the command
in the Deployment
revision history

Save configuration in
resource's annotations

Creating or Modifying a Deployment

Use the **kubectl apply** command along with the **--filename** or **-f** switch

Record the command
in the Deployment
revision history

```
# Create initial deployment
kubectl apply -f file.deployment.yml --record
```

Checking the Deployment Status

The **kubectl rollout status** command can be used to get information about a specific Deployment

```
# Get information about a Deployment  
kubectl rollout status deployment [deployment-name]
```

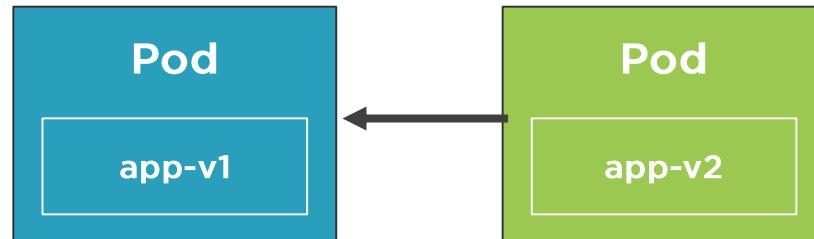
Rolling Update Deployment in Action



Rolling Back Deployments



Rolling Back Deployments



Rolling update revisions can be tracked using `--record`

If a Deployment has issues, a new Deployment can be applied, or you can revert to a previous revision

Several kubectl commands can be used for rollbacks:

- `kubectl rollout status`
- `kubectl rollout history`
- `kubectl rollout undo`



Checking Deployment History

The **kubectl rollout history** command can be used to view history of a Deployment

```
# Get information about a Deployment  
kubectl rollout history deployment [deployment-name]  
  
# Get information about a Deployment  
kubectl rollout history deployment [deployment-name] --revision=2
```

Rolling Back a Deployment

Use the **kubectl rollout undo** command to rollback to a specific Deployment revision

```
# Check status  
kubectl rollout status -f file.deployment.yml  
  
# Rollback a Deployment  
kubectl rollout undo -f file.deployment.yml  
  
# Rollback to a specific revision  
kubectl rollout undo deployment [deployment-name] --to-revision=2
```

Rolling Back Deployments in Action



Summary



Rolling updates are the default Deployment strategy used by Kubernetes

Ensures zero-downtime during a Deployment

Maximum and minimum Pods available during a Deployment can be defined

Deployments can be recorded and stored in history using --record

Deployments can be rolled back to a specific revision



Performing Canary Deployments



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

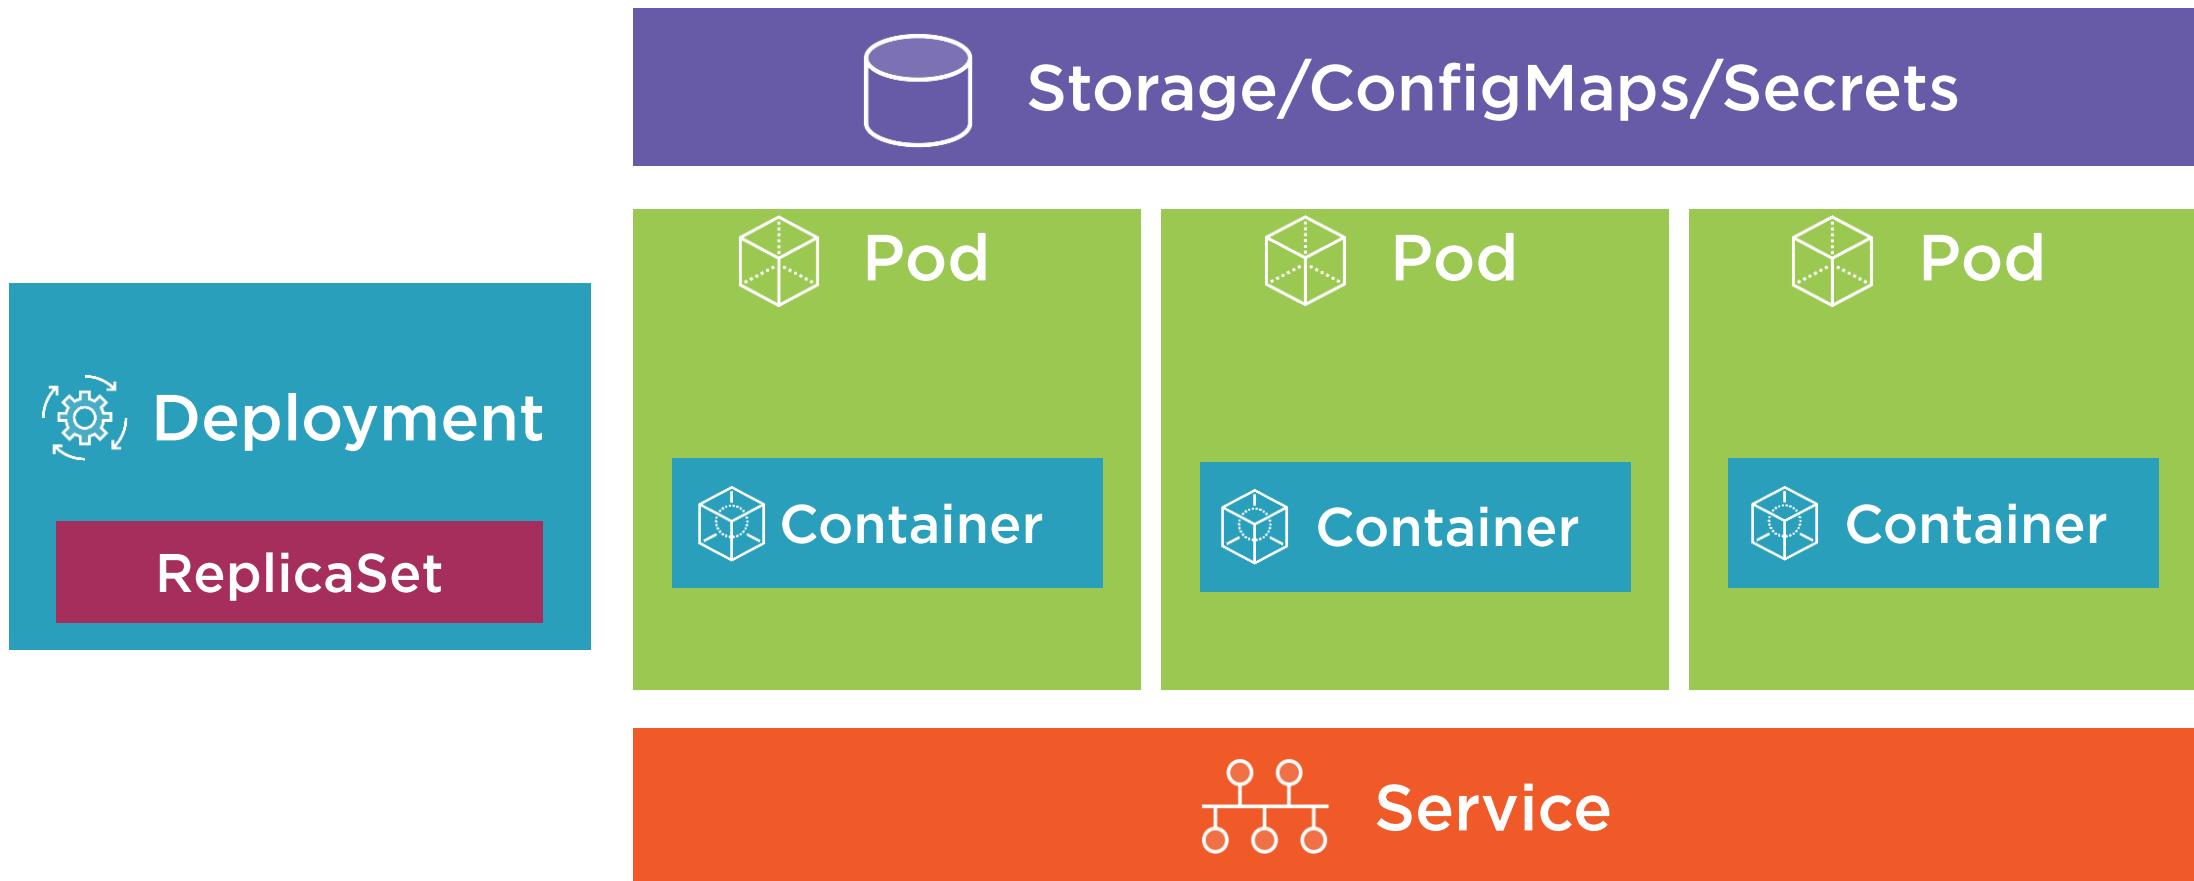
Understanding Canary Deployments

Creating a Canary Deployment

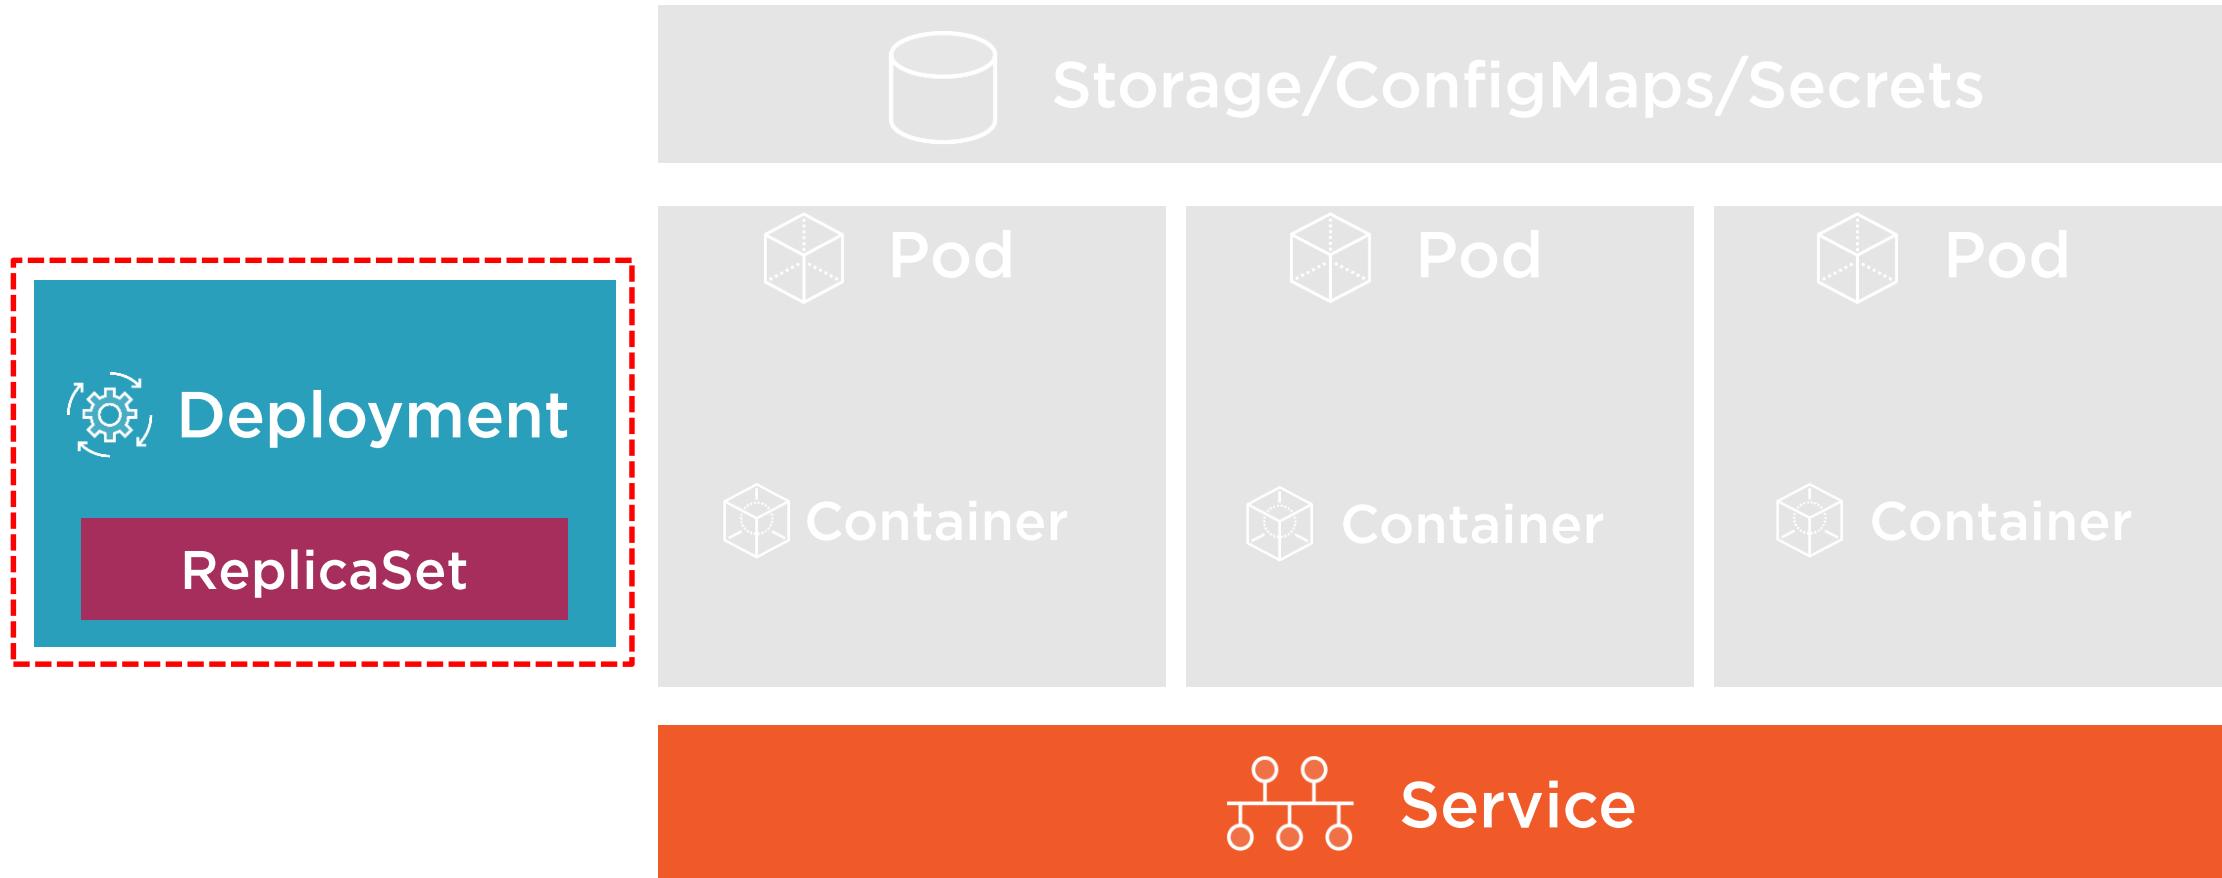
Canary Deployments in Action



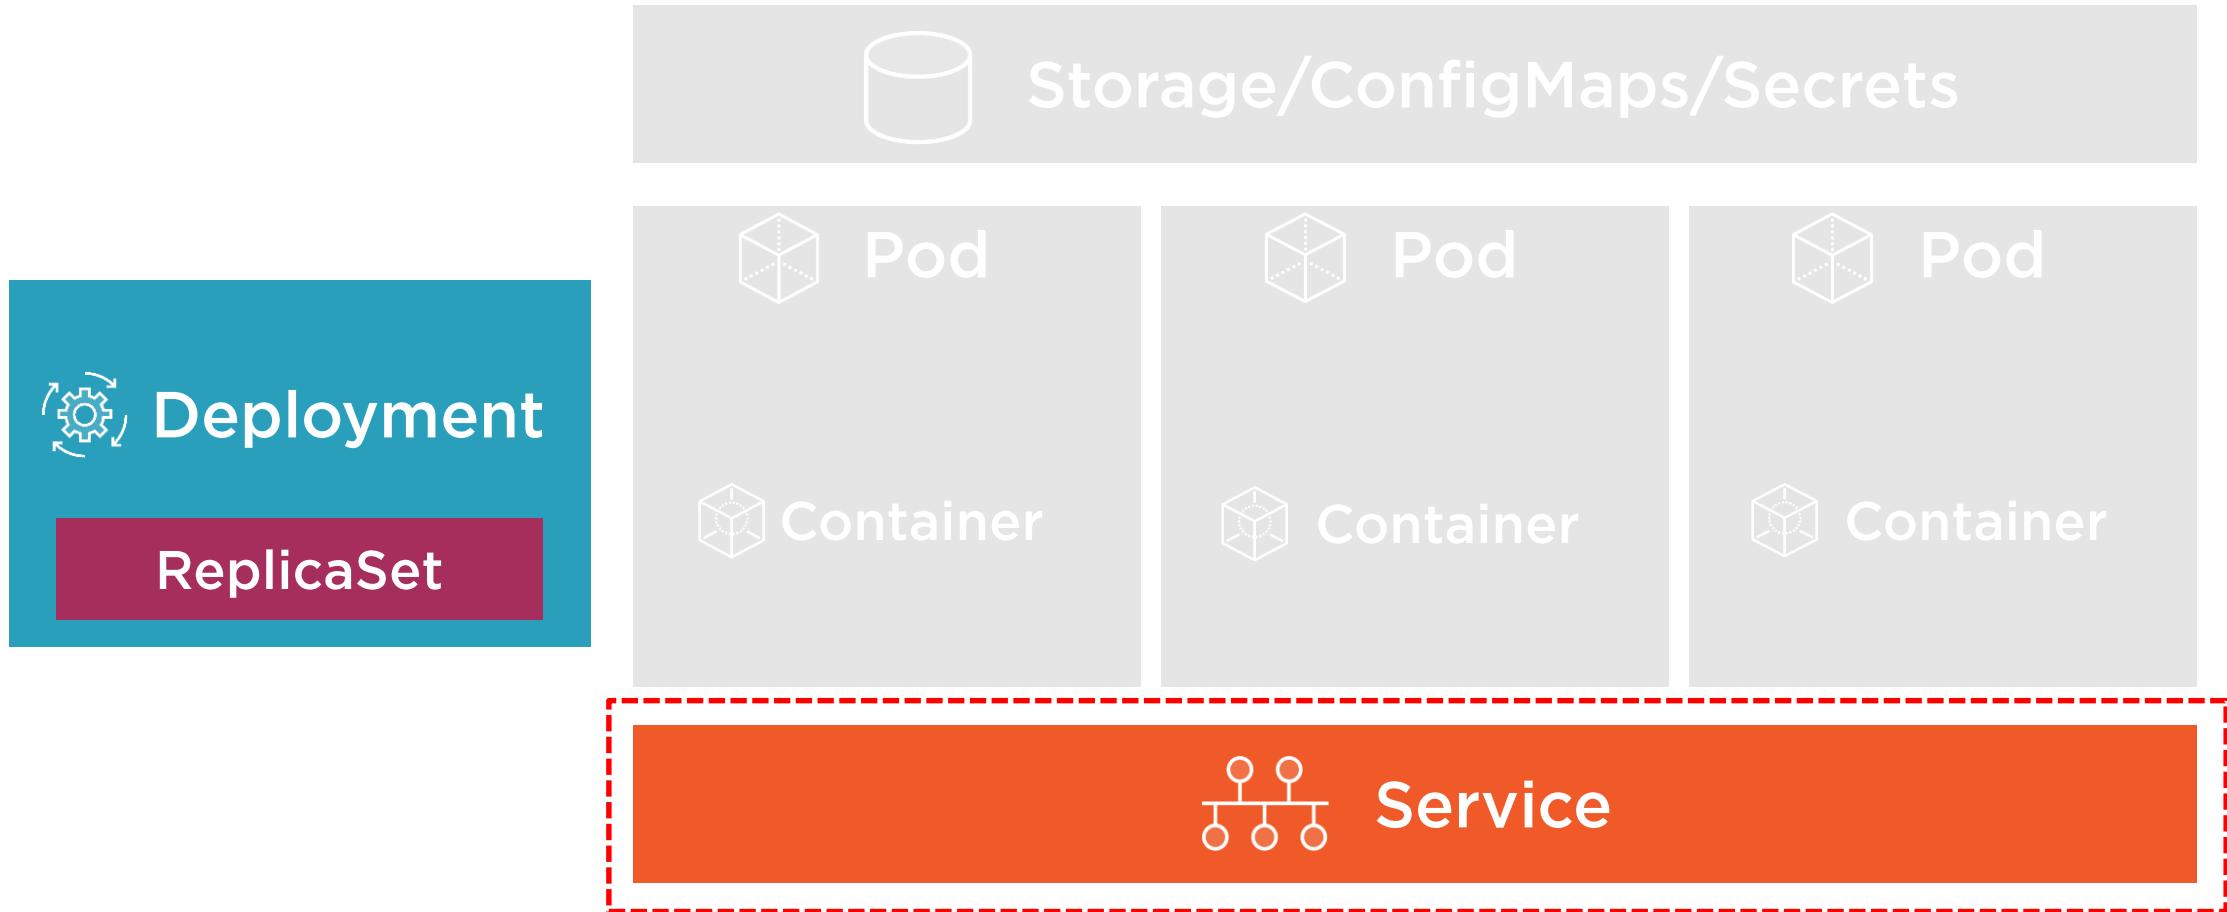
Kubernetes Resources



Kubernetes Resources



Kubernetes Resources



Understanding Canary Deployments





Wouldn't it be nice to rollout a new Deployment but only route a small percentage of the overall traffic to it to ensure it's working properly?

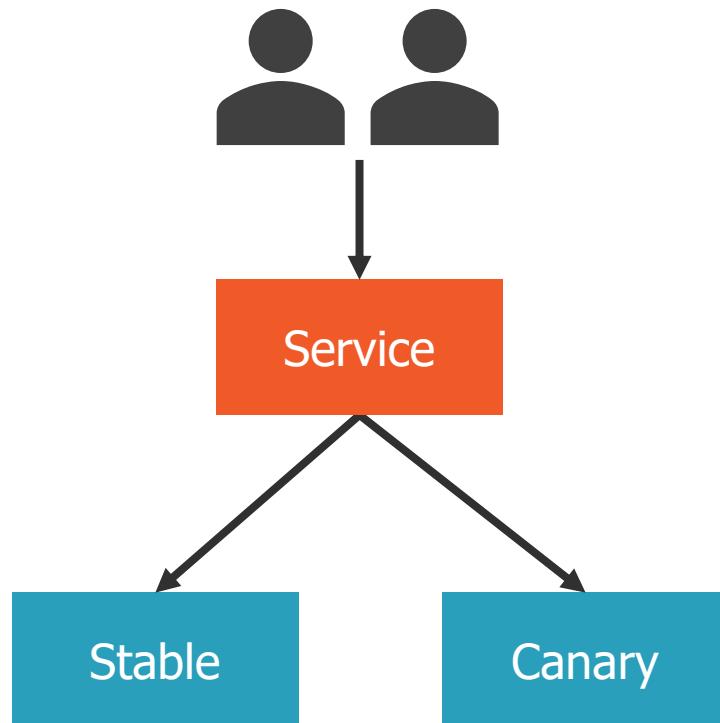


"Canary deployment strategy involves deploying new versions of applications next to stable production versions to see how the canary version compares against the baseline before promoting or rejecting the deployment."

~ <https://docs.microsoft.com>



Canary Deployments



Strategy for checking the viability of a deployment

Run two identical production environments at the same time

Canary Deployment runs alongside the existing stable Deployment

Canary Deployment is setup to receive minimal traffic

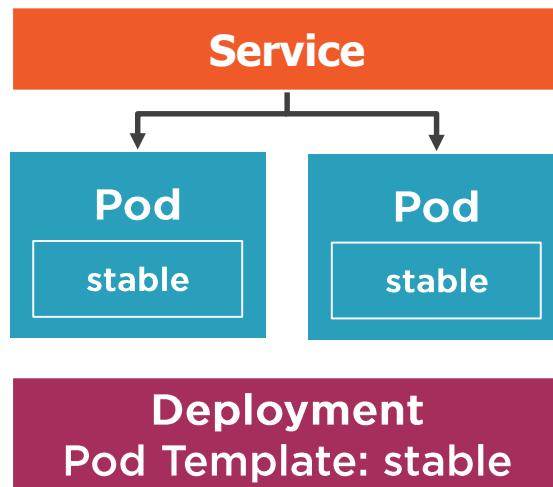


Canary Deployments

1

Create Stable Deployment and Service

Stable

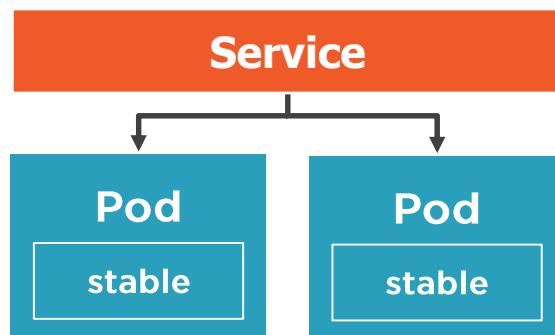


Canary Deployments

2

Create Canary Deployment

Stable



Deployment
track: stable

Canary



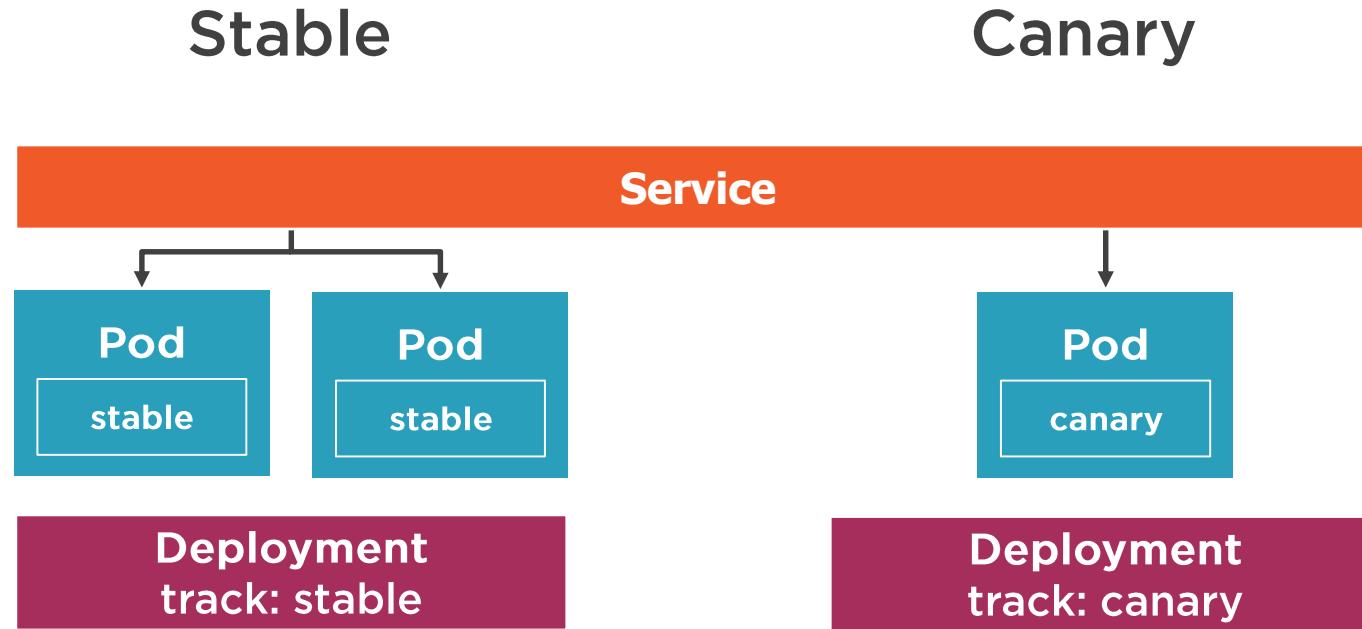
Deployment
track: canary



Canary Deployments

3

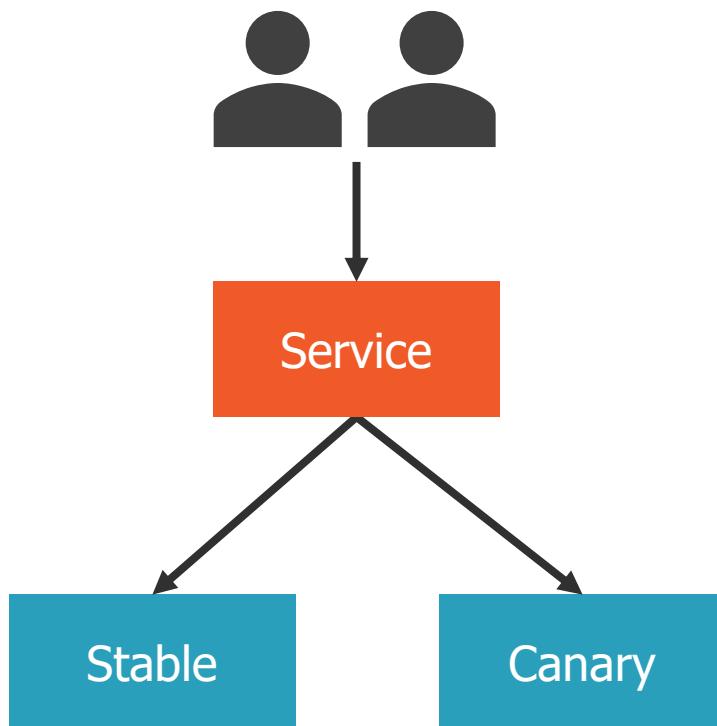
Service adds Canary Pod(s) and traffic is routed



Creating a Canary Deployment



Canary Deployment Resources



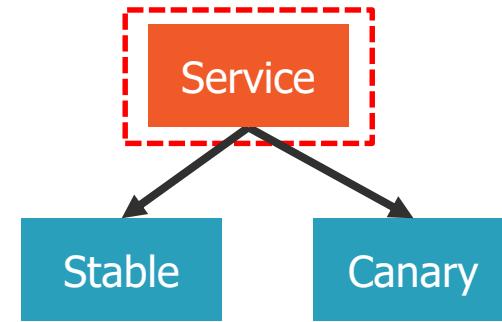
A Canary Deployment involves 3 main Kubernetes resources:

- Service
- Stable Deployment
- Canary Deployment



Defining a Service

```
kind: Service
apiVersion: v1
metadata:
  name: stable-service
  labels:
    app: aspnetcore
spec:
  type: LoadBalancer
  selector:
    app: aspnetcore
ports:
  - port: 80
    targetPort: 80
```

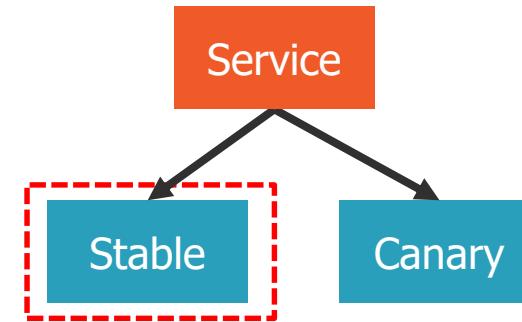


◀ Pod Label to select for Service



Defining a Stable Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: stable-deployment
spec:
  replicas: 4
  selector:
    matchLabels:
      app: aspnetcore
      track: stable
  template:
    metadata:
      labels:
        app: aspnetcore
        track: stable
  spec:
    containers:
      - name: stable-app
        image: stable-app
```



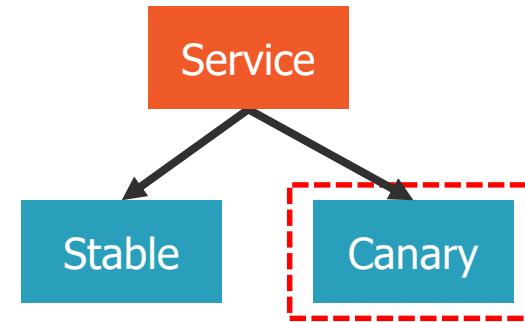
◀ Create stable replicas

◀ Pod labels (recall that
app:aspnetcore is targeted by
the Service)



Defining a Canary Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: canary-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: aspnetcore
      track: canary
  template:
    metadata:
      labels:
        app: aspnetcore
        track: canary
  spec:
    containers:
    - name: canary-app
      image: canary-app
```



◀ Create canary replicas (25% of stable in this example)

◀ Pod labels (recall that **app:aspnetcore** is targeted by the Service)



Creating the Stable and Canary Resources

Use `kubectl create` or `kubectl apply` commands to create the Service, Stable Deployment, and Canary Deployment

```
# Create Service, Stable Deployment, and Canary Deployment  
kubectl create -f [folder-name] --save-config --record
```

Canary Deployments in Action



Summary



Canary Deployments allows a new version to be deployed next to a stable version

Configured to only handle a small percentage of the traffic initially

Once the Canary Deployment is verified it can be scaled up and the existing stable Deployment can be scaled down



Performing Blue-Green Deployments



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

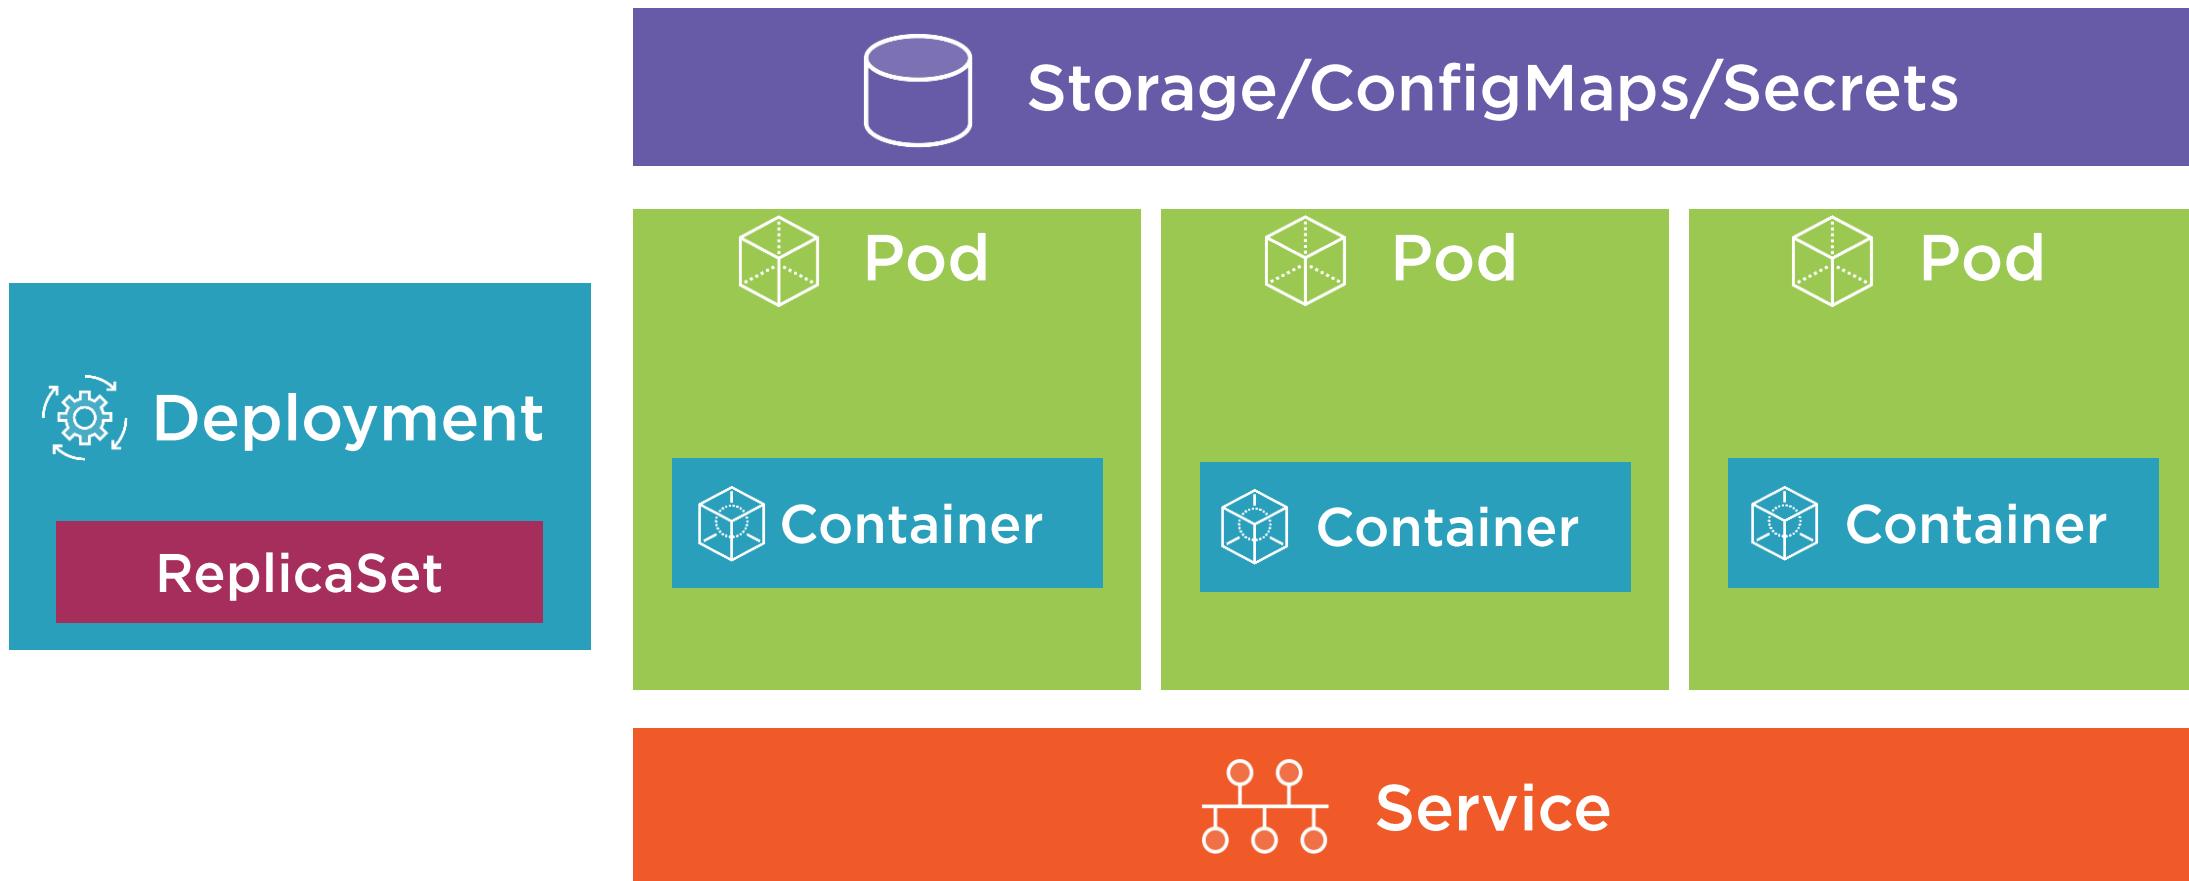
Understanding Blue-Green Deployments

Creating a Blue-Green Deployment

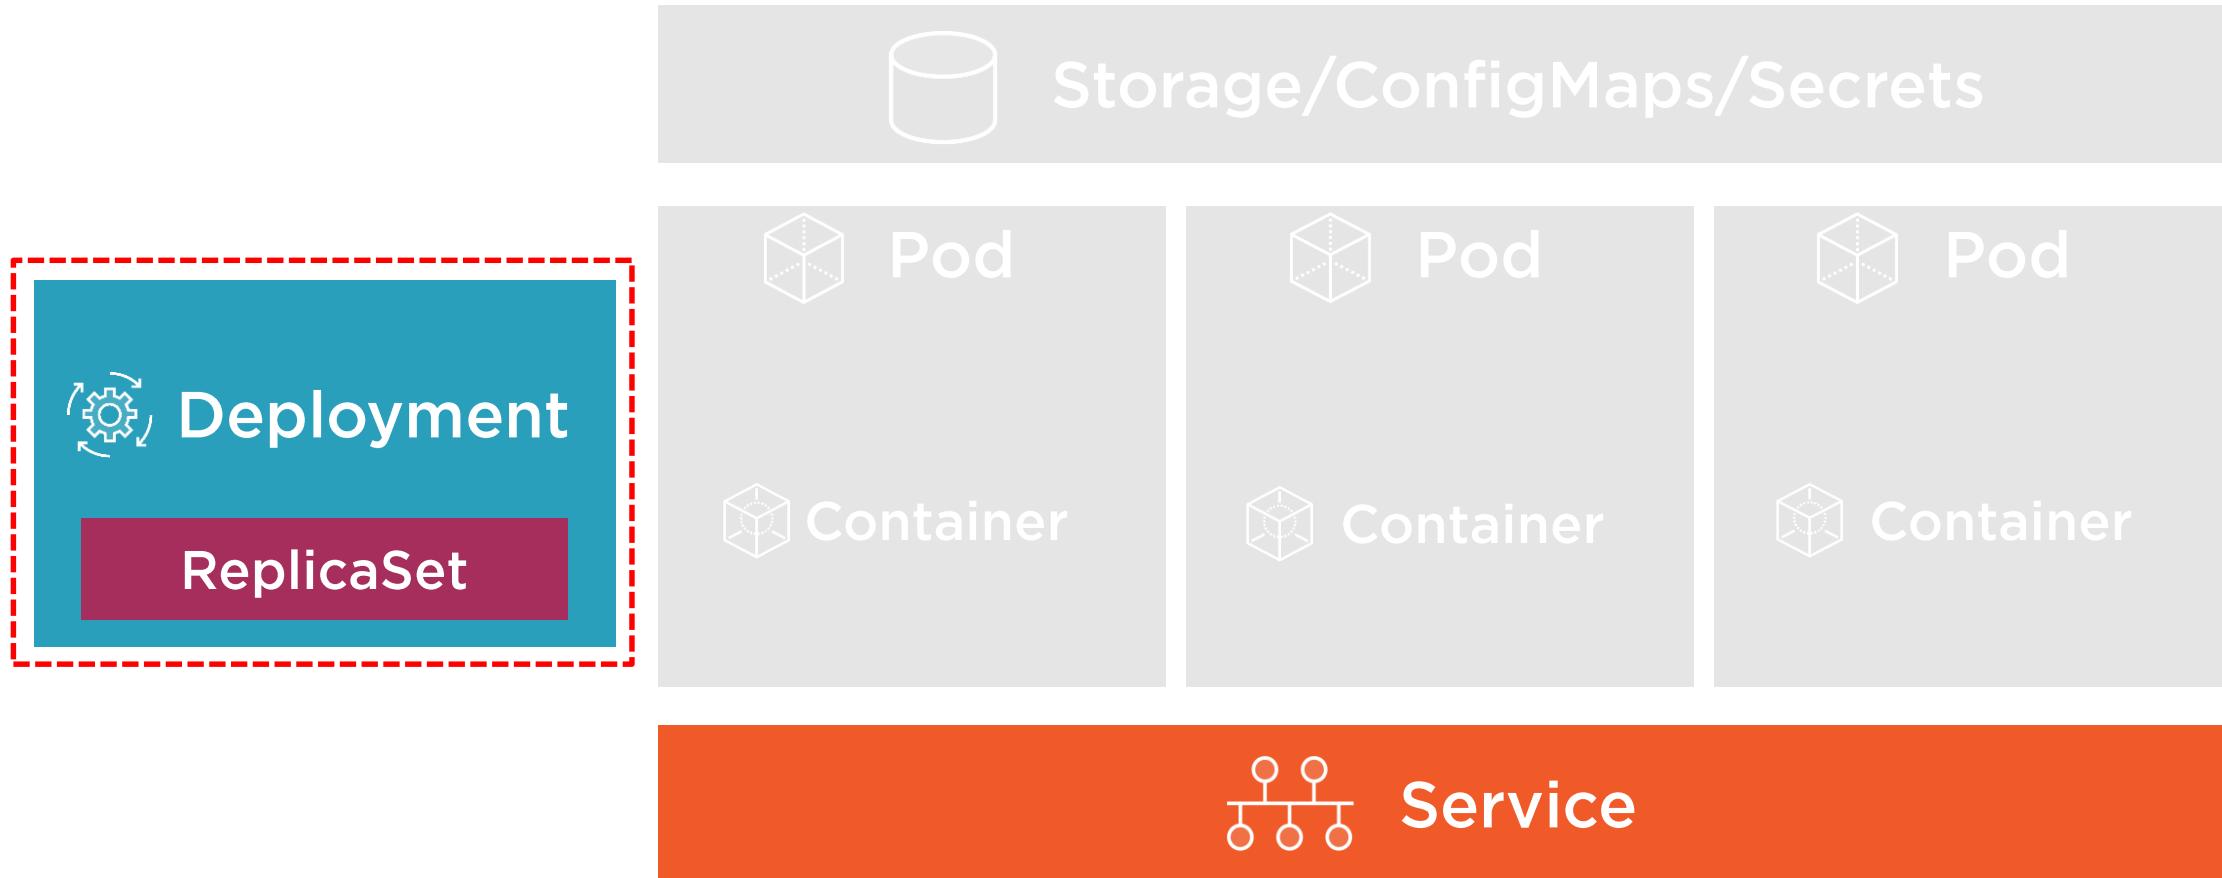
Blue-Green Deployments in Action



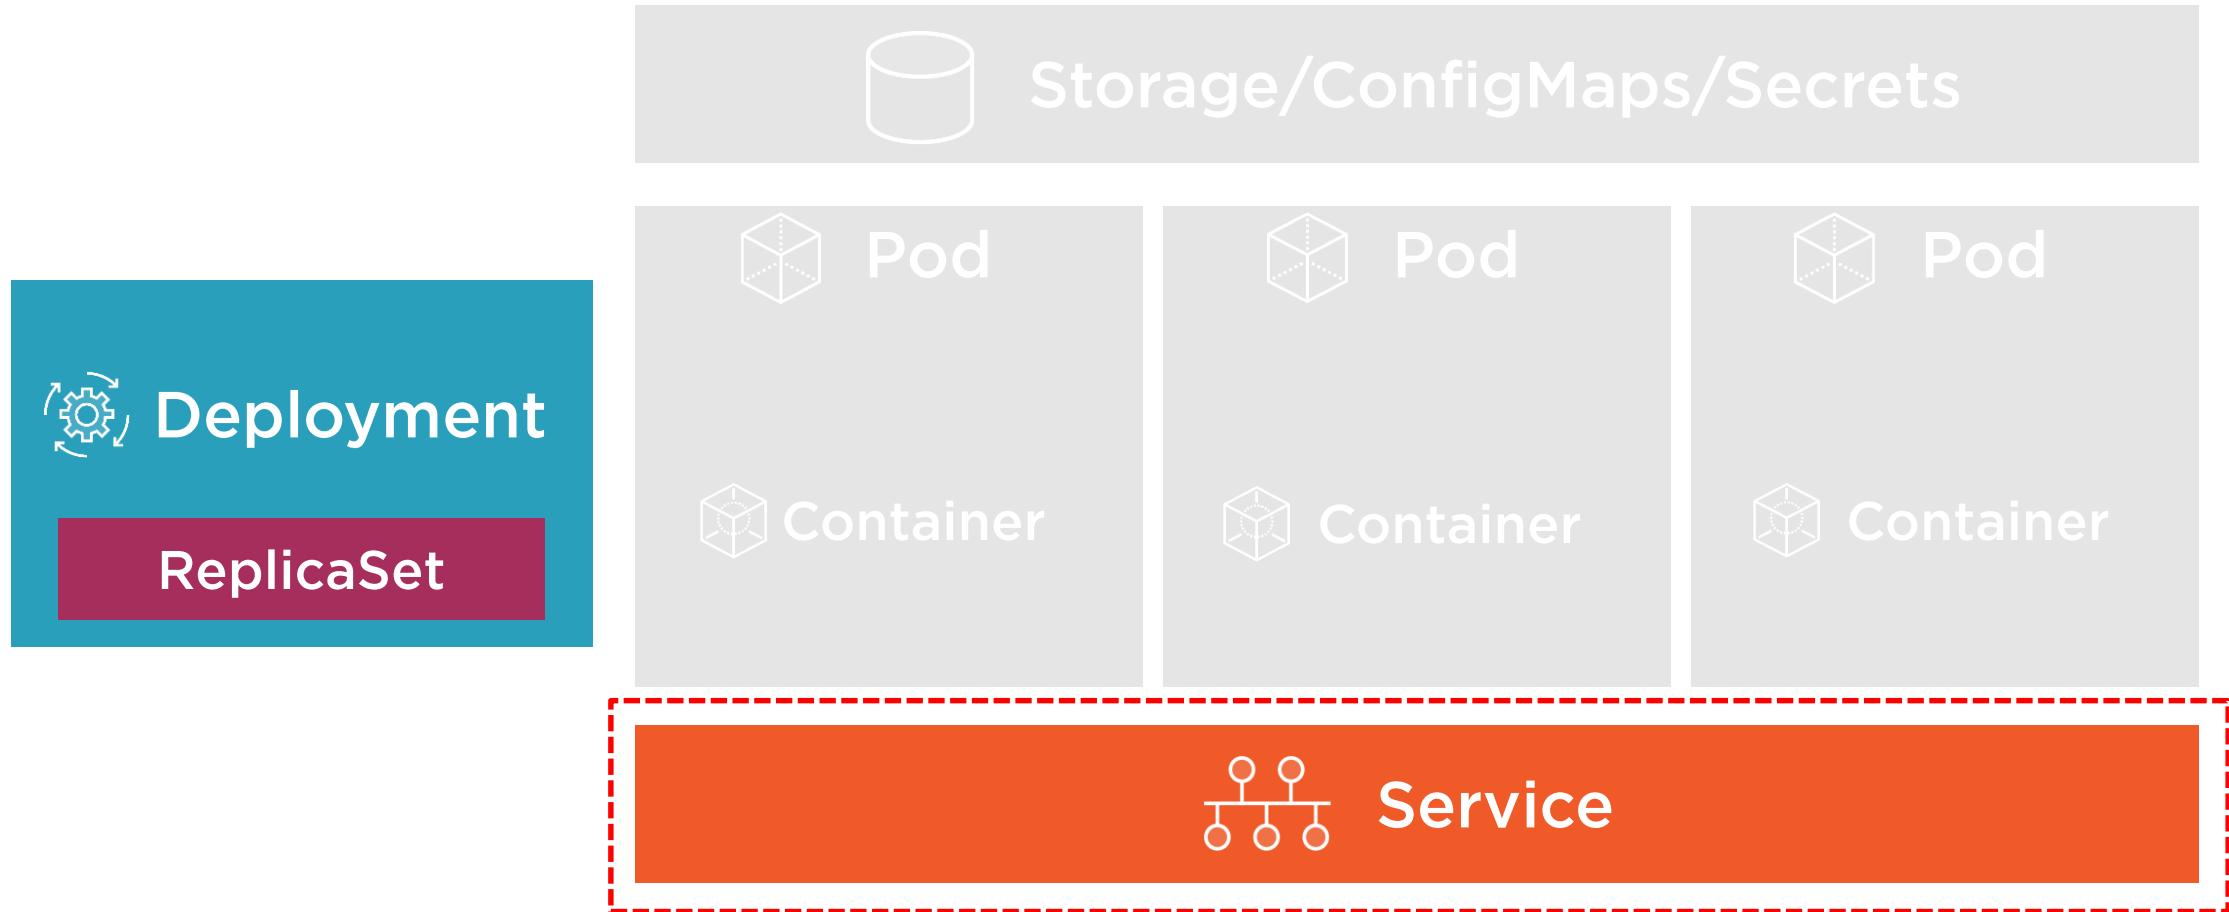
Kubernetes Resources



Kubernetes Resources



Kubernetes Resources



Understanding Blue-Green Deployments



Have you ever deployed an application to production and experienced problems?
(don't laugh too much ☺)



"A blue/green deployment is a change management strategy for releasing software code."

~ TechTarget.com



"A blue/green deployment is a change management strategy for releasing software code. Blue/green deployments, which may also be referred to as A/B deployments **require two identical hardware environments that are configured exactly the same way. While one environment is active and serving end users, the other environment remains idle."**

~ TechTarget.com

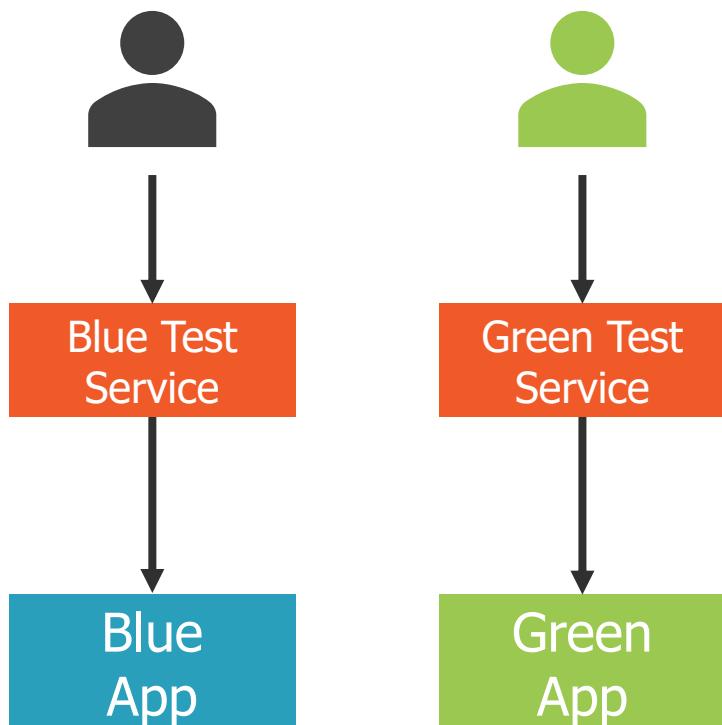


"A blue/green deployment is a change management strategy for releasing software code. Blue/green deployments, which may also be referred to as A/B deployments require two identical hardware environments that are configured exactly the same way. While one environment is active and serving end users, the other environment remains idle."

~ TechTarget.com



Blue-Green Deployments



Strategy for checking the viability of a deployment before it's publicly available

Run two identical production environments at the same time

New application (green) is deployed alongside the old application (blue)

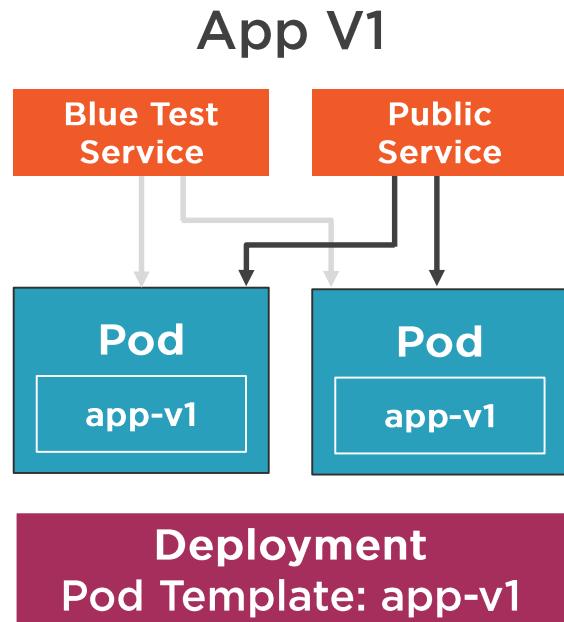
Traffic routed from blue to green when checks pass



Blue-Green Deployments

1

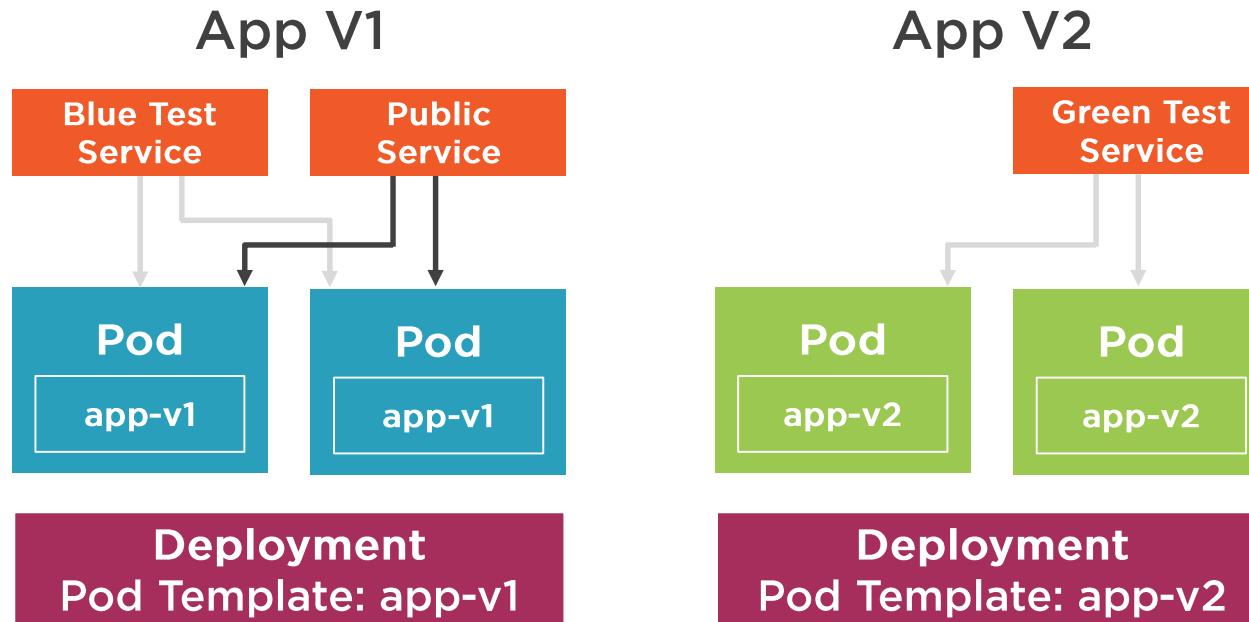
Create BLUE Deployment and Services



Blue-Green Deployments

2

Create GREEN Deployment and Service

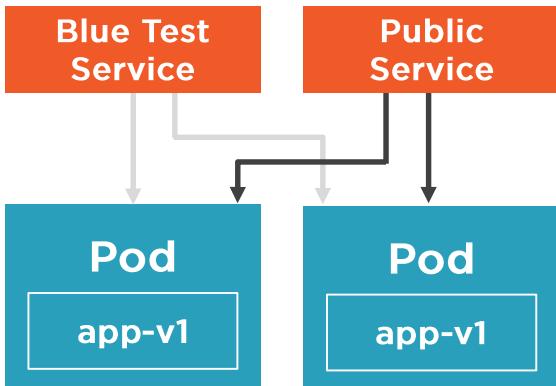


Blue-Green Deployments

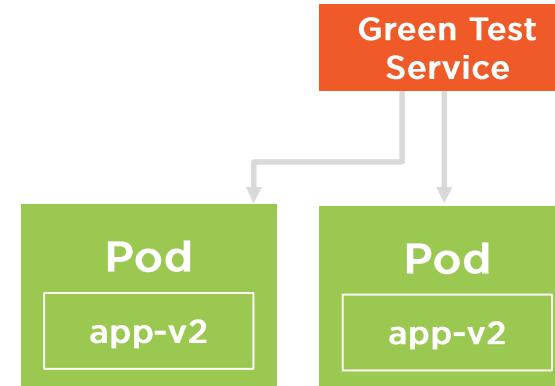
3

Test GREEN Pods

App V1



App V2



Deployment
Pod Template: app-v1

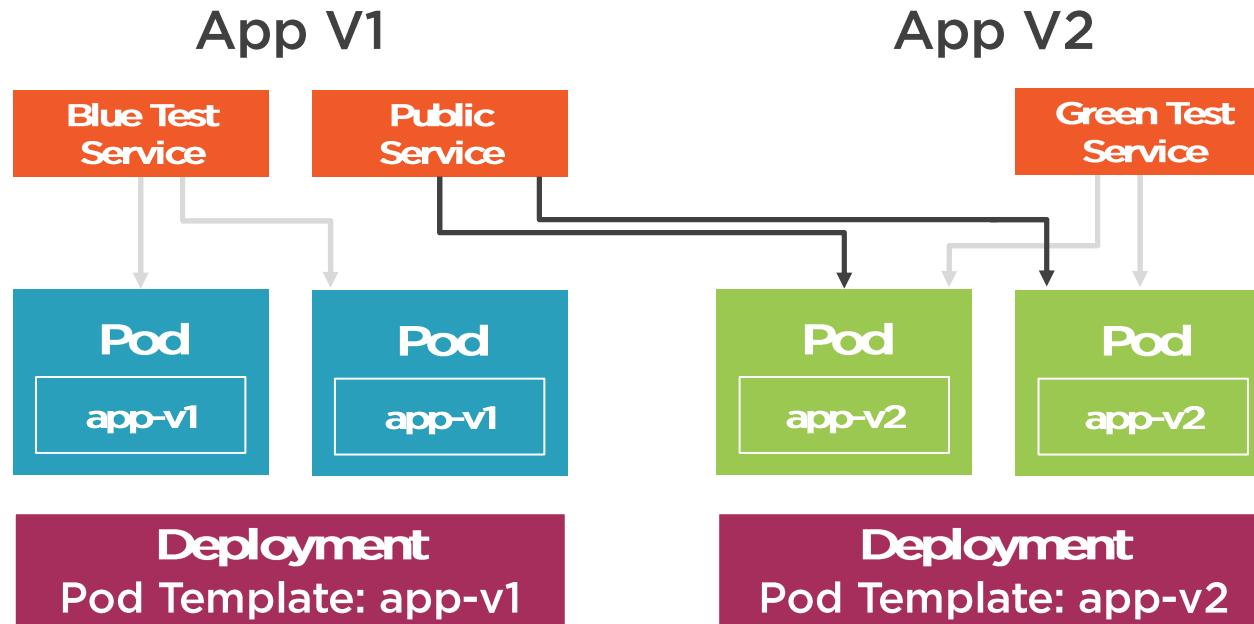
Deployment
Pod Template: app-v2



Blue-Green Deployments

4

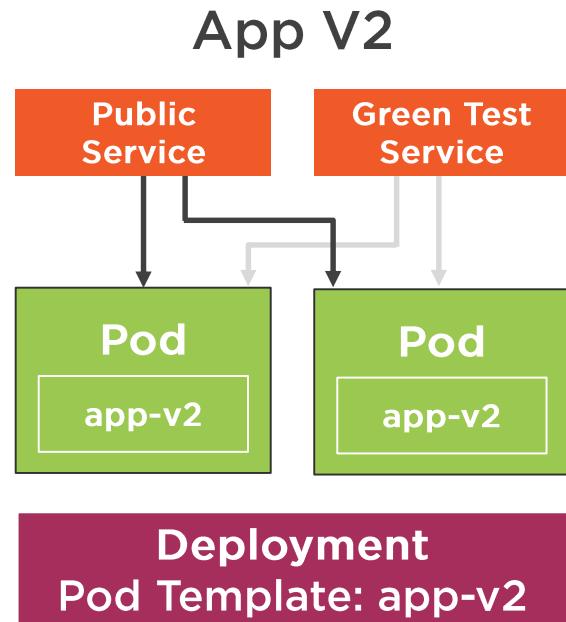
Change public Service from BLUE to GREEN



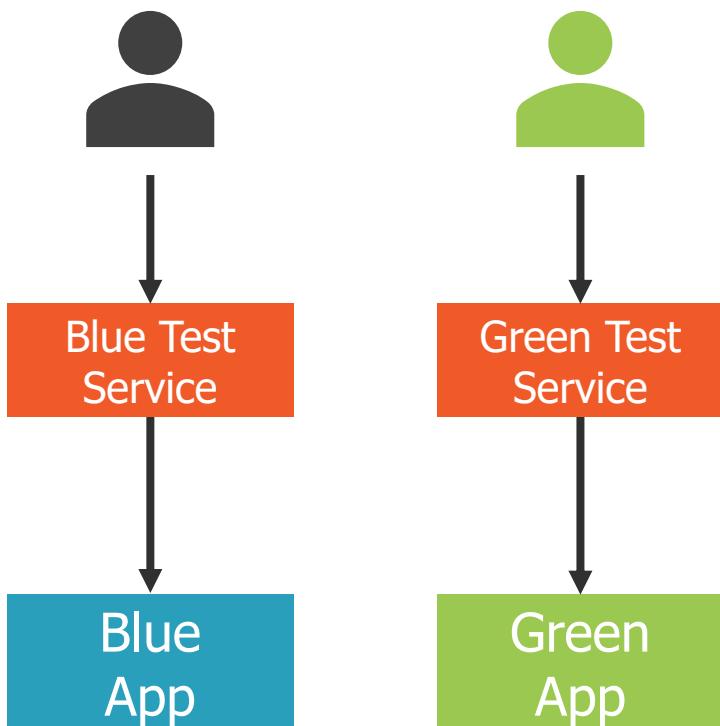
Blue-Green Deployments

5

Remove BLUE Deployment and Service



Blue-Green Deployment Considerations



Key considerations:

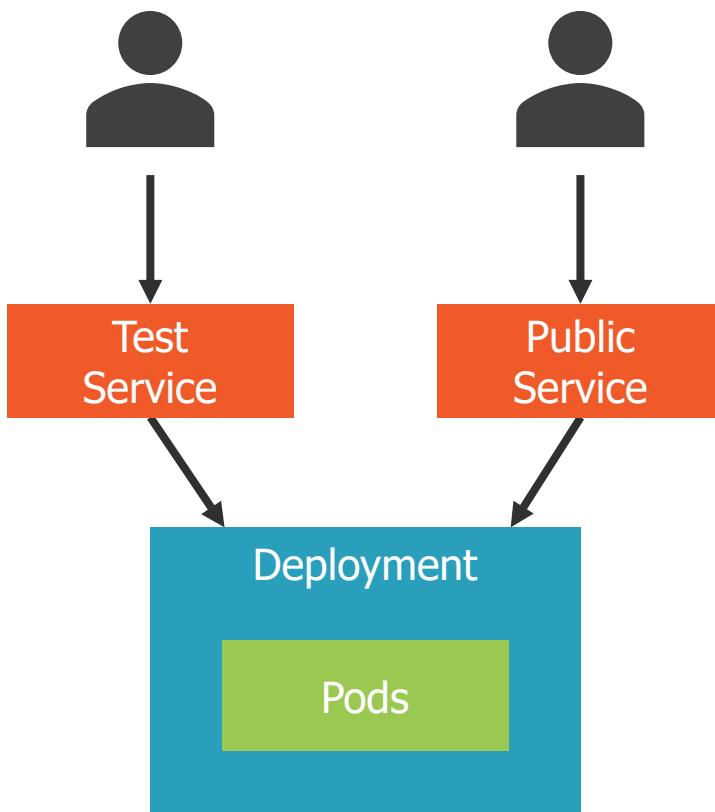
- How many Pods are being deployed to each environment?
- How much memory is required to run the Pods?
- What are the CPU requirements?
- Other considerations (volumes, sessions, node affinity, etc.)



Creating a Blue-Green Deployment



Blue-Green Deployment Resources



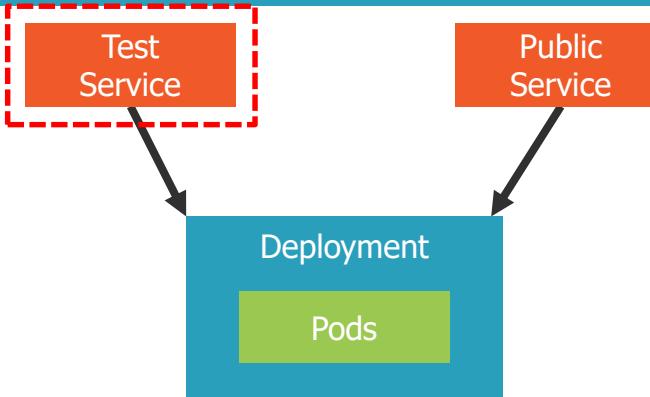
A Blue-Green Deployment involves 3 main Kubernetes resources:

- Test Service
- Public Service
- Deployment



Defining a Test Service

```
kind: Service
apiVersion: v1
metadata:
  name: nginx-blue-test
  labels:
    app: nginx
    role: blue-test
    env: test
spec:
  type: LoadBalancer
  selector:
    app: nginx
    role: blue
  ports:
    - port: 9000
      targetPort: 80
```



◀ Role that the service plays
(blue-test for test environment)

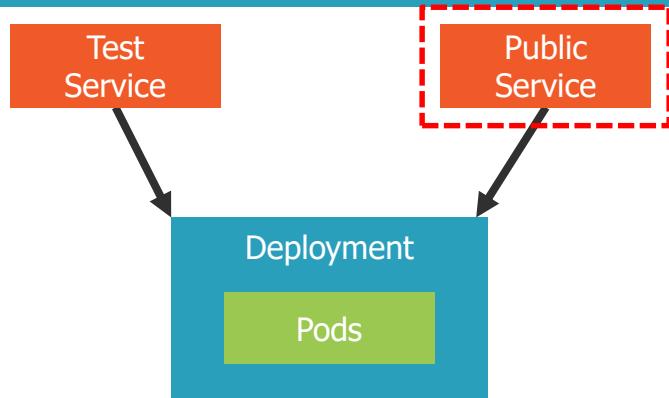
◀ Service will apply to Pods with
these labels (note the role: blue)

◀ Expose port 9000 (test port)



Defining a Public Service

```
kind: Service
apiVersion: v1
metadata:
  name: nginx-service
  labels:
    app: nginx
    role: blue
    env: prod
spec:
  type: LoadBalancer
  selector:
    app: nginx
    role: blue
  ports:
    - port: 80
      targetPort: 80
```



◀ Role that the service plays (blue for production environment)

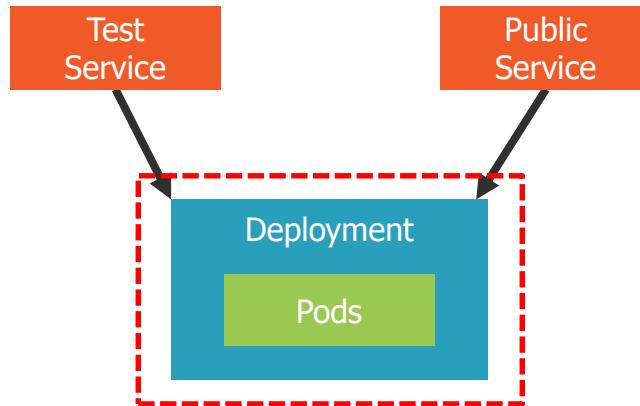
◀ Service will apply to Pods with these labels (note the role: blue)

◀ Expose port 80 (public port)



Defining a Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment-blue
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
      role: blue
  template:
    metadata:
      labels:
        app: nginx
        role: blue
  spec:
    containers:
      - name: nginx-blue
        image: nginx:1.x.x-alpine
    ports:
      - containerPort: 80
```



◀ Apply to Pods with these labels
(note the role: blue)

◀ Pod labels (note role: blue)

◀ Container image



Changing From Blue to Green

Once a **GREEN** deployment has been successfully rolled out and tested, change the public service's selector to "green"

```
selector:  
  app: nginx  
  role: green
```

Change role to green in
public service's selector

```
# Apply changes made to service's YAML (declarative)  
kubectl apply -f file.service.yml
```

```
# Change Service's selector to green (imperative)  
kubectl set selector svc [service-name] 'role=green'
```

Blue-Green Deployments in Action – The Blue Deployment



Blue-Green Deployments in Action – The Green Deployment



Summary



Blue-Green Deployments allow two environments to be deployed at the same time

Provides a way to test a new version of an application before switching over to it

Works by changing the "blue" Service's selector to point to the "green" Deployment



Running Jobs and CronJobs



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Understanding Jobs

Understanding CronJobs

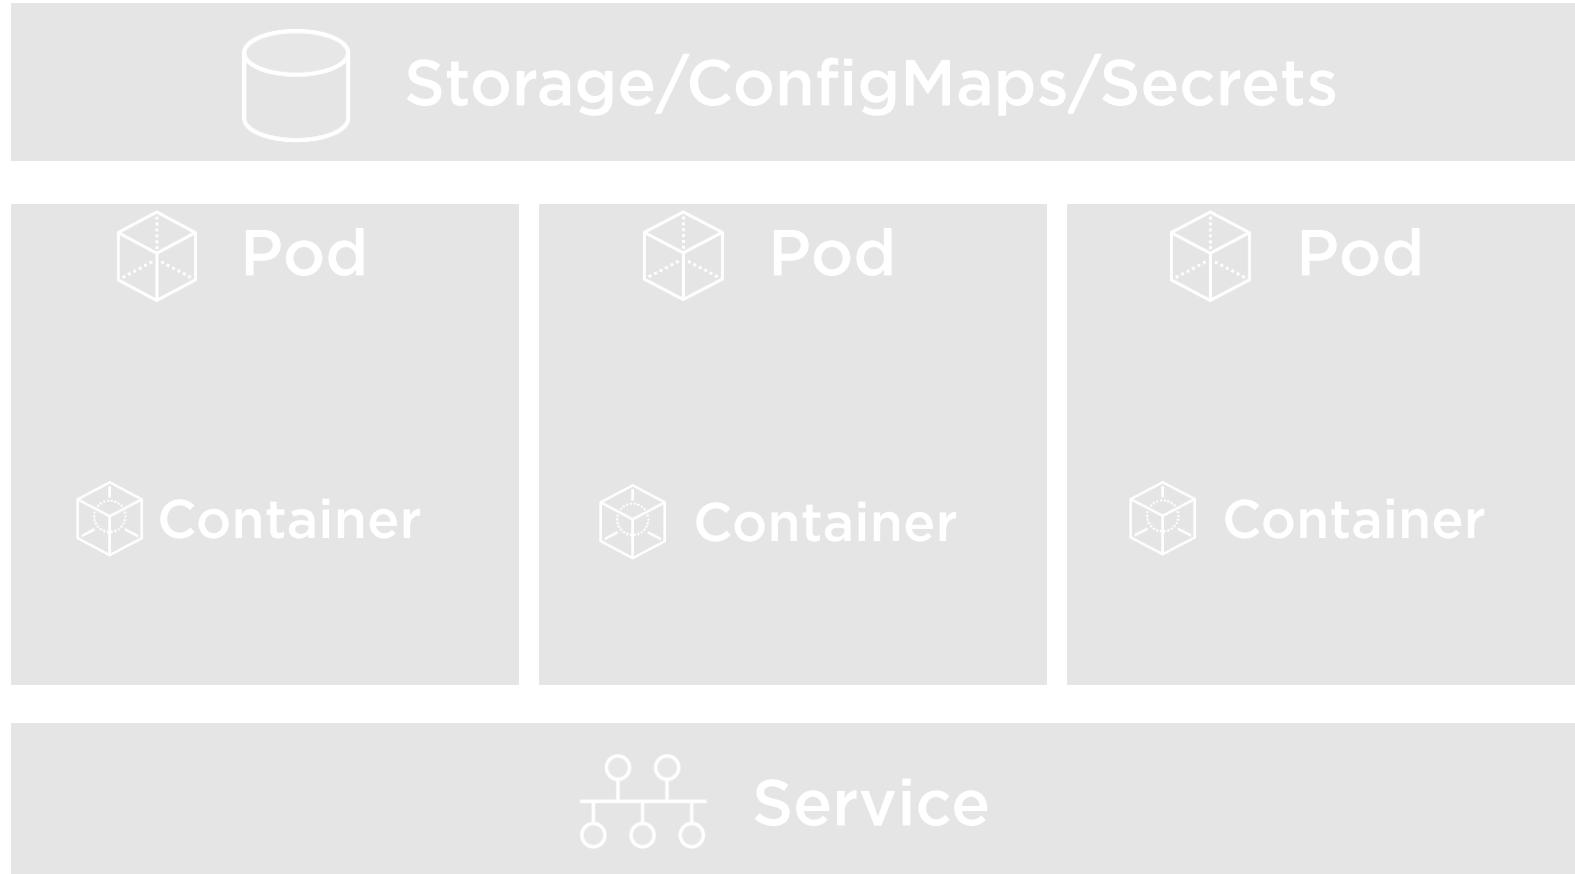
Creating a Job and CronJob

Jobs in Action

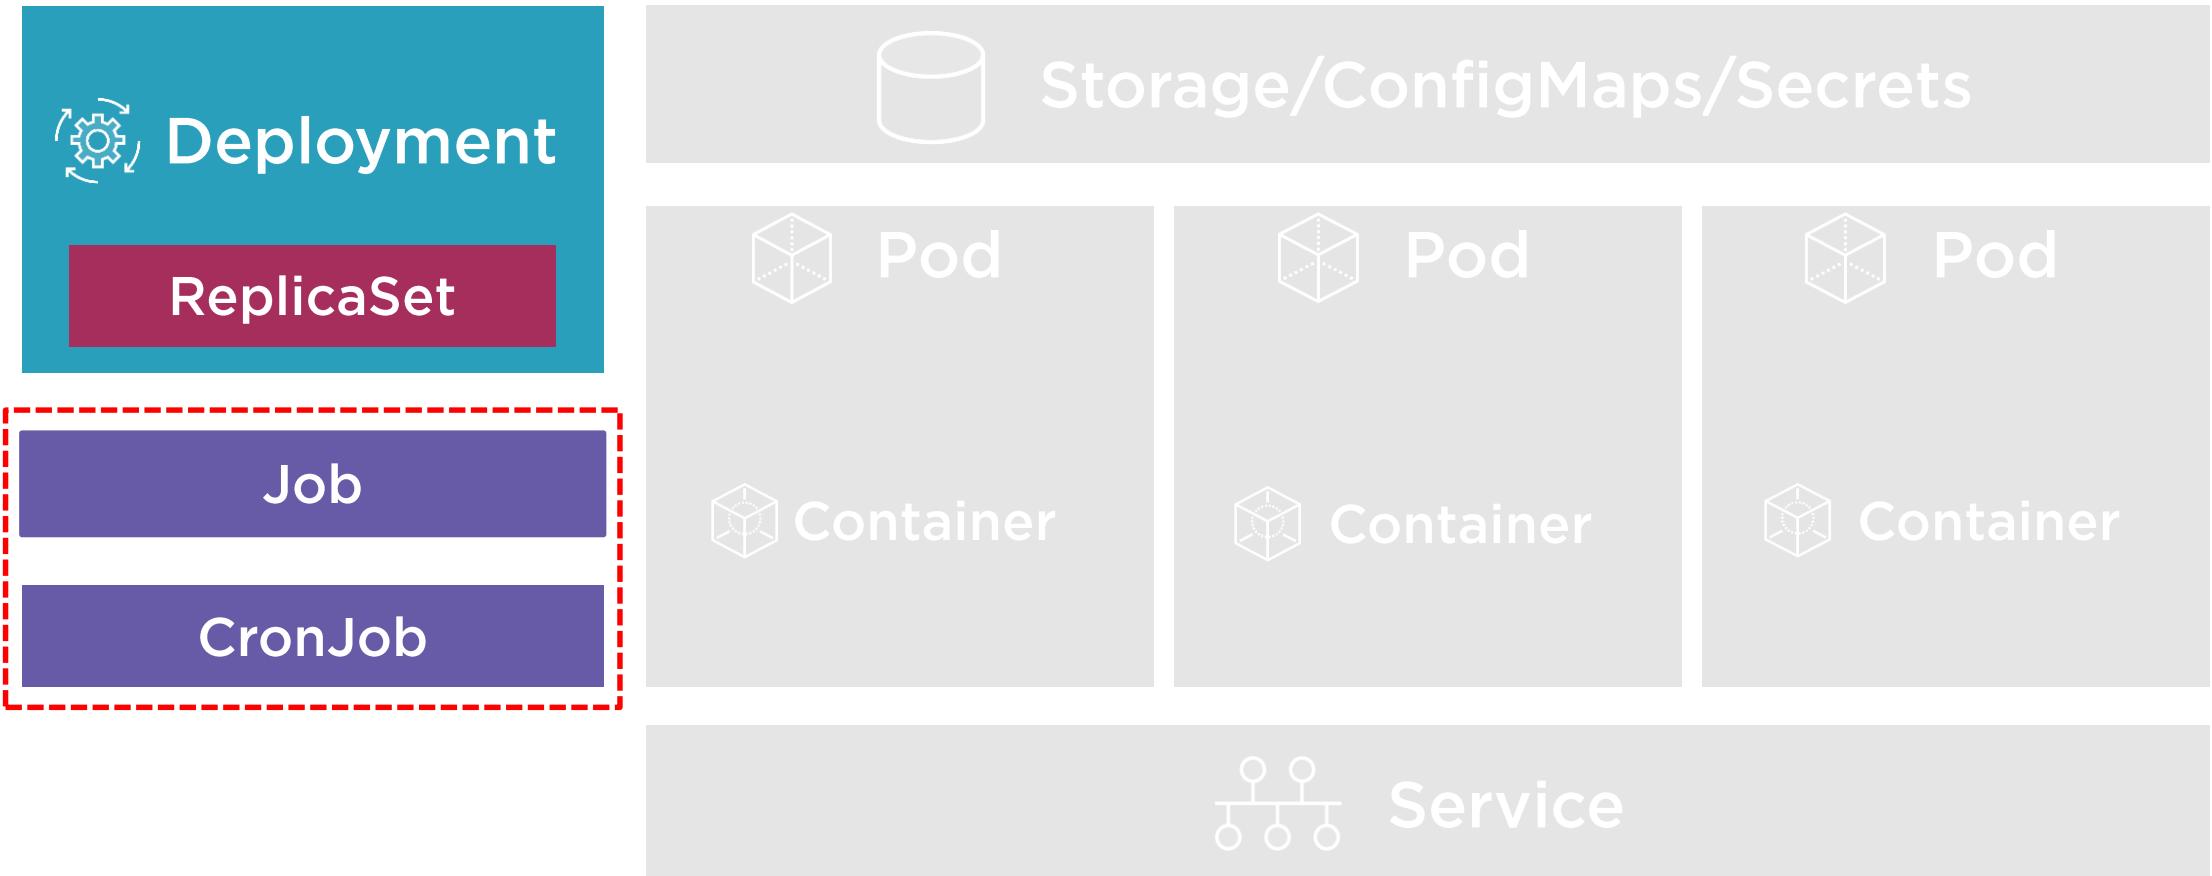
CronJobs in Action



Kubernetes Resources



Kubernetes Resources



Understanding Jobs



Have you ever needed to run a job that performs a task and then terminates?

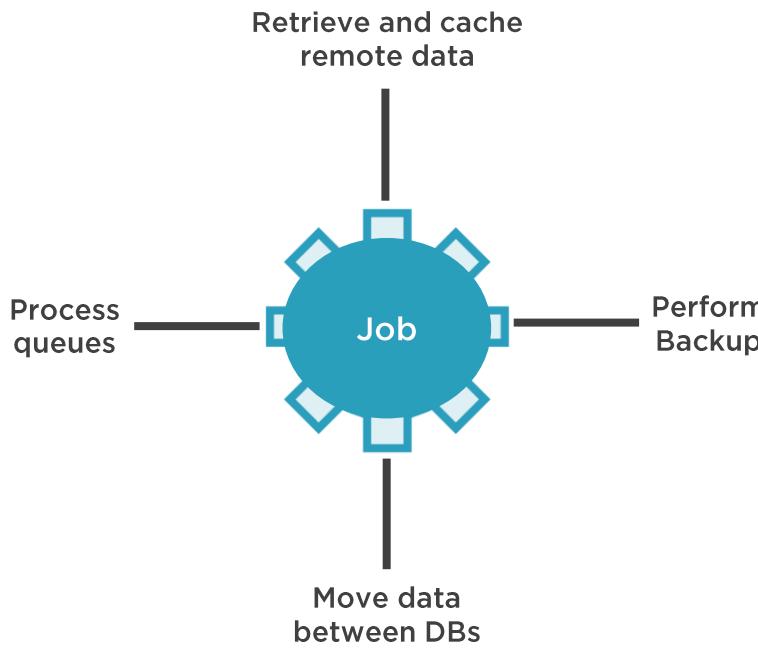


"A Job creates one or more Pods and ensures that a specified number of them successfully terminate."

~ Kubernetes Documentation



Understanding Jobs



A Job creates a Pod(s) that performs a task or batch process

Unlike standard Pods, a Job does not run indefinitely

A Job can be configured to run multiple Pods in parallel

Successful completions are tracked

Once a Job is deleted its Pods are removed



Understanding CronJobs



Have you ever needed to run a job on a scheduled basis?

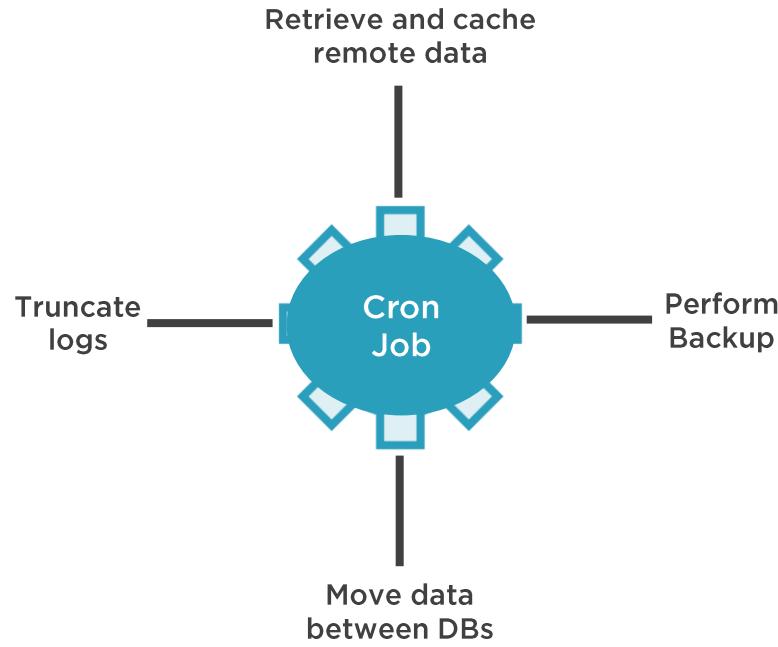


"A Cron Job creates Jobs on a time-based schedule."

~ Kubernetes Documentation



Understanding CronJobs



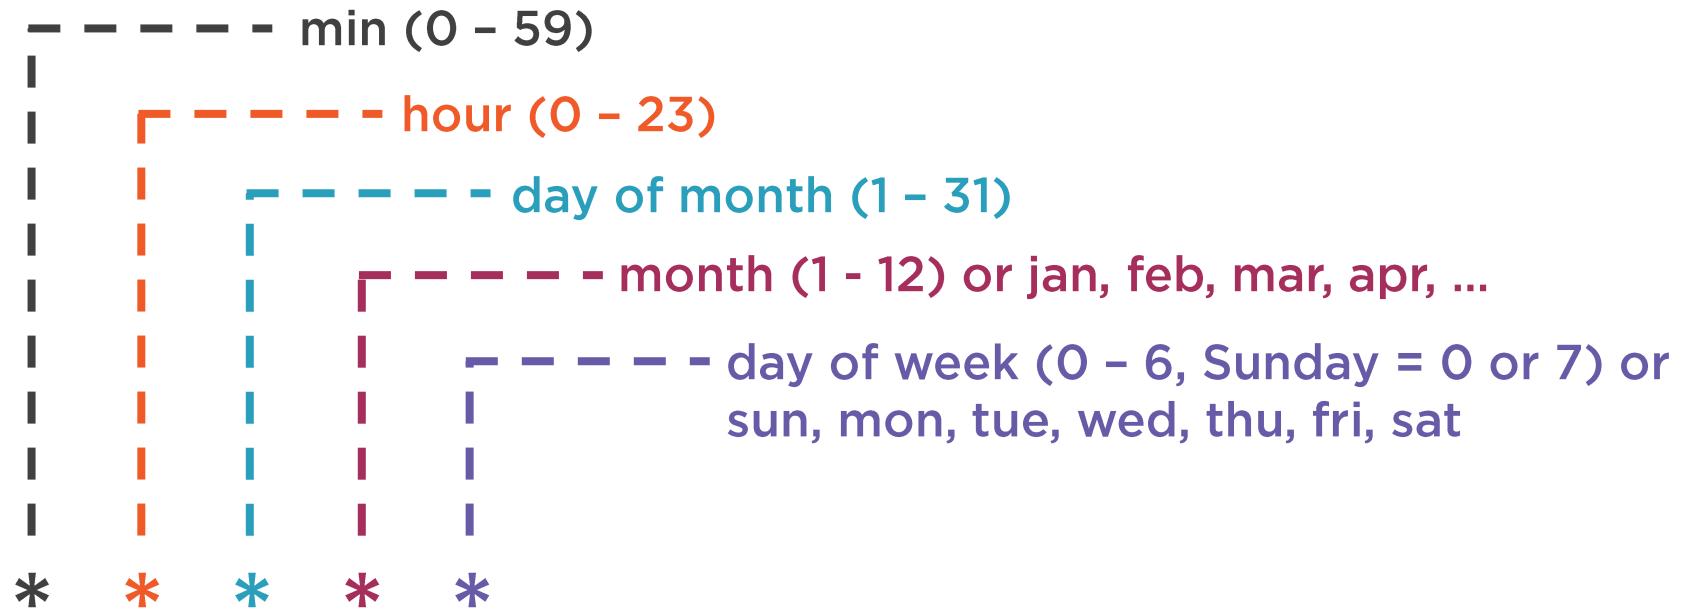
A CronJob is a Job that runs on a scheduled basis

Scheduled using the Cron format

CronJob names must be less than 52 characters

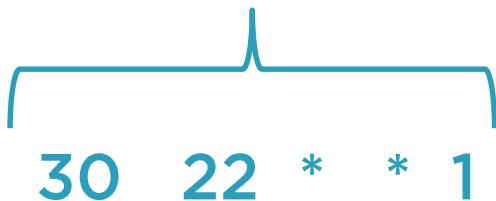


Understanding the Cron Format

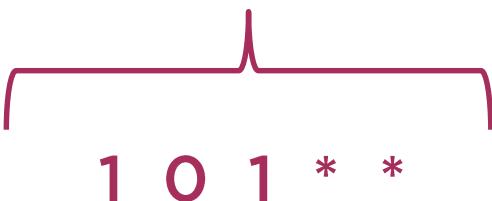


Cron Format Examples

Run at 22:30 every Monday



Run at 00:01 on the first day of each month



Additional Cron Formats

0 * * * *

0 0 * * *

0 0 * * 0

0 0 1 * *

0 0 1 1 *

*/1 * * * *

- ◀ **@hourly** – run once every hour
- ◀ **@daily** – run once every day at midnight
- ◀ **@weekly** – run once every week
- ◀ **@monthly** – run once every month
- ◀ **@yearly** – run once every year
- ◀ Run once every minute



Creating a Job and CronJob



Defining a Job

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pie-counter
spec:
  template:
    metadata:
      name: pie-counter
    spec:
      restartPolicy: Never
      containers:
        - name: pie-counter
          image: alpine
          command:
            - "sh"
            - "-c"
            - "echo 'scale=1000; 4*a(1)' ...;"
```

◀ Batch API
◀ Job kind

◀ Never try to restart (Never or OnFailure)

◀ Job command to run



Defining a Job that Requires Multiple Completions

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pie-counter
spec:
  completions: 4
  template:
    ...

```

◀ Run 4 Pods sequentially



Defining a Job that Can Run in Parallel

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pie-counter
spec:
  completions: 4
  parallelism: 2
  template:
    ...

```

- ◀ 4 Pods must complete successfully
- ◀ 2 Pods can run in parallel at a time



Creating a Job

A job can be created using the standard **kubectl create** or **kubectl apply** commands

```
# Create a new Job  
kubectl create -f file.job.yml --save-config
```

```
# Creating or modifying a Job  
kubectl apply -f file.job.yml
```

Defining a CronJob

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: pie-counter
spec:
  concurrencyPolicy: Allow
  # Run the job every minute
  schedule: "*/1 * * * *"
jobTemplate:
  spec:
    template:
      spec:
        restartPolicy: OnFailure
        containers:
        - name: pie-counter
          image: alpine
          command:
          - "sh"
          - "-c"
          - "echo 'scale=1000; 4*a(1)'"
```

◀ CronJob batch API

◀ CronJob kind

◀ Allow multiple Pods to run event if their scheduling overlaps

◀ Cron format to use for scheduling

◀ Restart if there's a failure

◀ Command to run



Creating a CronJob

A CronJob can be created using the standard **kubectl create** or **kubectl apply** commands

```
# Create a new CronJob
kubectl create -f file.cronjob.yml --save-config
```

```
# Creating or modifying a CronJob
kubectl apply -f file.cronjob.yml
```

Jobs in Action



CronJobs in Action



Summary



Jobs are used to run a task or batch process

Successful completions are tracked

CronJobs allow a task/batch process to be run on a scheduled basis

Relies on the Cron format



Performing Monitoring and Troubleshooting Tasks



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

**Monitoring and Troubleshooting
Overview**

Web UI Dashboard in Action

**Prometheus, Metrics Server, and
kube-state-metrics in Action**

Grafana in Action

**Troubleshooting Techniques
in Action**



Monitoring and Troubleshooting Overview



How important is it to monitor Deployments, Pods, and other resources in your Kubernetes cluster?



Q What should you monitor?

A Short answer...everything!



"Kubernetes makes managing a containerized infrastructure much easier by creating levels of abstractions such as pods and services. We no longer have to worry about where applications are running or if they have enough resources to work properly."

<https://kubernetes.io/blog>

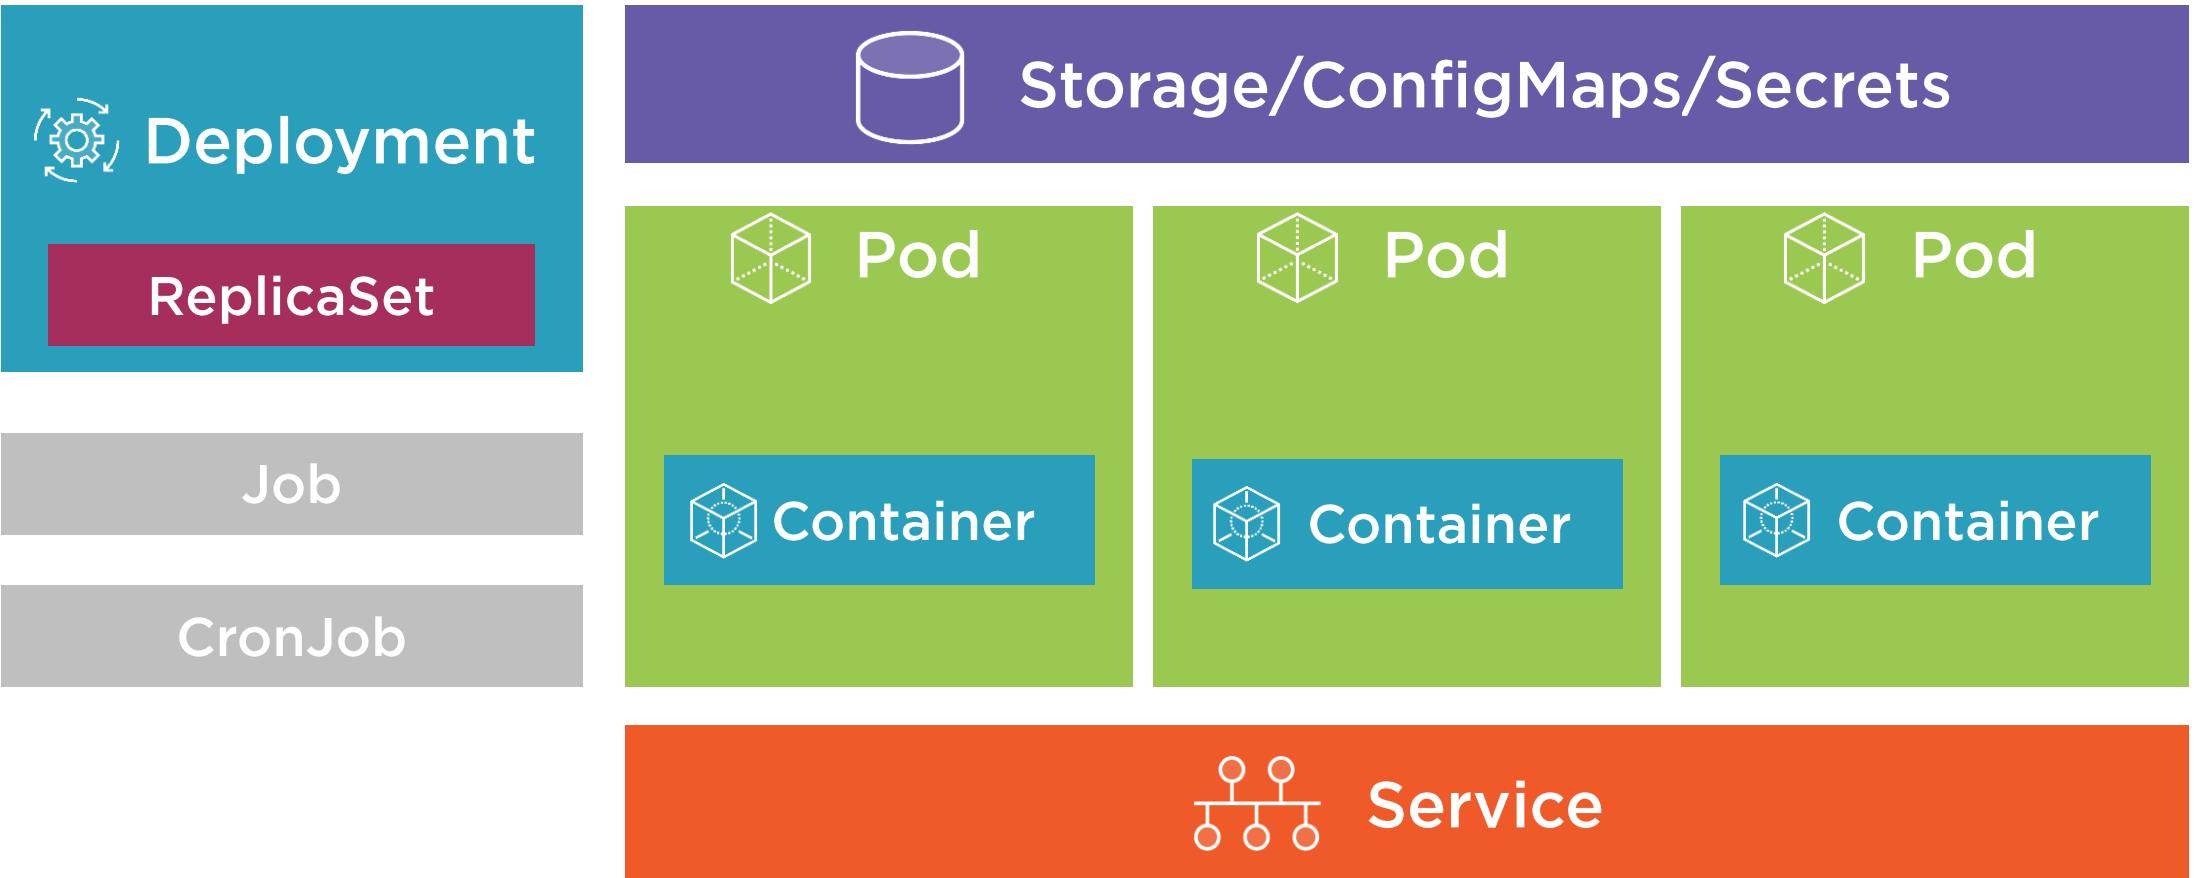


"Kubernetes makes managing a containerized infrastructure much easier by creating levels of abstractions such as pods and services. We no longer have to worry about where applications are running or if they have enough resources to work properly. But that doesn't change the fact that, in order to ensure good performance, we need to monitor our applications, the containers running them, and Kubernetes itself."

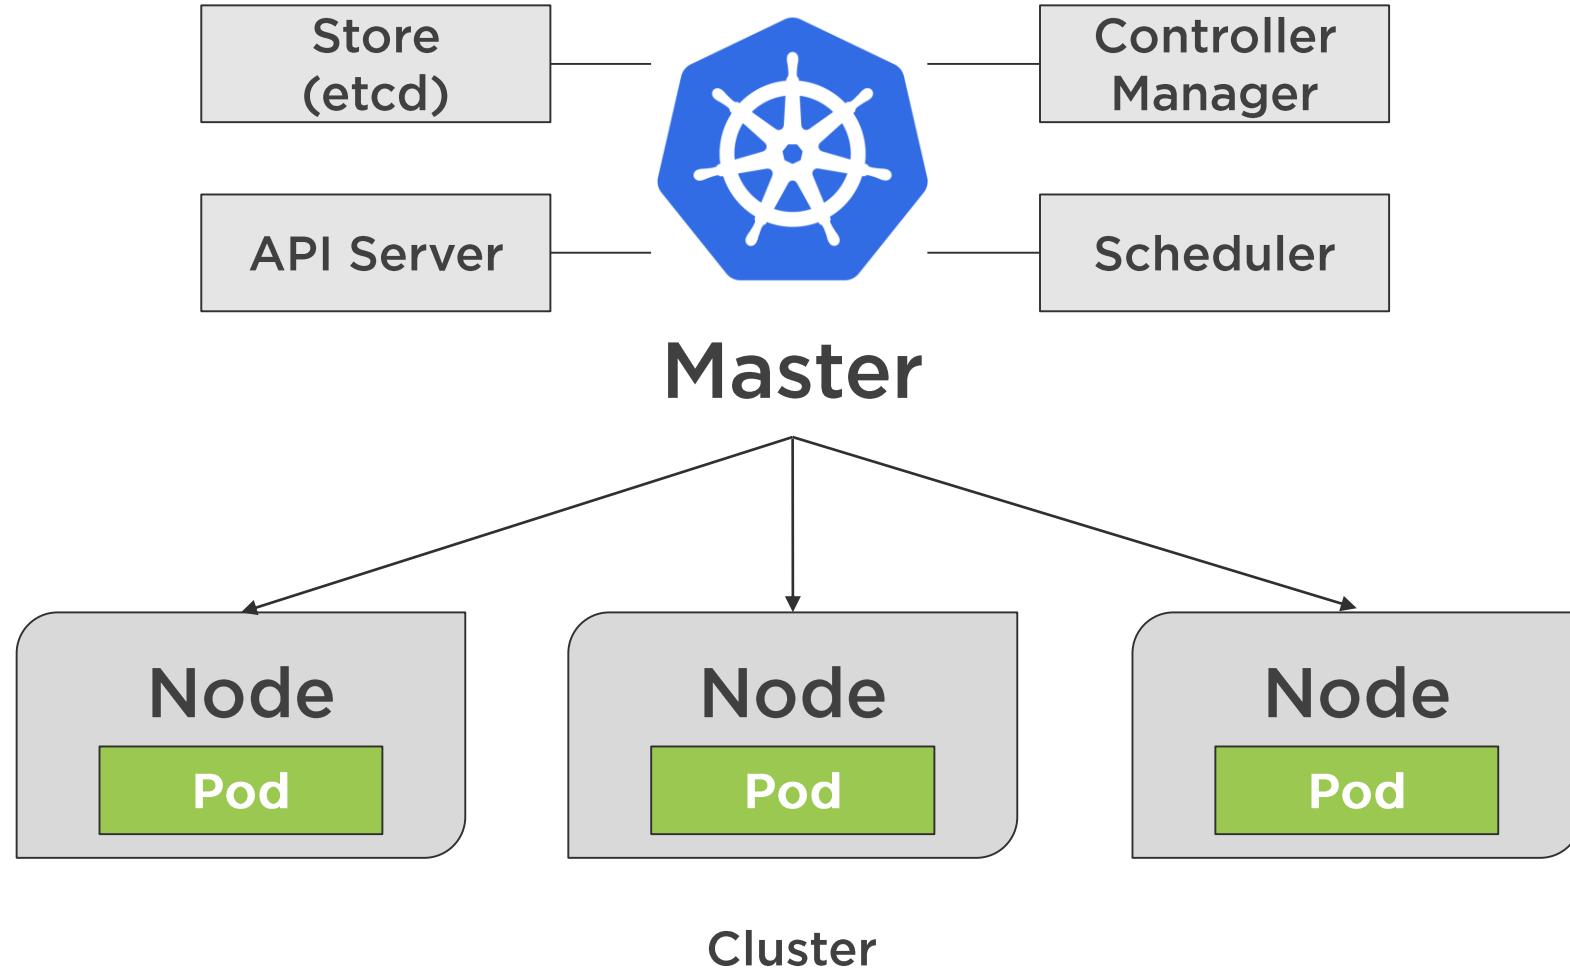
<https://kubernetes.io/blog>

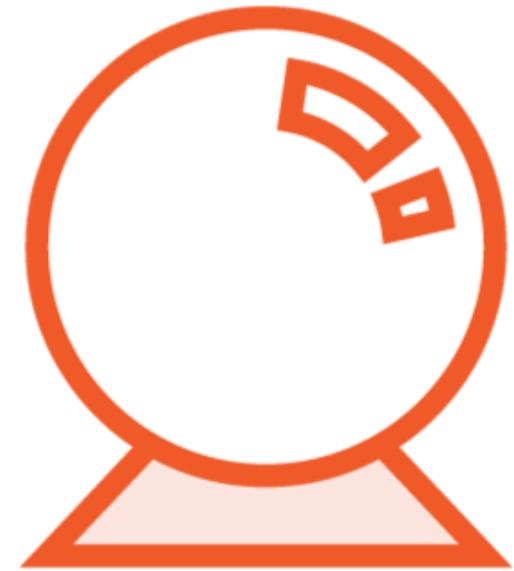


What Should You Monitor?



What Should You Monitor?





Monitoring and Troubleshooting Options



Key monitoring and alerting options:

- Web UI Dashboard
- Metrics Server
- kube-state-metrics
- Prometheus
- Grafana
- Many more...



Web UI Dashboard in Action



"Web UI (Dashboard) is a web-based Kubernetes user interface. You can use Dashboard to deploy containerized applications to a Kubernetes cluster, troubleshoot your containerized application, and manage the cluster resources."

<https://kubernetes.io>

<https://github.com/kubernetes/dashboard>



Metrics Server, kube-state-metrics, and Prometheus in Action



"Metrics Server is a cluster-wide aggregator of resource usage data. It is deployed by default in clusters created by kube-up.sh script as a Deployment object."

<https://kubernetes.io>

<https://github.com/kubernetes-sigs/metrics-server>



"kube-state-metrics is a simple service that listens to the Kubernetes API server and generates metrics about the state of the objects. It is not focused on the health of the individual Kubernetes components, but rather on the health of the various objects inside, such as deployments, nodes and pods."

<https://github.com/kubernetes/kube-state-metrics>

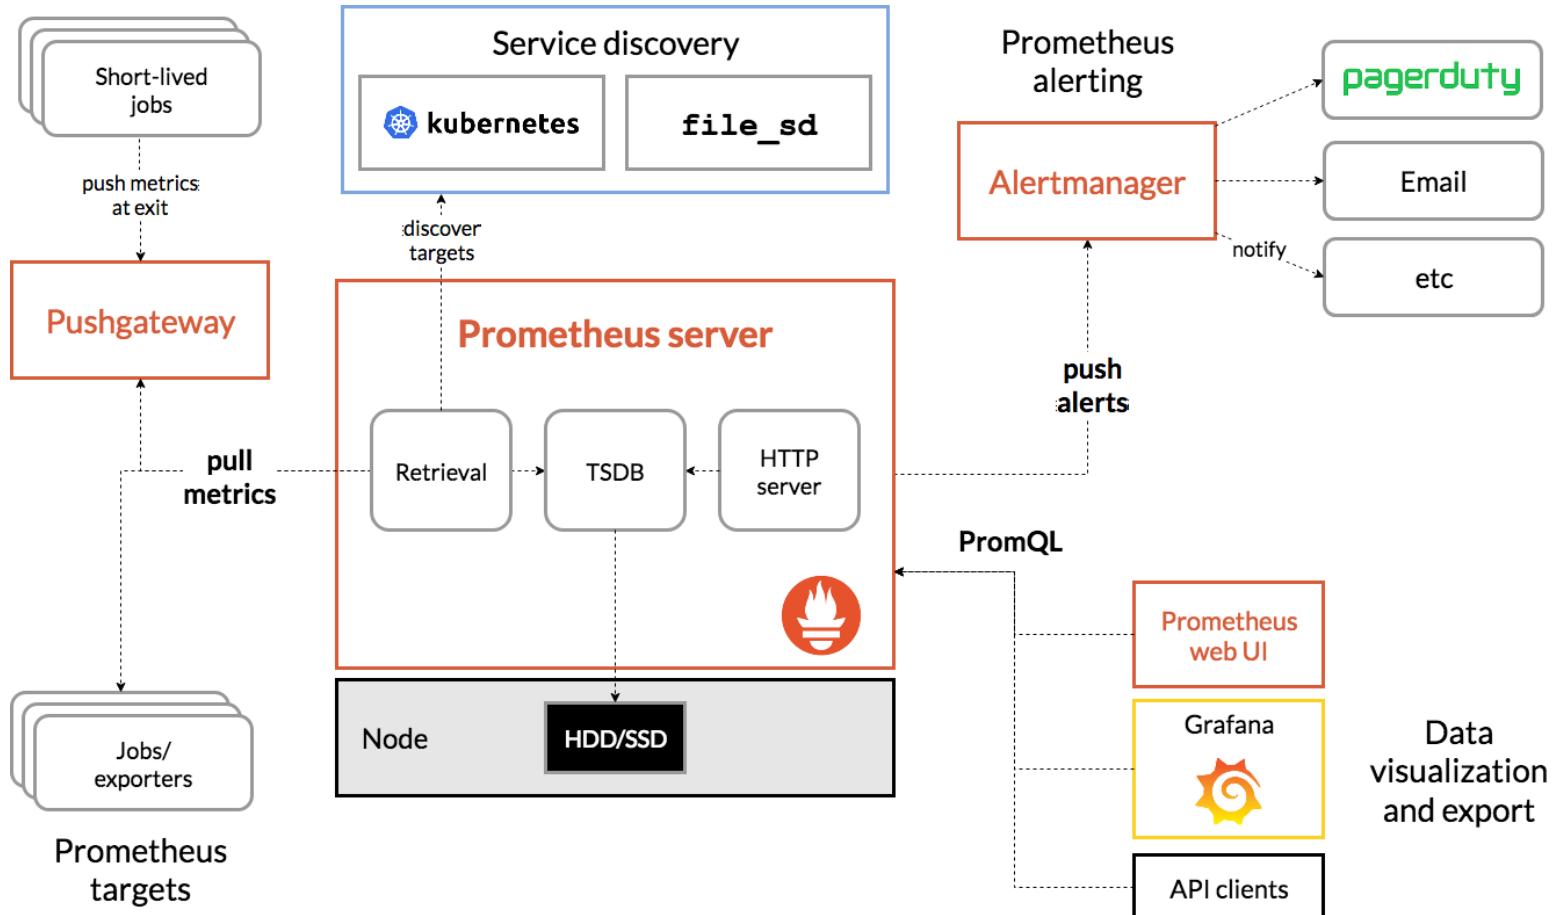


"Prometheus is an open-source systems monitoring and alerting toolkit originally built at SoundCloud....Prometheus joined the Cloud Native Computing Foundation in 2016 as the second hosted project, after Kubernetes."

<https://prometheus.io>



Prometheus Architecture



<https://prometheus.io/docs/introduction/overview>



Grafana in Action



"Grafana allows you to query, visualize, alert on and understand your metrics no matter where they are stored. Create, explore, and share dashboards with your team and foster a data driven culture."

<https://grafana.com>



Troubleshooting Techniques with kubectl



Key Troubleshooting Commands



Key troubleshooting commands:

- `kubectl get pod [pod-name] -o yaml`
- `kubectl describe pod [pod-name]`
- `kubectl exec [pod-name] -it sh`

Viewing Pod logs

- `kubectl logs [pod-name]`
- `kubectl logs [pod-name] -c [container-name]`
- `kubectl logs -p [pod-name]`
- `kubectl logs -f [pod-name]`



Troubleshooting Techniques in Action



Summary



Monitoring should be a critical part of any Kubernetes rollout plan

Several tools can be used to provide monitoring and alerts:

- Web UI Dashboard
- Prometheus
- Grafana
- Many more...

Master kubectl commands that can help troubleshoot issues with Kubernetes resources



Putting It All Together



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Thank You!



Dan Wahlin
Wahlin Consulting

@DanWahlin
<https://codewithdan.com>

