

Kubernetes for Developers: Core Concepts

KUBERNETES FROM A DEVELOPER PERSPECTIVE



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Course Overview

Kubernetes from a
Developer Perspective

Understanding Storage
Options

Creating Pods

Creating ConfigMaps and
Secrets

Creating Deployments

Putting It All Together

Creating Services

Course Summary



Target Audience



**Developers looking to understand
Kubernetes core concepts**



Course Prereqs



Comfortable using command-line tools and virtual machines

General familiarity with software development

Understanding of Docker containers and how they work



Introduction



Module Overview

Kubernetes Overview

The Big Picture

Benefits and Use Cases

Running Kubernetes Locally

Getting Started with kubectl

Web UI Dashboard



Kubernetes Overview

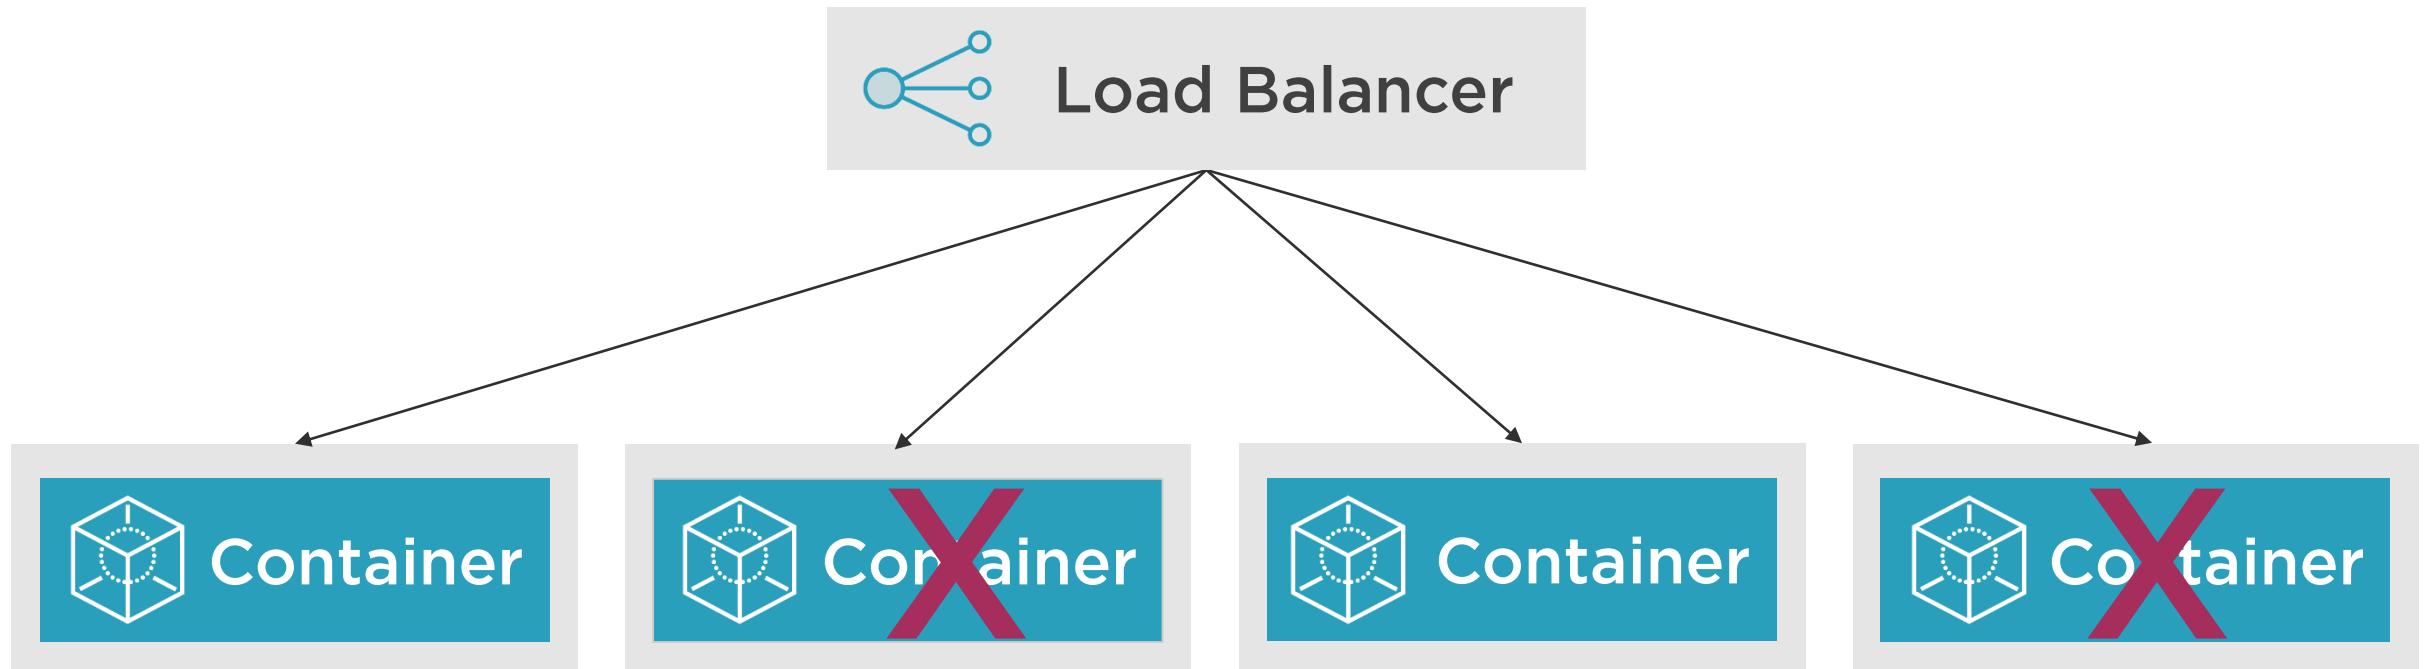


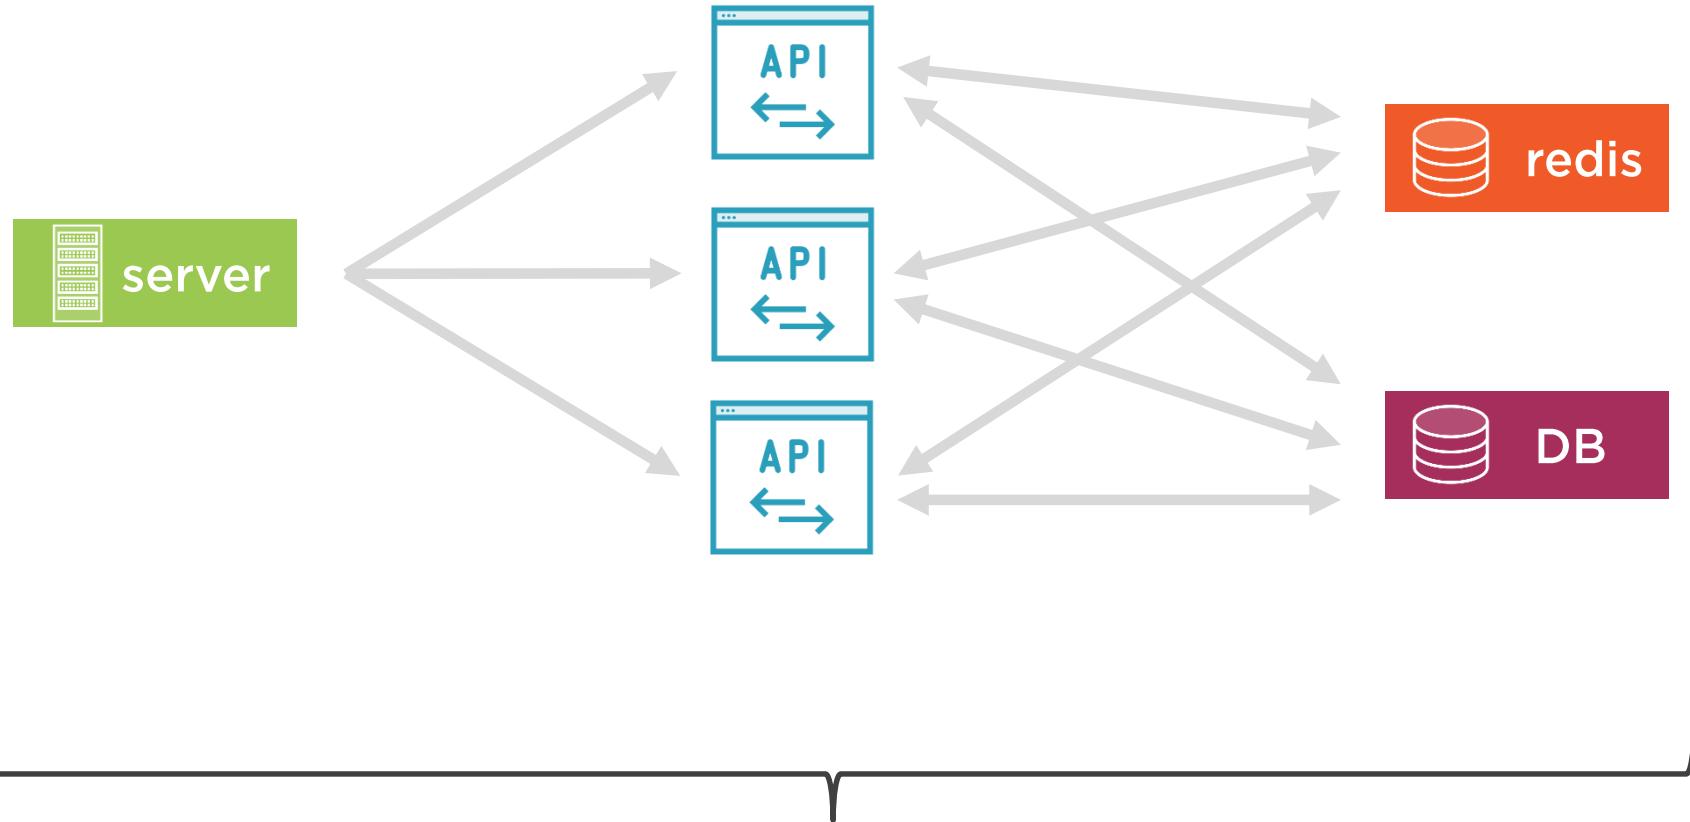
"Kubernetes (K8s) is an open-source system for automating deployment, scaling, and management of containerized applications."

<https://kubernetes.io>



How Are You Managing Containers?





How do you manage all of these containers?



It Would Be
Nice if We Could...



Package up an app and let something else manage it for us

Not worry about the management of containers

Eliminate single points of failure

Scale containers

Update containers without bringing down the application

Have robust networking and persistent storage options





Container

Container

Kubernetes
(conductor)

Container

Kubernetes is the conductor of a container orchestra.



Key Kubernetes Features

Service Discovery/
Load Balancing

Storage
Orchestration

Automate
Rollouts/Rollbacks

Self-healing

Secret and
Configuration
Management

Horizontal Scaling



The Big Picture



Kubernetes



Container and cluster management

Open source project

**Used internally by Google for 15+ years and
donated to the Cloud Native Computing
Foundation**

Supported by all major cloud platforms

**Provides a "declarative" way to define a
cluster's state**



Kubernetes Moves You to a Desired State

Current State



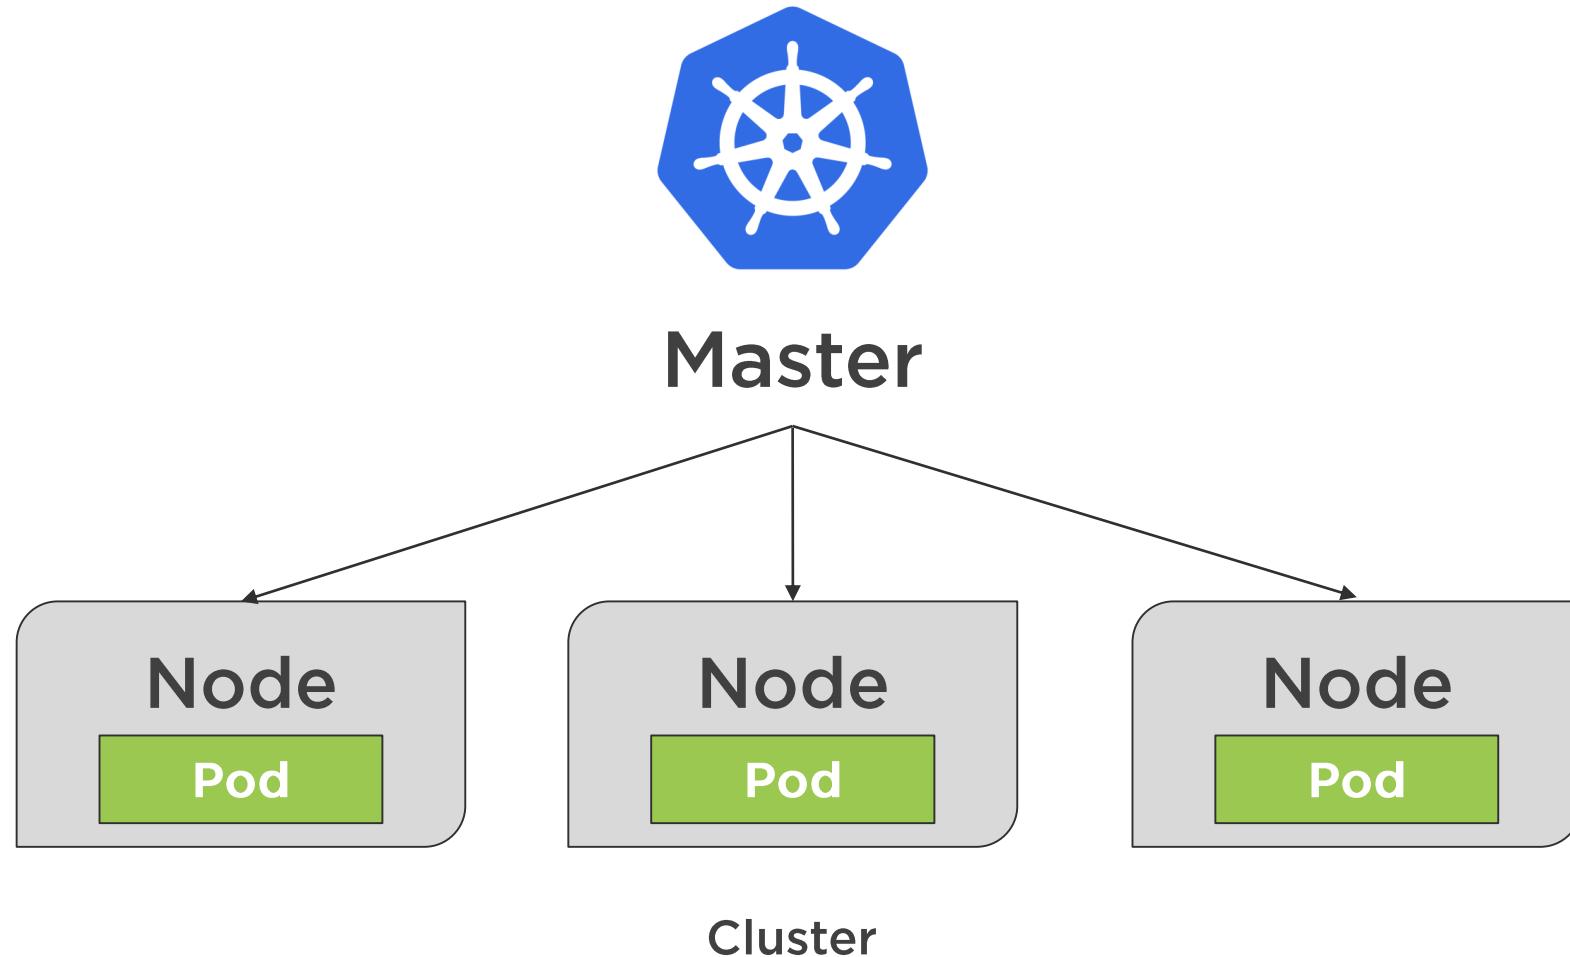
Kubernetes



Desired State



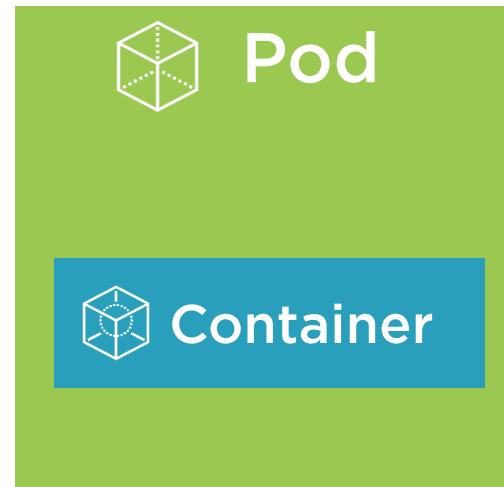
The Big Picture



Pods and Containers

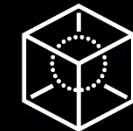
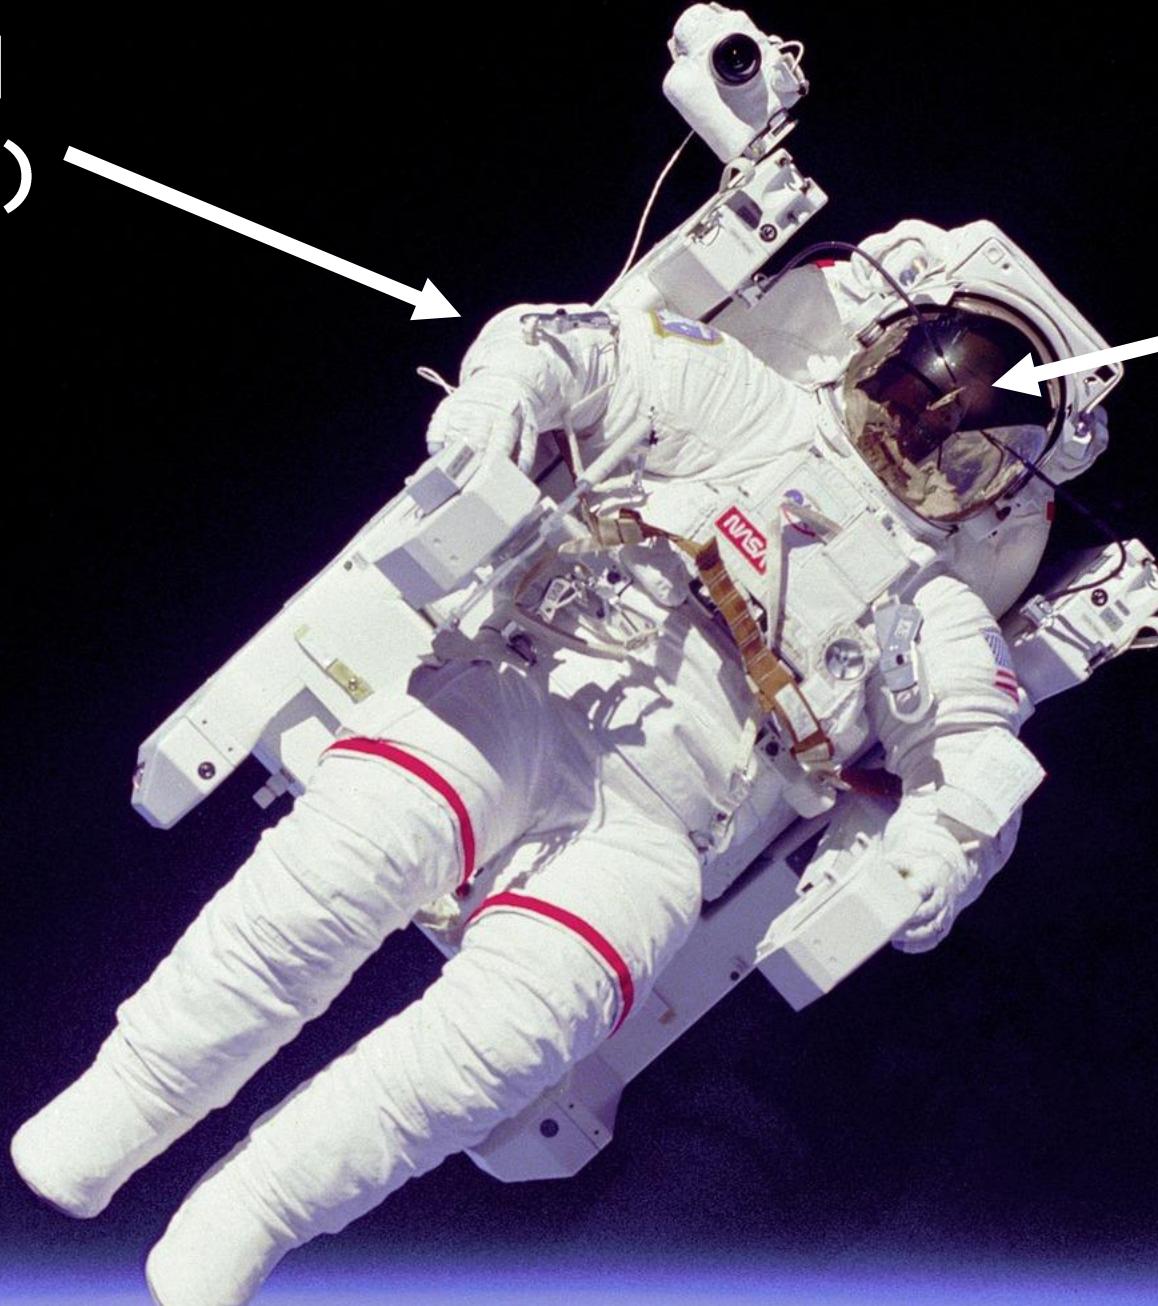


Pods and Containers



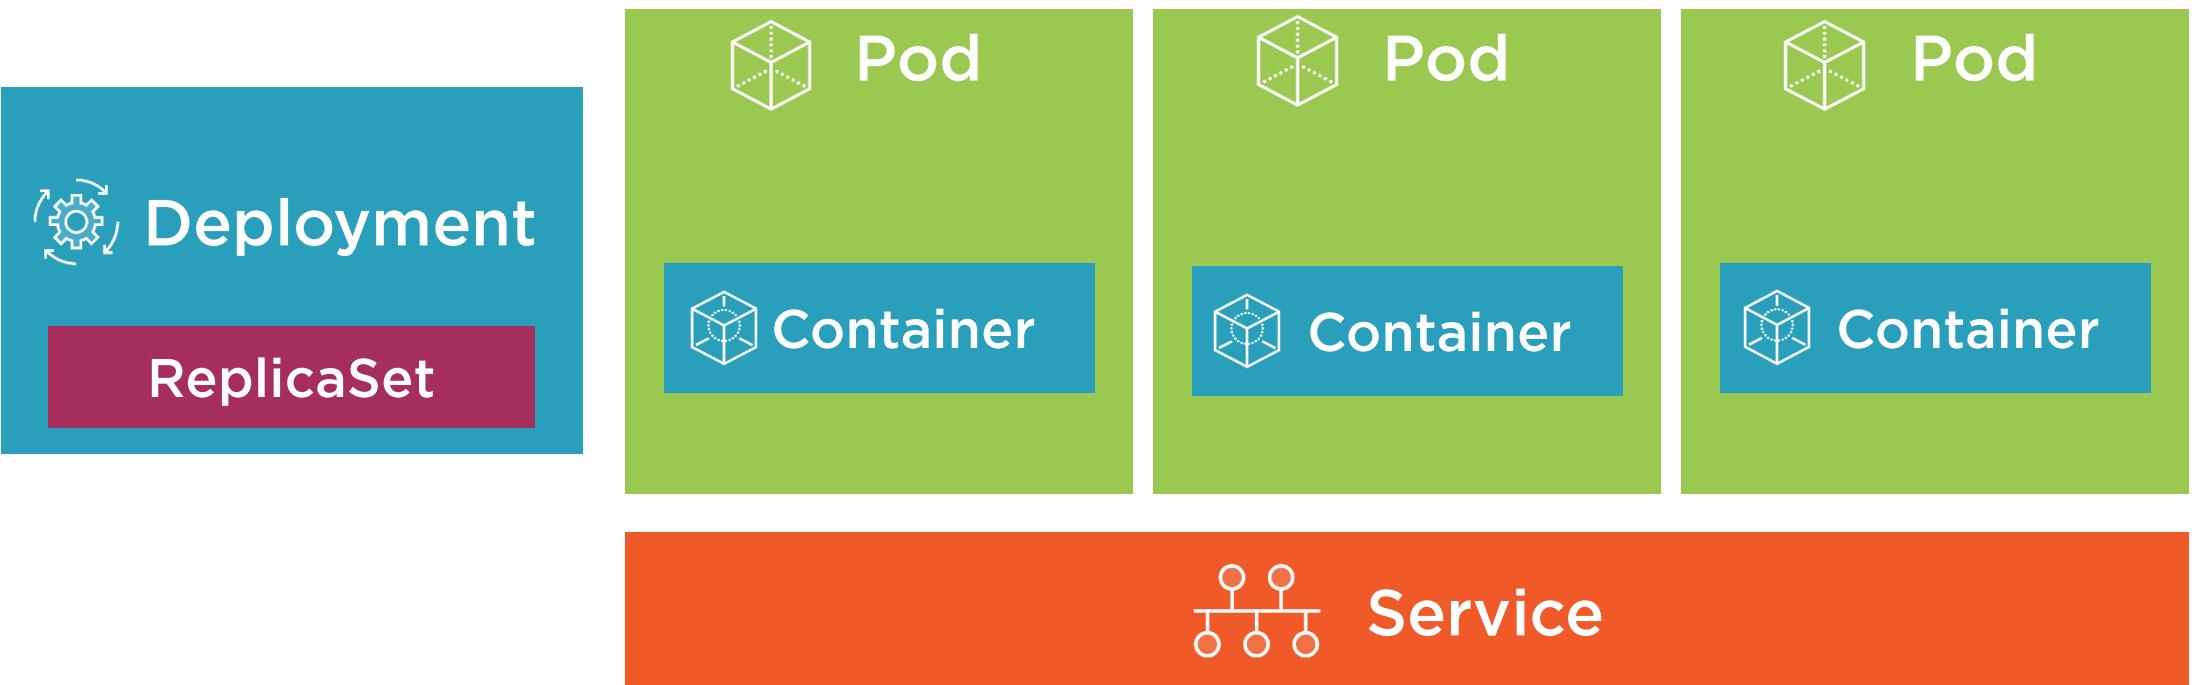


Pod
(suit)

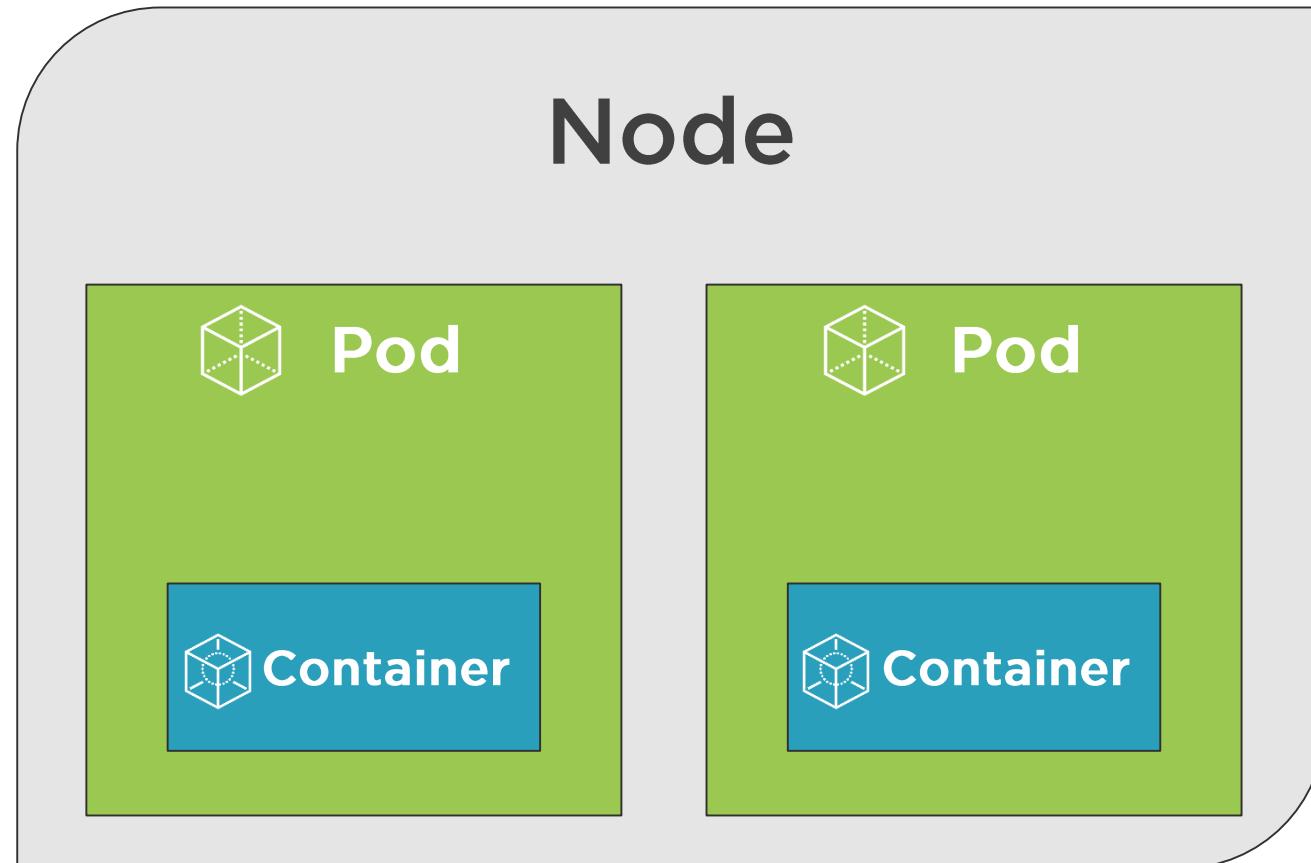


Container
(person)

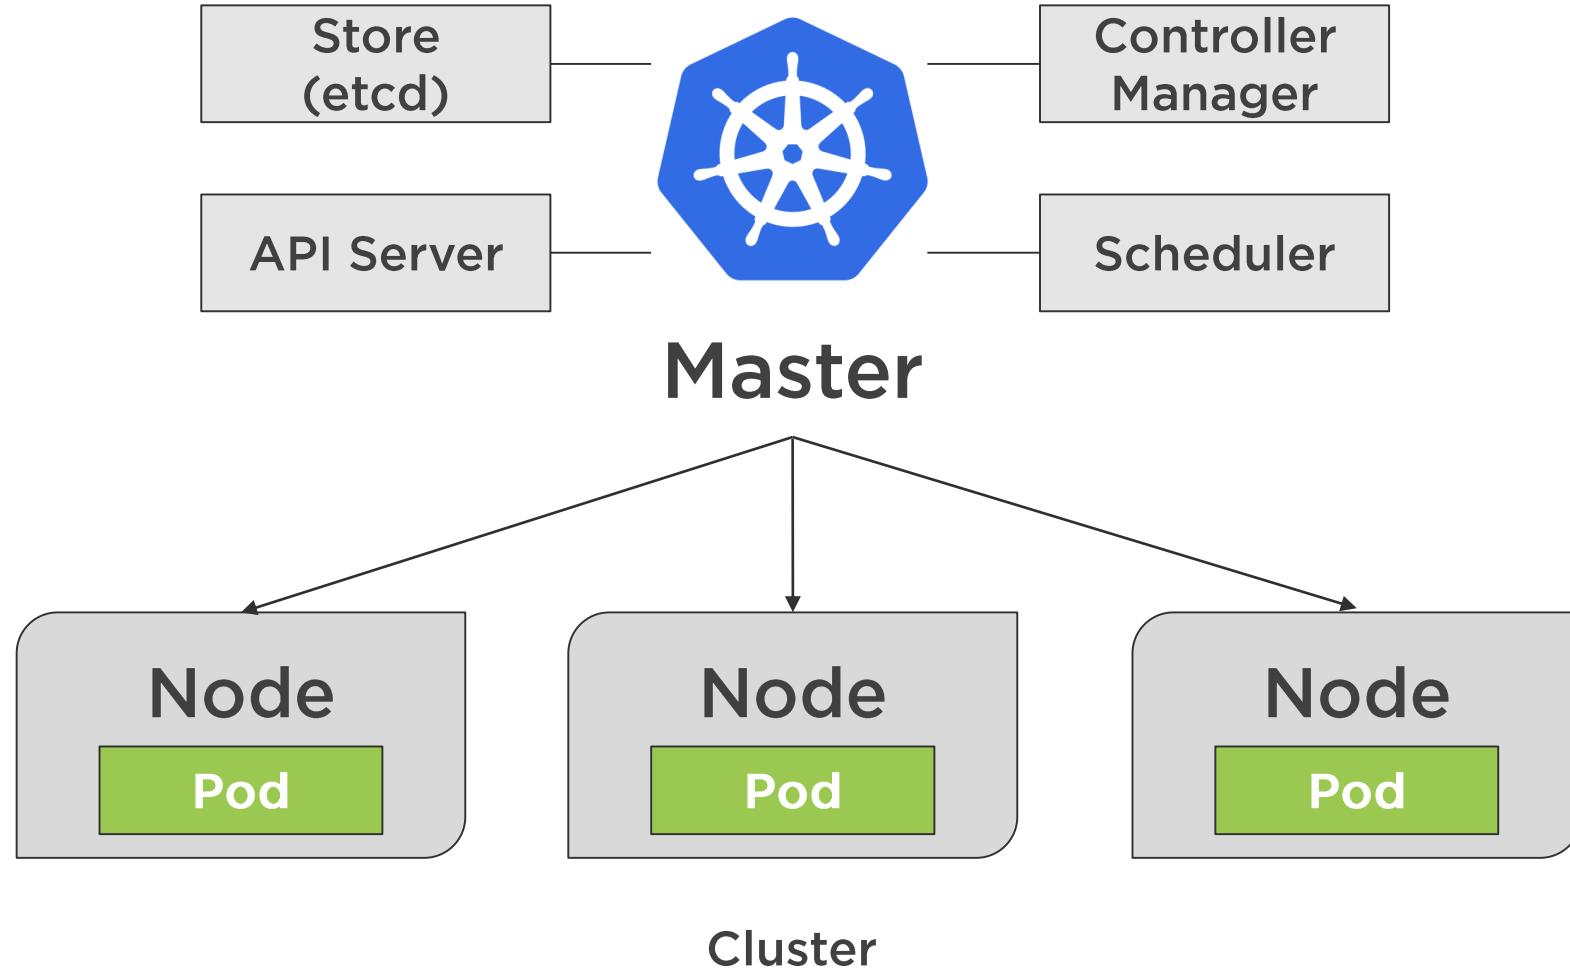
Kubernetes Building Blocks



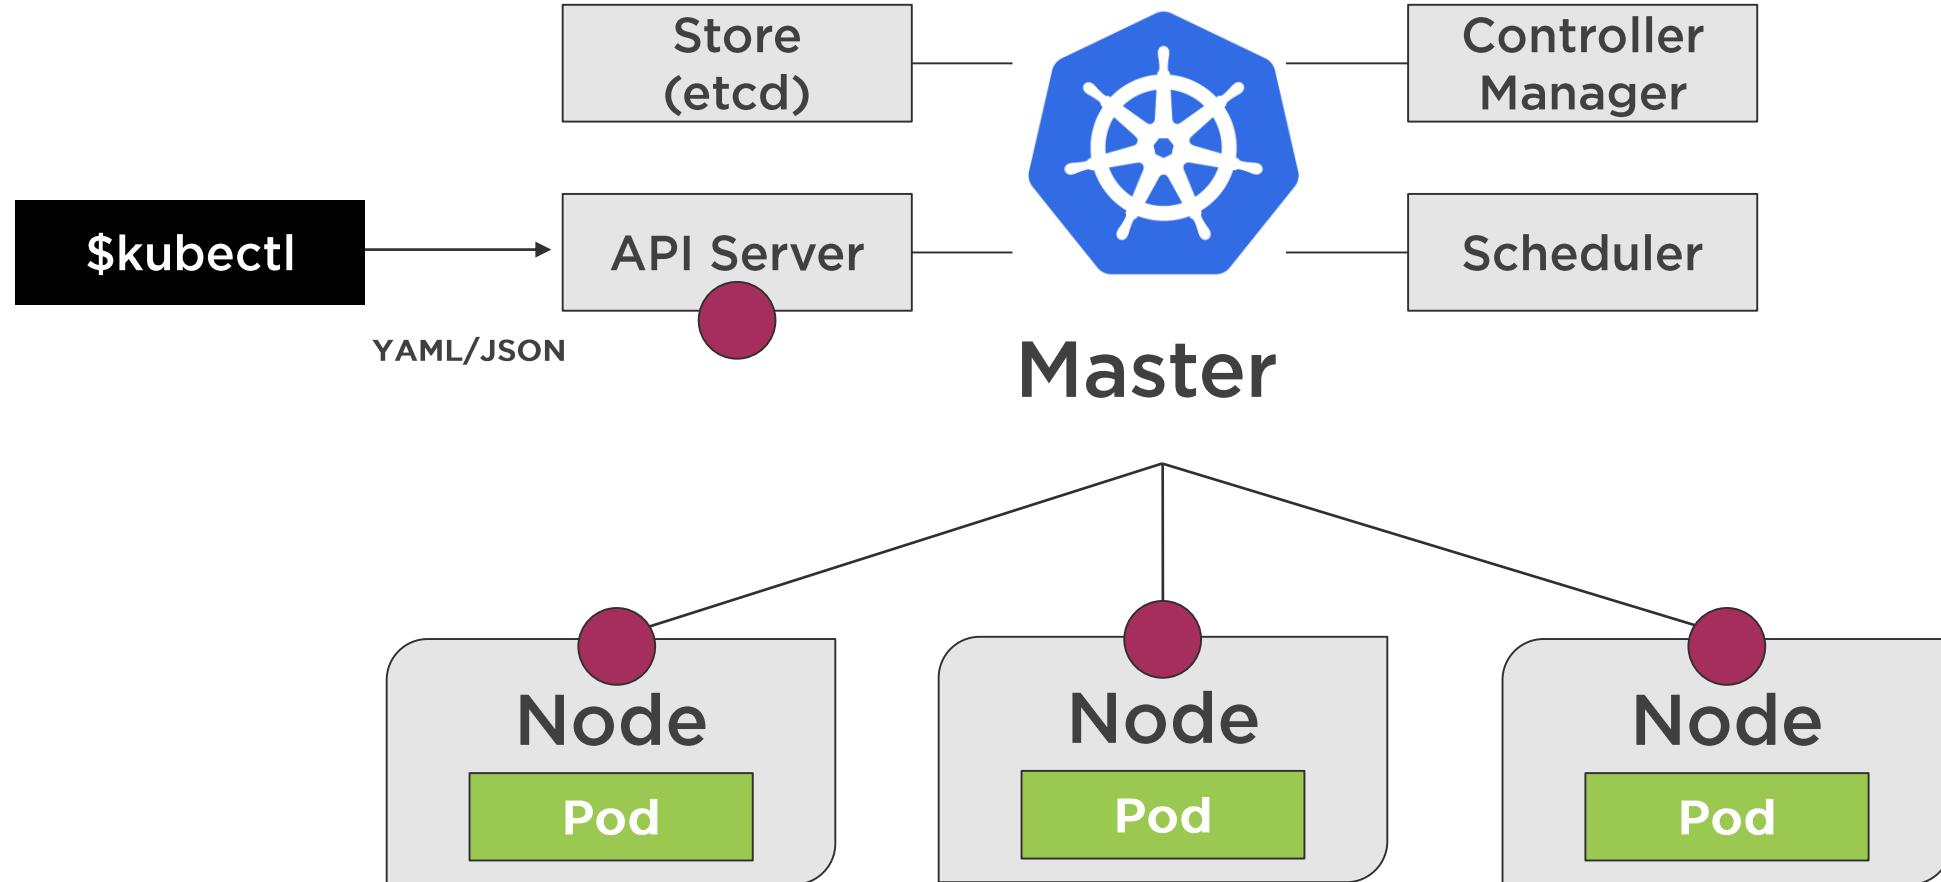
Pods and Nodes



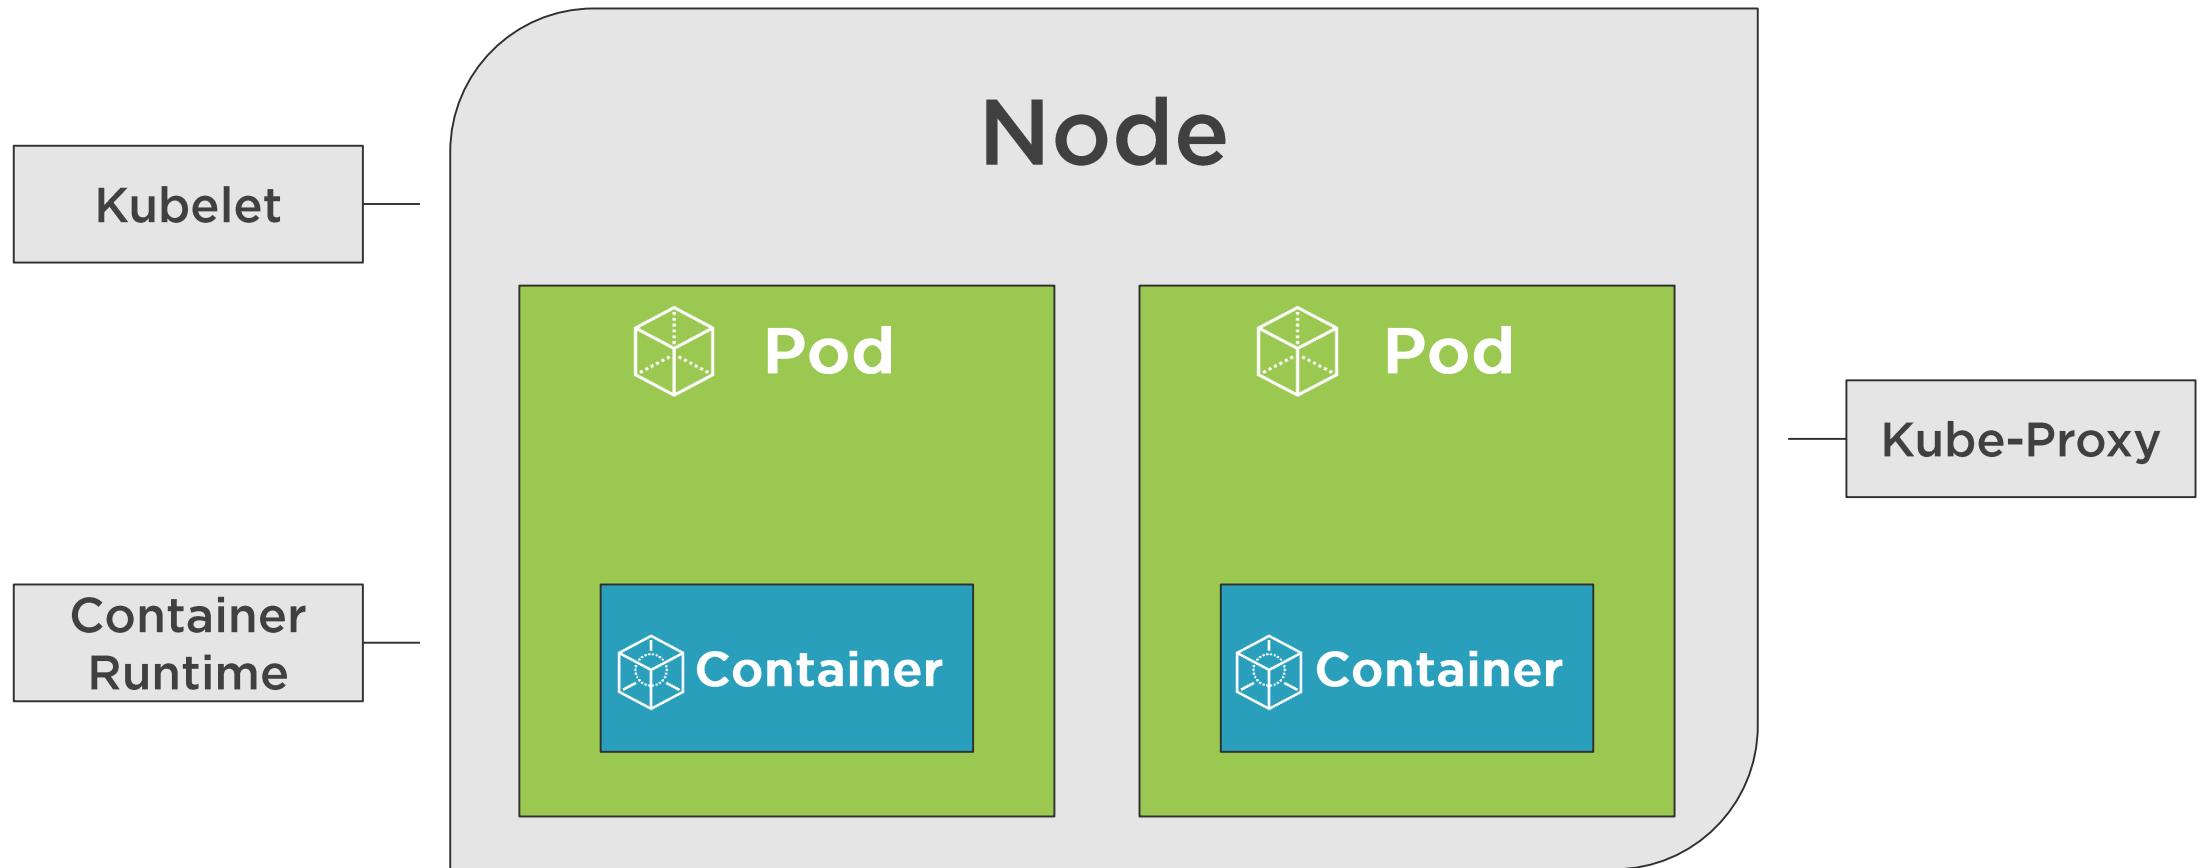
The Master Node



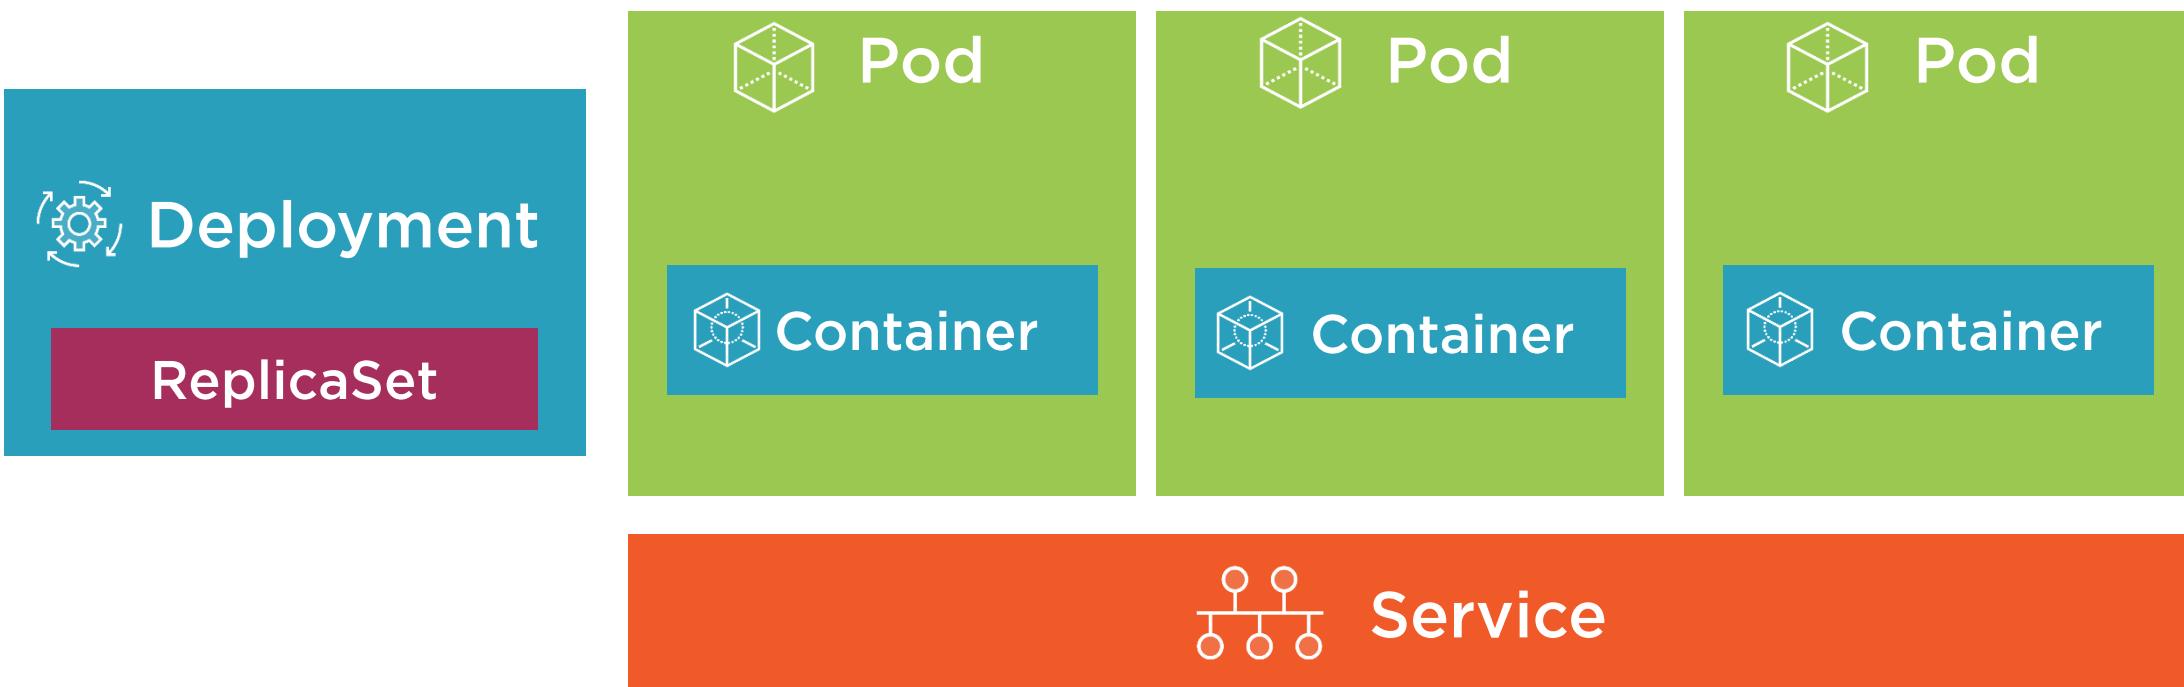
Communicating with kubectl



Kubernetes Nodes



The Big Picture



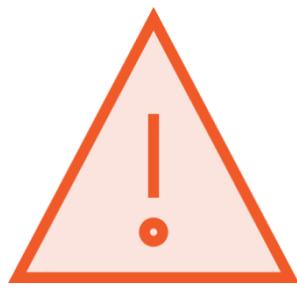
Benefits and Use Cases



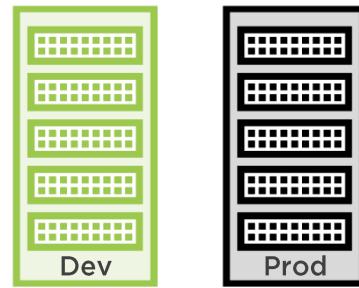
Key Container Benefits



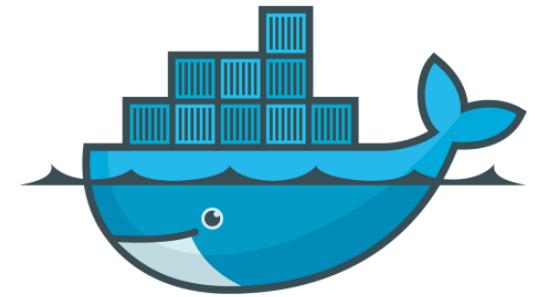
Accelerate
Developer
Onboarding



Eliminate App
Conflicts



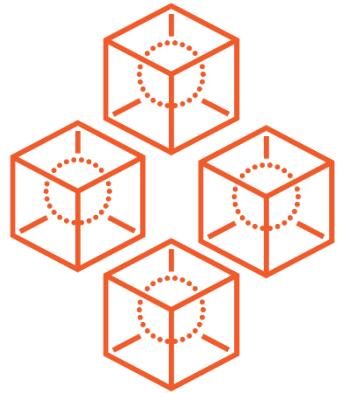
Environment
Consistency



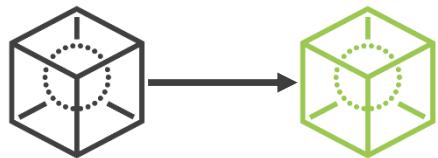
Ship Software
Faster



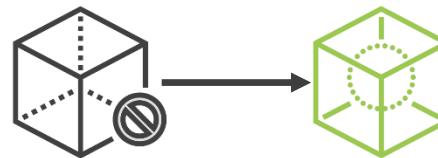
Key Kubernetes Benefits



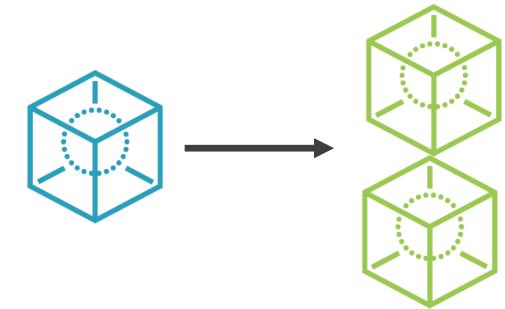
Orchestrate
Containers



Zero Downtime
Deployments



Self-healing
(Superpowers)



Scale
Containers



Developer Use Cases

Emulate production locally

Move from Docker Compose to Kubernetes

Create an end-to-end testing environment

Ensure application scales properly

Ensure secrets/config are working properly

Performance testing scenarios

Workload scenarios (CI/CD and more)

Learn how to leverage deployment options

Help DevOps create resources and solve problems



Running Kubernetes Locally



Installing and Running Kubernetes

Minikube

<https://github.com/kubernetes/minikube>

Docker
Desktop

<https://www.docker.com/products/docker-desktop>



Installing and Running Kubernetes

kind

<https://kind.sigs.k8s.io>

kubeadm

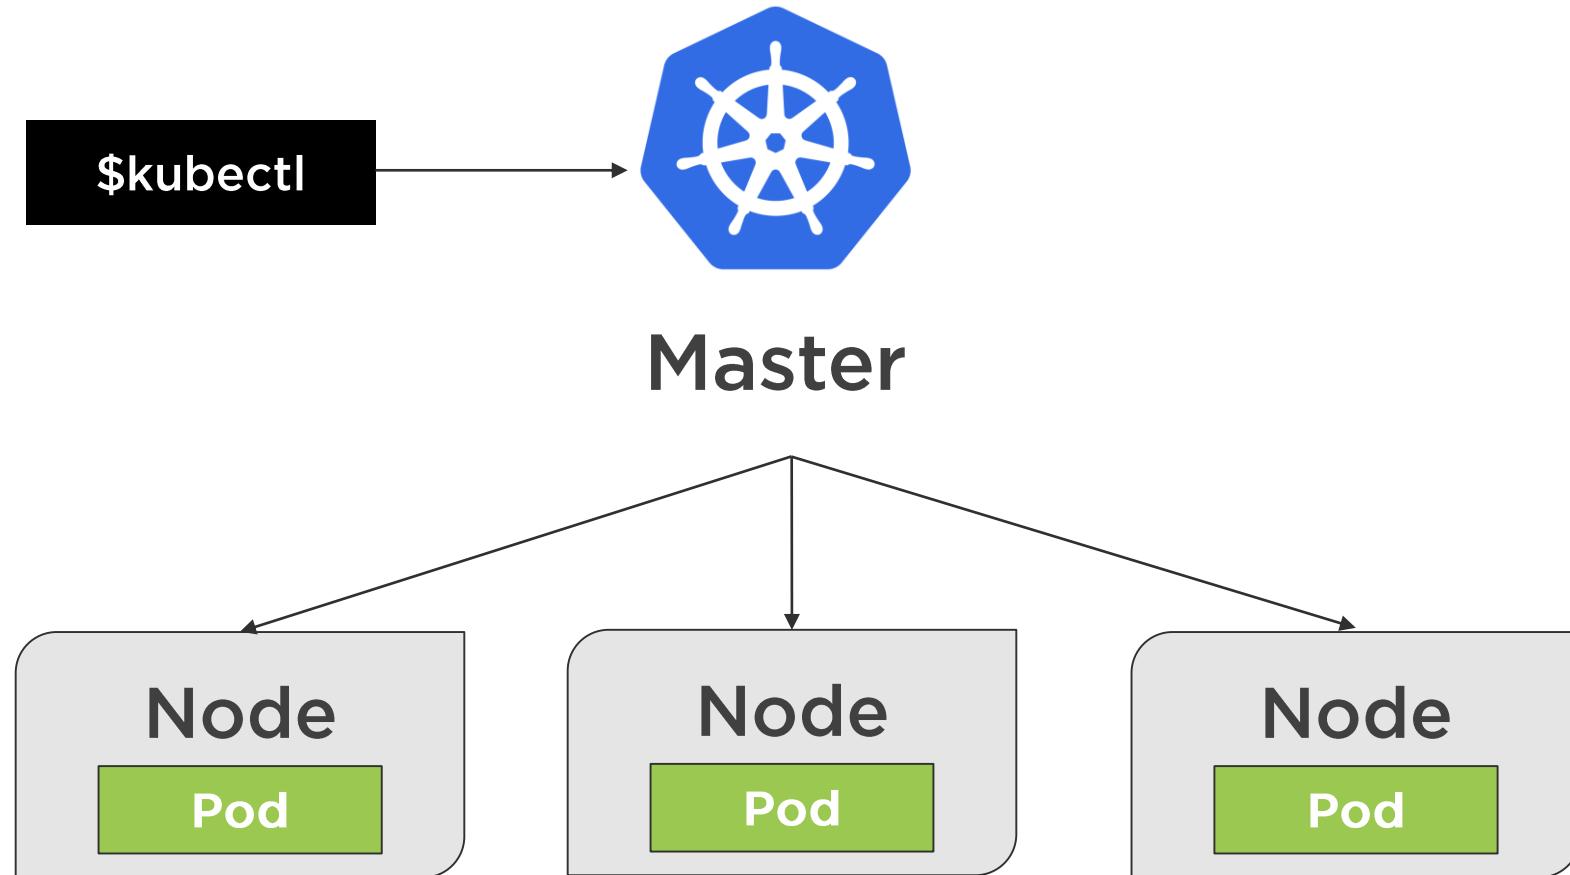
[https://kubernetes.io/docs/reference/
setup-tools/kubeadm/kubeadm](https://kubernetes.io/docs/reference/setup-tools/kubeadm/kubeadm)



Getting Started with kubectl



Using kubectl



Getting Started with kubectl Commands

kubectl version

kubectl cluster-info

kubectl get all

kubectl run [container-name]
--image=[image-name]

kubectl port-forward [pod] [ports]

kubectl expose ...

kubectl create [resource]

kubectl apply [resource]

- ◀ Check Kubernetes version
- ◀ View cluster information
- ◀ Retrieve information about Kubernetes Pods, Deployments, Services, and more
- ◀ Simple way to create a Deployment for a Pod
- ◀ Forward a port to allow external access
- ◀ Expose a port for a Deployment/Pod
- ◀ Create a resource
- ◀ Create or modify a resource



Aliasing kubectl (to save on typing)

```
# PowerShell  
Set-Alias -Name k -Value kubectl
```

◀ Create alias for PowerShell

```
# Mac/Linux  
alias k="kubectl"
```

◀ Create alias for Mac/Linux shell



Web UI Dashboard



The Web UI Dashboard

 kubernetes

Search

+  

Overview

Cluster

- Cluster Roles
- Namespaces
- Nodes
- Persistent Volumes
- Storage Classes

Namespace

- default

Overview

Workloads

Workload Status

 100.0%	 100.0%	 100.0%
Deployments	Pods	Replica Sets

Deployments

Name	Namespace	Labels	Pods	Age	Images
my-nginx	default	run: my-nginx	1 / 1	30 seconds	nginx:alpine

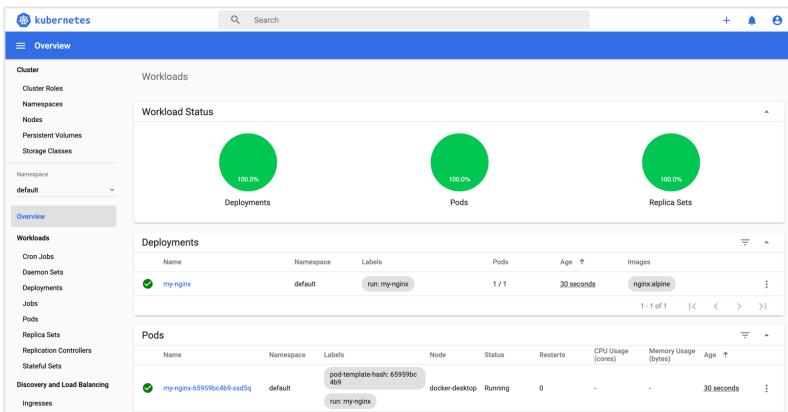
1 - 1 of 1 | < < > >|

Pods

Name	Namespace	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Age
my-nginx-65959bc4b9-xsd5q	default	pod-template-hash: 65959bc4b9 run: my-nginx	docker-desktop	Running	0	-	-	30 seconds



Enabling the Web UI Dashboard



Web UI dashboard provides a user interface to view Kubernetes resources

Steps to enable the UI Dashboard:

- kubectl apply [dashboard-yaml-url]
- kubectl describe secret -n kube-system
- Locate the kubernetes.io/service-account-token and copy the token
- kubectl proxy
- Visit the dashboard URL and login using the token



Summary



Kubernetes provides container orchestration capabilities

Use for production, emulating production, testing, and more

Several options are available to run Kubernetes locally

Interact with Kubernetes using kubectl



Creating Pods



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Pod Core Concepts

Creating a Pod

kubectl and Pods

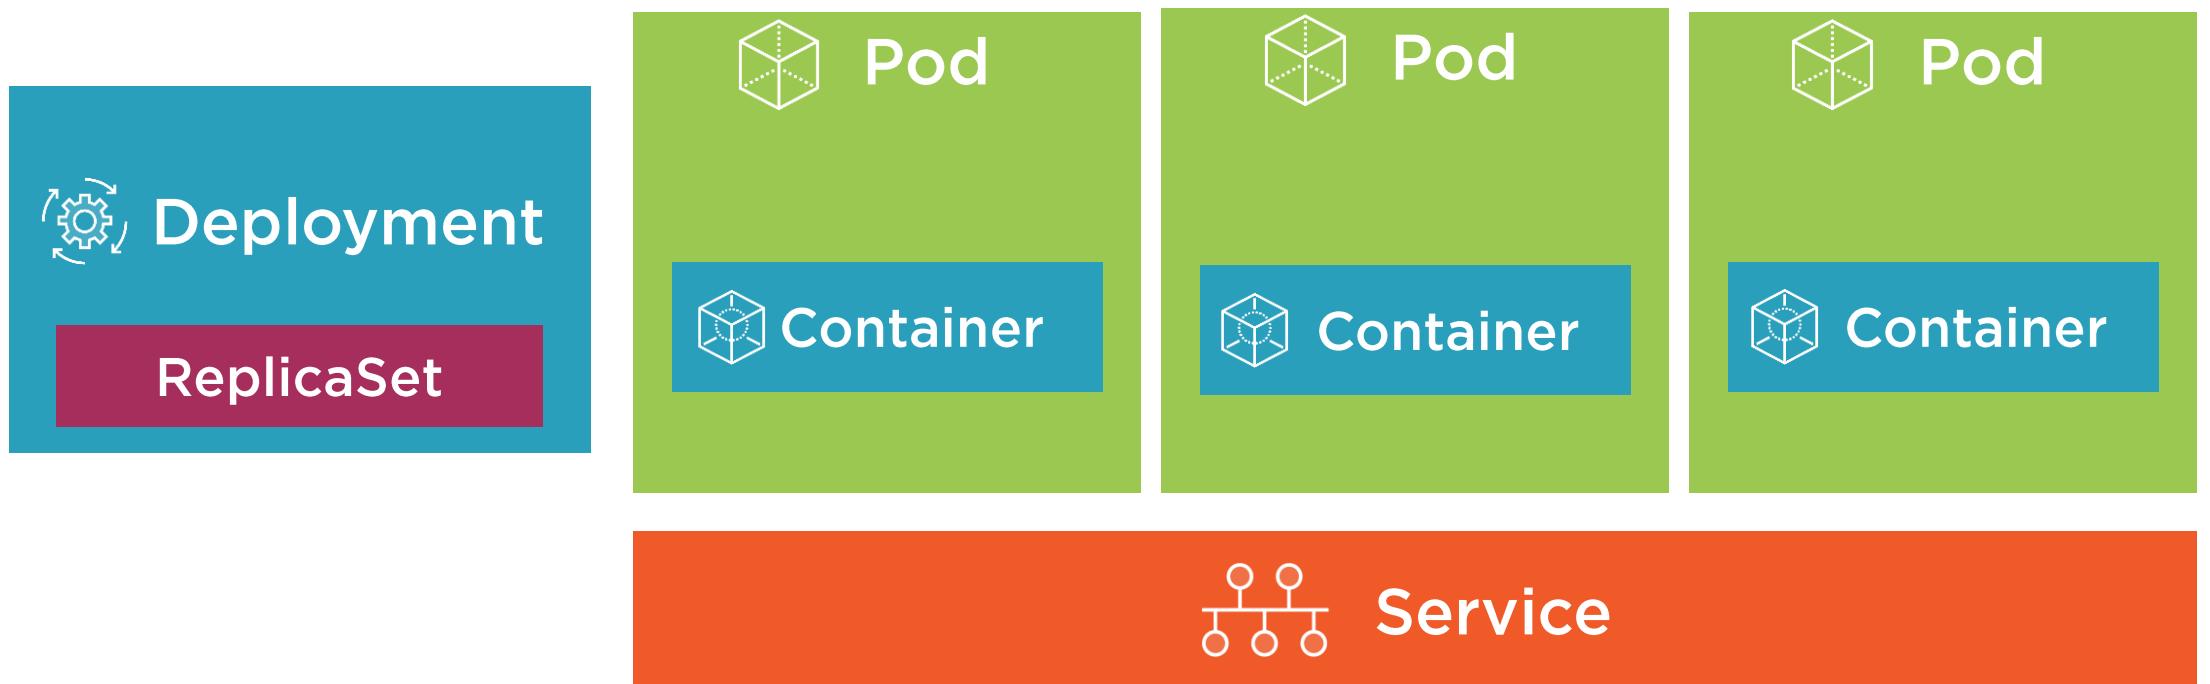
YAML Fundamentals

Defining a Pod with YAML

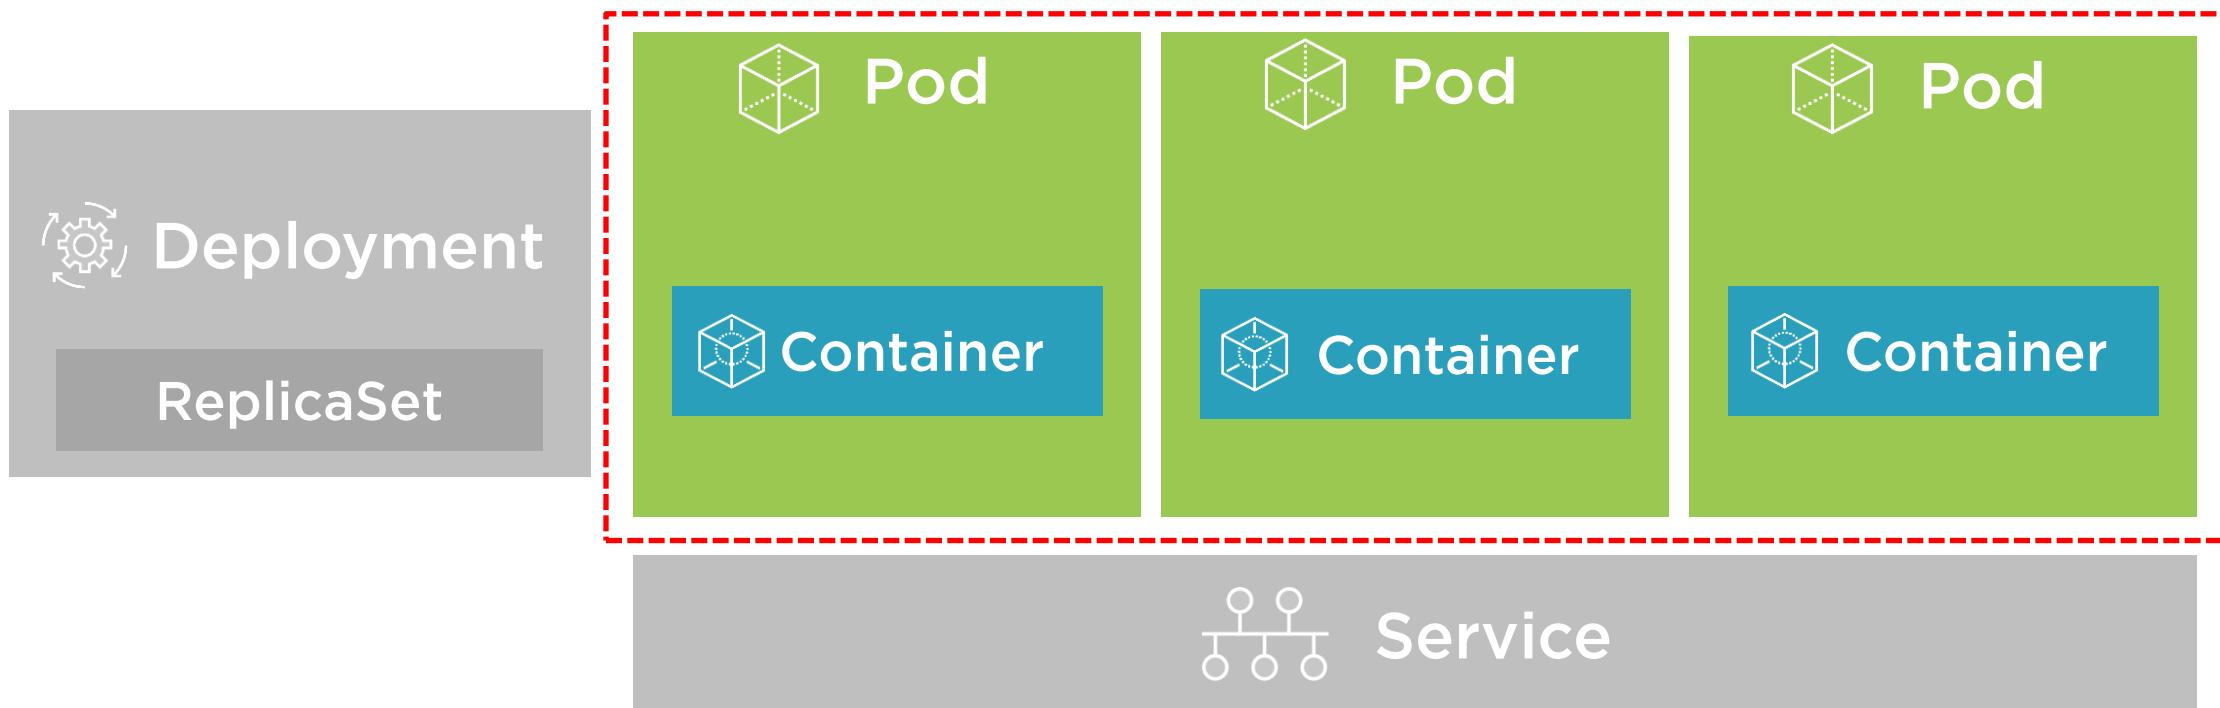
Pod Health



You Are Here



You Are Here



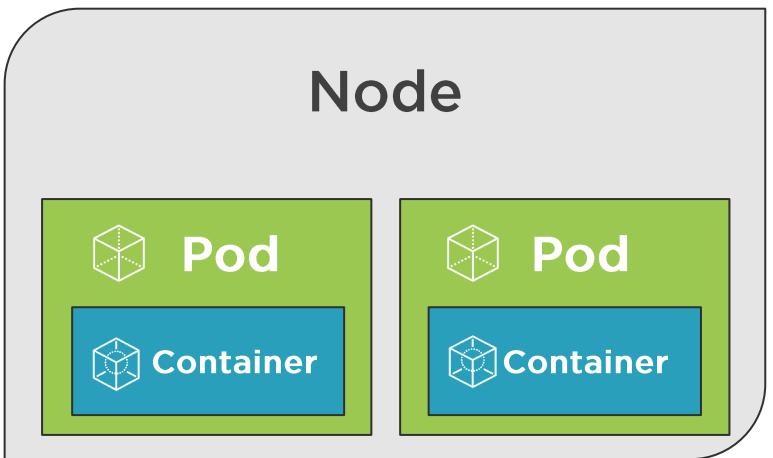
Pod Core Concepts



A Pod is the basic execution unit of a Kubernetes application—the smallest and simplest unit in the Kubernetes object model that you create or deploy.



Kubernetes Pods



Smallest object of the Kubernetes object model

Environment for containers

**Organize application "parts" into Pods
(server, caching, APIs, database, etc.)**

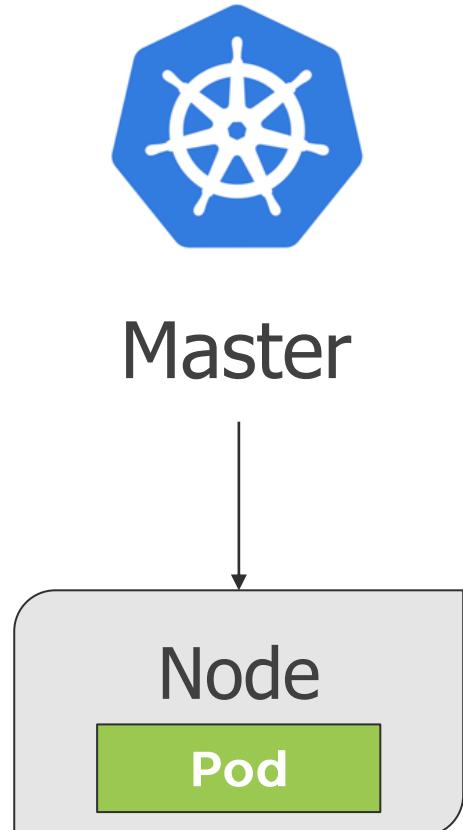
**Pod IP, memory, volumes, etc. shared
across containers**

Scale horizontally by adding Pod replicas

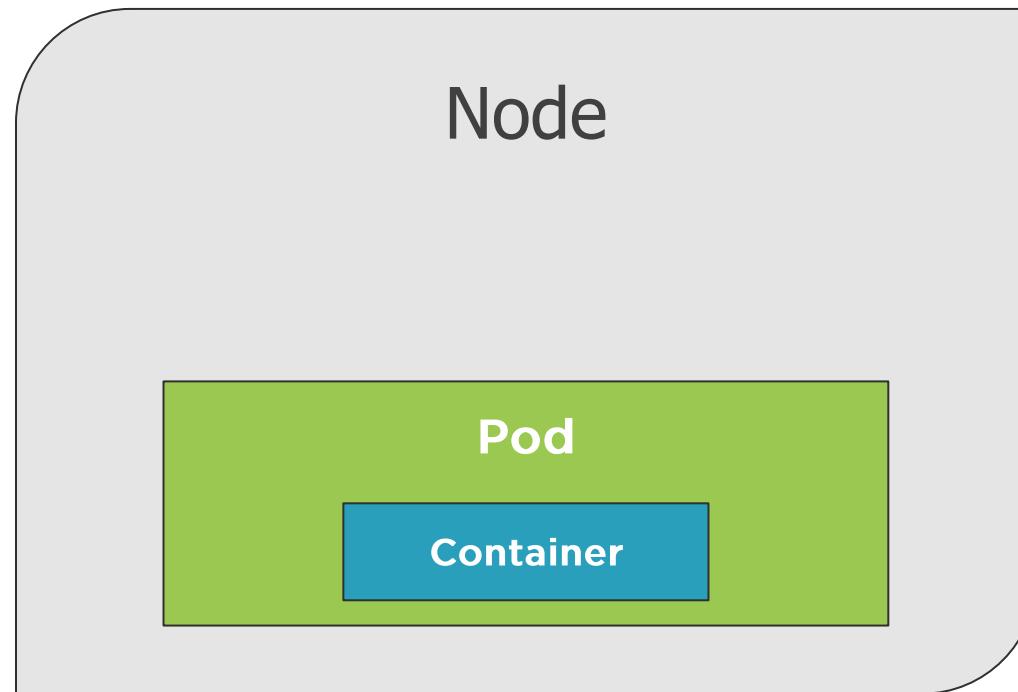
Pods live and die but never come back to life



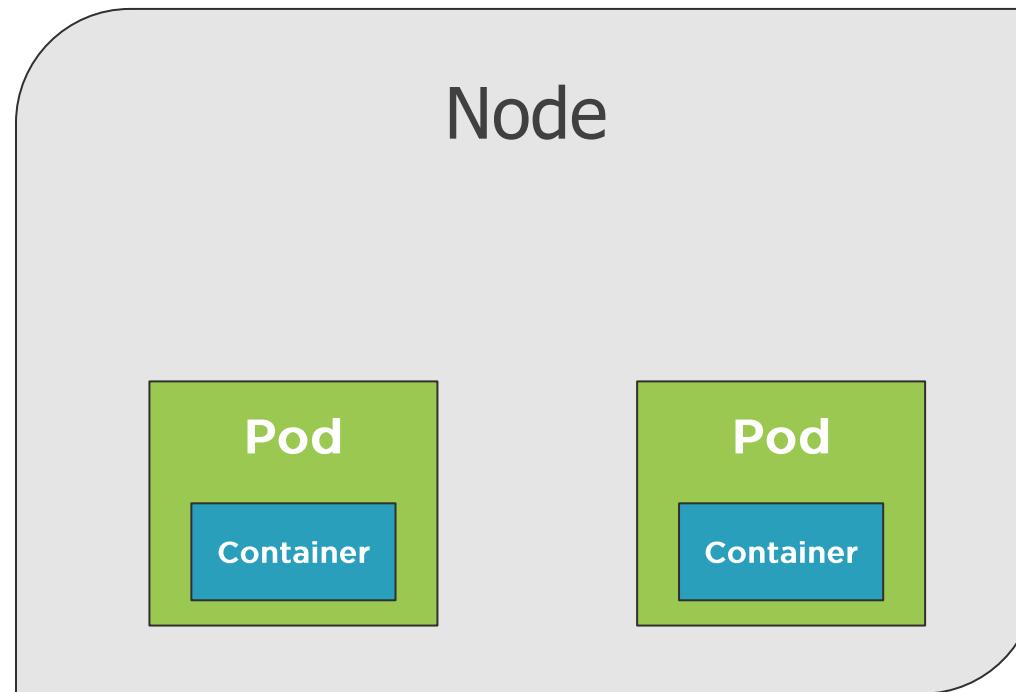
The Role of Pods



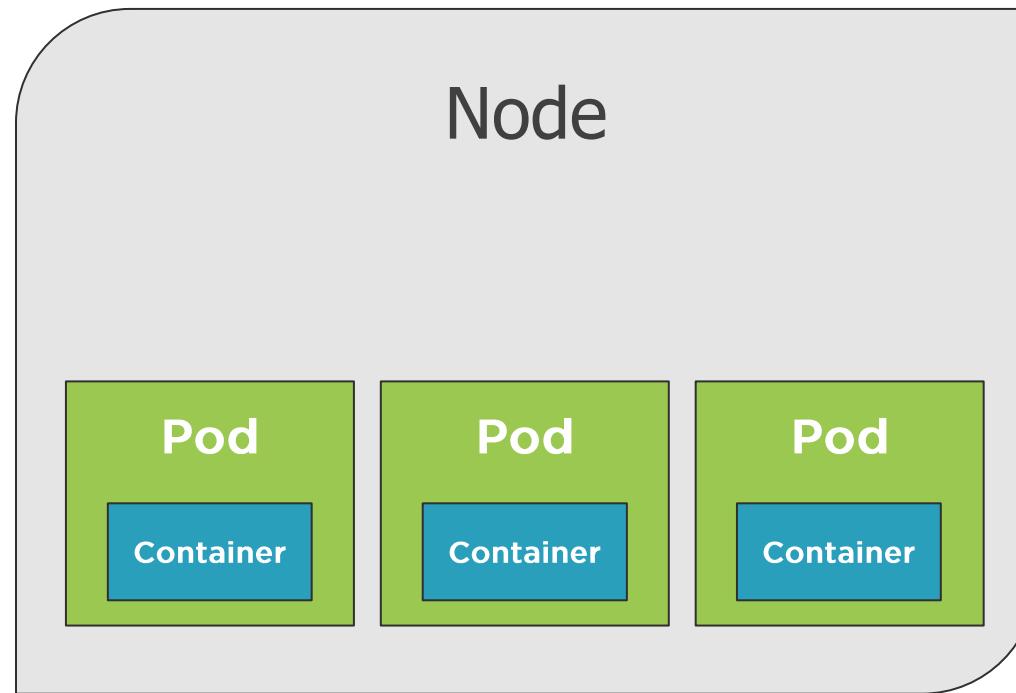
The Role of Pods



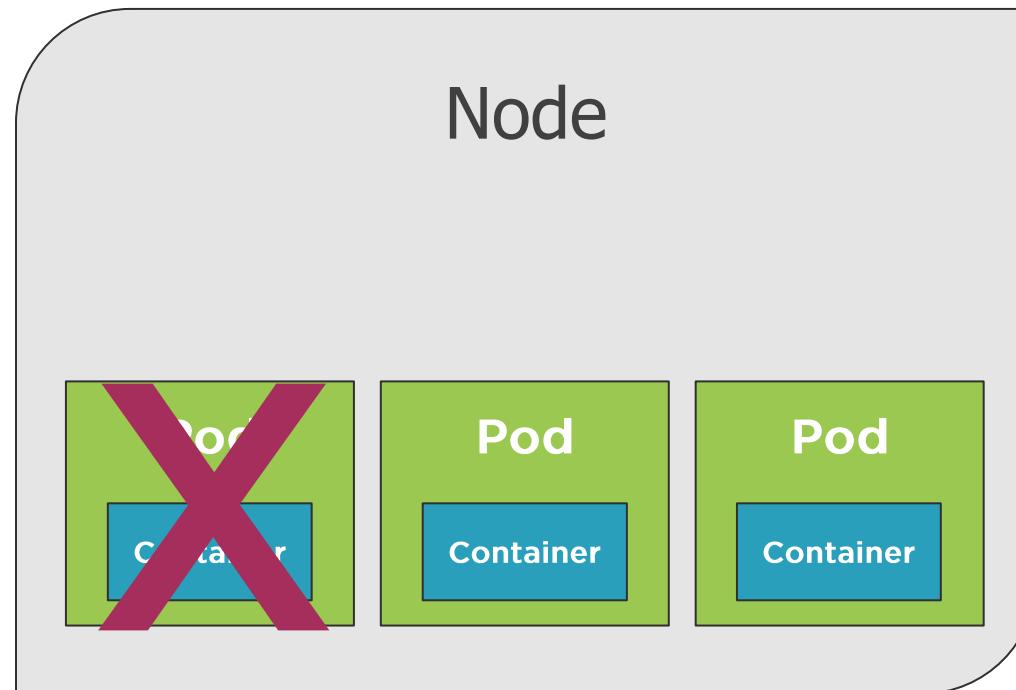
The Role of Pods



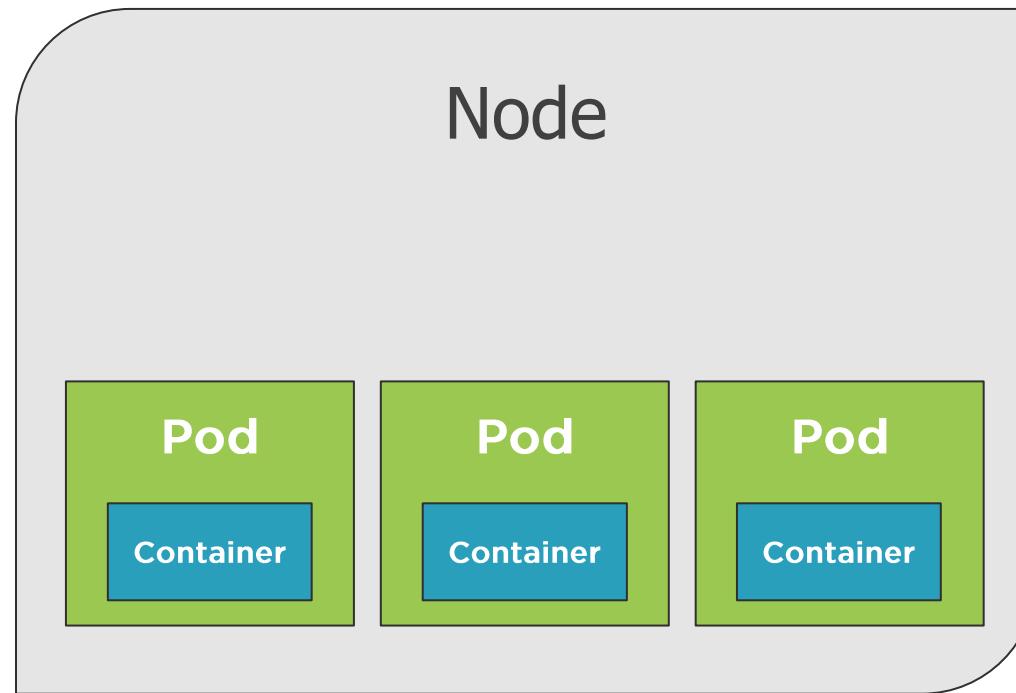
The Role of Pods



The Role of Pods



The Role of Pods



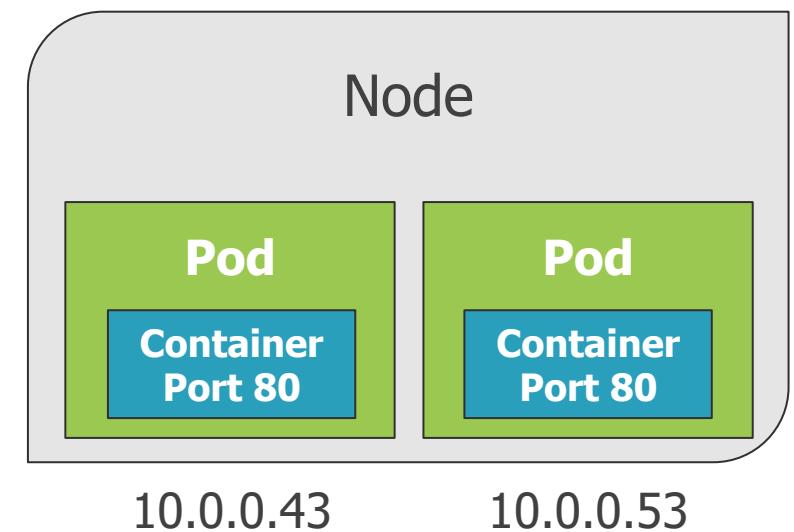
Pod containers share the same Network namespace (share IP/port)

Pod containers have the same loopback network interface (localhost)

Container processes need to bind to different ports within a Pod

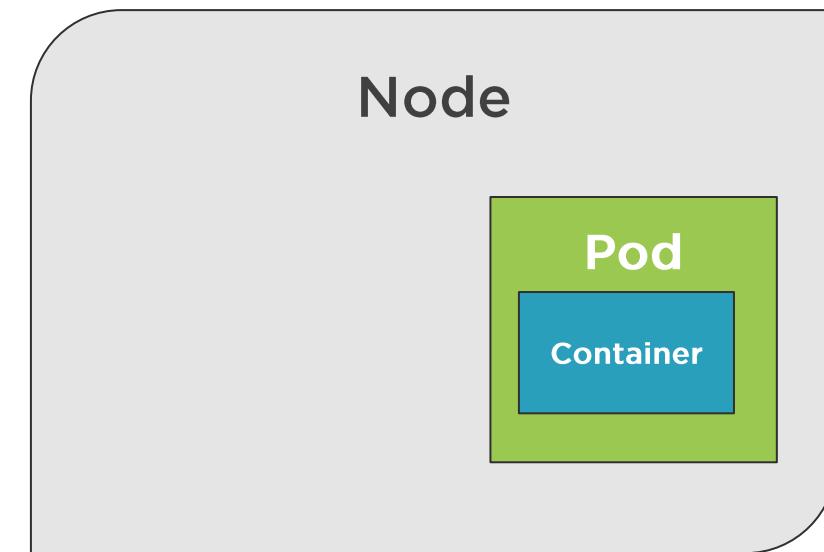
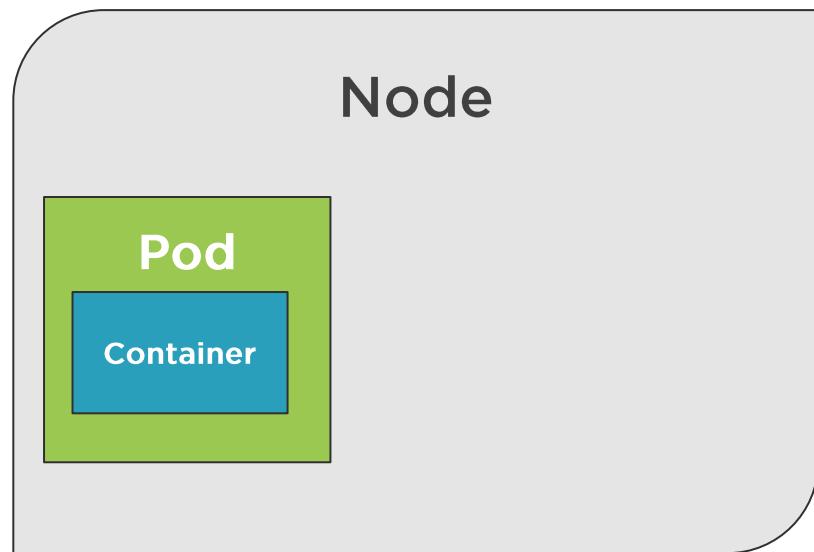
Ports can be reused by containers in separate Pods

Pods, IPs, and Ports



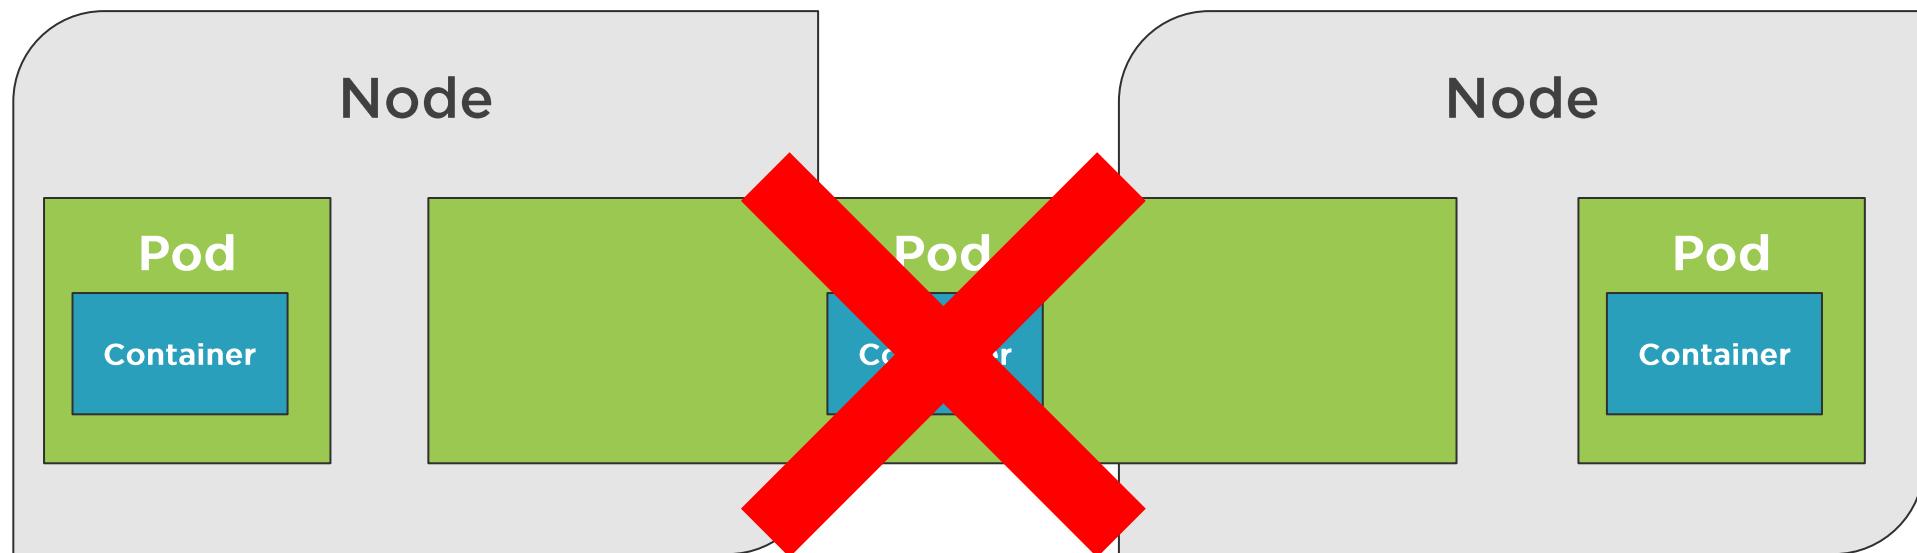
Nodes and Pods

Pods do not span nodes



Nodes and Pods

Pods do not span nodes



Creating a Pod



Running a Pod

There are several different ways to schedule a Pod:

kubectl run command

kubectl create/apply command with a yaml file

```
# Run the nginx:alpine container in a Pod  
kubectl run [podname] --image=nginx:alpine
```

```
# List only Pods
```

```
kubectl get pods
```

```
# List all resources
```

```
kubectl get all
```

Get Information about a Pod

The **kubectl get** command can be used to retrieve information about Pods and many other Kubernetes objects



Expose a Pod Port

Pods and containers are only accessible within the Kubernetes cluster by default

One way to expose a container port externally: **kubectl port-forward**

```
# Enable Pod container to be  
# called externally  
kubectl port-forward [name-of-pod] 8080:80
```

Internal port



External port

```
# Will cause pod to be recreated  
kubectl delete pod [name-of-pod]
```

```
# Delete Deployment that manages the Pod  
kubectl delete deployment [name-of-deployment]
```

Deleting a Pod

Running a Pod will cause a Deployment to be created

To delete a Pod use **kubectl delete pod** or find the deployment and use **kubectl delete deployment**



kubectl and Pods



```
kubectl run [pod-name] --image=nginx:alpine
```

```
kubectl get pods
```

```
kubectl port-forward [pod-name] 8080:80
```

```
kubectl delete pod [pod-name]
```

Working with Pods using kubectl

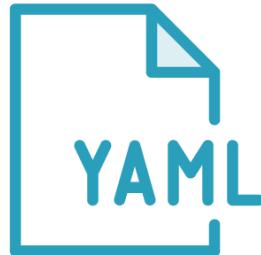
Different kubectl commands can be used to run, view, and delete Pods



YAML Fundamentals



YAML Review



YAML files are composed of maps and lists

Indentation matters (be consistent!)

Always use spaces

Maps:

- name: value pairs
- Maps can contain other maps for more complex data structures

Lists:

- Sequence of items
- Multiple maps can be defined in a list



Introduction to YAML

```
key: value  
  
complexMap:  
  
    key1: value  
  
    key2:  
  
        subKey: value  
  
items:  
  
    - item1  
  
    - item2  
  
itemsMap:  
  
    - map1: value  
  
        map1Prop: value  
  
    - map2: value  
  
        map2Prop: value
```

- ◀ YAML maps define a key and value
- ◀ More complicated map structures can be defined using a key that references another map
- ◀ YAML lists can be used to define a sequence of items

- ◀ YAML lists can define a sequence maps

Note:

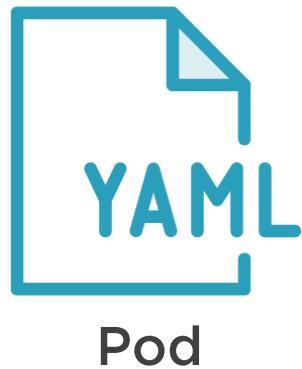
- Indentation matters
- Use spaces NOT tabs



Defining a Pod with YAML



Defining a Pod with YAML



+ **kubectl** =



nginx.pod.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: my-nginx
spec:
  containers:
    - name: my-nginx
      image: nginx:alpine
```

- ◀ Kubernetes API version
- ◀ Type of Kubernetes resource
- ◀ Metadata about the Pod
- ◀ The spec/blueprint for the Pod
- ◀ Information about the containers that will run in the pod



Creating a Pod Using YAML

To create a pod using YAML use the **kubectl create** command along with the **--filename** or **-f** switch

```
# Perform a "trial" create and also validate the YAML  
kubectl create -f file.pod.yml --dry-run --validate=true
```

```
# Create a Pod from YAML  
# Will error if Pod already exists  
kubectl create -f file.pod.yml
```



Default value

Creating or Applying Changes to a Pod

To create or apply changes to a pod using YAML use the **kubectl apply** command along with the **--filename** or **-f** switch

```
# Alternate way to create or apply changes to a  
# Pod from YAML  
kubectl apply -f file.pod.yml
```

Store current
properties in
resource's annotations

```
# Use --save-config when you want to use  
# kubectl apply in the future  
kubectl create -f file.pod.yml --save-config
```

Using kubectl create --save-config

```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubectl.kubernetes.io/
    last-applied-configuration:
    {"apiVersion":"v1","kind":"Pod",
     "metadata":{  
       "name": "my-nginx"  
     ...  
   }  
}  
...
```

- ◀ **--save-config causes the resource's configuration settings to be saved in the annotations**
- ◀ **Example of saved configuration**
- ◀ **Having this allows in-place changes to be made to a Pod in the future using kubectl apply**

In-place/non-disruptive changes can also be made to a Pod using **kubectl edit** or **kubectl patch**.



Deleting a Pod

To delete a Pod use **kubectl delete**

```
# Delete Pod  
kubectl delete pod [name-of-pod]
```

```
# Delete Pod using YAML file that created it  
kubectl delete -f file.pod.yml
```

kubectl and YAML



```
kubectl create -f nginx.pod.yml --save-config  
kubectl describe pod [pod-name]  
kubectl apply -f nginx.pod.yml  
kubectl exec [pod-name] -it sh  
kubectl edit -f nginx.pod.yml  
kubectl delete -f nginx.pod.yml
```

Creating and Inspecting Pods with kubectl

Several different commands can be used to create and modify Pods



Pod Health



Kubernetes relies on Probes to determine the health of a Pod container.



A Probe is a diagnostic performed periodically by the kubelet on a Container.



Types of Probes



Liveness
Probe



Readiness
Probe



Types of Probes



Liveness probes determine if a Pod is healthy and running as expected

Readiness probes determine if a Pod should receive requests

Failed Pod containers are recreated by default (`restartPolicy` defaults to `Always`)



ExecAction – Executes an action inside the container

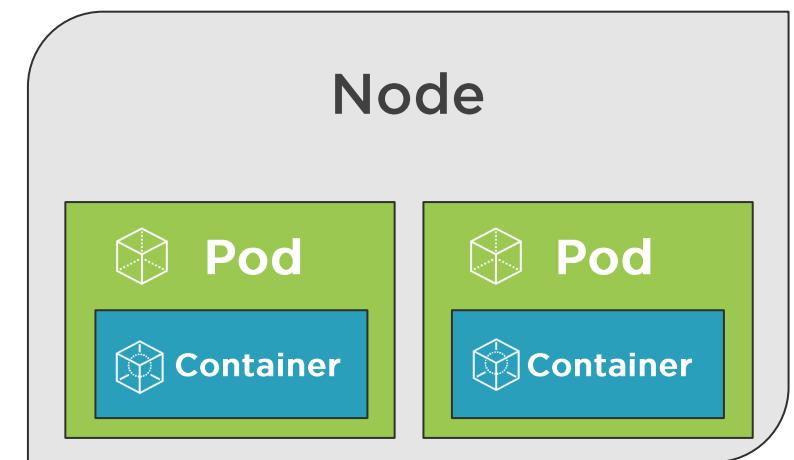
TCPSocketAction – TCP check against the container's IP address on a specified port

HTTPGetAction – HTTP GET request against container

Probes can have the following results:

- Success
- Failure
- Unknown

Probe Types



Defining an HTTP Liveness Probe

```
apiVersion: v1
kind: Pod
...
spec:
  containers:
    - name: my-nginx
      image: nginx:alpine
      livenessProbe:
        httpGet:
          path: /index.html
          port: 80
        initialDelaySeconds: 15
        timeoutSeconds: 2
        periodSeconds: 5
        failureThreshold: 1
```

- ◀ Define liveness probe
- ◀ Check /index.html on port 80
- ◀ Wait 15 seconds
- ◀ Timeout after 2 seconds
- ◀ Check every 5 seconds
- ◀ Allow 1 failure before failing Pod



Defining an ExecAction Liveness Probe

```
apiVersion: v1
kind: Pod
...
spec:
  containers:
    - name: liveness
      image: k8s.gcr.io/busybox
      args:
        - /bin/sh
        - -c
        - touch /tmp/healthy; sleep 30;
          rm -rf /tmp/healthy; sleep 600
  livenessProbe:
    exec:
      command:
        - cat
        - /tmp/healthy
  initialDelaySeconds: 5
  periodSeconds: 5
```

◀ Define args for container

◀ Define liveness probe

◀ Define action/command to execute



Defining a Readiness Probe

```
apiVersion: v1
kind: Pod
...
spec:
  containers:
    - name: my-nginx
      image: nginx:alpine
      readinessProbe:
        httpGet:
          path: /index.html
          port: 80
        initialDelaySeconds: 2
        periodSeconds: 5
    ◀ Define readiness probe
    ◀ Check /index.html on port 80
    ◀ Wait 2 seconds
    ◀ Check every 5 seconds
```



Readiness Probe:
When should a container start
receiving traffic?

Liveness Probe:
When should a container restart?



Pod Health in Action



Summary



Pods are the smallest unit of Kubernetes

Containers run within Pods and share a Pod's memory, IP, volumes, and more

Pods can be started using different kubectl commands

YAML can be used to create a Pod

Health checks provide a way to notify Kubernetes when a Pod has a problem



Creating Deployments



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Deployments Core Concepts

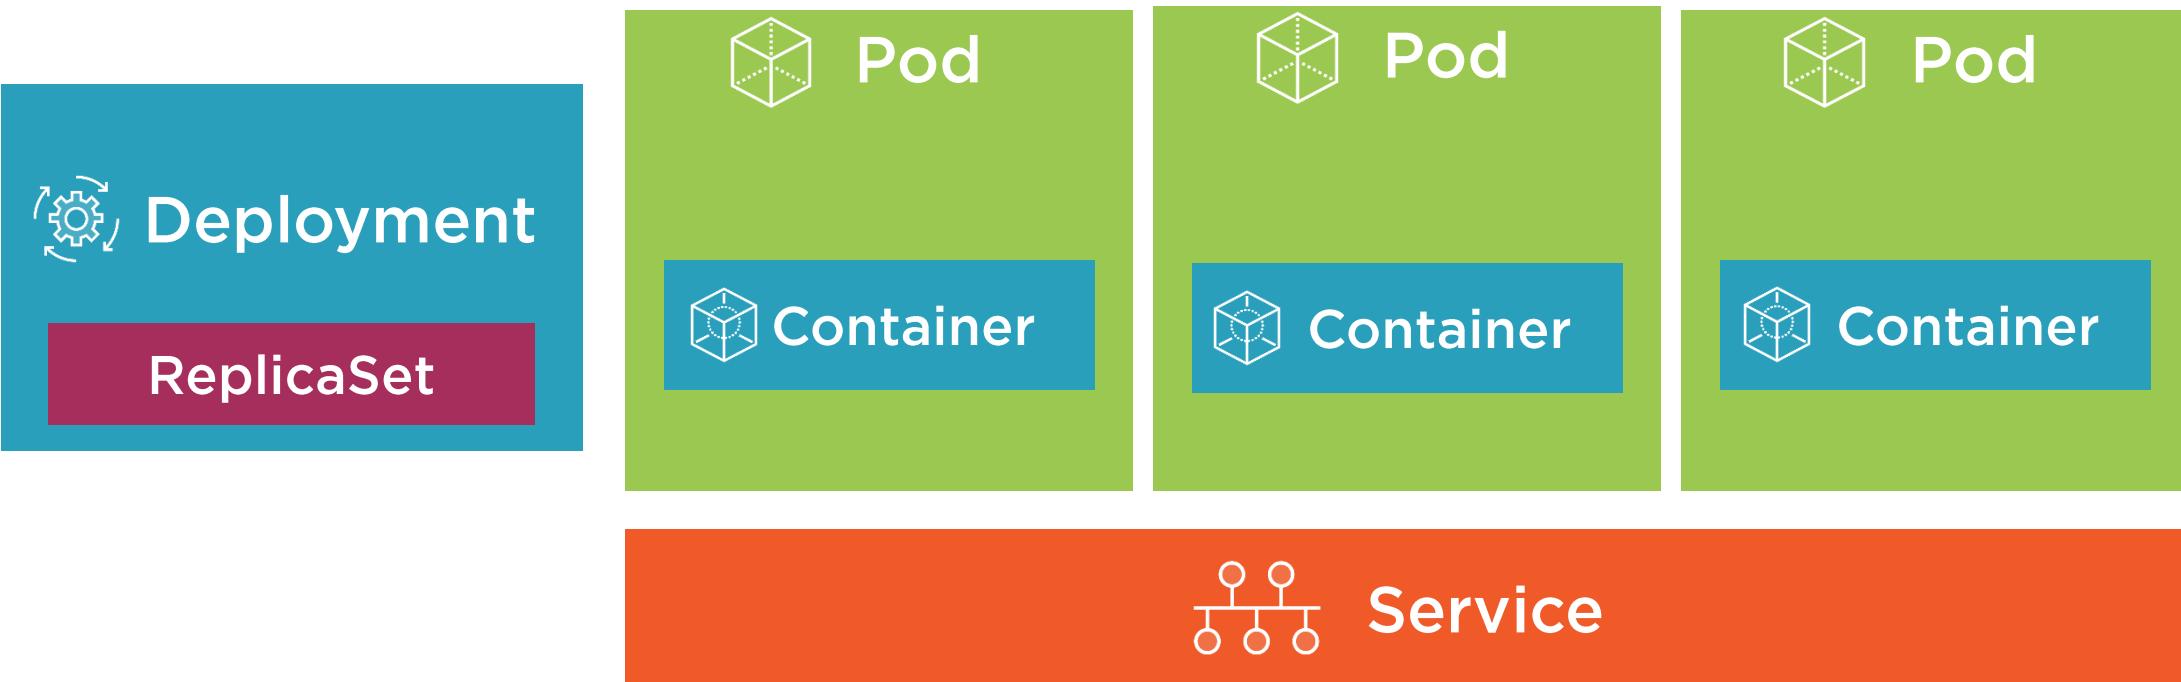
Creating a Deployment

kubectl and Deployments

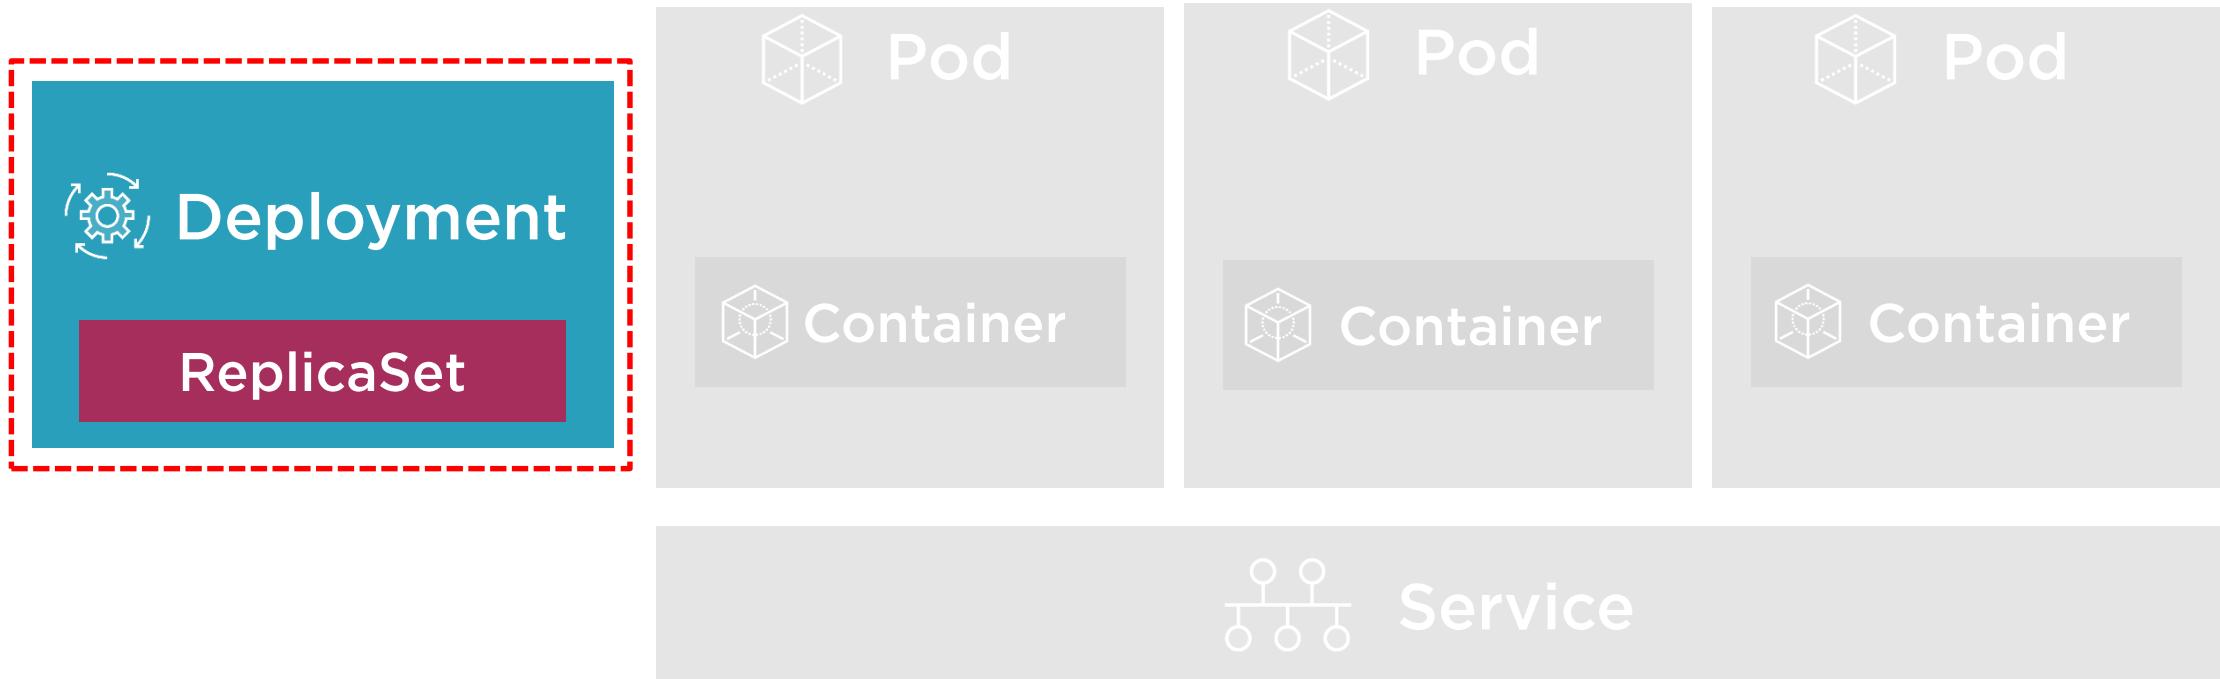
Deployment Options



You Are Here



You Are Here



Deployments Core Concepts



A ReplicaSet is a declarative way to manage Pods



A Deployment is a declarative
way to manage Pods using a
ReplicaSet



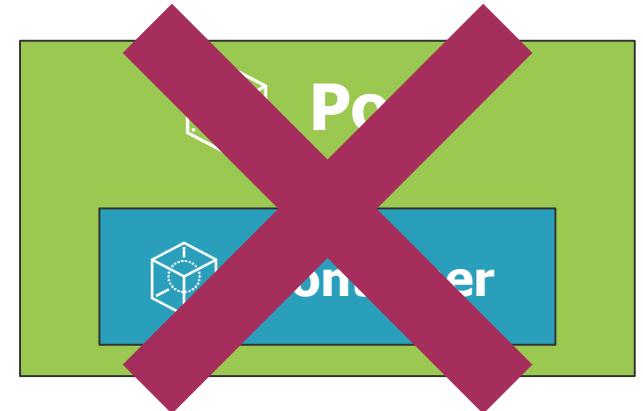
Pods, Deployments, and ReplicaSets

Pods represent the most basic resource in Kubernetes

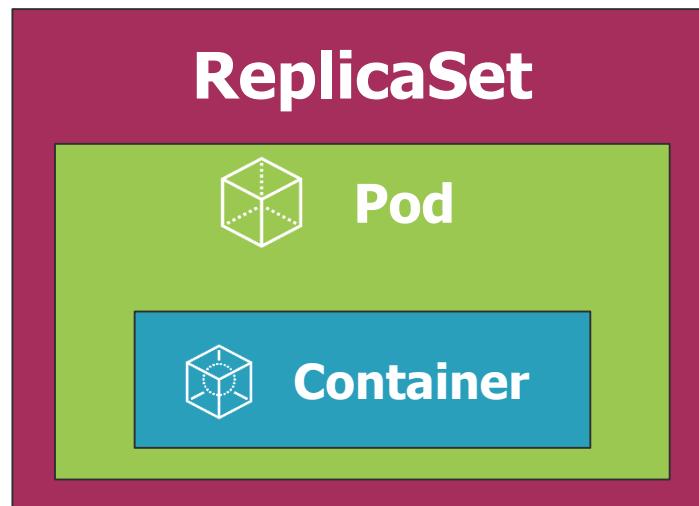
Can be created and destroyed but are never re-created

What happens if a Pod is destroyed?

Deployments and ReplicaSets ensure Pods stay running and can be used to scale Pods



The Role of ReplicaSets

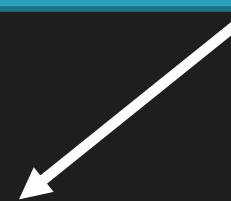


ReplicaSets act as a Pod controller:

- Self-healing mechanism
- Ensure the requested number of Pods are available
- Provide fault-tolerance
- Can be used to scale Pods
- Relies on a Pod template
- No need to create Pods directly!
- Used by Deployments



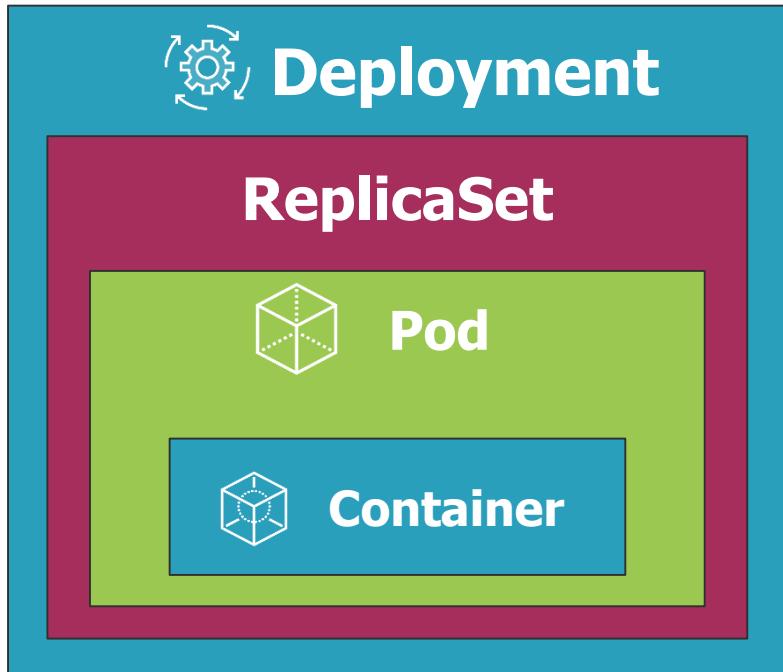
Result of Creating a ReplicaSet

2 Pods created					
iMac-3:replicaSets danwahlin\$ k get all + kubectl get all					
NAME	READY	STATUS	RESTARTS	AGE	
pod/frontend-8rslg	1/1	Running	0	8s	
pod/frontend-zc8gl	1/1	Running	0	8s	
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	41d
NAME	DESIRED	CURRENT	READY	AGE	
replicaset.apps/frontend	2	2	2	8s	

ReplicaSet created



The Role of Deployments



A Deployment manages Pods:

- Pods are managed using ReplicaSets
- Scales ReplicaSets, which scale Pods
- Supports zero-downtime updates by creating and destroying ReplicaSets
- Provides rollback functionality
- Creates a unique label that is assigned to the ReplicaSet and generated Pods
- YAML is very similar to a ReplicaSet



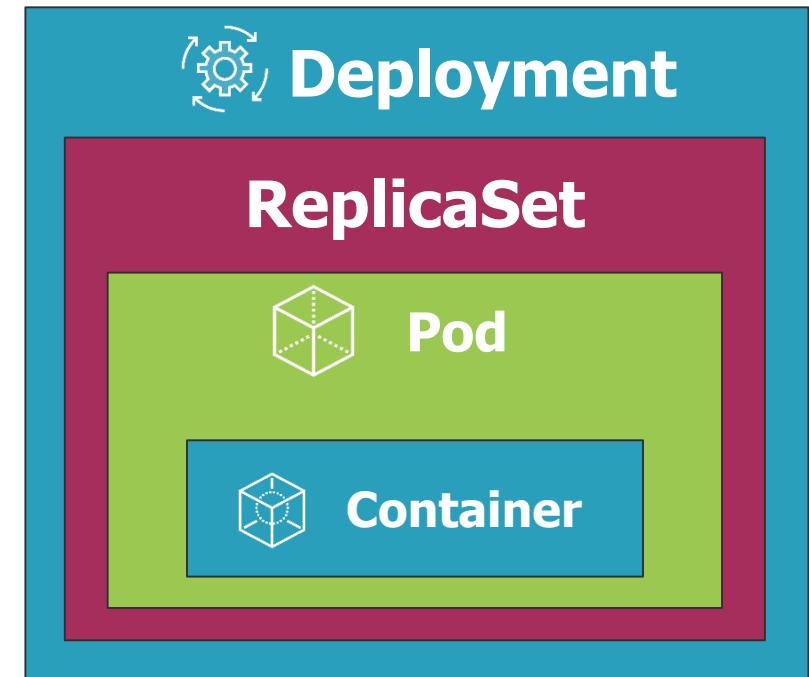
Creating a Deployment



Defining a Deployment with YAML



+ kubectl =



Defining a Deployment (From a High-Level)

```
apiVersion: apps/v1
kind: Deployment
metadata:
spec:
  selector:
  template:
    spec:
      containers:
        - name: my-nginx
          image: nginx:alpine
```

- ◀ **Kubernetes API version and resource type (Deployment)**
- ◀ **Metadata about the Deployment**
- ◀ **Select Pod template label(s)**
- ◀ **Template used to create the Pods**
- ◀ **Containers that will run in the Pod**



Defining a Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
  labels:
    app: my-nginx
    tier: frontend
spec:
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: my-nginx
          image: nginx:alpine
```

- ◀ **Kubernetes API version and resource type (Deployment)**
- ◀ **Metadata about the Deployment**
- ◀ **The selector is used to "select" the template to use (based on labels)**
- ◀ **Template to use to create the Pod/Containers (note that the selector matches the label)**



Defining Probes in a Deployment

```
apiVersion: apps/v1
kind: Deployment
...
template:
  spec:
    containers:
      - name: my-nginx
        image: nginx:alpine
    livenessProbe:
      httpGet:
        path: /index.html
        port: 80
      initialDelaySeconds: 15
      timeoutSeconds: 2
      periodSeconds: 5
      failureThreshold: 1
```

- ◀ Define liveness probe (readiness probes can also be defined)
- ◀ Check /index.html on port 80



kubectl and Deployments



```
# Create a Deployment
```

```
kubectl create -f file.deployment.yml
```

Creating a Deployment

Use the `kubectl create` command along with the `--filename` or `-f` switch



Creating or Applying Changes

Use the **kubectl apply** command along with the **--filename** or **-f** switch

```
# Alternate way to create or apply changes to a  
# Deployment from YAML  
kubectl apply -f file.deployment.yml  
  
# Use --save-config when you want to use  
# kubectl apply in the future  
kubectl create -f file.deployment.yml --save-config
```

Store current
properties in
resource's annotations

```
kubectl get deployments
```

Getting Deployments

List all Deployments



```
# List all Deployments and their labels  
kubectl get deployment --show-labels
```

```
# Get all Deployments with a specific label  
kubectl get deployment -l app=nginx
```

Deployments and Labels

List the labels for all Deployments using the **--show-labels** switch

To get information about a Deployment with a specific label, use the **-l** switch



Deleting a Deployment

To delete a Deployment use kubectl delete

Will delete the Deployment and all associated Pods/Containers

```
# Delete Deployment  
kubectl delete deployment [deployment-name]
```

```
# Scale the Deployment Pods to 5  
kubectl scale deployment [deployment-name] --replicas=5
```

```
# Scale by refencing the YAML file  
kubectl scale -f file.deployment.yml --replicas=5
```

Scaling Pods Horizontally

Update the YAML file or use the **kubectl scale** command

```
spec:  
  replicas: 3  
  selector:  
    tier: frontend
```



kubectl Deployments in Action



```
kubectl create -f nginx.deployment.yml --save-config
```

```
kubectl describe [pod | deployment] [pod-name | deployment-name]
```

```
kubectl apply -f nginx.pod.yml
```

```
kubectl get deployments --show-labels
```

```
kubectl get deployments -l app=my-nginx
```

```
kubectl scale -f nginx.deployment.yml --replicas=4
```

kubectl Deployments Commands

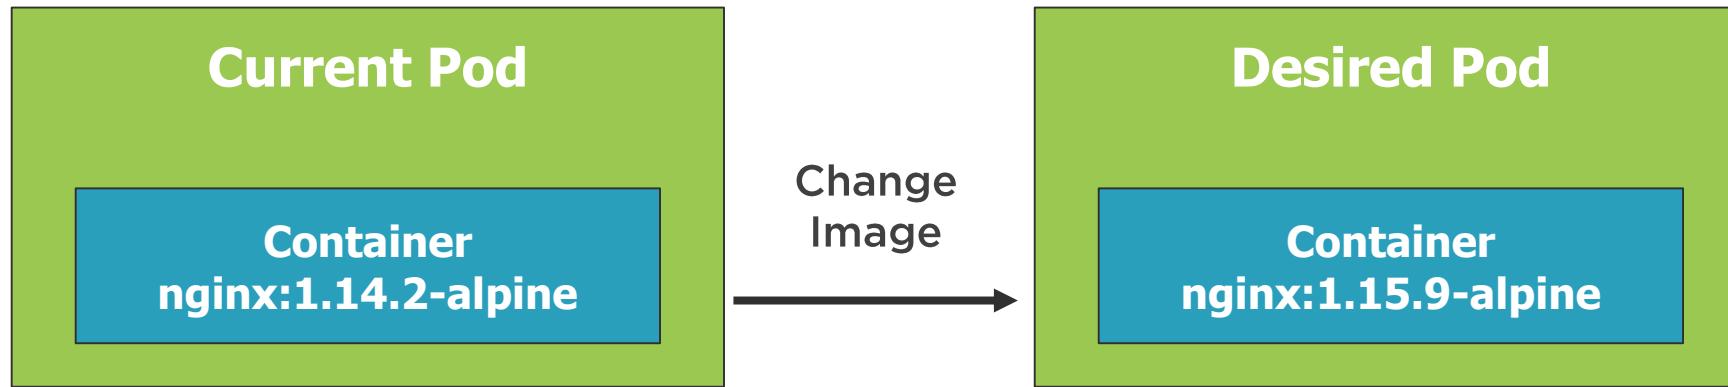
Several different kubectl commands can be used to create and work with Deployments



Deployment Options



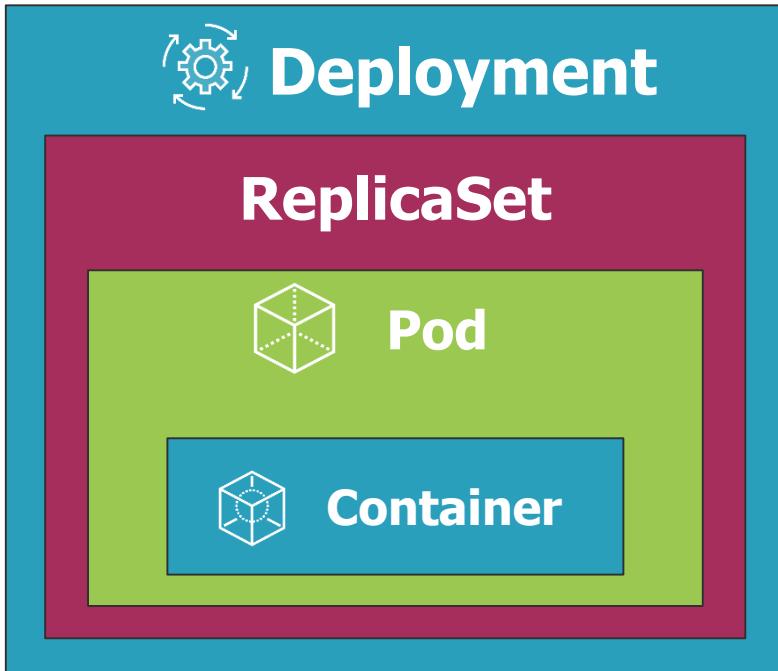
How Do You Update Existing Pods?



Zero downtime deployments
allow software updates to be
deployed to production without
impacting end users



Deployment Options



One of the strengths of Kubernetes is zero downtime deployments

Update an application's Pods without impacting end users

Several options are available:

- Rolling updates
- Blue-green deployments
- Canary deployments
- Rollbacks



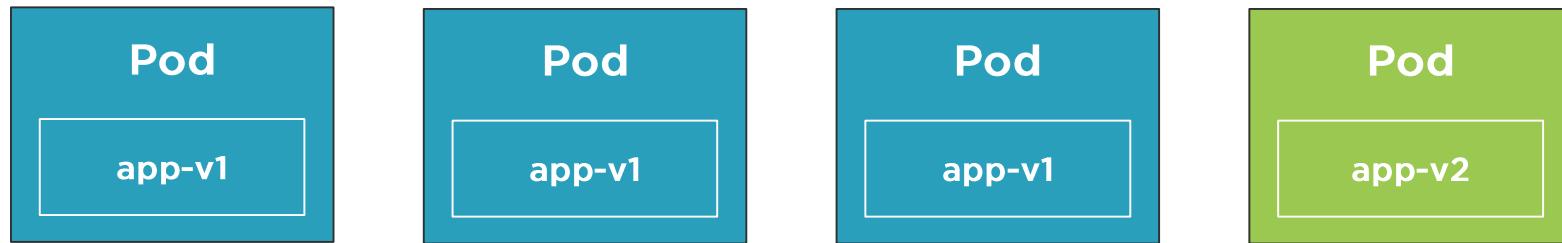
Rolling Deployments

Initial Pod State



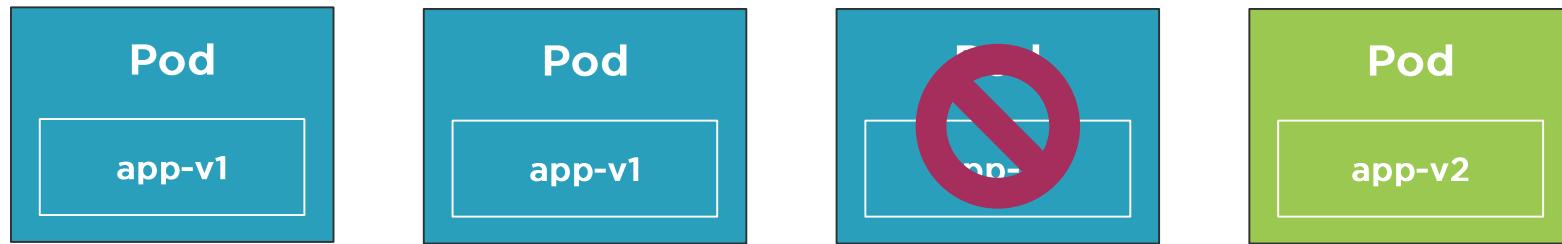
Rolling Deployments

Rollout New Pod



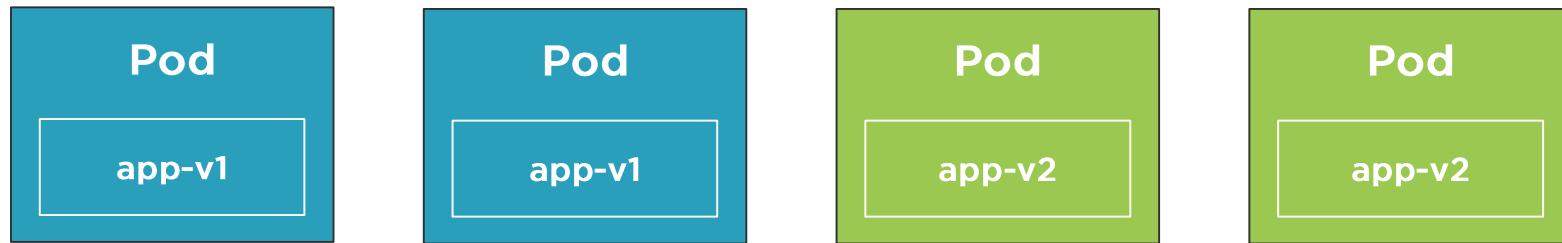
Rolling Deployments

Delete Pod



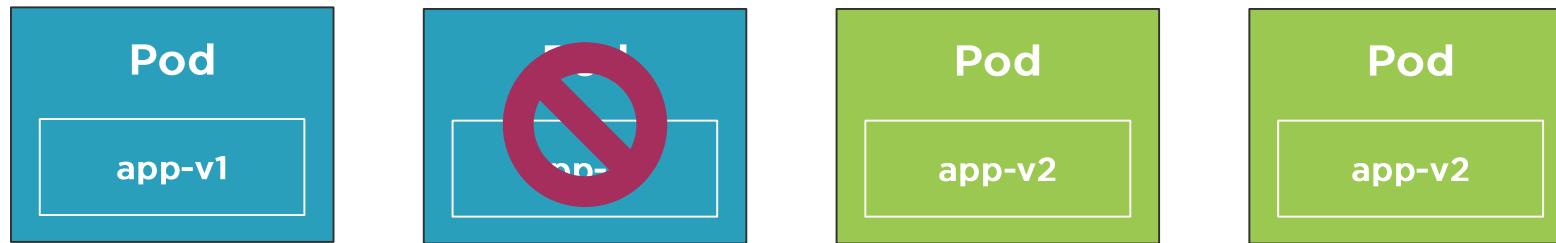
Rolling Deployments

Rollout New Pod



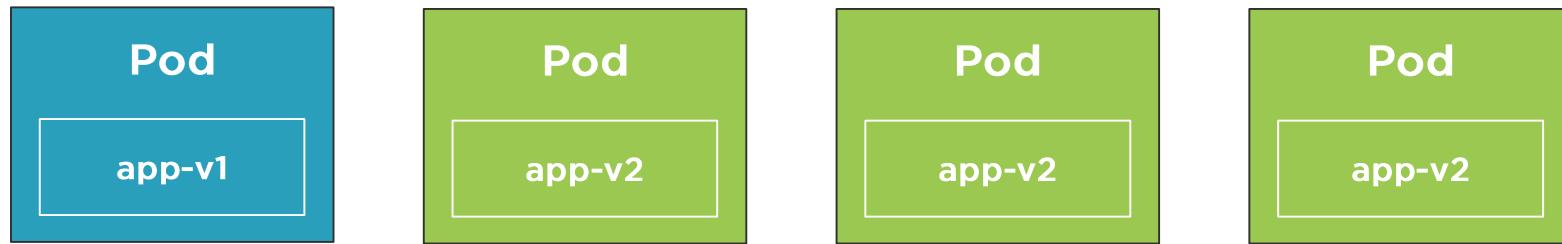
Rolling Deployments

Delete Pod



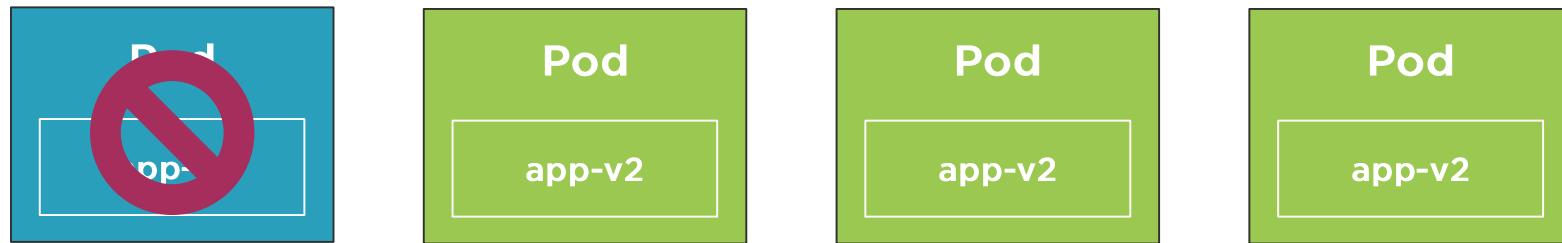
Rolling Deployments

Rollout New Pod



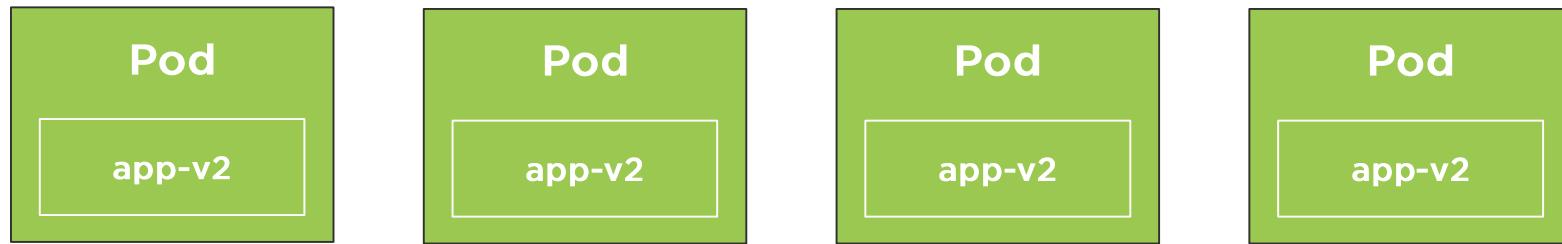
Rolling Deployments

Delete Pod



Rolling Deployments

Rollout New Pod



Updating a Deployment

Update a deployment by changing the YAML and applying changes to the cluster with kubectl apply

```
# Apply changes made in a YAML file  
kubectl apply -f file.deployment.yml
```

Zero Downtime Deployments in Action



Summary



Pods are deployed, managed, and scaled using deployments and ReplicaSets

Deployments are a higher-level resource that define one or more Pod templates

The `kubectl create` or `apply` commands can be used to run a deployment

Kubernetes supports zero downtime deployments



Creating Services



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Services Core Concepts

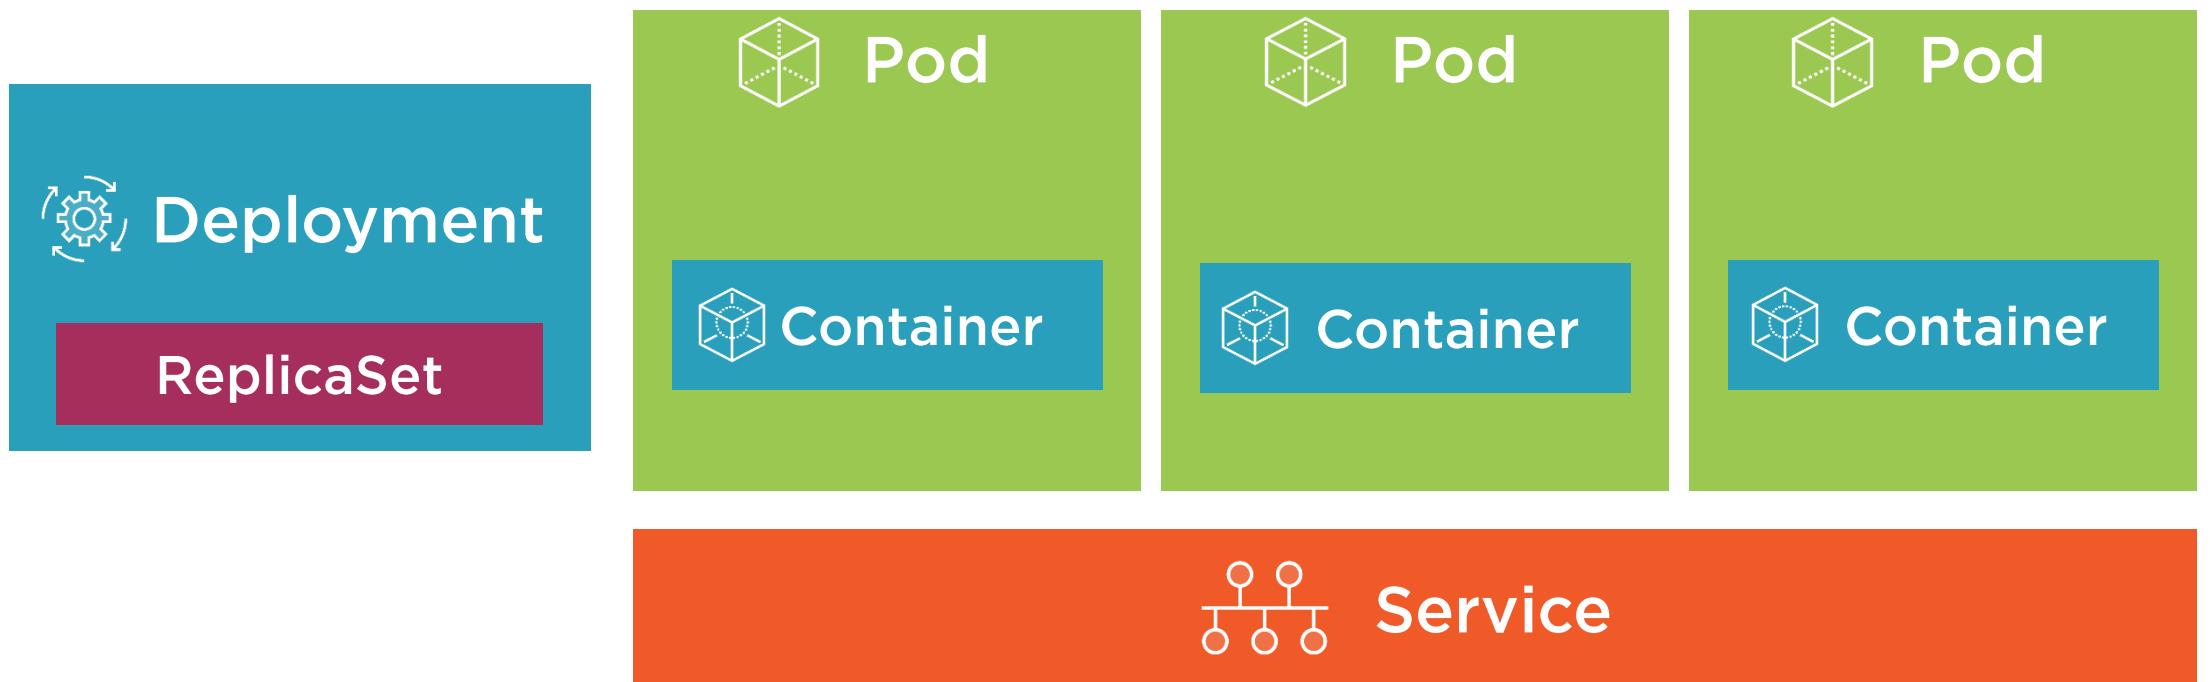
Service Types

Creating a Service with kubectl

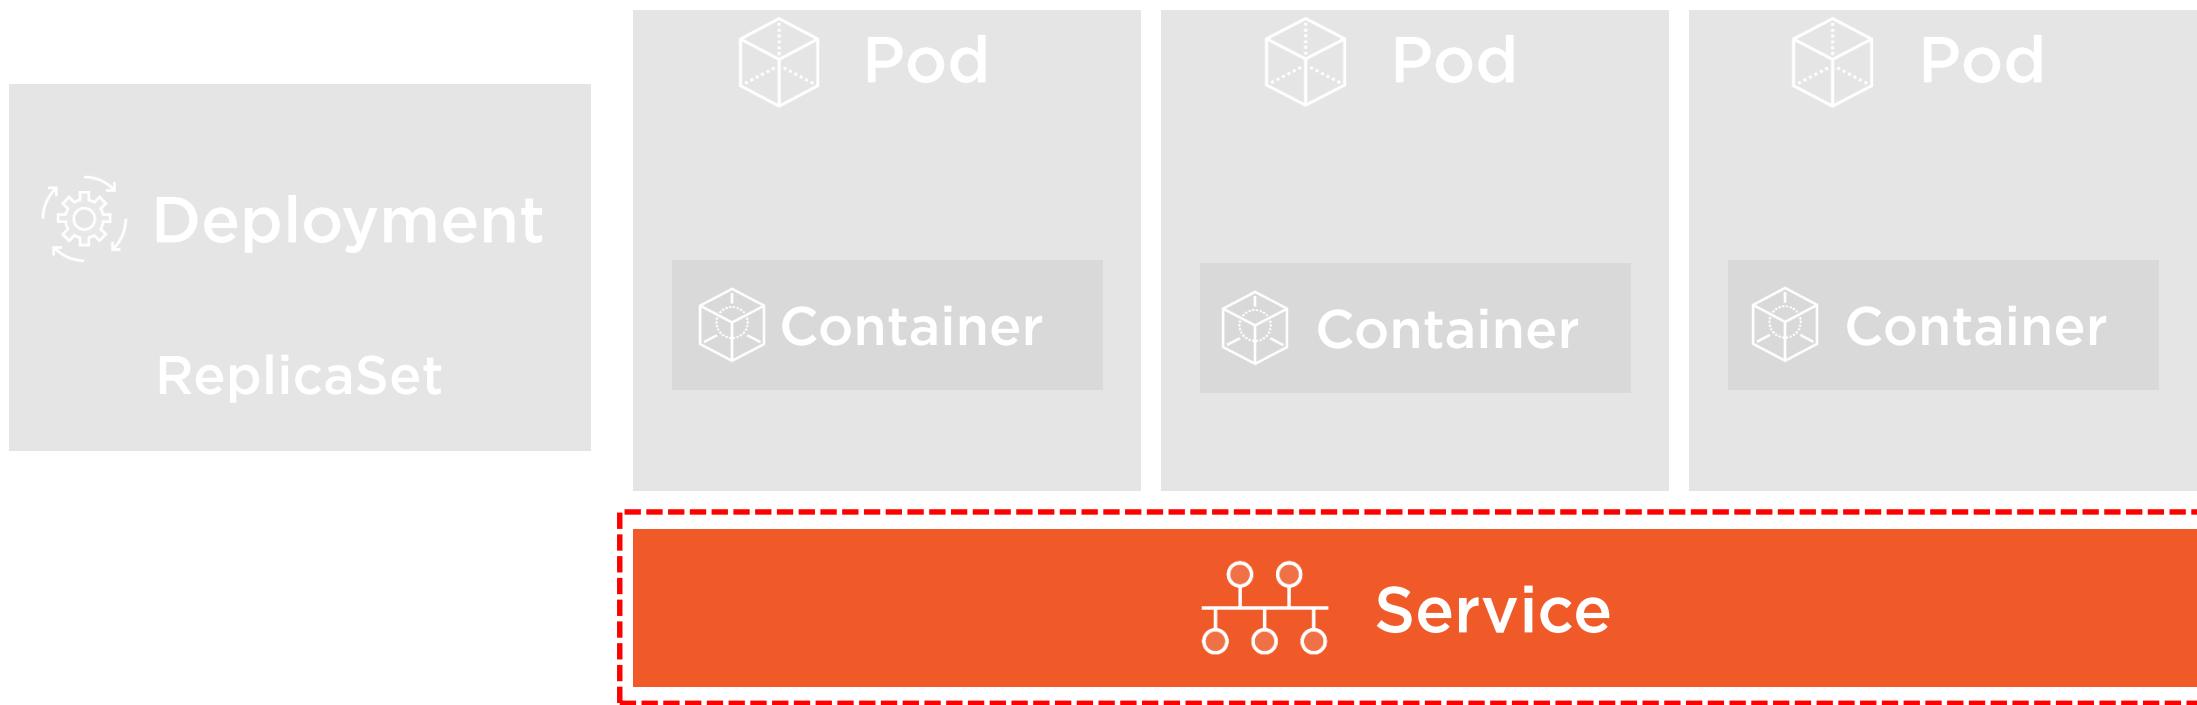
Creating a Service with YAML



You Are Here



You Are Here



Services Core Concepts

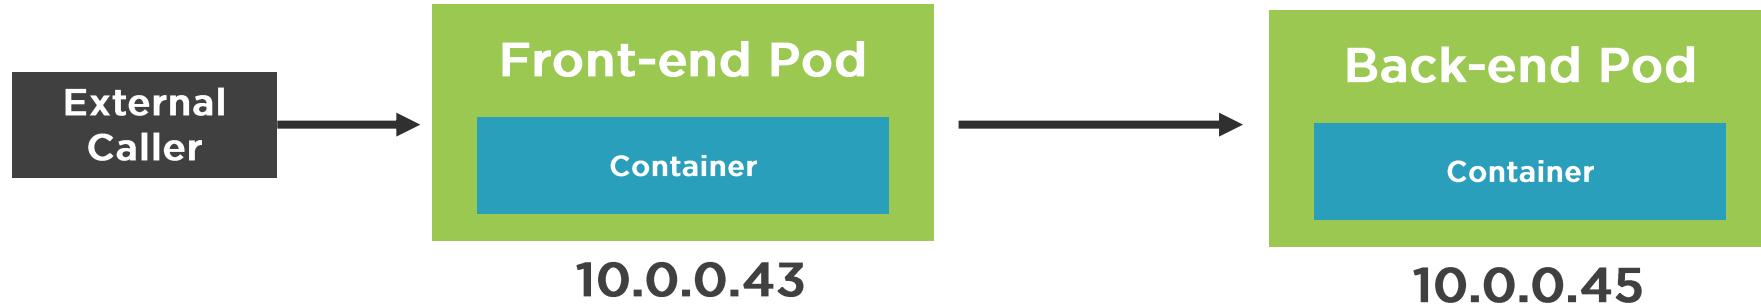


A Service provides a single point
of entry for accessing one or
more Pods



Review Question:

Since Pods live and die, can you rely on their IP?



Answer:

No! That's why we need Services - IPs change a lot!



Pods are "mortal" and may only live a short time (ephemeral)

You can't rely on a Pod IP address staying the same

Pods can be horizontally scaled so each Pod gets its own IP address

A Pod gets an IP address after it has been scheduled (no way for clients to know IP ahead of time)

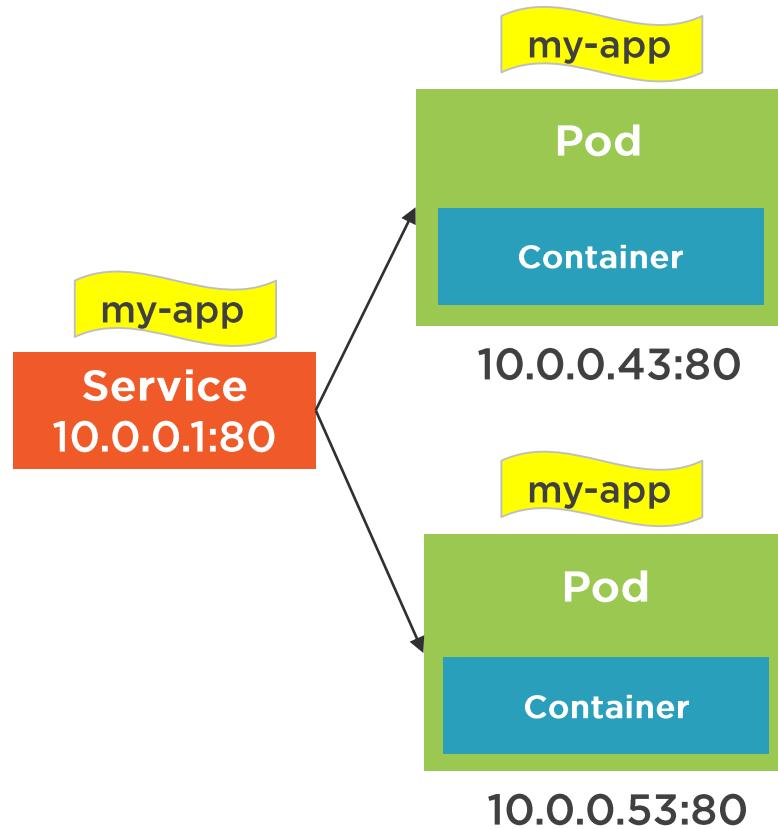
The Life of a Pod



10.0.0.43



The Role of Services



Services abstract Pod IP addresses from consumers

Load balances between Pods

Relies on labels to associate a Service with a Pod

Node's kube-proxy creates a virtual IP for Services

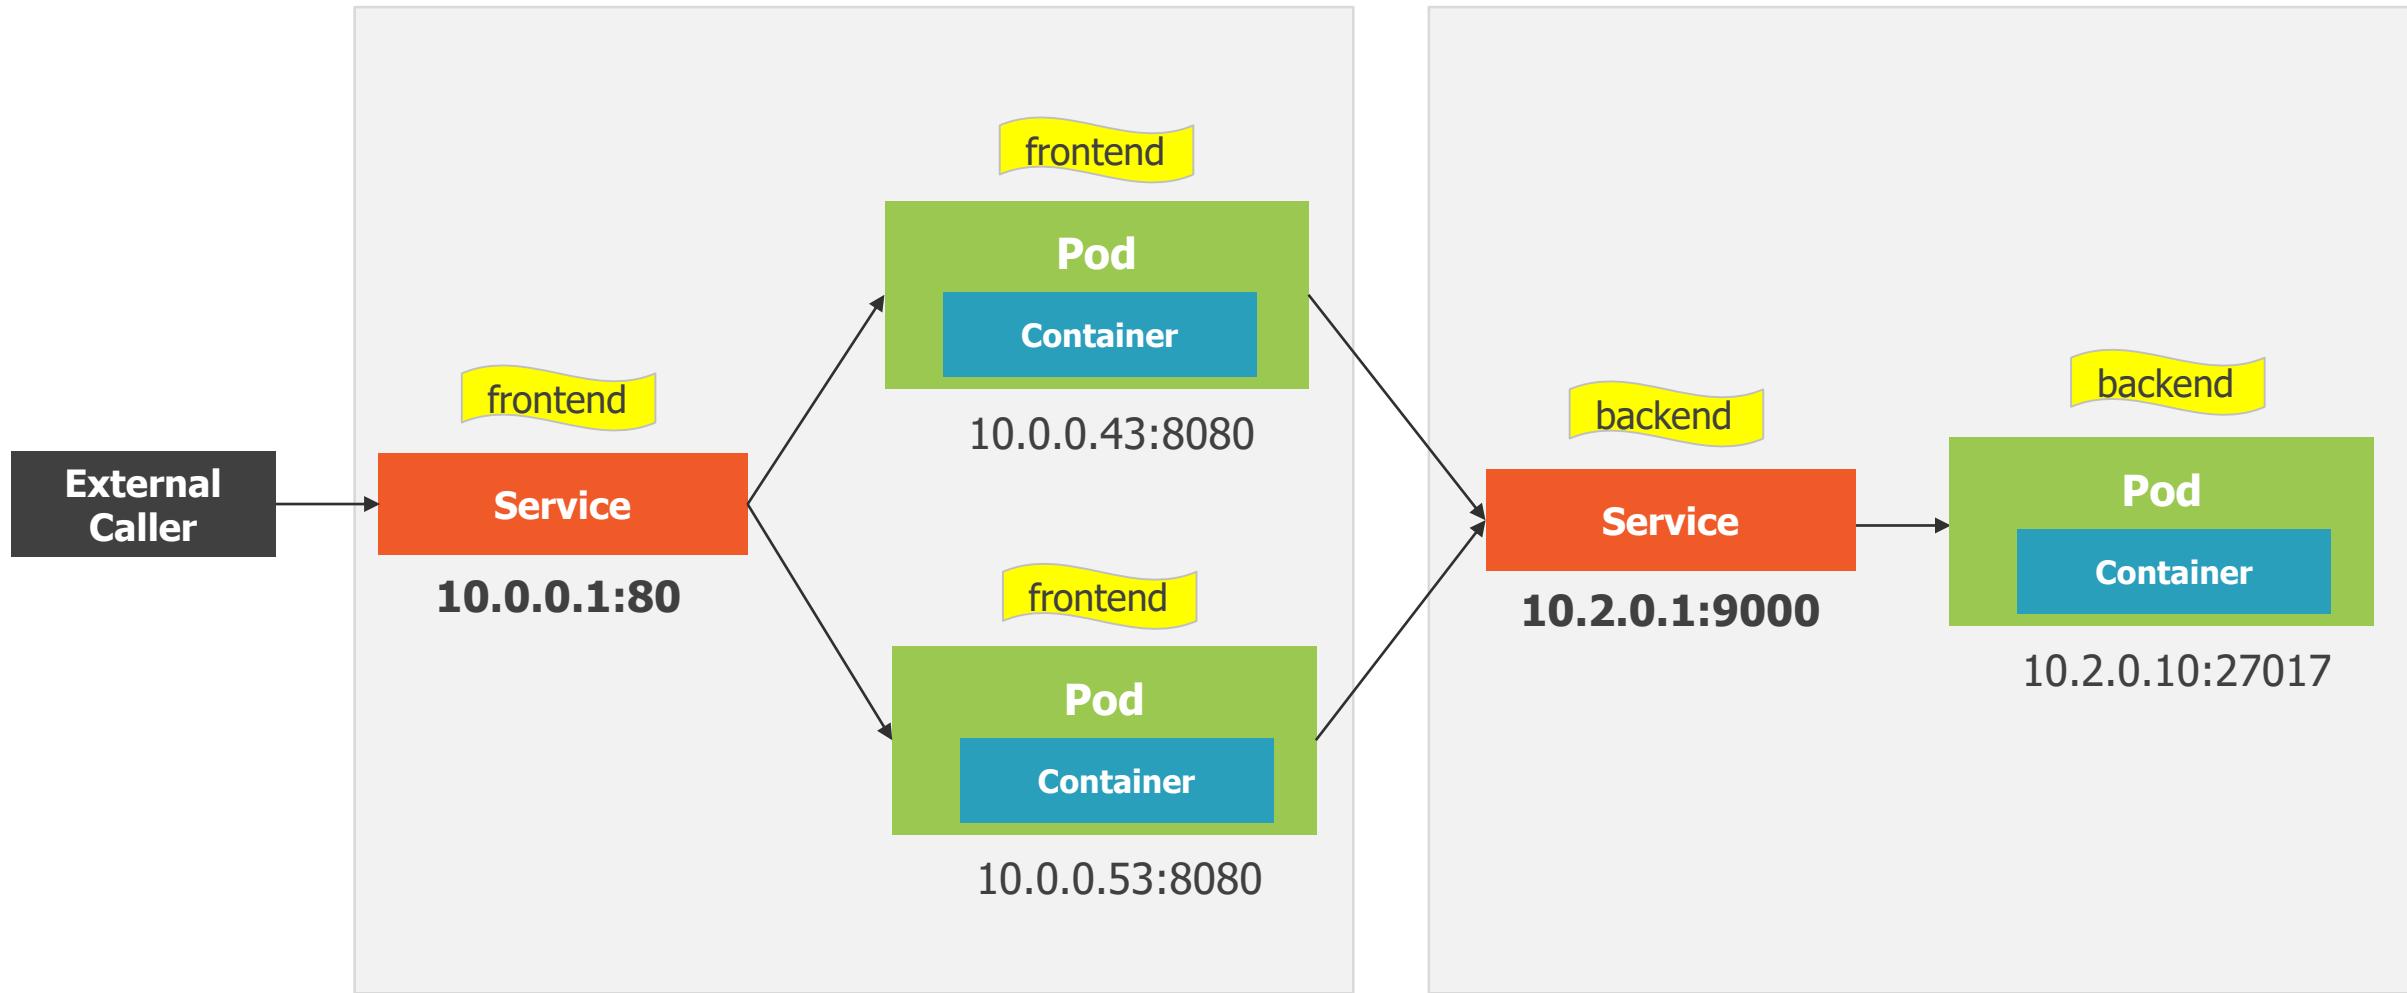
Layer 4 (TCP/UDP over IP)

Services are not ephemeral

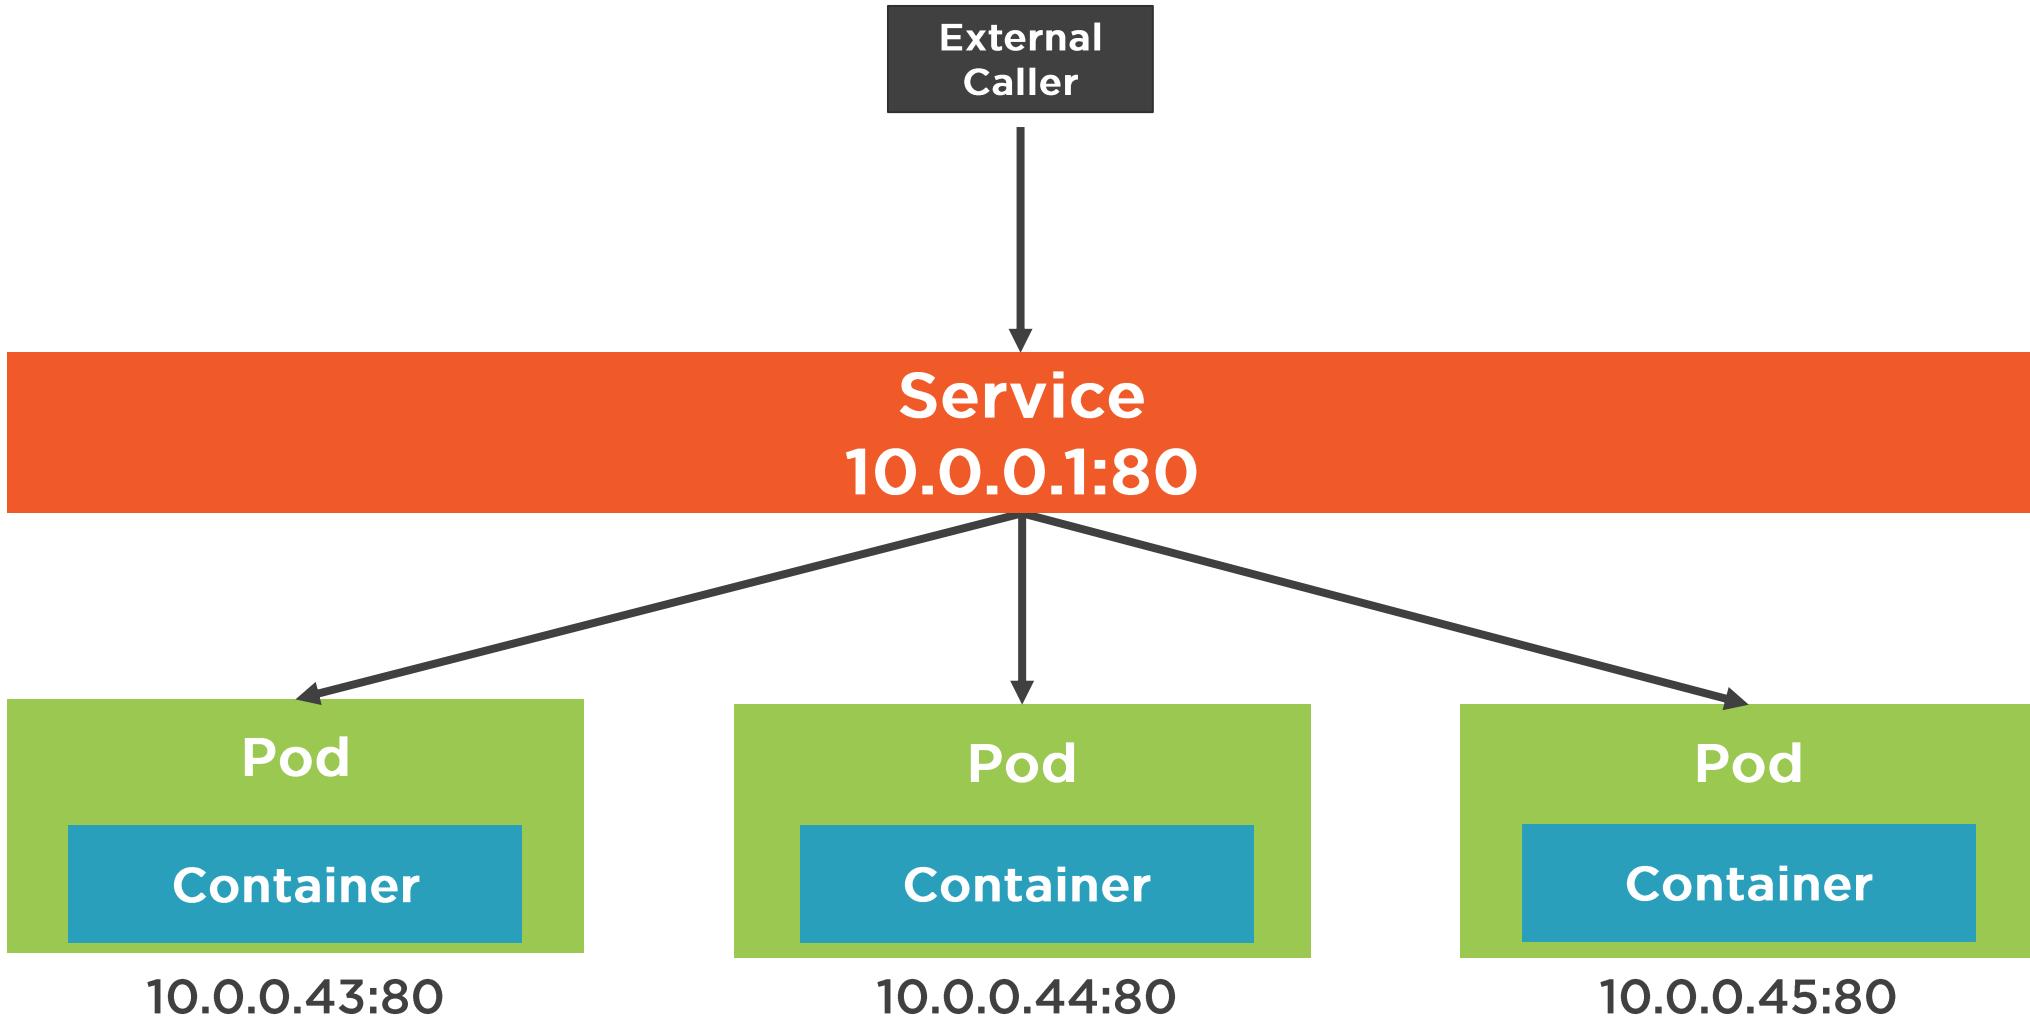
Creates endpoints which sit between a Service and Pod



Calling Services



Services and Pod Load Balancing



Service Types



Service Types



Service

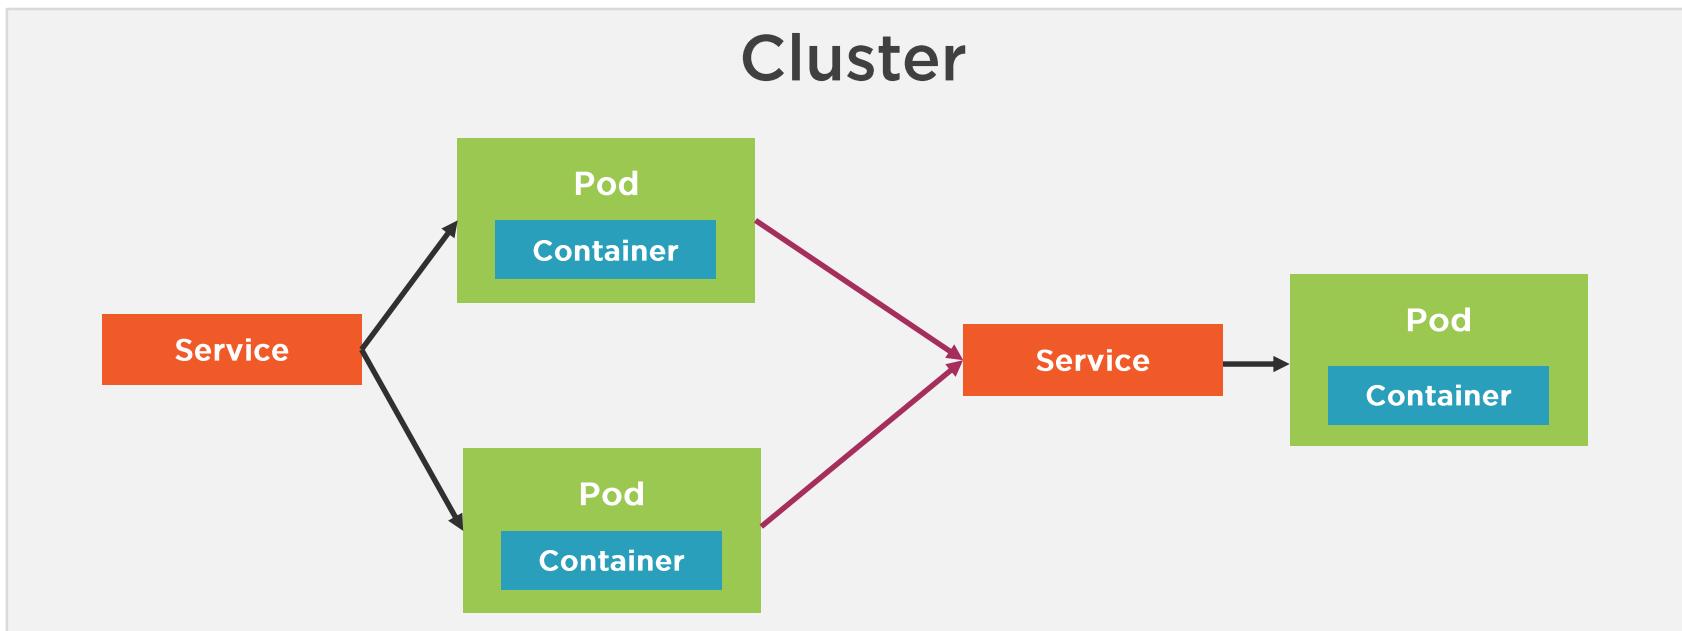
Services can be defined in different ways:

- ClusterIP – Expose the service on a cluster-internal IP (default)
- NodePort – Expose the service on each Node's IP at a static port.
- LoadBalancer – Provision an external IP to act as a load balancer for the service
- ExternalName – Maps a service to a DNS name



ClusterIP Service

Service IP is exposed internally within the cluster
Only Pods within the cluster can talk to the Service
Allows Pods to talk to other Pods

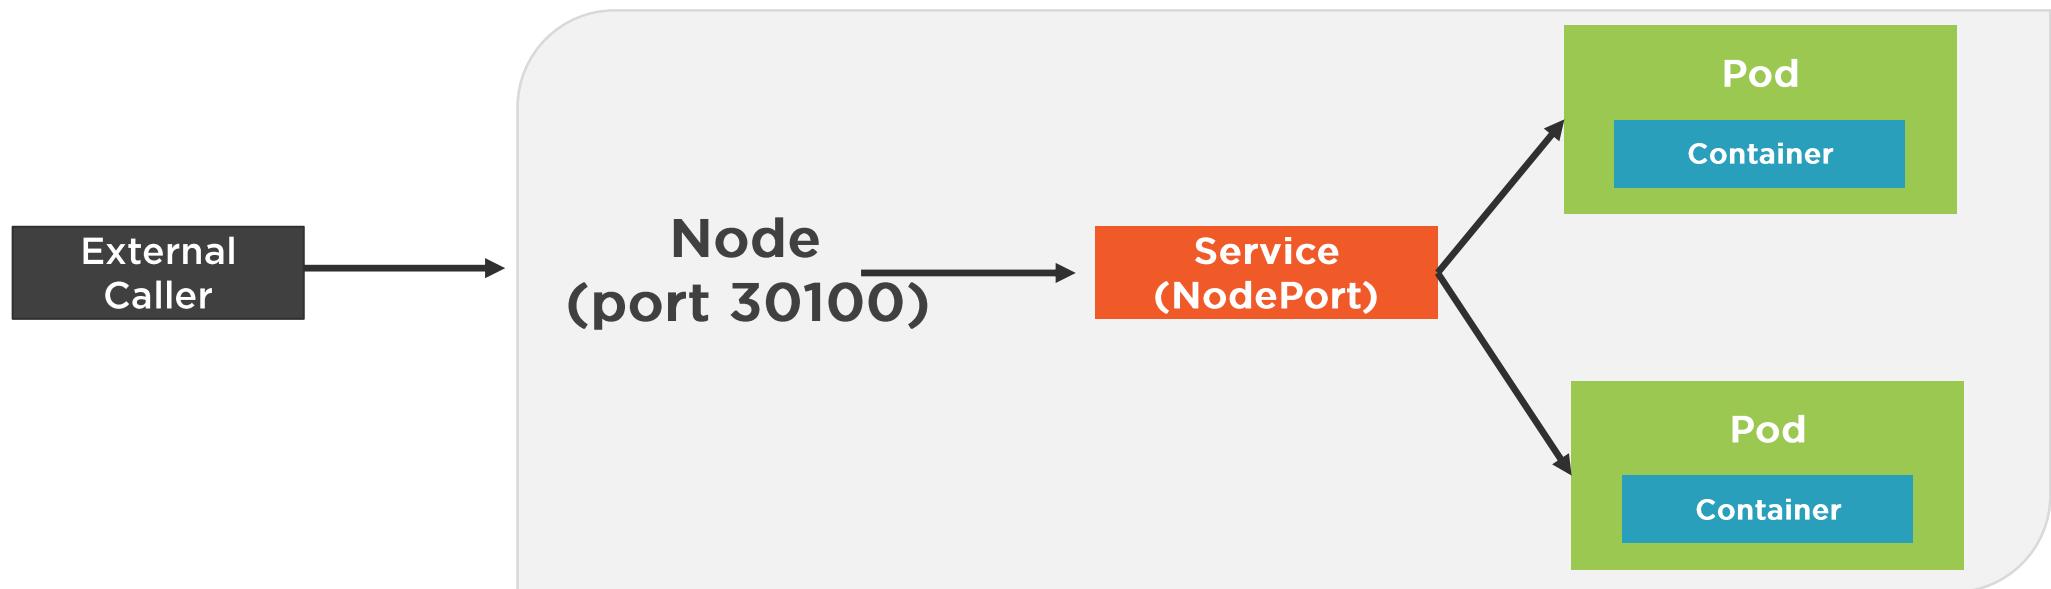


NodePort Service

Exposes the Service on each Node's IP at a static port

Allocates a port from a range (default is 30000-32767)

Each Node proxies the allocated port



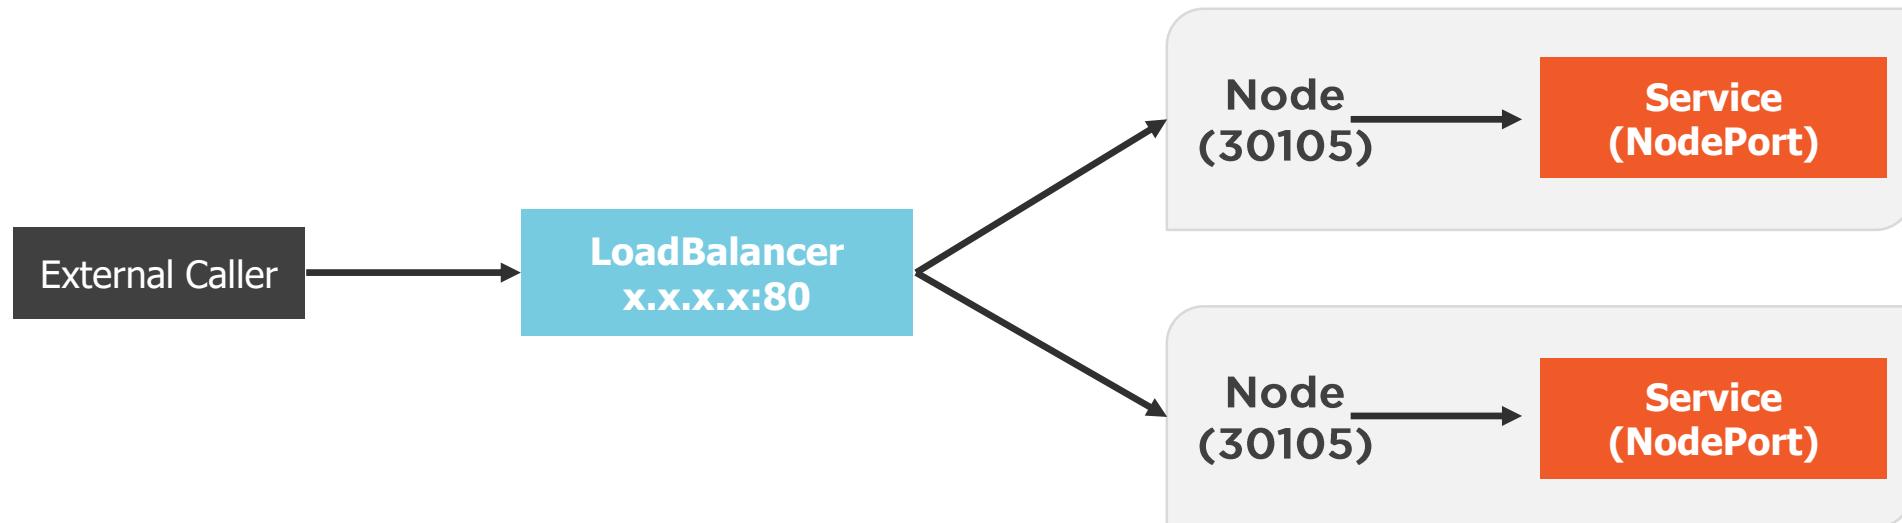
LoadBalancer Service

Exposes a Service externally

Useful when combined with a cloud provider's load balancer

NodePort and ClusterIP Services are created

Each Node proxies the allocated port

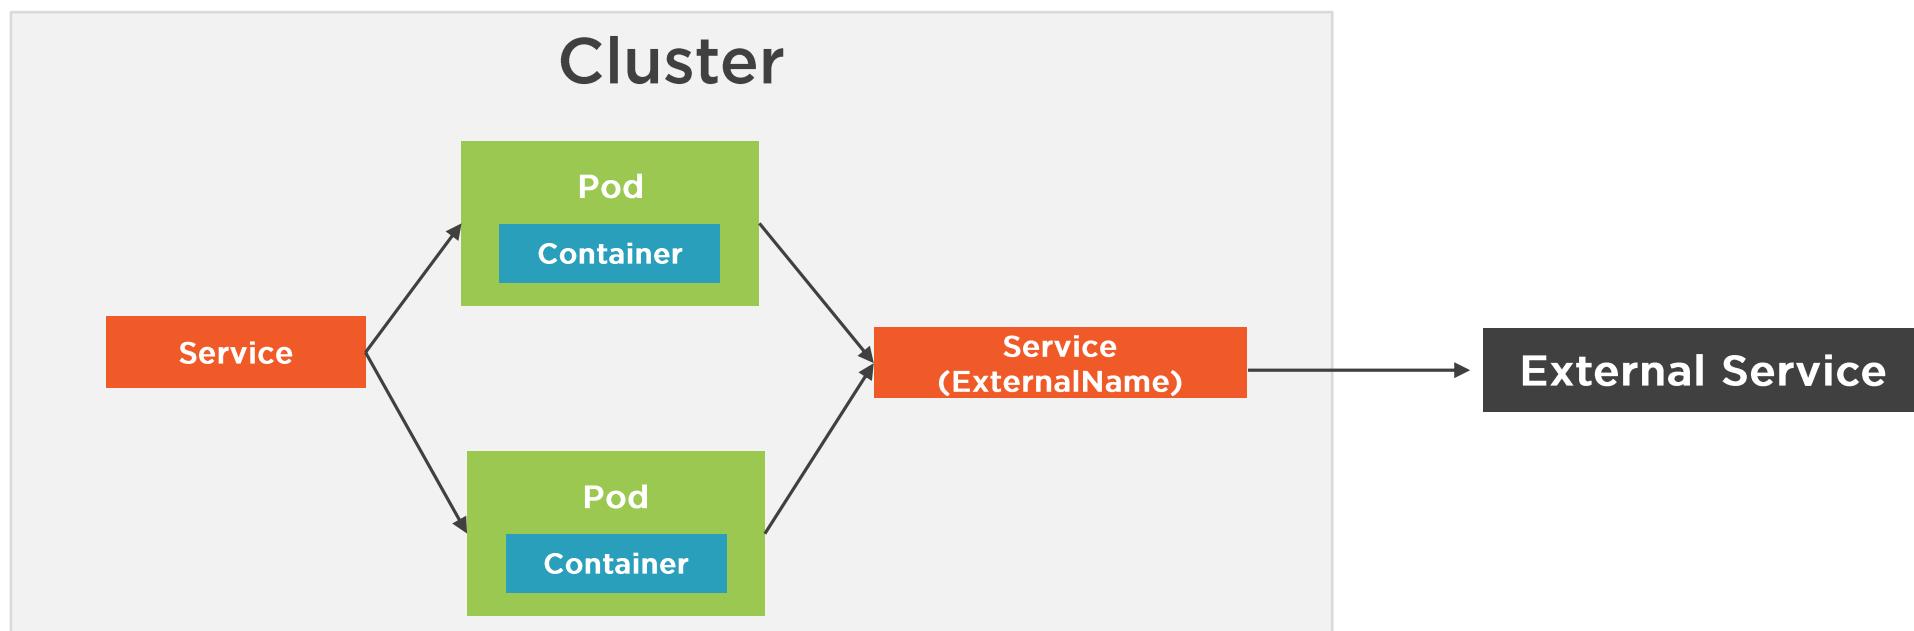


ExternalName Service

Service that acts as an alias for an external service

Allows a Service to act as the proxy for an external service

External service details are hidden from cluster (easier to change)



Creating a Service with kubectl



Port Forwarding

Q. How can you access a Pod from outside of Kubernetes?

A. Port forwarding

Use the kubectl port-forward to forward a local port to a Pod port

```
# Listen on port 8080 locally and forward to port 80 in Pod  
kubectl port-forward pod/[pod-name] 8080:80
```

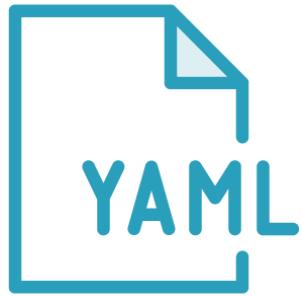
```
# Listen on port 8080 locally and forward to Deployment's Pod  
kubectl port-forward deployment/[deployment-name] 8080
```

```
# Listen on port 8080 locally and forward to Service's Pod  
kubectl port-forward service/[service-name] 8080
```

Creating a Service with YAML

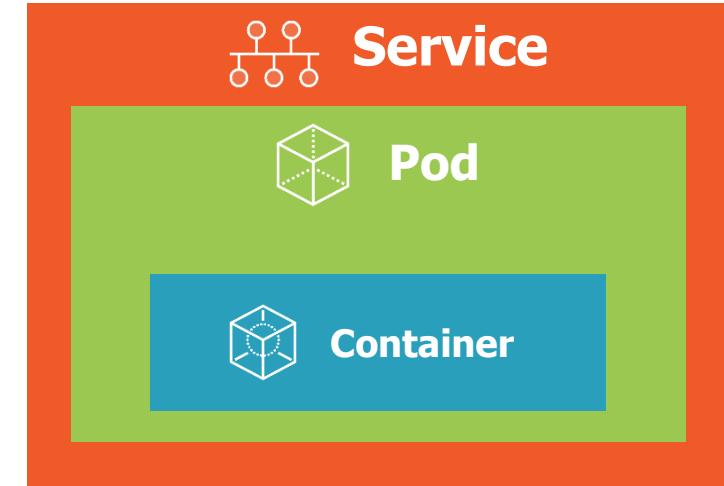


Defining a Service with YAML



Service

+ kubectl =



```
apiVersion: v1
kind: Service
metadata:
spec:
  type:
  selector:
  ports:
```

- ◀ **Kubernetes API version and resource type (Service)**
- ◀ **Metadata about the Service**
- ◀ **Type of service (ClusterIP, NodePort, LoadBalancer) – defaults to ClusterIP**
- ◀ **Select Pod template label(s) that service will apply to**
- ◀ **Define container target port and the port for the service**



```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  selector:
    app: nginx
  ports:
  - name: http
    port: 80
    targetPort: 80
```

◀ **Kubernetes API version and resource type (Service)**

◀ **Metadata about the Service**

◀ **Service will apply to resources with a label of app: nginx**

◀ **Define container target port(s) and the port(s) for the Service**



Connecting to a Service by It's DNS Name

```
apiVersion: v1
kind: Service
metadata:
  name: frontend
...
...
```

```
apiVersion: v1
kind: Service
metadata:
  name: backend
...
...
```

◀ **Name of Service (each Service gets a DNS entry)**

◀ **A frontend Pod can access a backend Pod using backend:port**



Creating a NodePort Service

```
apiVersion: v1
kind: Service
metadata:
  ...
spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - port: 80
      targetPort: 80
      nodePort: 31000
```

◀ Set Service *type* to NodePort

◀ Optionally set NodePort value
(defaults between 30000-32767)



Creating a LoadBalancer Service

```
apiVersion: v1
kind: Service
metadata:
...
spec:
  type: LoadBalancer
  selector:
    app: nginx
  ports:
  - port: 80
    targetPort: 80
```

◀ Set Service *type* to LoadBalancer
(normally used with cloud providers)



Creating an ExternalName Service

```
apiVersion: v1
kind: Service
metadata:
  name: external-service
spec:
  type: ExternalName
  externalName: api.acmecorp.com
  ports:
  - port: 9000
```

- ◀ Other Pods can use this FQDN to access the external service
- ◀ Set type to ExternalName
- ◀ Service will proxy to FQDN



kubectl and Services



Creating a Service

Use the `kubectl create` command along with the `--filename` or `-f` switch



```
# Create a Service
kubectl create -f file.service.yml
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	55d
nginx-clusterip	ClusterIP	10.102.26.70	<none>	8080/TCP	6s

```
# Update a Service  
# Assumes --save-config was used with create  
kubectl apply -f file.service.yml
```

Updating or Creating a Service

Use the kubectl apply command along with the --filename or -f switch



Deleting a Service

Use the `kubectl delete` command along with the `--filename` or `-f` switch

```
# Delete a Service  
kubectl delete -f file.service.yml
```

```
# Shell into a Pod and test a URL. Add -c [containerID]  
# in cases where multiple containers are running in the Pod  
kubectl exec [pod-name] -- curl -s http://podIP  
  
# Install and use curl (example shown is for Alpine Linux)  
kubectl exec [pod-name] -it sh  
> apk add curl  
> curl -s http://podIP
```

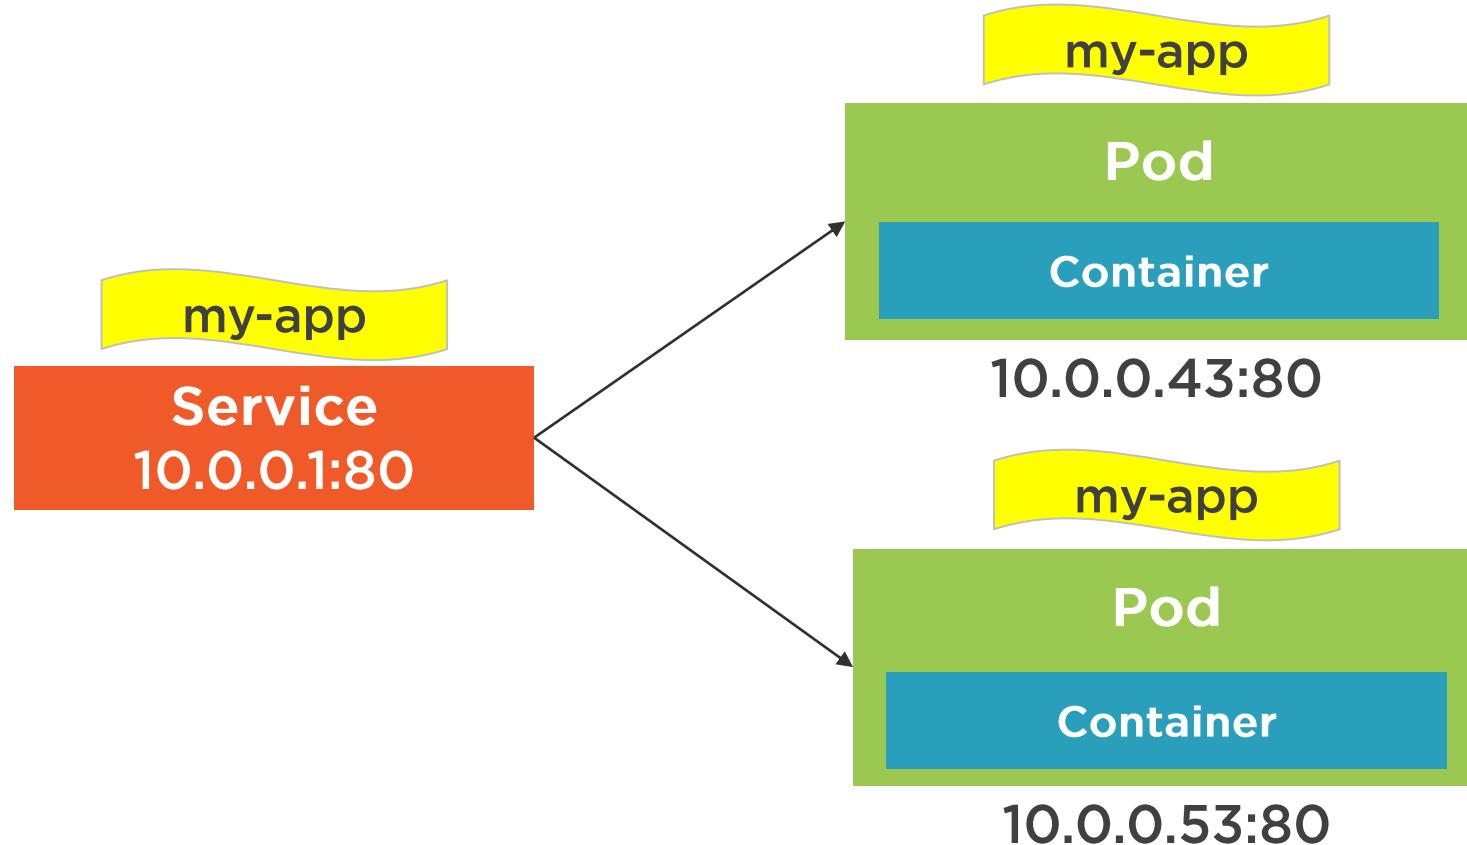
Testing a Service and Pod with curl

How can you quickly test if a Service and Pod is working?

Use kubectl exec to shell into a Pod/Container



The Role of Services



kubectl Services in Action



Summary



Pods live and die so their IP address can change

Services abstract Pod IP addresses from consumers

Labels associate a Service with a Pod

Service types:

- ClusterIP (internal to cluster - default)
- NodePort (exposes Service on each's Node's IP)
- LoadBalancer (exposes a Service externally)
- ExternalName (proxies to an external service)



Understanding Storage Options



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Storage Core Concepts

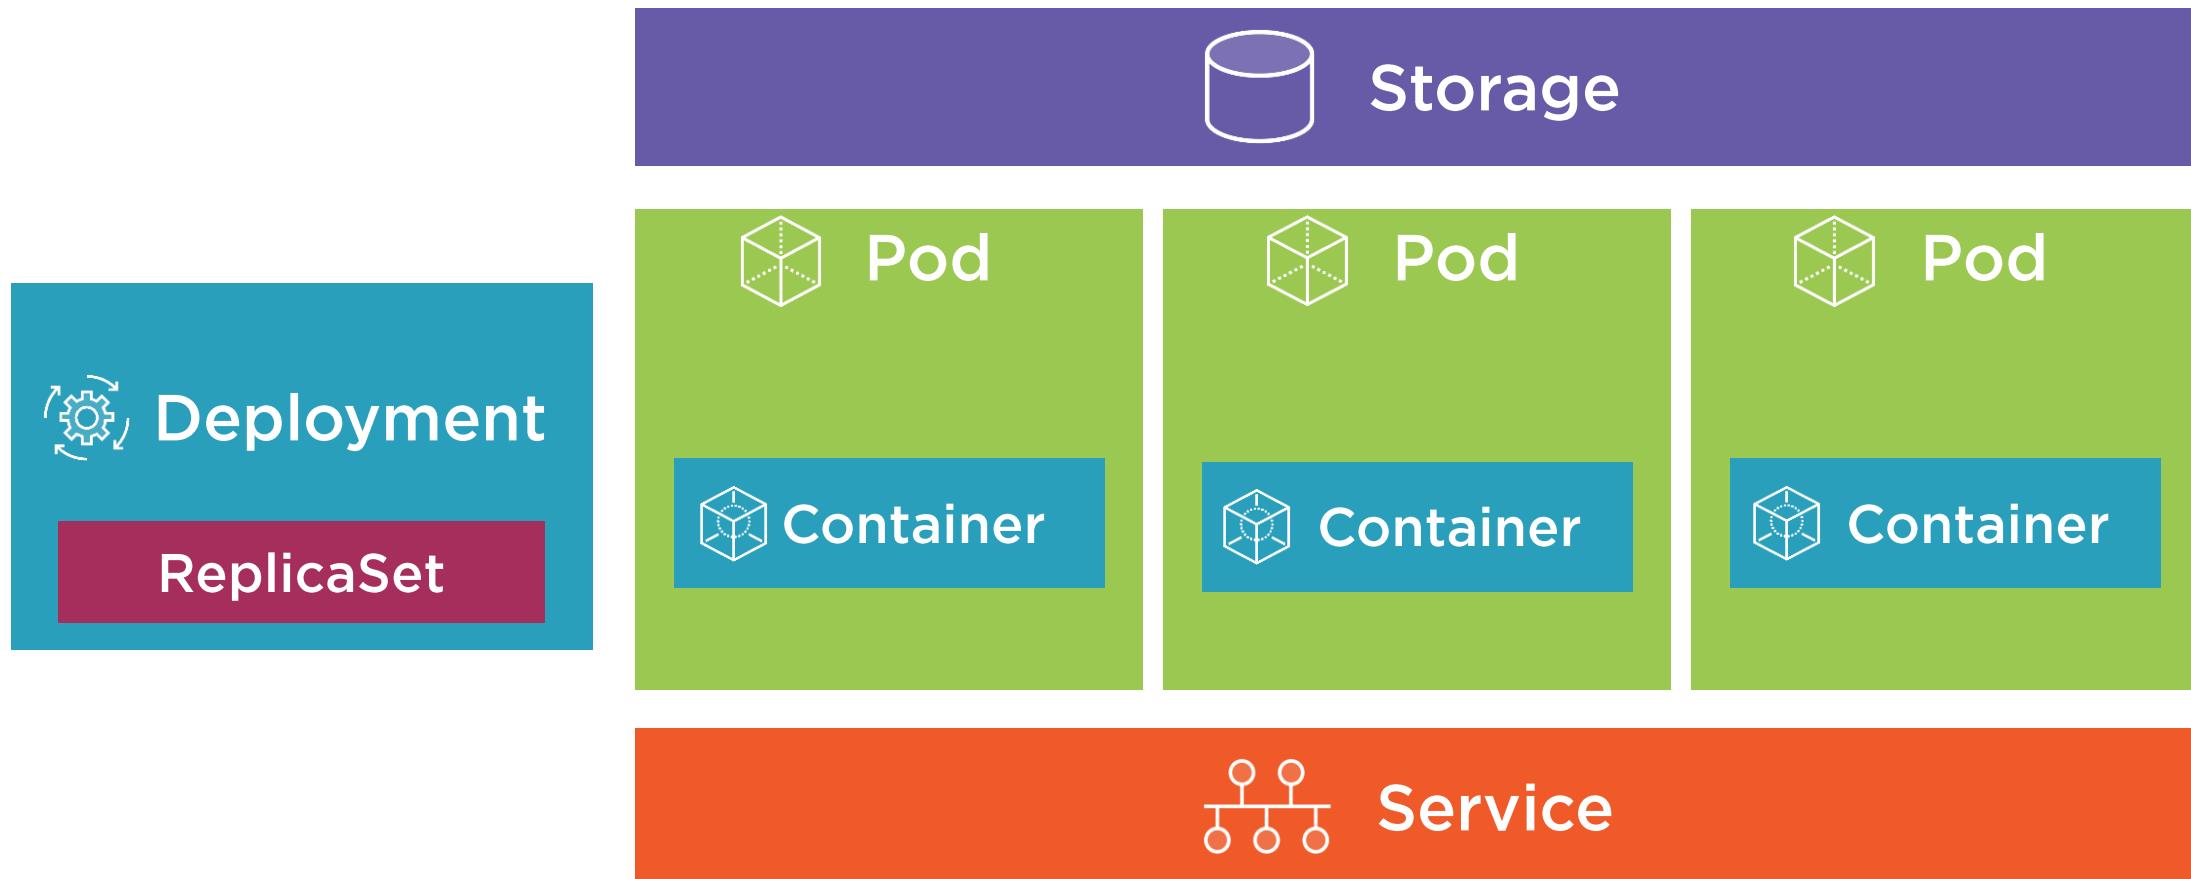
Volumes

**PersistentVolumes and
PersistentVolumeClaims**

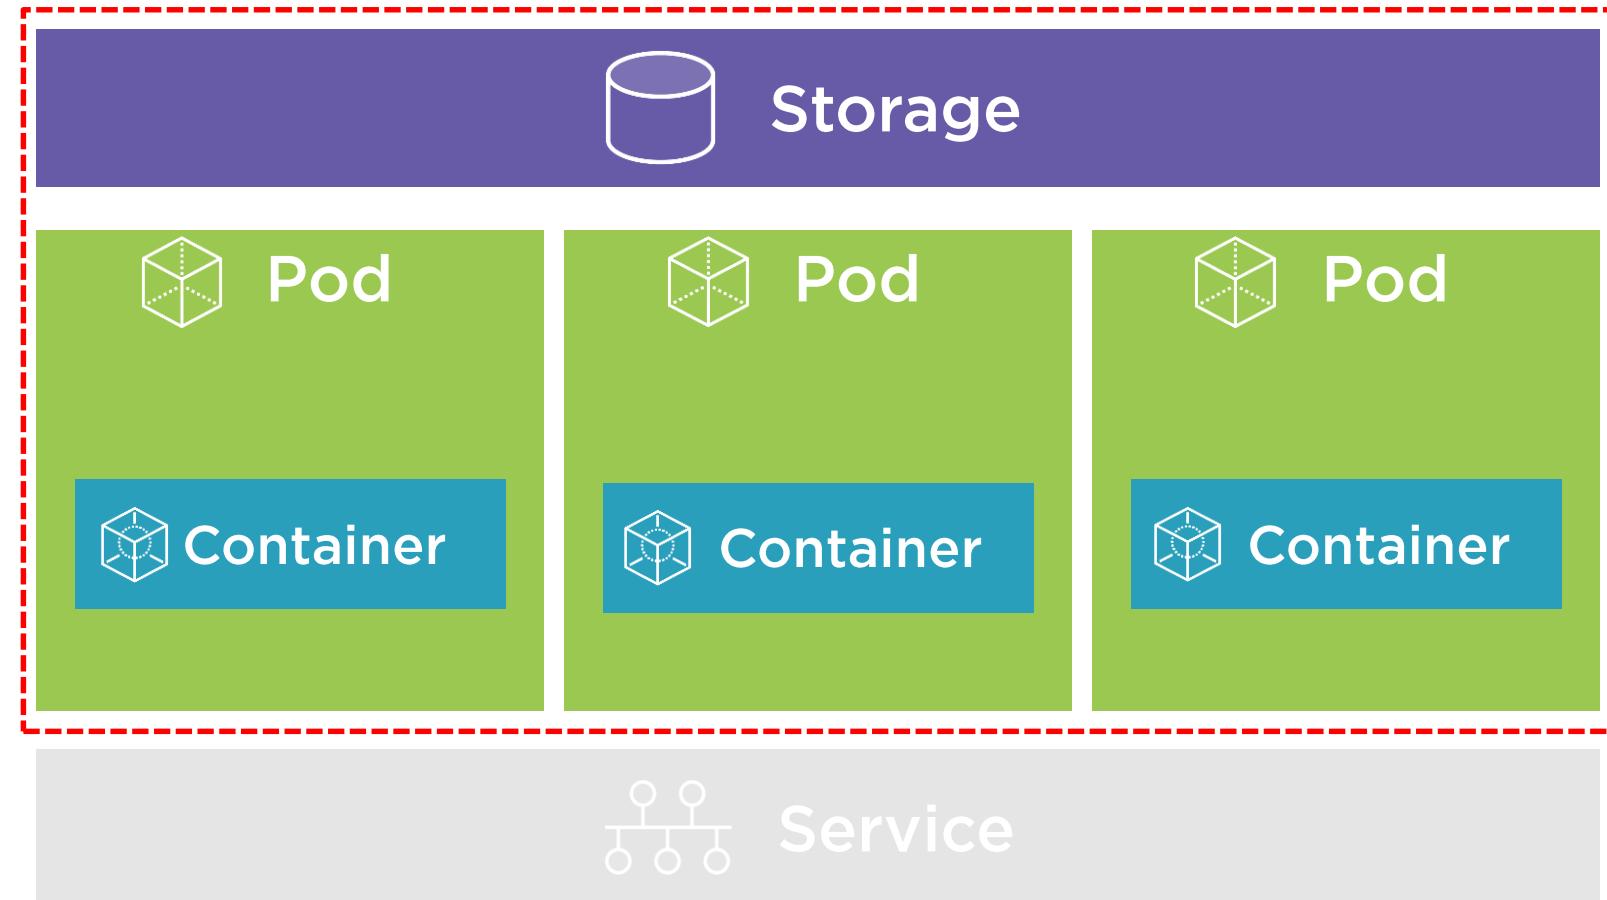
StorageClasses



You Are Here



You Are Here



Storage Core Concepts



Question:

How do you store application state/data and exchange it between Pods with Kubernetes?

Answer:

Volumes (although other data storage options exist)



A Volume can be used to hold data and state for Pods and containers.



Pods live and die so their file system is short-lived (ephemeral)

Volumes can be used to store state/data and use it in a Pod

A Pod can have multiple Volumes attached to it

Containers rely on a mountPath to access a Volume

Kubernetes supports:

- Volumes
- PersistentVolumes
- PersistentVolumeClaims
- StorageClasses

Pod State and Data



Volumes



Volumes and Volume Mounts



A Volume references a storage location

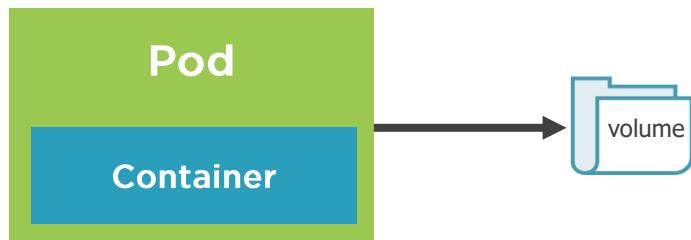
Must have a unique name

Attached to a Pod and may or may not be tied to the Pod's lifetime (depending on the Volume type)

A Volume Mount references a Volume by name and defines a mountPath



Volumes Type Examples



emptyDir – Empty directory for storing "transient" data (shares a Pod's lifetime) useful for sharing files between containers running in a Pod

hostPath – Pod mounts into the node's filesystem

nfs – An NFS (Network File System) share mounted into the Pod

configMap/secret – Special types of volumes that provide a Pod with access to Kubernetes resources

persistentVolumeClaim – Provides Pods with a more persistent storage option that is abstracted from the details

Cloud – Cluster-wide storage



Volume Types

awsElasticBlockStore	azureDisk	azureFile	cephfs	configMap
csi	downwardAPI	emptyDir	fc	flexVolume
flocker	gcePersistentDisk	glusterfs	hostPath	iscsi
local	nfs	persistentVolumeClaim	projected	portworxVolume
quobyte	rbd	scaleIO	secret	storageos
vsphereVolume				



Defining an emptyDir Volume

```
apiVersion: v1
kind: Pod
spec:
  volumes:
    - name: html
      emptyDir: {}
  containers:
    - name: nginx
      image: nginx:alpine
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
          readOnly: true
    - name: html-updater
      image: alpine
      command: ["/bin/sh", "-c"]
      args:
        - while true; do date >> /html/index.html;
          sleep 10; done
      volumeMounts:
        - name: html
          mountPath: /html
```

- ◀ Define initial Volume named "html" that is an empty directory (lifetime of the Pod)
- ◀ Reference "html" Volume and define a mountPath
- ◀ Update file in Volume mount /html path with latest date every 10 seconds
- ◀ Reference "html" Volume (defined above) and define a mountPath



Defining a hostPath Volume

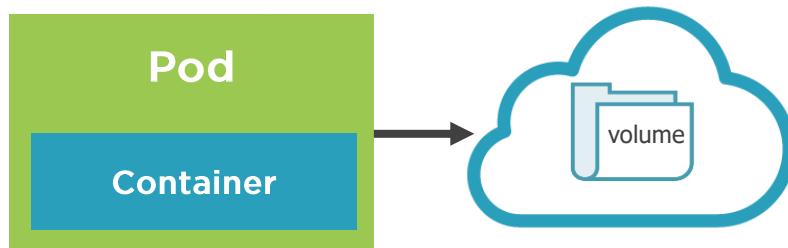
```
apiVersion: v1
kind: Pod
spec:
  volumes:
    - name: docker-socket
      hostPath:
        path: /var/run/docker.sock
        type: Socket
  containers:
    - name: docker
      image: docker
      command: ["sleep"]
      args: ["100000"]
  volumeMounts:
    - name: docker-socket
      mountPath: /var/run/docker.sock
```

◀ Define a socket volume on host that points to /var/run/docker.sock

◀ Reference "docker-socket" Volume and define mountPath



Cloud Volumes



Cloud providers (Azure, AWS, GCP, etc.) support different types of Volumes:

- Azure – Azure Disk and Azure File
- AWS – Elastic Block Store
- GCP – GCE Persistent Disk



Defining an Azure File Volume

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  volumes:
    - name: data
      azureFile:
        secretName: <azure-secret>
        shareName: <share-name>
        readOnly: false
  containers:
    - image: someimage
      name: my-app
      volumeMounts:
        - name: data
          mountPath: /data/storage
```

◀ Define initial Volume named "data" that is Azure File storage

◀ Reference "data" Volume and define a mountPath



Defining an AWS Volume

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  volumes:
    - name: data
      awsElasticBlockStore:
        volumeID: <volume_ID>
        fsType: ext4
  containers:
    - image: someimage
      name: my-app
      volumeMounts:
        - name: data
          mountPath: /data/storage
```

◀ Define initial Volume named "data" that is a awsElasticBlockStore

◀ Reference "data" Volume and define a mountPath



Defining a Google Cloud gcePersistentDisk Volume

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  volumes:
    - name: data
      gcePersistentDisk:
        pdName: datastorage
        fsType: ext4
  containers:
    - image: someimage
      name: my-app
      volumeMounts:
        - name: data
          mountPath: /data/storage
```

◀ Define initial Volume named "data" that is a gcePersistentDisk

◀ Reference "data" Volume and define a mountPath



Viewing a Pod's Volumes

Several different techniques can be used to view a Pod's Volumes

```
# Describe Pod
```

```
kubectl describe pod [pod-name]
```

```
Volumes:
```

```
  html:
```

```
    Type:  EmptyDir (a temporary directory that shares a pod's lifetime)
```

```
    Medium:
```

```
# Get Pod YAML
```

```
kubectl get pod [pod-name] -o yaml
```

```
volumeMounts:
```

```
- mountPath: /html  
  name: html
```

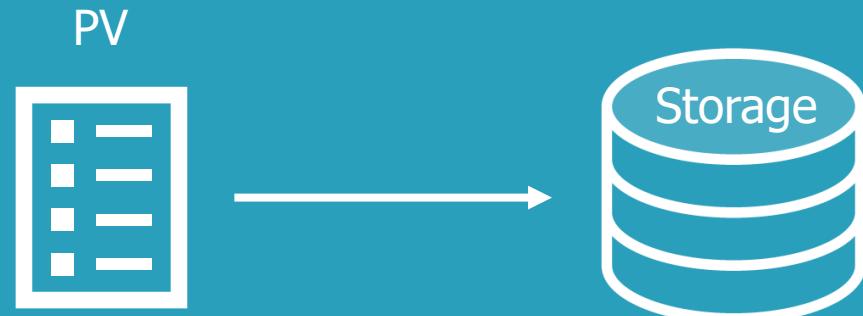
Volumes in Action



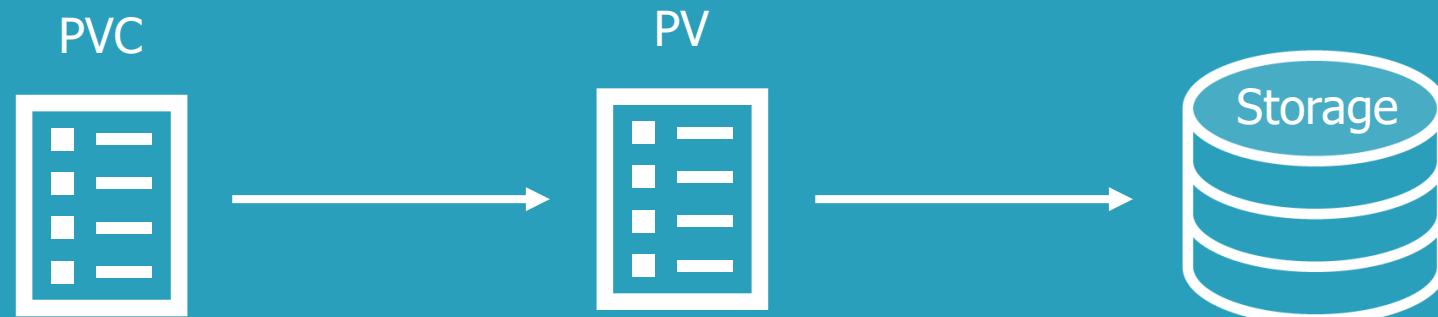
PersistentVolumes and PersistentVolumeClaims



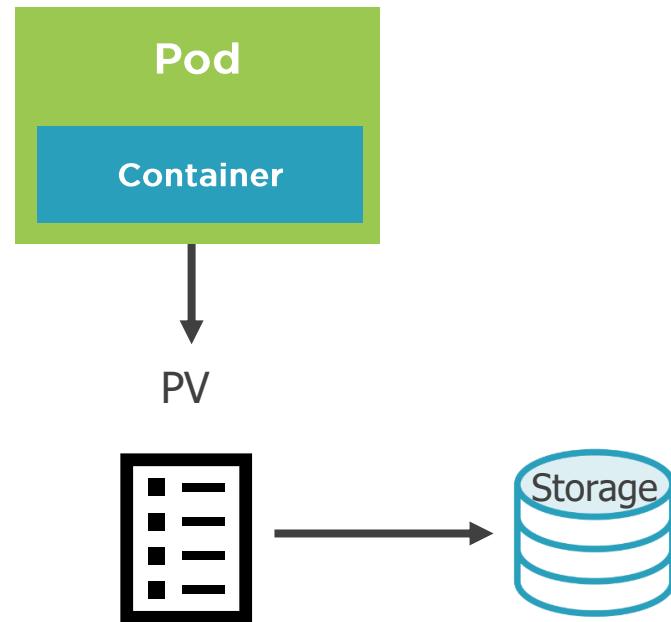
A PersistentVolume (PV) is a cluster-wide storage unit provisioned by an administrator with a lifecycle independent from a Pod.



A PersistentVolumeClaim (PVC) is a request for a storage unit (PV).



PersistentVolume



A **PersistentVolume** is a cluster-wide storage resource that relies on network-attached storage (NAS)

Normally provisioned by a cluster administrator

Available to a Pod even if it gets rescheduled to a different Node

Rely on a storage provider such as NFS, cloud storage, or other options

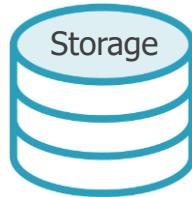
Associated with a Pod by using a **PersistentVolumeClaim (PVC)**



PersistentVolume Workflow

1

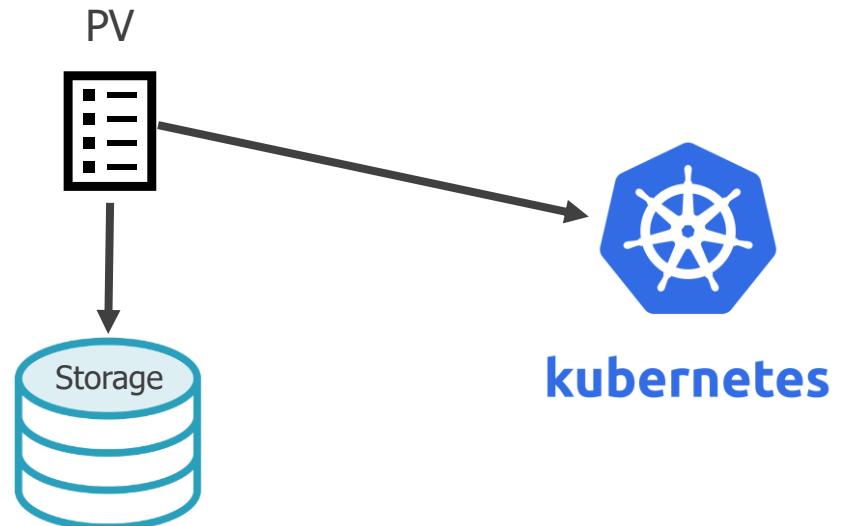
Create network storage resource (NFS, cloud, etc.)



PersistentVolume Workflow

2

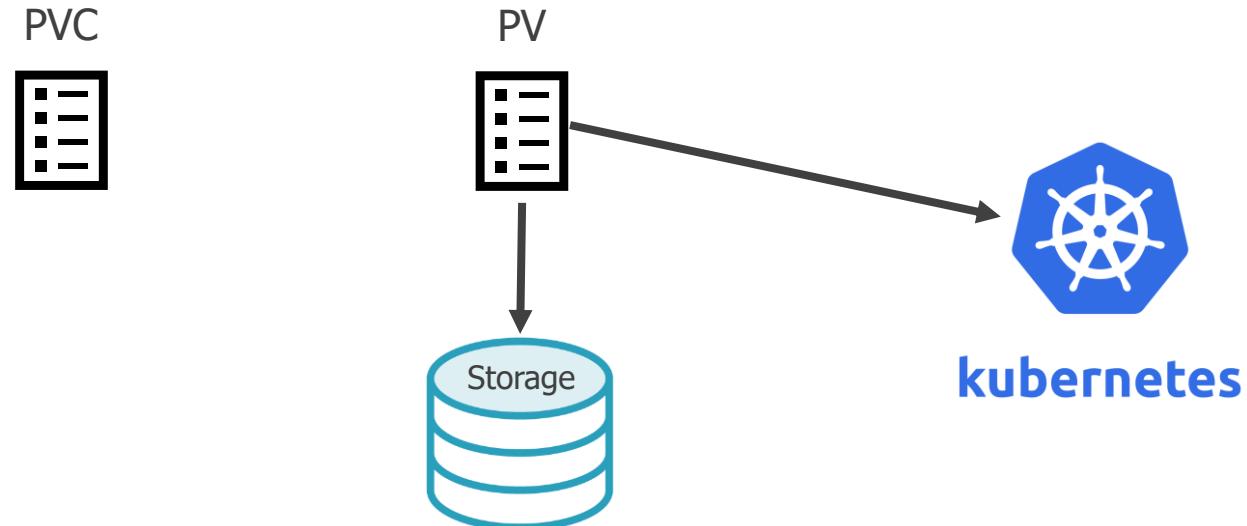
**Define a Persistent Volume (PV)
and send to the Kubernetes API**



PersistentVolume Workflow

3

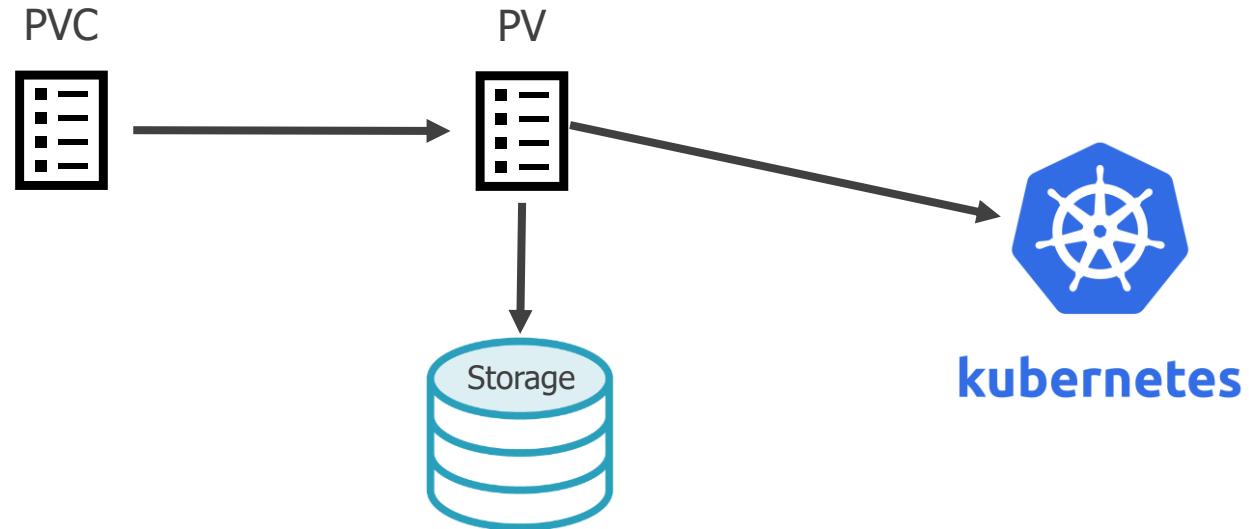
Create a PersistentVolumeClaim (PVC)



PersistentVolume Workflow

4

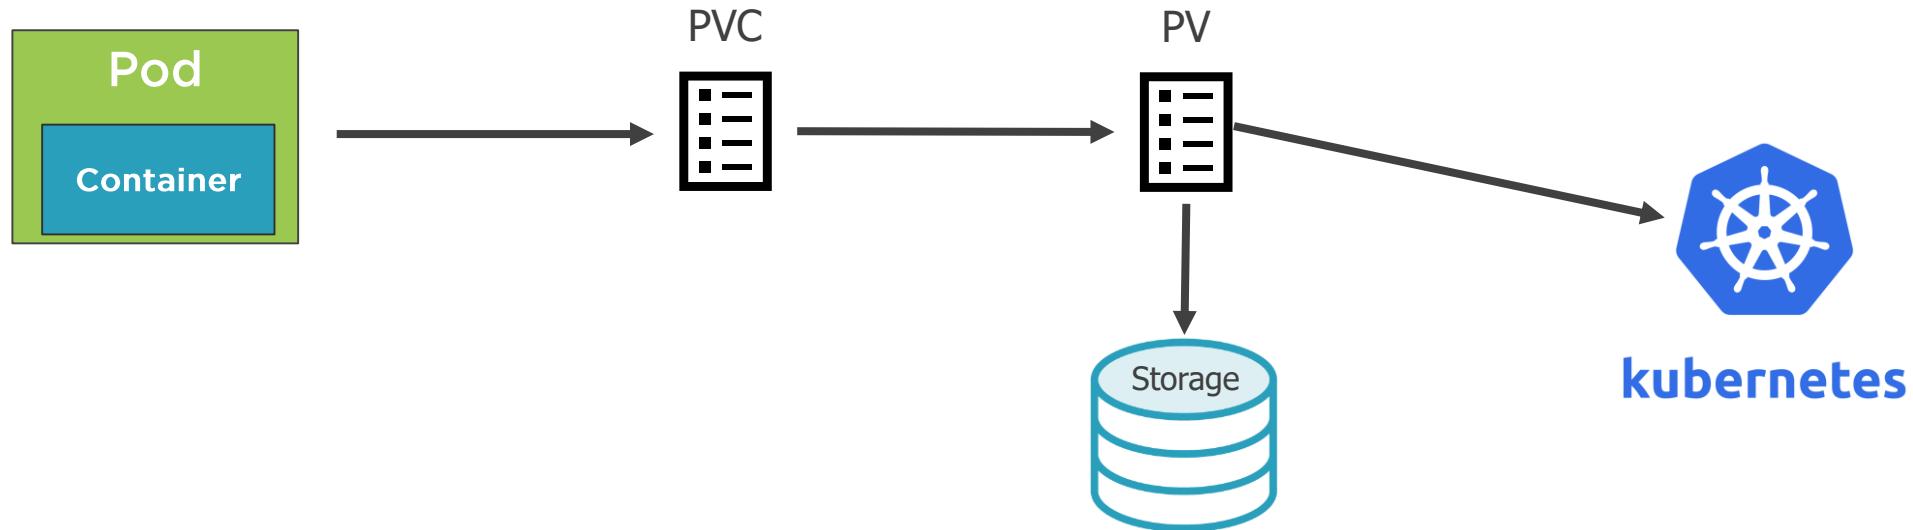
Kubernetes binds the PVC to the PV



PersistentVolume Workflow

5

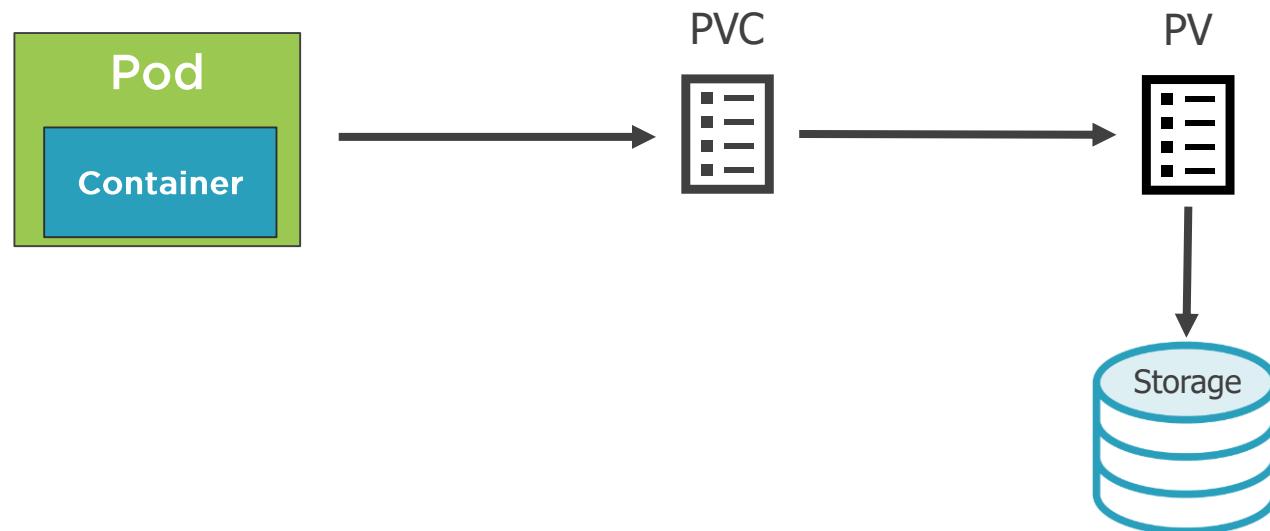
Pod Volume references the PVC



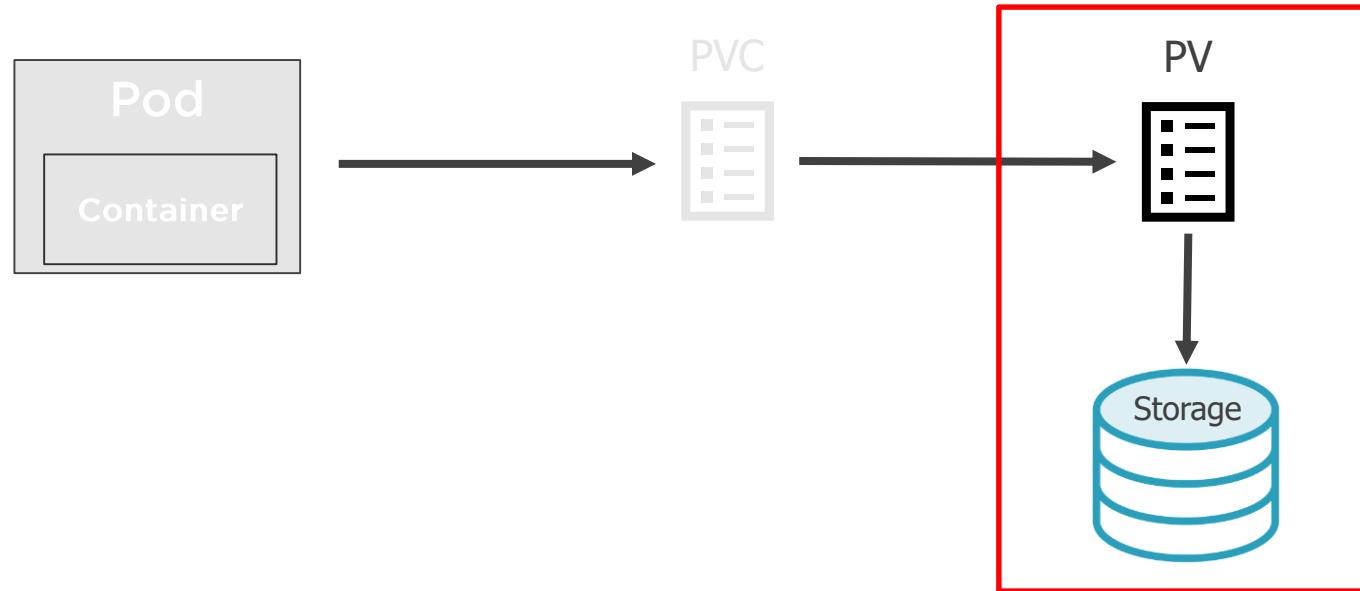
PersistentVolume and PersistentVolumeClaim YAML



Defining a PV and PVC



Creating a PersistentVolume



Defining a PersistentVolume for Azure

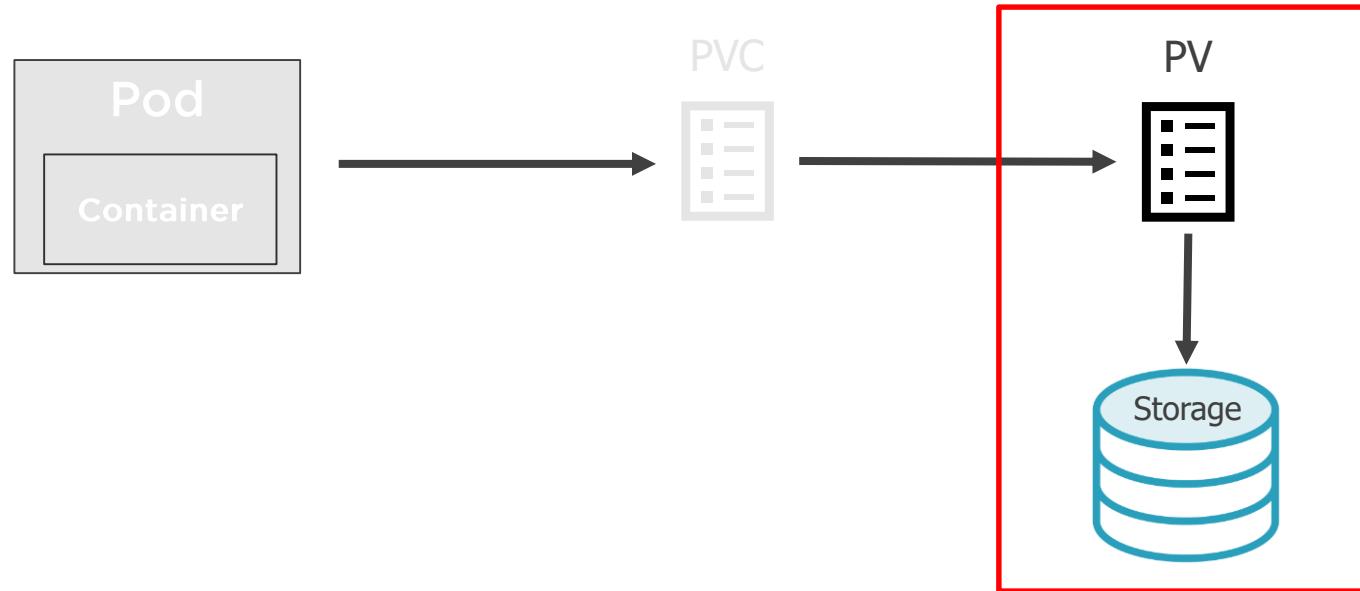
```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  capacity: 10Gi
  accessModes:
    - ReadWriteOnce
    - ReadOnlyMany
  persistentVolumeReclaimPolicy: Retain
  azureFile:
    secretName: <azure-secret>
    shareName: <name_from_azure>
    readOnly: false
```

- ◀ **Create PersistentVolume kind**
- ◀ **Define storage capacity**
- ◀ **One client can mount for read/write**
- ◀ **Many clients can mount for reading**
- ◀ **Retain even after claim is deleted (not erased/deleted)**
- ◀ **Reference storage to use (specific to Cloud provider, NFS setup, etc.)**

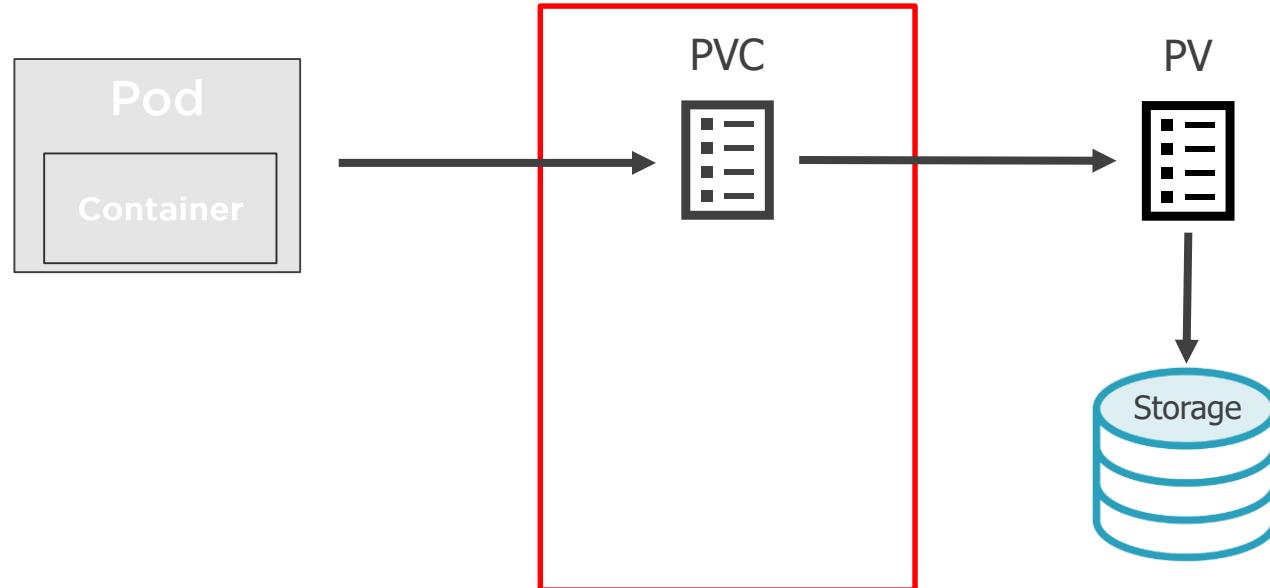
<https://github.com/kubernetes/examples>



Creating a PersistentVolume



Creating a PersistentVolumeClaim



Defining a PersistentVolumeClaim

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: pv-dd-account-hdd-5g
  annotations:
    volume.beta.kubernetes.io/storage-class: accounthdd
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
```

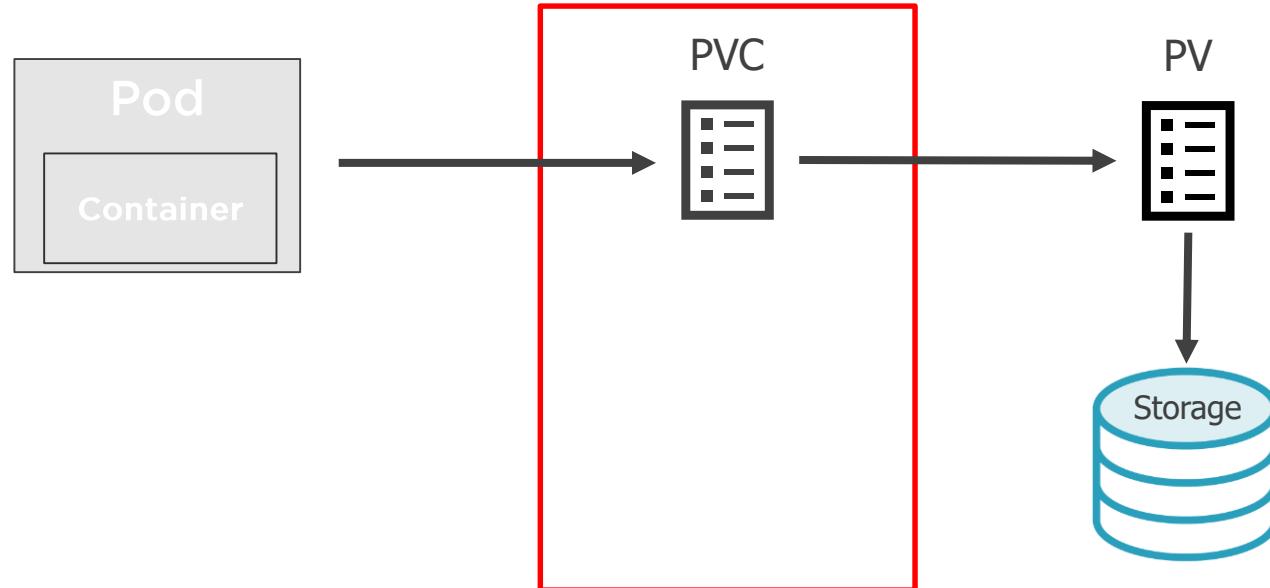
◀ Define a PersistentVolumeClaim (PVC)

◀ Define access mode

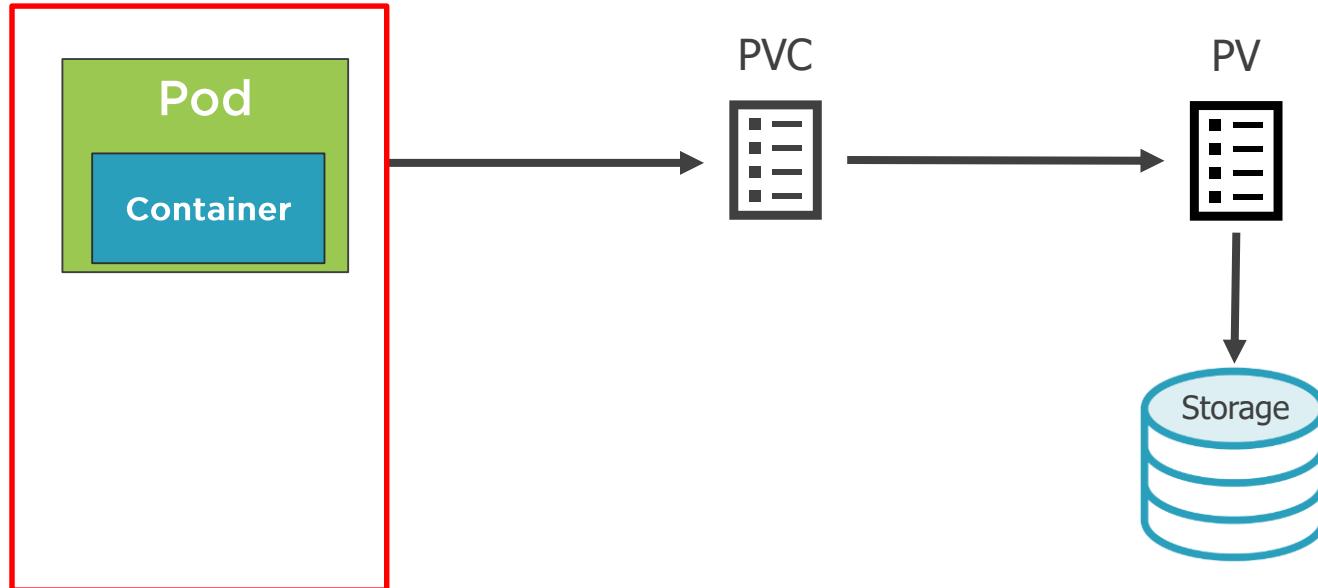
◀ Request storage amount



Creating a PersistentVolumeClaim



Defining a Volume that Uses a PVC



Using a PersistentVolumeClaim

```
kind: Pod
apiVersion: v1
metadata:
  name: pod-uses-account-hdd-5g
  labels:
    name: storage
spec:
  containers:
  - image: nginx
    name: az-c-01
    command:
    - /bin/sh
    - -c
    - while true; do echo $(date) >> /mnt/blobdisk/outfile; sleep 1; done
  volumeMounts:
  - name: blobdisk01
    mountPath: /mnt/blobdisk
volumes:
- name: blobdisk01
  persistentVolumeClaim:
    claimName: pv-dd-account-hdd-5g
```

◀ Mount to Volume

◀ Create Volume that binds to
PersistentVolumeClaim



StorageClasses

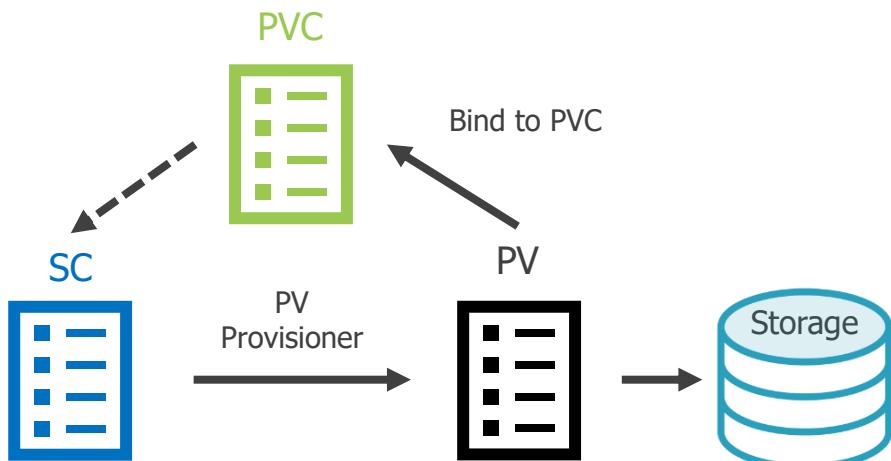


A StorageClass (SC) is a type of storage template that can be used to dynamically provision storage.

PVC



StorageClass



Used to define different "classes" of storage

Act as a type of storage template

Supports dynamic provisioning of PersistentVolumes

Administrators don't have to create PVs in advance

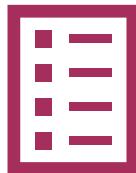


StorageClass Workflow

1

Create Storage Class

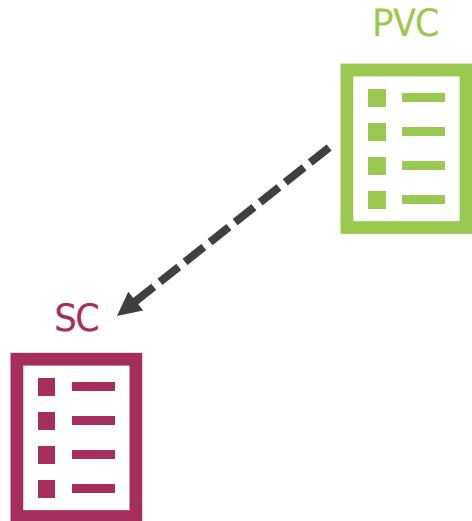
SC



StorageClass Workflow

2

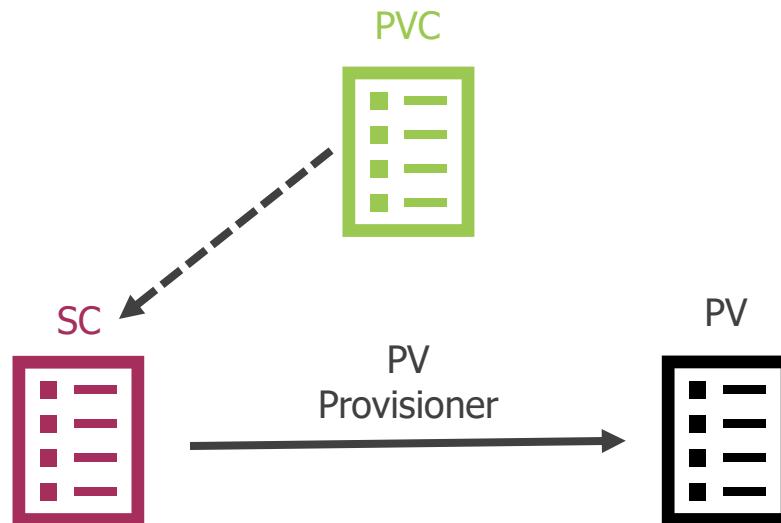
Create PersistentVolumeClaim that references StorageClass



StorageClass Workflow

3

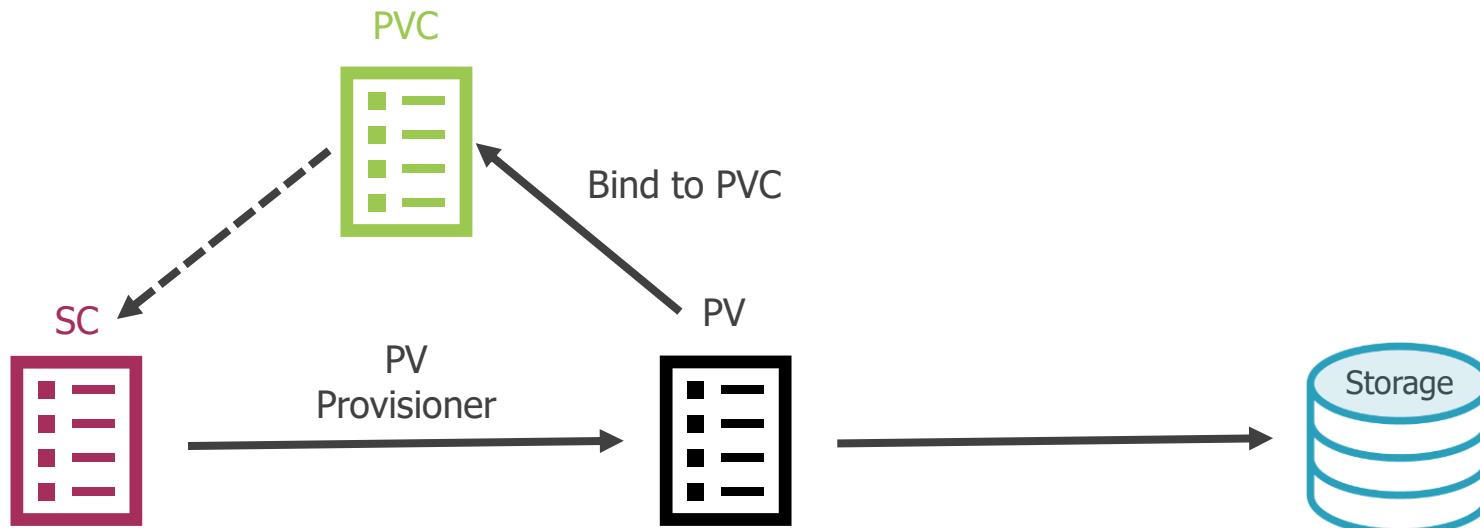
Kubernetes uses **StorageClass** provisioner
to provision a **PersistentVolume**



StorageClass Workflow

4

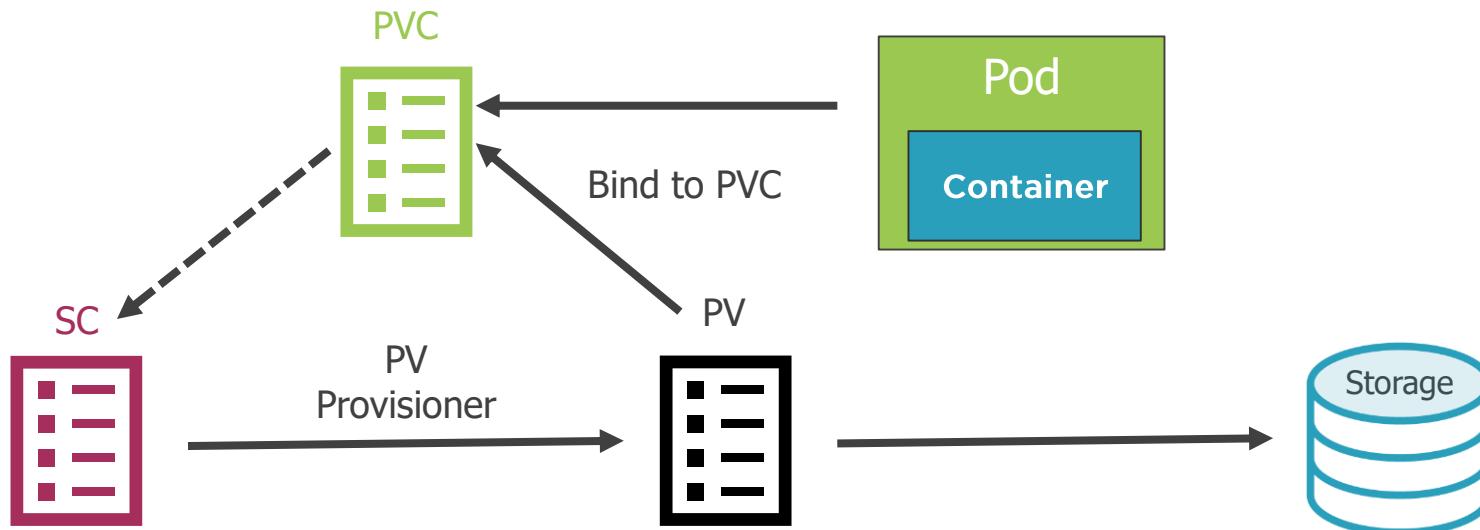
Storage provisioned, PersistentVolume created and bound to PersistentVolumeClaim



StorageClass Workflow

5

Pod volume references PersistentVolumeClaim



Defining a Local Storage StorageClass

```
apiVersion: storage.k8s.io/v1
```

```
kind: StorageClass
```

```
metadata:
```

```
  name: local-storage
```

```
reclaimPolicy: Retain
```

```
provisioner: kubernetes.io/no-provisioner
```

```
volumeBindingMode: WaitForFirstConsumer
```

- ◀ API version

- ◀ A StorageClass resource

- ◀ Retain storage or Delete (default) after PVC is released

- ◀ Provisioner (volume plugin) that will be used to create PersistentVolume resource.

- ◀ Wait to create until Pod making PVC is created. Default is Immediate (create once PVC is created)



Defining a Local Storage PersistentVolume

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  capacity:
    storage: 10Gi
  volumeMode: Block
  accessModes:
    - ReadWriteOnce
  storageClassName: local-storage
  local:
    path: /data/storage
  nodeAffinity:
    required:
      nodeSelectorTerms:
        - matchExpressions:
          - key: kubernetes.io/hostname
            operator: In
            values:
              - <node-name>
```

- ◀ One client can mount for read/write
- ◀ Reference StorageClass
- ◀ Path where data is stored on Node
- ◀ Select the Node where the local storage PV is created



Defining a PersistentVolumeClaim

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
  - ReadWriteOnce
  storageClassName: local-storage
resources:
  requests:
    storage: 1Gi
```

◀ Define a PersistentVolumeClaim (PVC)

◀ Access Mode and storage classification
PV needs to support

◀ Storage request information



Using a PersistentVolumeClaim

```
apiVersion: apps/v1  
kind: [Pod | StatefulSet | Deployment]  
...  
spec:  
volumes:  
- name: my-volume  
persistentVolumeClaim:  
claimName: my-pvc
```

◀ Define a Volume

◀ Use a PVC to claim the required storage



PersistentVolumes in Action



Summary



Kubernetes supports several different types of storage:

- Ephemeral storage (emptyDir)
- Persistent storage (many options)
- PersistentVolumes,
PersistentVolumeClaims, and
StorageClasses
- ConfigMaps (key/value pairs)
- Secrets



Creating ConfigMaps and Secrets



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

ConfigMaps Core Concepts

Creating a ConfigMap

Using a ConfigMap

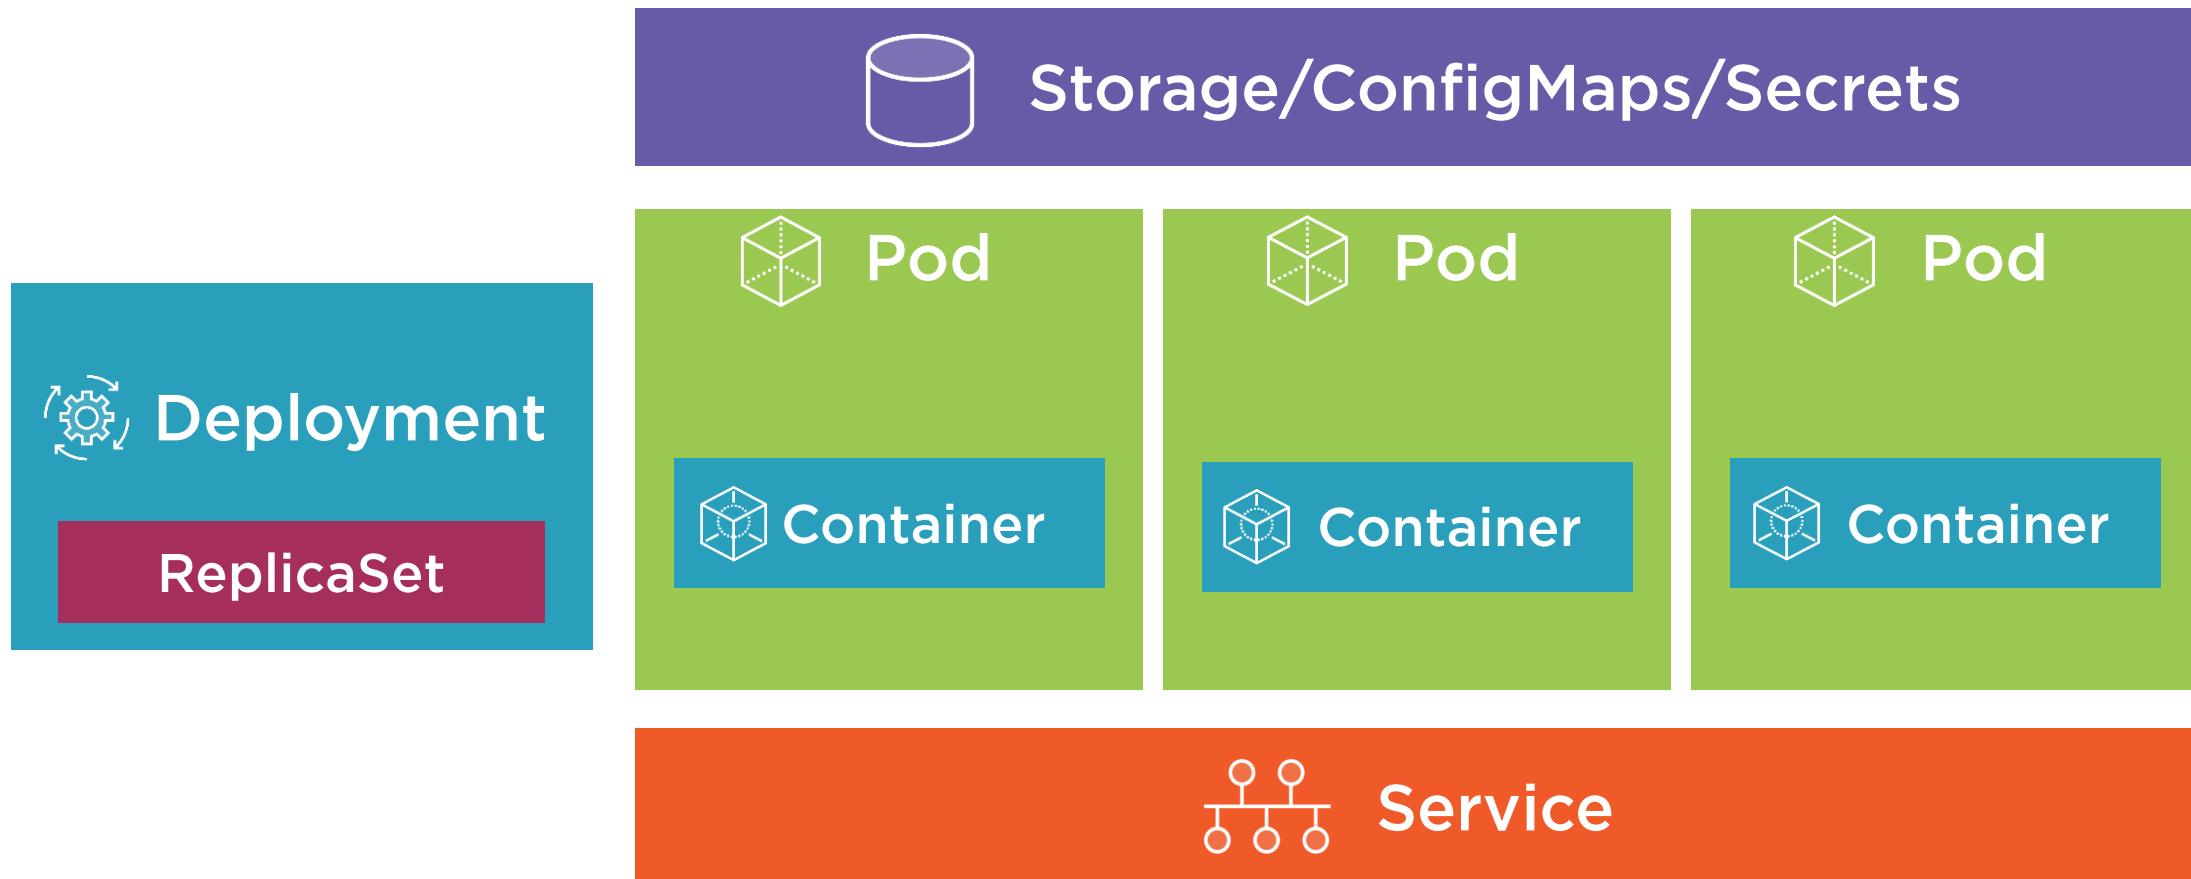
Secrets Core Concepts

Creating a Secret

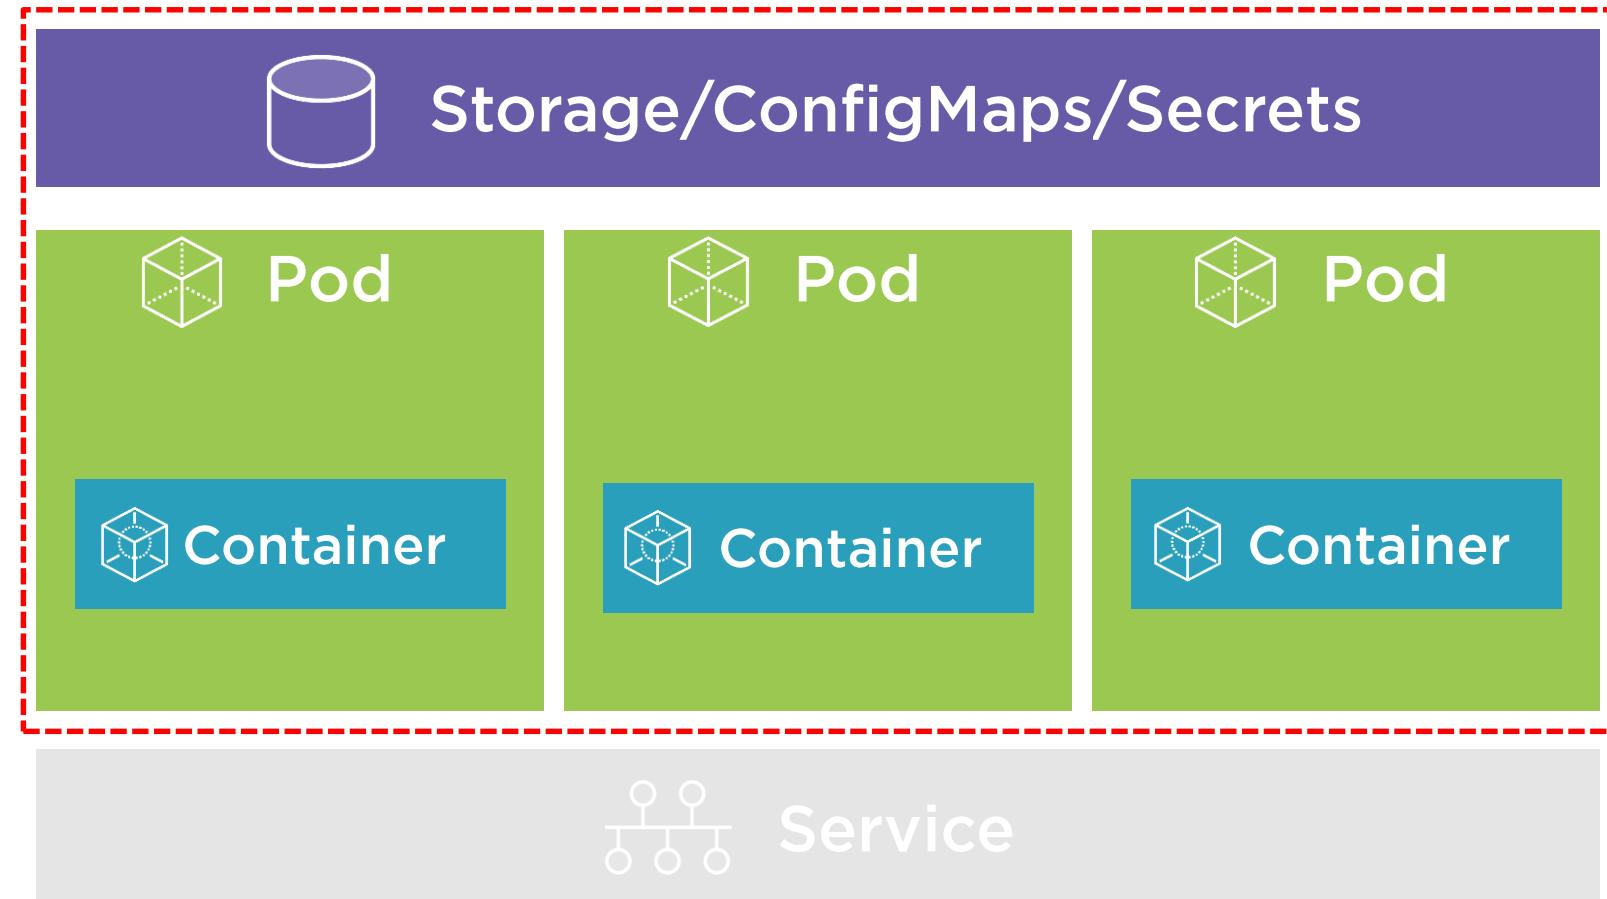
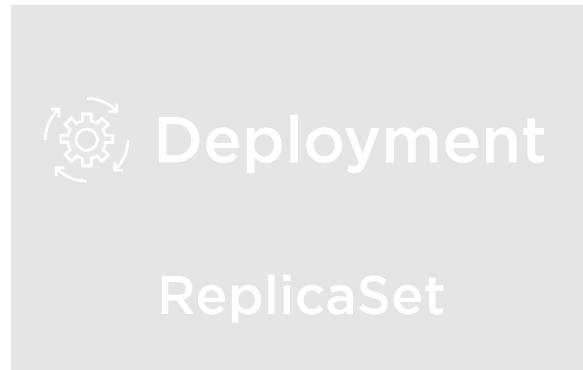
Using a Secret



You Are Here



You Are Here



ConfigMaps Core Concepts



ConfigMaps provide a way to store configuration information and provide it to containers.



Provides a way to inject configuration data into a container

Can store entire files or provide key/value pairs:

- Store in a File. Key is the filename, value is the file contents (can be JSON, XML, keys/values, etc.).
- Provide on the command-line
- ConfigMap manifest

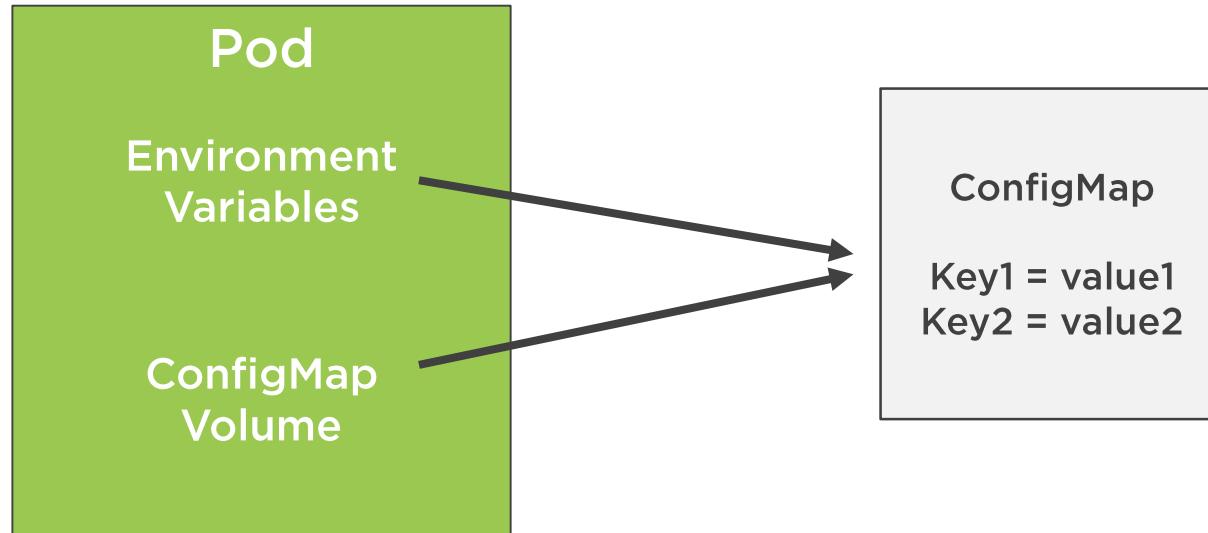
ConfigMaps



Accessing ConfigMap Data in a Pod

ConfigMaps can be accessed from a Pod using:

- **Environment variables (key/value)**
- **ConfigMap Volume (access as files)**



Creating a ConfigMap



Defining Values in a ConfigMap Manifest

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-settings
  labels:
    app: app-settings
data:
  enemies: aliens
  lives: "3"
  enemies.cheat: "true"
  enemies.cheat.level=noGoodRotten
```

```
# Create from a ConfigMap manifest
kubectl create -f file.configmap.yml
```

◀ A ConfigMap resource

◀ Name of ConfigMap

◀ ConfigMap data



Defining Key/Value Pairs in a File

```
enemies=aliens  
lives=3  
enemies.cheat=true  
enemies.cheat.level=noGoodRotten
```

```
# Create a ConfigMap using data from a file  
kubectl create configmap [cm-name]  
--from-file=[path-to-file]
```

```
apiVersion: v1  
kind: ConfigMap  
data:  
  game.config: |-  
    enemies=aliens  
    lives=3  
    enemies.cheat=true  
    enemies.cheat.level=noGoodRotten
```

- ◀ Key/value pairs defined in a file named game.config
- ◀ Nested properties can be defined and assigned a value

- ◀ Note that the file name is used as the key for the values
- ◀ Your application can now work with the content just as it would a normal configuration file (JSON, XML, keys/values, could be used)



Defining Key/Value Pairs in an Env File

```
enemies=aliens  
lives=3  
enemies.cheat=true  
enemies.cheat.level=noGoodRotten
```

- ◀ Key/value pairs can be defined in an "environment" variables file (game-config.env)
- ◀ Nested properties can be defined and assigned a value

```
# Create a env ConfigMap using data from a file  
kubectl create configmap [cm-name]  
--from-env-file=[path-to-file]
```

```
apiVersion: v1  
kind: ConfigMap  
data:  
  enemies=aliens  
  lives=3  
  enemies.cheat=true  
  enemies.cheat.level=noGoodRotten
```

- ◀ Note that the file name is NOT included as a key



```
# Create a ConfigMap using data from a config file
kubectl create configmap [cm-name] --from-file=[path-to-file]

# Create ConfigMap from an env file
kubectl create configmap [cm-name] --from-env-file=[path-to-file]

# Create a ConfigMap from individual data values
kubectl create configmap [cm-name]
  --from-literal=apiUrl=https://my-api
  --from-literal=otherKey=otherValue

# Create from a ConfigMap manifest
kubectl create -f file.configmap.yml
```

Creating a ConfigMap

A ConfigMap can be created using **kubectl create**

Key command-line switches include:

--from-file

--from-env-file

--from-literal



Using a ConfigMap



```
# Get a ConfigMap  
kubectl get cm [cm-name] -o yaml
```

Getting a ConfigMap

kubectl get cm can be used to get a ConfigMap and view its contents



Accessing a ConfigMap: Environment Vars

**Pods can access ConfigMap values through environment vars
ENEMIES environment variable created (value=aliens)**

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-settings
data:
  enemies: aliens
  lives: "3"
  enemies.cheat: "true"
  enemies.cheat.level=noGoodRotten
```

```
apiVersion: apps/v1
...
spec:
  template:
    ...
      spec:
        containers: ...
        env:
          - name: ENEMIES
            valueFrom:
              configMapKeyRef:
                name: app-settings
                key: enemies
```

**Environment
variable name**



Accessing a ConfigMap: Environment Vars

envFrom can be used to load all ConfigMap keys/values into environment variables

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-settings
data:
  enemies: aliens
  lives: "3"
  enemies.cheat: "true"
  enemies.cheat.level=noGoodRotten
```

```
apiVersion: apps/v1
...
spec:
  template:
    ...
      spec:
        containers: ...
          envFrom:
            - configMapRef:
                name: app-settings
```

Environment
variables created
for all data keys



Accessing a ConfigMap: Volume

ConfigMap values can be loaded through a Volume

Each key is converted to a file - value is added into the file

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-settings
data:
  enemies: aliens
  lives: "3"
  enemies.cheat: "true"
  enemies.cheat.level=noGoodRotten
```

```
apiVersion: apps/v1
...
spec:
  template:
    ...
      spec:
        volumes:
          - name: app-config-vol
            configMap:
              name: app-settings
        containers:
          volumeMounts:
            - name: app-config-vol
              mountPath: /etc/config
```

ConfigMap values stored
at /etc/config



ConfigMaps in Action



Secrets Core Concepts



A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key.



Secrets



Kubernetes can store sensitive information (passwords, keys, certificates, etc.)

Avoids storing secrets in container images, in files, or in deployment manifests

Mount secrets into pods as files or as environment variables

Kubernetes only makes secrets available to Nodes that have a Pod requesting the secret

Secrets are stored in tmpfs on a Node (not on disk)



**Enable encryption at rest for cluster data
(<https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data>)**

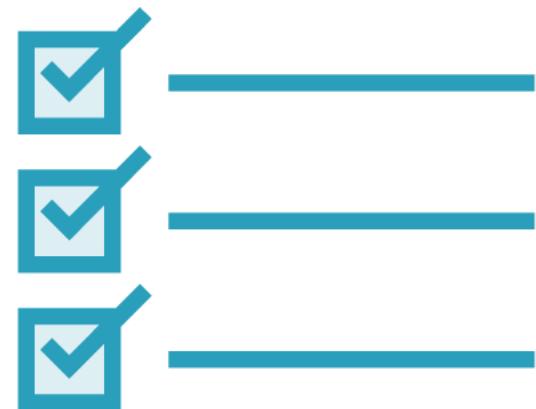
Limit access to etcd (where Secrets are stored) to only admin users

Use SSL/TLS for etcd peer-to-peer communication

Manifest (YAML/JSON) files only base64 encode the Secret

Pods can access Secrets so secure which users can create Pods. Role-based access control (RBAC) can be used.

Secrets Best Practices



Creating a Secret



Creating a Secret

Secrets can be created using kubectl create secret

```
# Create a secret and store securely in Kubernetes
kubectl create secret generic my-secret
--from-literal=pwd=my-password

# Create a secret from a file
kubectl create secret generic my-secret
--from-file=ssh-privatekey=~/.ssh/id_rsa
--from-file=ssh-publickey=~/.ssh/id_rsa.pub

# Create a secret from a key pair
kubectl create secret tls tls-secret --cert=path/to/tls.cert
--key=path/to/tls.key
```

Question:

Can I declaratively define secrets using YAML?

Answer:

**Yes – but any secret data is only base64 encoded
in the manifest file!**



Defining a Secret in YAML

```
apiVersion: v1
kind: Secret
metadata:
  name: db-passwords
type: Opaque
data:
  app-password: cGFzc3dvcmQ=
  admin-password: dmVyeV9zZWNyZXQ=
```

◀ Define a Secret

◀ Secret name

◀ Keys/values for Secret



Using a Secret



```
# Get secrets  
kubectl get secrets
```



A terminal window titled "danwahlin — bash — 73x6" showing the output of the command "k get secrets". The output lists two secrets: "db-passwords" and "default-token-rxmjb". The "db-passwords" secret is of type "Opaque" and has 2 data bytes, created 34m ago. The "default-token-rxmjb" secret is of type "kubernetes.io/service-account-token" and has 3 data bytes, created 66d ago.

NAME	TYPE	DATA	AGE
db-passwords	Opaque	2	34m
default-token-rxmjb	kubernetes.io/service-account-token	3	66d

```
# Get YAML for specific secret  
kubectl get secrets db-passwords -o yaml
```



A terminal window titled "danwahlin — bash — 73x16" showing the output of the command "k get secrets db-passwords -o yaml". The output is a YAML representation of the "db-passwords" secret. It includes fields for apiVersion, data (containing two password entries), kind, metadata (creationTimestamp, name, namespace, resourceVersion, selfLink, uid), and type.

```
apiVersion: v1  
data:  
  mongodb-password: cGFzc3dvcmQ=  
  mongodb-root-password: cGFzc3dvcmQ=  
kind: Secret  
metadata:  
  creationTimestamp: "2019-03-22T00:40:05Z"  
  name: db-passwords  
  namespace: default  
  resourceVersion: "3481795"  
  selfLink: /api/v1/namespaces/default/secrets/db-passwords  
  uid: 0982413e-4c3b-11e9-b7f0-025000000001  
type: Opaque
```

Listing Secret Keys

A list of secrets can be retrieved using **kubectl get secrets**



Accessing a Secret: Environment Vars

Pods can access Secret values through environment vars

DATABASE_PASSWORD environment var created

```
apiVersion: v1
```

```
kind: Secret
```

```
metadata:
```

```
  name: db-passwords
```

```
type: Opaque
```

```
data:
```

```
  db-password: cGFzc3dvcmQ=
```

```
  admin-password: dmVyeV9zZWNyZXQ=
```

```
apiVersion: apps/v1
```

```
...
```

```
spec:
```

```
  template:
```

```
  ...
```

```
spec:
```

```
  containers: ...
```

```
env:
```

```
  - name: DATABASE_PASSWORD
```

```
    valueFrom:
```

```
      secretKeyRef:
```

```
        name: db-passwords
```

```
        key: db-password
```



Accessing a Secret: Volumes

Pods can access secret values through a volume

Each key is converted to a file - value is added into the file

```
apiVersion: v1
kind: Secret
metadata:
  name: db-passwords
type: Opaque
data:
  db-password: cGFzc3dvcmQ=
  admin-password: dmVyeV9zZWNyZXQ=
```

```
apiVersion: apps/v1
...
spec:
  template:
    ...
spec:
  volumes:
    - name: secrets
      secret:
        secretName: db-passwords
  containers:
    volumeMounts:
      - name: secrets
        mountPath: /etc/db-passwords
        readOnly: true
```



Secrets in Action



Summary



ConfigMaps provide a way to store configuration data

Secrets provide a way to store sensitive data or files

Access key/value pairs using environment variables or volumes

Use caution when working with Secrets and ensure proper security is in place



Putting It All Together



Dan Wahlin

WAHLIN CONSULTING

@danwahlin www.codewithdan.com



Module Overview

Application Overview

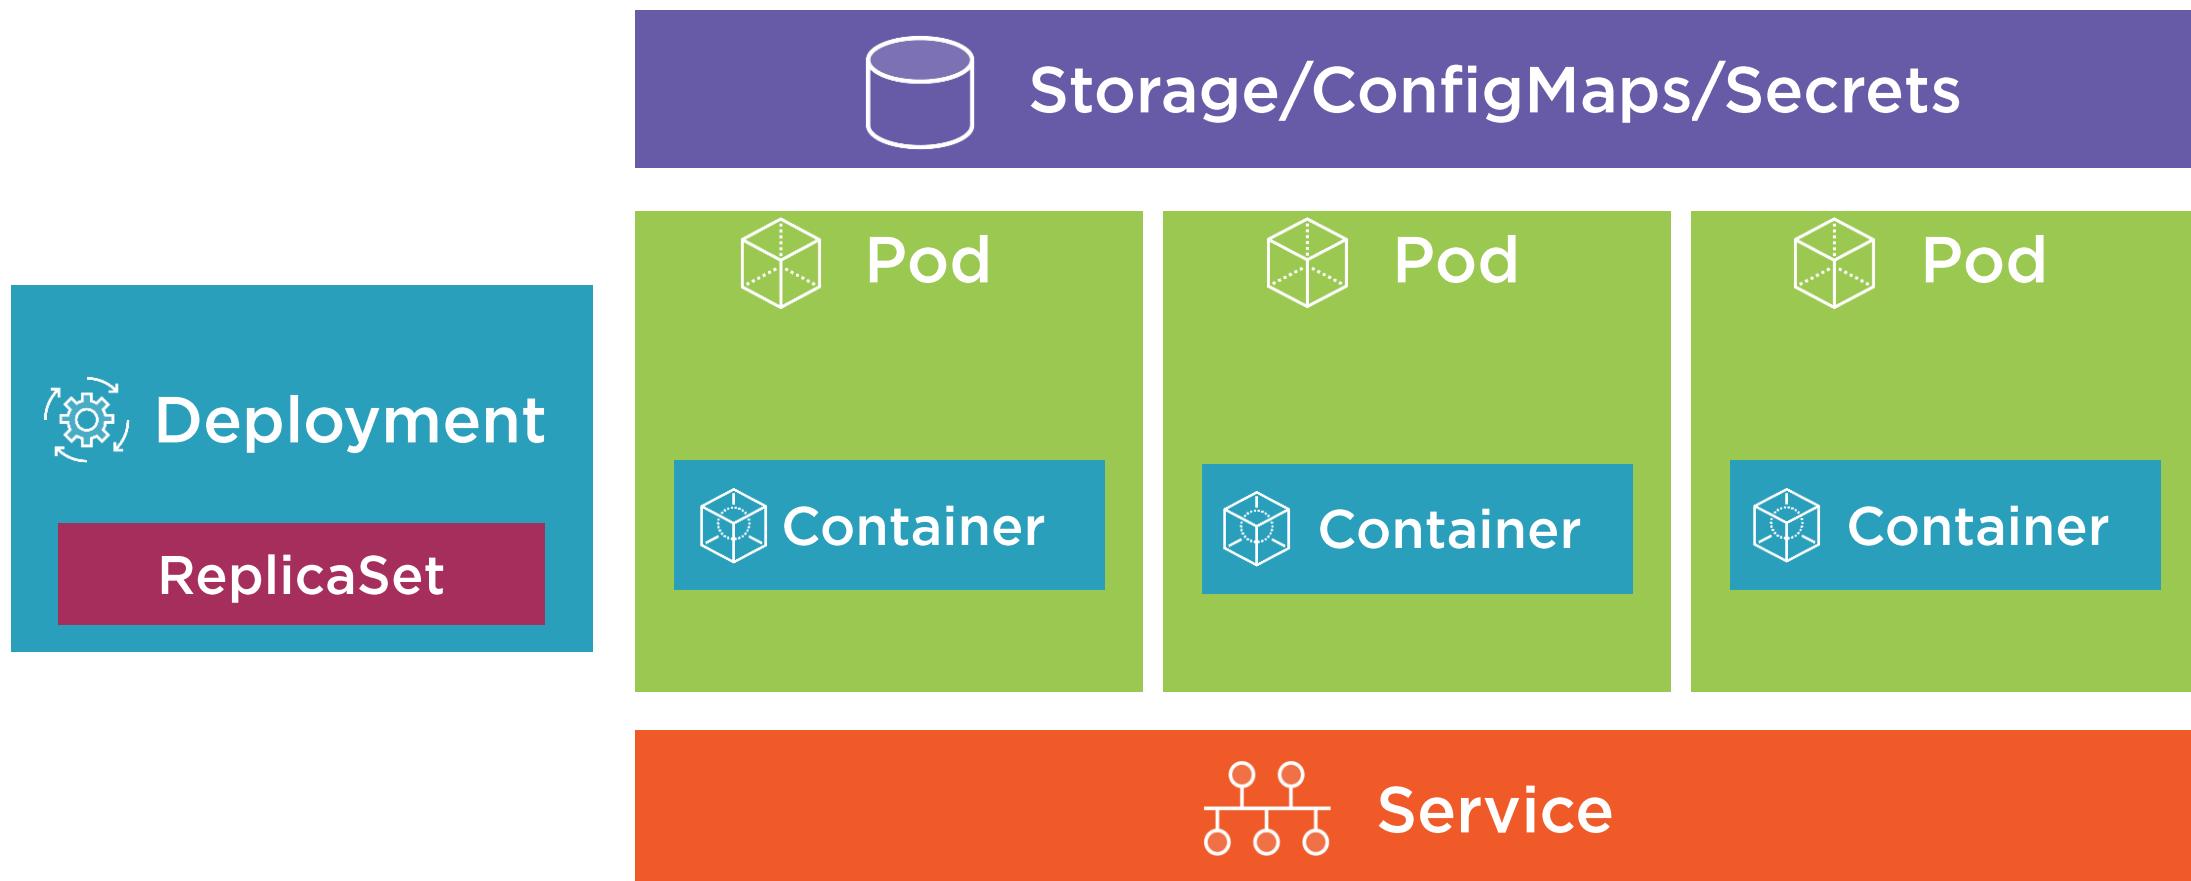
YAML Manifests

Running the Application

Troubleshooting Techniques



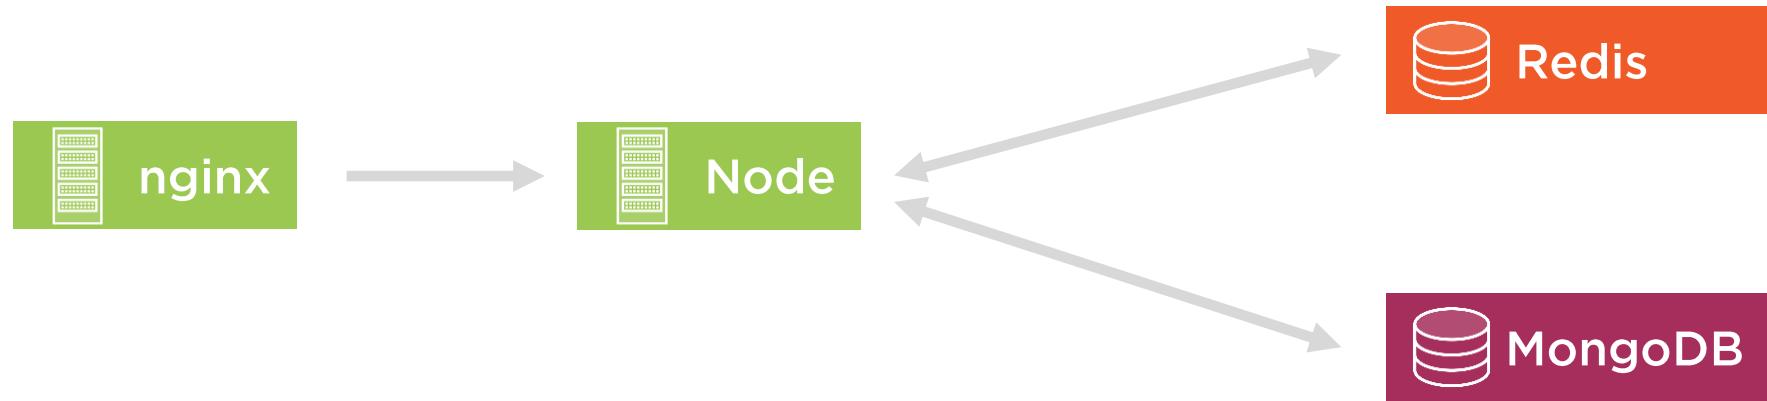
You Are Here



Application Overview



Application Overview



YAML Manifests



Running the Application



Troubleshooting Techniques



```
# View the logs for a Pod's container  
kubectl logs [pod-name]  
  
# View the logs for a specific container within a Pod  
kubectl logs [pod-name] -c [container-name]  
  
# View the logs for a previously running Pod  
kubectl logs -p [pod-name]  
  
# Stream a Pod's logs  
kubectl logs -f [pod-name]
```

Accessing Logs

Pod container logs can be viewed using the kubectl logs command



```
# Describe a Pod  
kubectl describe pod [pod-name]
```

```
# Change a Pod's output format  
kubectl get pod [pod-name] -o yaml
```

```
# Change a Deployment's output format  
kubectl get deployment [deployment-name] -o yaml
```

Getting Details About a Pod

Get details about a Pod using kubectl describe or kubectl get with -o



```
# Shell into a Pod container  
kubectl exec [pod-name] -it sh
```

Shell into a Pod Container

Shell into a Pod container using kubectl exec



Summary



YAML manifest files are used to define different Kubernetes resources

kubectl create or apply can be used with -f to deploy multiple manifest files

Learning different Pod troubleshooting commands and techniques is important to resolve issues

